Universitat de les Illes Balears

MASTER'S THESIS

PARAMETER ESTIMATION AND SKY LOCALIZATION OF LENSED IMAGE PAIRS OF BINARY BLACK HOLE MERGERS

Xavier Garcia Sàbat

Master's Degree in Advanced Physics and Applied Mathematics

Centre for Postgraduate Studies

Academic Year 2025-26

PARAMETER ESTIMATION AND SKY LOCALIZATION OF LENSED IMAGE PAIRS OF BINARY BLACK HOLE MERGERS

Xavier Garcia Sàbat

Master's Thesis

Centre for Postgraduate Studies

University of the Balearic Islands

Academic Year 2025-26

Keywords:

Gravitational waves, Strong gravitational lensing, Parameter estimation, Sky localization, Binary black holes, LIGO

Thesis Supervisor's Name: David Benjamin Keitel

Acknowledgments

Aquests agraïments són per totes i cadascuna de les persones que han contribuït al fet que hagi pogut arribar fins a on he arribat.

Aquest últim any he tingut un debat amb diverses persones, debatíem sobre quan pots considerar que ets 'físic', i en concret, astrofísic. Quan et gradues? Quan acabes el màster en una especialització? O quan treballes de 'físic'? Les opinions són variades, però el fet de contribuir amb la comunitat científica fa que no et calgui cap títol per saber fins on has arribat. I això és el que he fet aquest últim any. Sigui astrofísic (computacional) o no encara, sé que he complert el meu somni, aquell somni idílic d'un Xavier il·lusionat, innocent i molt motivat, el d'investigar en astrofísica.

Però aquest somni, que s'ha convertit en el meu present, viu única i exclusivament gràcies a la meva família, la meva parella, amics, tutors, i totes les persones que m'han donat força i motivació per seguir cada dia, i en especial m'agradaria mencionar-ne unes quantes.

Mama, Papa, Sabina, Martí, sou el meu pilar, el meu exemple, i els meus guies.

Mama, no sé si recordes que tu ets la culpable que sigui (o estigui a punt de ser) astrofísic. Quan tenia 16 anys i no sabia què volia estudiar a la universitat (però ja m'interessava una mica l'univers), em vas regalar un llibre del National Geographic sobre forats negres. Un llibre que va ser el meu preferit durant molt de temps i sempre guardaré com un tresor. Gràcies a aquell llibre, gràcies a tu, em vaig apassionar per l'astrofísica, i des d'aquell moment no ha passat ni un sol dia en què no m'hagis donat suport en les meves decisions i no t'hagis interessat pels meus interessos. No oblidaré mai veure la teva cara d'il·lusió quan t'explicava els meus èxits i somnis. Gràcies mama, gràcies per tot el que m'has donat sense demanar-me res a canvi. Només puc desitjar que estiguis orgullosa de fins a on he arribat i que algun dia et pugui compensar per tot el que m'has donat.

Papa, el meu exemple, tant de bo fos una desena part del que ets tu. Igual que la mama, el teu suport sempre ha estat més que incondicional, molt més del que a vegades em mereixia. La teva confiança en mi ha superat, la majoria de cops, la que jo tenia per mi. I això m'ha obligat a demostrar-me que tenies raó i que podia aconseguir el que vulgués. D'això ja fa anys, però mai oblidaré el temps en què et preguntaves si eres el culpable que treiés males notes de mates i física (de llengua sempre m'ha anat bé, si no no seria fill teu), no només mai m'has pressionat per les notes sinó que a més et feies responsable del que no et corresponia. Avui li dic al meu pare de fa anys i al d'ara que sí, que ell és el responsable, el responsable que hagi aconseguit el meu somni. Gràcies papa, algun dia trobaré la manera de compensar-te com a la mama.

Martí, Sabina, qui seria sense vosaltres... Gràcies per donar-me humanitat en un món que sovint et fa ser fred i calculador, i fer-me recordar perquè vivim. També m'heu donat identitat, valors, felicitat, humiltat... No heu deixat mai que em desviés del camí i que oblidés qui sóc. Necessitaria tres vides senceres per poder-vos tornar tot el que m'heu donat i les infinites vegades que m'heu donat ànims o m'heu felicitat per qualsevol èxit. Tenir germans sempre valdrà molt més que qualsevol somni. Tan sols me'n penedeixo de no haver-vos dedicat prou temps i escoltar-vos més, espero ser-hi a temps. Gràcies germans, per ser qui sou, i per fer que sigui qui sóc.

Paula, si hi ha una persona a la qual dec haver aconseguit arribar fins a on sóc aquesta ets tu. Des de gairebé el primer dia que ens vam conèixer, em vas ajudar amb els incomptables problemes que tenia per entendre totes les assignatures de física. Sempre has estat molt pacient, agradable, comprensible i, per descomptat, molt més intel·ligent que jo. Sense cap dubte, sense la teva ajuda no hagués estat capaç de graduar-me o com a mínim hagués tardat molt més temps i m'hagués desesperat molt més. També has fet que viure lluny de la família durant 5 anys s'hagi fet molt menys dur, el fet de poder estar amb tu i parlar del que fos m'ha fet ser molt feliç. Sempre estaré infinitament agraït per tot el que m'has donat i per permetre'm complir dos somnis, el de ser físic, i el de conèixer una persona com tu. T'estimo molt.

Tiets, avis, cosins, padrins, tots formeu part del meu passat, present i futur, i no podria estar més orgullós de tenir-vos a tots i cadascun de vosaltres. A part de complir somnis i objectius, la vida també és passar temps amb la teva família: àpats familiars, excursions, viatges... qualsevol moment amb vosaltres formarà part dels meus records per sempre. Sense vosaltres hagués estat impossible seguir endavant cada dia. Gràcies família, per interessar-vos sempre pels meus èxits i no tan èxits, i pel temps que m'heu dedicat.

Amics, si sense família complir un somni seria impossible, sense amics no tindria sentit. Us dec molt per les infinites (o no tan infinites últimament per culpa meva) hores de converses i celebracions. Poder comptar amb vosaltres per qualsevol problema o dubte és i serà sempre el millor tresor. Tornar a Catalunya de tant en tant i desconnectar de la física ha estat un suport molt gran i us dec molt. Perdoneu per no dedicar-vos el temps que us mereixeu, però no hi ha dia que passi que no pensi en vosaltres.

Gràcies també en especial als meus companys de viatge i de grup d'investigació, sense vosaltres he de dir que seria molt més avorrit. Poder parlar cada dia amb vosaltres i desfogar-nos del que sigui ha estat sense cap dubte una gran motivació. No hi ha res que em motivi més que poder començar el doctorat juntament amb vosaltres. Gràcies Raúl, Sofía, Unai, Andrés, Marian, Jesús, Ángel, Alejandro, Samuel, Joan-René, Alicia C., Felip, Joan, Maria, Jorge i Lorenzo. Us desitjo tota la sort del món. And thanks of course to all members of the GRAVITY Física Gravitacional group that I haven't mentioned, you have all helped me in one way or another and I am very grateful.

Voglio anche ringraziare in modo speciale un compagno di gruppo 'adottato', Francesco. Questi tre mesi che hai passato qui sono volati come tre settimane grazie a te, ma allo stesso tempo sono sembrati tre anni di amicizia. Hai sopportato mille volte le conversazioni in cui ti raccontavo tutti i problemi e i dubbi che avevo, e mi hai sempre ascoltato, sei stato comprensivo e mi hai dato ottimi consigli. Posso solo augurarmi che la vita mi porti presto a Tivoli o che ci rivediamo a qualche congresso. A presto, Francesco.

Por descontado también debo un gran agradecimiento a mis compañeros de máster del lado oscuro, del IFISC. Gracias por adoptarme aunque fuera de otro máster y por pasar tanto tiempo juntos. No me importaría hacer un año más de máster si fuera con vosotros. Sea donde sea que vayáis ahora os deseo la mejor de las suertes. Y que el futuro nos vuelva a juntar.

També es mereixen un agraïment especial els meus companys d'equip del Raimon. Heu fet que sigui un any molt especial, ple d'entrenaments, partits i estones úniques. Gràcies per tot.

M'agradaria ara poder donar els agraïments que no vaig poder donar pel meu treball final de grau als meus tutors juntament amb els agraïments del treball final de màster.

Ramon, el meu primer tutor i mentor de la universitat que va confiar plenament en mi. Encara em pregunto per què vas decidir escoltar i ajudar aquell estudiant de tercer promig que estava perdut i que només tenia motivació. Sobretot després de demanar-te 6 cartes de recomanació i que no m'agafessin enlloc... hahah. Tot i així l'any següent vas acceptar tutoritzar-me el treball de final de grau juntament amb l'Iñigo i enviar-me a l'Institut d'Astrofísica de les Illes Canàries (IAC), el qual era un dels meus somnis, poder veure els millors observatoris d'Europa i fins i tot del món. I no va ser fàcil, mai oblidaré que vau acabar signant juntament amb la Maria Antònia (a la qual també li dec moltíssim) i l'Iñigo un conveni entre la UIB i l'IAC que es va publicar al BOE... Hagués estat molt fàcil per a tu tirar-me per la finestra del teu despatx i no buscar-te problemes. Però no va ser així, vas confiar en mi i em vas donar l'oportunitat més gran que m'havien donat mai i em vas buscar el millor tutor de l'IAC i de les Canàries, l'Iñigo. Estaré sempre en deute amb tu, Ramon.

Iñigo, qué decirte que no te haya dicho ya... Otro tutor que decidió meterse en problemas burocráticos en vez de vivir tranquilo. Y no podría estar más agradecido por esta decisión y por tu esfuerzo y buena voluntad. Sin ti, sin duda, tampoco habría llegado donde he llegado. Poder cumplir el sueño de ir a las Islas Canarias y ver los mejores observatorios y además poder hacer prácticas y el TFG en el IAC con un tutor como tú sin duda fue infinitamente más de lo que podía desear o merecer. No me cansaré de darte las gracias por todo, Iñigo, sólo me queda pendiente ir a un congreso contigo, escuchar una charla y poder decir que es science fiction... Hasta pronto, Iñigo.

Vull fer també una menció especial a tots els membres del grup de física solar de la UIB, SolPhys. Em vau acollir molt ràpid i amb una qualitat humana inimaginable, aquell congrés en què em vau donar l'oportunitat de participar (IX Reunión Española de Física Solar y Heliosférica) va ser el primer congrés de la meva vida i estic segur que un dels millors. Gràcies Llorenç, Miquel, Varsha, Guillem i també Sushree i Víctor. Moltes gràcies també a en Jaume, Roberto i Manuel per tot.

Dec ara uns agraïments molt sincers a en David Keitel i l'Alicia Sintes.

David, I have nothing but good words for you. You are the third supervisor in a row who decided to trust me without me understanding why. It would have been very easy for you to set a higher grade threshold and choose someone more intelligent, but thanks to your trust and for allowing me to be your student collaborator, later giving me a paid contract in your project, and now being my Master's thesis supervisor, I am now halfway to achieving another one of my dreams: getting into a PhD. Everything I have learned this last year is thanks to you. I have

learned an enormous number of concepts and techniques at a speed that I would never have imagined, and even though I still have certain limitations in terms of intelligence, I also owe you thanks for your patience. Even though you are the busiest person I know after Alicia ;) , one thing that can never be held against you is that one hour in a meeting with you is like 100 hours with anyone else. Your efficiency is something to admire. But I don't want you to think that I will lower my guard just because I have made it this far. Now is when I have to prove that you were not wrong and give it my very best.

Alicia, com comença a ser costum a la meva vida, quarta tutora que decideix confiar en mi tenint infinits candidats millors que jo per poder triar. Començo a pensar que la meva família us envia sobrets amb diners... Quan semblava que no podria fer doctorat aquest any, de cop vaig rebre un missatge inesperat d'en David dient que tu em podries tutoritzar el doctorat. Gràcies a això, les meves possibilitats de fer doctorat han augmentat exponencialment. La teva confiança em dona una motivació més alta de la que ja tinc (si és possible). Tal com li he dit a en David, no vull que et pensis que ara em relaxaré, al contrari. Ara el que em toca fer és demostrar que no us heu equivocat de decisió i ho penso aconseguir. Gràcies Alicia, per la teva humanitat i sinceritat.

A totes les persones que no he mencionat, que sapigueu que penso en vosaltres cada dia, i tinc molt present tot el que m'heu aportat, i estaré eternament agraït.

This work was supported by the Universitat de les Illes Balears (UIB) with funds from the Programa de Foment de la Recerca i la Innovació de la UIB 2024-2026 (supported by the yearly plan of the Tourist Stay Tax ITS2023-086); the Spanish Agencia Estatal de Investigación grants PID2022-138626NB-I00, RED2024-153978-E, RED2024-153735-E, funded by MICIU/AEI/10.13039/501100011033 and the ERDF/EU; and the Comunitat Autònoma de les Illes Balears through the Conselleria d'Educació i Universitats with funds from the European Union – European Regional Development Fund (ERDF) (SINCO2022/18146 - Plataforma HiTech-IAC3-BIO).

Xavier Garcia Sàbat is supported by the Govern de les Illes Balears through the TEC 2025 grant call for the recruitment of technical research support personnel in the Balearic Islands (reference TEC2025_13); earlier stages of this work were supported by the project "Novel techniques for the next first detections in gravitational wave astrophysics" (reference CNS2022-135440), funded by MICINN and the European Union through the NextGenerationEU / Plan de Recuperación, Transformación y Resiliencia.

This work used the computing resources of the LIGO Laboratory computing cluster at Caltech (CIT), supported by National Science Foundation Grants PHY-0757058 and PHY-0823459.

This research has made use of data or software obtained from the Gravitational Wave Open Science Center (https://gwosc.org [1]), a service of the LIGO Scientific Collaboration, the Virgo Collaboration, and KAGRA. This material is based upon work supported by NSF's LIGO Laboratory which is a major facility fully funded by the National Science Foundation, as well as the Science and Technology Facilities Council (STFC) of the United Kingdom, the Max-Planck-Society (MPS), and the State of Niedersachsen/Germany for support of the construction of Advanced LIGO and construction and operation of the GEO600 detector. Additional support for Advanced LIGO was provided by the Australian Research Council. Virgo is funded, through the European Gravitational Observatory (EGO), by the French Centre National de Recherche

Scientifique (CNRS), the Italian Istituto Nazionale di Fisica Nucleare (INFN) and the Dutch Nikhef, with contributions by institutions from Belgium, Germany, Greece, Hungary, Ireland, Japan, Monaco, Poland, Portugal, Spain. KAGRA is supported by Ministry of Education, Culture, Sports, Science and Technology (MEXT), Japan Society for the Promotion of Science (JSPS) in Japan; National Research Foundation (NRF) and Ministry of Science and ICT (MSIT) in Korea; Academia Sinica (AS) and National Science and Technology Council (NSTC) in Taiwan.

The authors thankfully acknowledge the computer resources at Picasso and the technical support provided by the Barcelona Supercomputing Center (BSC) through grants No. AECT-2025-1-0018, AECT-2025-2-0030 and AECT-2025-3-0007 from the Red Española de Supercomputación (RES). The authors thank the Supercomputing and Bioinnovation Center (SCBI) of the University of Malaga for their provision of computational resources and technical support (https://www.scbi.uma.es/site).

We thank Neha Singh, Hemantakumar Phurailatpam, and Angel Garrón for their generous help with cluster usage, code implementation, and troubleshooting the many errors encountered along the way. We also thank Alvin K. Y. Li and Otto A. Hannuksela for clarifying aspects of [2].

DCC: LIGO-P2600393

Summary

Around 1920, a journalist asked the British astrophysicist Arthur Eddington whether it was true that only three people in the world understood general relativity (GR). Eddington paused for a moment and replied: "I am trying to think who the third person is." The anecdote may not be true, but Eddington's claim was not. Within months of Einstein publishing his field equations in 1915, Schwarzschild's black hole solution was already found; the de Sitter and anti-de Sitter cosmologies (1917) and Friedmann's expanding universe (1922) followed within a few years. The theory proved far more accessible than its reputation suggested [3].

That reputation, however, is not entirely undeserved. GR rests on tensor analysis and differential geometry, mathematics most physics students only meet when they need it. Yet beneath the formalism lies a surprisingly simple physical idea: spacetime is a dynamical entity, shaped by matter and energy, and what we call gravity is nothing more than the curvature of that entity. As John Wheeler famously put it, "spacetime tells matter how to move; matter tells spacetime how to curve" [4].

One hundred and ten years after the publication of the field equations, GR has passed every experimental test devised. Among its most striking predictions are gravitational waves (GWs), ripples in the fabric of spacetime produced by the most violent events in the universe. First detected in September 2015 by the Laser Interferometer Gravitational-wave Observatory (LIGO) detectors and announced jointly by the LIGO Scientific Collaboration and the Virgo Collaboration [5], they opened an entirely new observational window onto the cosmos, one based not on light but on the oscillation of spacetime itself. In the decade since, more than three hundred compact binary mergers have been detected, and the field is entering what can only be described as its golden age.

This thesis sits at the intersection of two of the richest areas in modern GW astronomy: strong gravitational lensing and sky localization. When a GW passes near a massive galaxy on its way to Earth, its path is deflected, and it can arrive at the detector as multiple copies of the same signal, separated in time and amplitude. No confirmed lensed GW event has been observed to date, but forecasts for the upcoming O5 observing run predict more than one such event per year. Identifying these lensed image pairs, and analyzing them jointly rather than in isolation, opens the possibility of dramatically improving our knowledge of where on the sky the source lies, a measurement that with only two detectors operating is often poorly constrained.

This is the question this thesis sets out to answer: can the joint analysis of two lensed images

of the same binary black hole (BBH) merger meaningfully improve sky localization, and by how much? To answer it, a population of 500 strongly lensed BBH events detectable by the H1 and L1 detectors under projected O5a sensitivity is generated with the LeR framework, and a realistic detector duty cycle is applied to obtain a final injection set of 157 events, of which 150 were successfully carried through the full analysis. Each image is analyzed individually with Bilby using the `IMRPhenomXPNR` approximant, and the two images are then analyzed together with Hanabi, sharing a coherent likelihood across both. The joint analysis reduces the median 90% credible sky area from 571 deg^2, obtained from the better of the two single-image analyses, to 99.9 deg^2, a typical improvement by a factor of ~ 5, and brings 3.3% of events below the 10 deg^2 threshold, a regime that no single-image analysis reaches in this two-detector configuration.

This thesis is organized as follows:

Chapter 1: Introduction to general relativity. Develops the theoretical foundations of GR, from the equivalence principle and the mathematics of curved spacetime to the Einstein field equations and the Schwarzschild solution.

Chapter 2: Gravitational wave theory. Derives GWs from the linearized field equations, characterizes their polarization states and effect on matter, and introduces the detection principle and Bayesian parameter estimation framework.

Chapter 3: Gravitational lensing of gravitational waves. Develops the strong lensing formalism and characterizes the signatures lensing imprints on GW signals, motivating the joint analysis of lensed image pairs for sky localization.

Chapter 4: Simulation setup and implementation. Describes the construction of the population of strongly lensed BBH events with LeR, including the detectability and duty-cycle selection process used to obtain the final injection set.

Chapter 5: Parameter estimation pipeline and results. Presents the parameter estimation pipeline: single-image analysis of each lensed image with Bilby, coherent joint analysis of the image pair with Hanabi, the automation infrastructure deployed on the Picasso HPC cluster, and a validation of the recovered parameters.

Chapter 6: Sky localization. Presents the sky map generation and area statistics, comparing the 50% and 90% credible sky areas and the constellation localization fraction obtained from single-image and joint analyses across the injected population, and places these results in the context of real GWTC-5 events and prior theoretical work.

Chapter 7: Conclusions. Summarizes the main findings of this work and discusses possible directions for future research.

Appendix A: Parameters of compact binary coalescences. Complete description of the parameter space of precessing binary black hole systems.

Appendix B: Practical guide to HPC workflows. Day-to-day commands, configuration patterns, and troubleshooting notes for working on the Picasso and CIT clusters, intended for future students.

Appendix C: Detailed implementation of the pipeline. Complete configuration files of the single-image and joint analyses.

Appendix D: Full posterior comparison for all events. Recovered joint and single-image posteriors for all 150 events, for the masses, effective distance, and spin parameters.

Contents

List of Figures

Abbreviations

BBH	**B**inary **B**lack **H**ole
BH	**B**lack **H**ole
BNS	**B**inary **N**eutron **S**tar
CBC	**C**ompact **B**inary **C**oalescence
CE	**C**osmic **E**xplorer
CW	**C**ontinuous gravitational **W**ave
EFE	**E**instein **F**ield **E**quations
EM	**E**lectro**m**agnetic
EOB	**E**ffective **O**ne **B**ody
EPL	**E**lliptical **P**ower **L**aw
ET	**E**instein **T**elescope
FLRW	**F**riedmann-**L**emaître-**R**obertson-**W**alker
GR	**G**eneral **R**elativity
GW	**G**ravitational **W**ave
GWTC	**G**ravitational-**W**ave **T**ransient **C**atalog
HLV	**H**anford-**L**ivingston-**V**irgo (three-detector network)
HPC	**H**igh **P**erformance **C**omputing
IGWN	**I**nternational **G**ravitational-**W**ave Observatory **N**etwork
IMR	**I**nspiral-**M**erger-**R**ingdown
IR1	**I**ntermediate **R**un **1** (between O4 and O5)
LHO	**L**IGO **H**anford **O**bservatory
LIGO	**L**aser **I**nterferometer **G**ravitational-wave **O**bservatory
LISA	**L**aser **I**nterferometer **S**pace **A**ntenna
LLO	**L**IGO **L**ivingston **O**bservatory
LSC	**L**IGO **S**cientific **C**ollaboration
LVK	**L**IGO-**V**irgo-**K**AGRA
MC	**M**onte **C**arlo
MCMC	**M**arkov **C**hain **M**onte **C**arlo
NFW	**N**avarro-**F**renk-**W**hite
NR	**N**umerical **R**elativity
NS	**N**eutron **S**tar
NSBH	**N**eutron **S**tar-**B**lack **H**ole (binary)
O1, O2, O3, O4, O5	**O**bservation run 1, 2, 3, 4, 5 of the LVK network
PE	**P**arameter **E**stimation
PN	**P**ost-**N**ewtonian
PSD	**P**ower **S**pectral **D**ensity
SGWB	**S**tochastic **G**ravitational **W**ave **B**ackground
SIE	**S**ingular **I**sothermal **E**llipsoid
SIS	**S**ingular **I**sothermal **S**phere
SNR	**S**ignal-to-**N**oise **R**atio
TDI	**T**ime **D**elay **I**nterferometry
TT	**T**ransverse-**T**raceless
WD	**W**hite **D**warf

Part I

Theoretical introduction

The first three chapters provide the theoretical foundation of this thesis. Chapter 1 introduces general relativity, developing the mathematical framework of curved spacetime and the Einstein field equations. Chapter 2 derives gravitational waves from the linearized field equations and describes the current network of interferometric detectors and the Bayesian parameter estimation framework. Chapter 3 introduces gravitational lensing, develops the strong lensing formalism for gravitational waves, and motivates the joint analysis of lensed image pairs for sky localization.

Chapter 1

Introduction to General Relativity

Nowadays, general relativity (GR) is fundamental to our understanding of modern astrophysics. Whether we are studying black holes (BHs), quasars, or gravitational lensing, Einstein's theory provides the essential framework. Efforts to detect gravitational waves (GWs) are a direct test of one of his key theoretical predictions, marking a remarkable chapter in the story of one of the most influential scientific achievements of the 20th century.

1.1 Motivation and Historical Background

The centenary of GR marks not only one of the greatest theoretical advances in physics, but also the foundation of modern astrophysical exploration. From its inception in 1915, Einstein's theory has evolved far beyond corrections to planetary orbits, becoming a central pillar in our understanding of the cosmos [6].

1.1.1 Key Historical Milestones

- **1905–1915:** Following his groundbreaking special relativity paper [7], Einstein seeks a relativistic framework for gravitation, which results in the formulation of GR in 1915.
- **1916:** Within months, Karl Schwarzschild finds the first exact solution to Einstein's equations [8], revealing the theoretical existence of BHs, a concept reinforced in 1917 when Einstein applies his theory to cosmology.
- **1919:** Arthur Eddington's eclipse expedition confirms the bending of starlight by the Sun's gravity [9], providing one of the first key empirical validations of the theory.
- **1927–1929:** Georges Lemaître [10] and Edwin Hubble [11] independently establish the expanding universe paradigm, linking GR to the observable evolution of the cosmos.
- **1960s–1970s:** After decades of theoretical neglect, interest in GR resurges. Joseph Weber pioneers early experimental searches for GWs [12], and the discovery of the Hulse-Taylor pulsar in 1974 offers indirect but compelling evidence that energy is indeed carried away by gravitational radiation.
- **1980s–1990s:** Interferometric efforts gain traction on both sides of the Atlantic, with prototype detectors developed in Europe and the United States; the resulting collaborations produce GEO600 and pave the way for the Laser Interferometer Gravitational-wave Observatory (LIGO) and Virgo [13].

- **2015–2017:** The detection of GW150914 by LIGO in September 2015 [5], publicly announced in February 2016, marks the first direct observation of GWs, fulfilling a century-old prediction and inaugurating the era of GW astronomy.

This historical path underscores GR's pivotal role beyond theoretical elegance. The final formulation emerged through a series of four weekly submissions to the Prussian Academy of Sciences in November 1915 [14, 15, 16, 17], culminating in the field equations presented on November 25th. Among the key results was the correct prediction of Mercury's perihelion shift of 43" per century [16], which had previously eluded Newtonian mechanics.

1.1.2 Equivalence principle

The equivalence principle played a central role in the development of GR. In 1907, Einstein realized that a freely falling observer does not locally experience gravity, an insight he later described as "the happiest thought of my life" [18, 19].

1.1.2.1 The weak equivalence principle

The roots of GR reach back to the 17$^{\text{th}}$ century, when Galileo Galilei and Simon Stevin (1548–1620) began to recognize that the velocity reached by freely falling objects is independent of their mass [3], a realization that contradicted the Aristotelian view of physics. Galileo first argued against this view in his early manuscript De Motu (ca. 1590) [20], and later established the result experimentally by rolling spheres down inclined planes during his time in Padua, published in his 1638 Discorsi [21].

The universality of free fall reflects a striking feature of Newtonian mechanics: the equivalence between inertial mass m_i and gravitational mass m_g, also called the weak equivalence principle. In Newtonian theory, inertial mass appears in Newton's second law,

$$\vec{F} = m_i \, \vec{a} \ , \tag{1.1}$$

while gravitational mass governs the gravitational potential energy between two bodies,

$$V = -\frac{G \, m_\mathrm{g} \, M_\mathrm{g}}{r} \ , \tag{1.2}$$

where M_g is the gravitational mass of the second body and $r > 0$ is the separation between them. The motion of an object in free fall is then governed by

$$-m_\mathrm{g} g = m_\mathrm{i} \ddot{z} \ , \tag{1.3}$$

whose solution,

$$z(t) = -\frac{1}{2}\frac{m_\mathrm{g}}{m_\mathrm{i}} g t^2 + v_0 t + z_0 \ , \tag{1.4}$$

is independent of mass only if $m_g/m_i =$ constant, which experiments confirm to be universal and set to 1 by an appropriate choice of units.

This equivalence has been tested with remarkable precision. In 1889, Eötvös used a torsion balance to reach a precision of $|(m_g - m_i)/m_g| \sim 5 \times 10^{-8}$, later refined with Pekár and Fekete to $\sim 10^{-9}$, published in 1922 [22], and later improved by Dicke in 1964 to $< 10^{-11}$ [23]. Currently, the MICROSCOPE satellite mission has achieved a precision of 10^{-15} [24].

1.1.2.2 The strong equivalence principle

Although Newtonian mechanics offered no fundamental reason for $m_g = m_i$, Einstein believed this could not be a mere coincidence. The breakthrough came when he realized that an observer in free fall does not feel their own weight and cannot distinguish, through any local experiment, whether they are in a gravitational field or drifting freely in empty space.

This reasoning extends to the case of uniform acceleration: an observer inside an enclosed rocket accelerating at $a = g$ cannot tell whether the apparent downward pull is due to gravity or to the rocket's acceleration. This leads to two equivalent formulations:

> **Equivalence Principle** (for uniform gravitational fields): *An observer in free fall within a constant gravitational field is equivalent to an inertial observer in the absence of gravity. It is impossible to distinguish between these two situations by means of physical experiments.*

> **Equivalence Principle** (for constant accelerations): *An observer undergoing uniform acceleration is equivalent to an inertial observer in a uniform gravitational field. It is impossible to distinguish between these two situations through physical experiments.*

However, this equivalence breaks down for general gravitational fields, where tidal forces (arising from field inhomogeneities) allow a free-falling observer to detect the presence of gravity [3]. For instance, two freely falling objects in Earth's radial field slowly converge toward the center rather than following parallel paths. Since tidal forces are proportional to the gradient of the field, they vanish at sufficiently small scales, and the equivalence principle recovers its validity locally.

> **Equivalence Principle** (general formulation): *Free-falling observers in a gravitational field are locally indistinguishable from inertial observers. There are no local experiments that can differentiate between these two scenarios.*

The local validity of the equivalence principle has a precise mathematical counterpart: at any point of spacetime it is always possible to choose coordinates in which the metric tensor (introduced formally in §1.2.4) reduces locally to the Minkowski form $\eta_{\mu\nu}$. This idea, made rigorous through the concept of normal coordinates, underlies the geometric interpretation of gravity developed in the following sections.

1.2 The description of curved spacetime

After his work on special relativity and learning about the mathematical structure introduced by Minkowski [25], Einstein realized that spacetime itself possesses a geometry. Instead of thinking of space and time as separate and absolute entities, special relativity showed that they are part of a unified, four-dimensional spacetime fabric.

Before introducing gravity, it is crucial to recall this idea: spacetime itself has a geometry. In special relativity, this geometry is flat (Minkowskian); however, the presence of mass and energy changes this picture.

To describe this curved geometry, Euclidean tools are no longer sufficient. Instead, we require a more general mathematical framework: Riemannian geometry, capable of capturing the curvature induced by mass and energy. In this section, we will extend that concept further.

1.2.1 From Euclidean to non-Euclidean geometry

The geometry most familiar to us, Euclidean geometry, describes flat space: parallel lines never meet, and the angles of a triangle always sum to 180°. These properties stem from the work of the Greek mathematician Euclid, who around 300 BC formulated five postulates from which he derived hundreds of theorems [26]. Among these, the fifth, the parallel postulate, stood out for its complexity, prompting centuries of efforts to prove it from the others [27].

Eventually, mathematicians like Nikolai I. Lobachevsky [28] and János Bolyai [29] explored what would happen if the parallel postulate were replaced, discovering entirely new self-consistent geometries (see Fig. 1.1). On a sphere, parallel lines do not exist and triangle angles exceed 180°; in hyperbolic geometry, infinitely many parallel lines pass through a point and triangle angles sum to less than 180°.

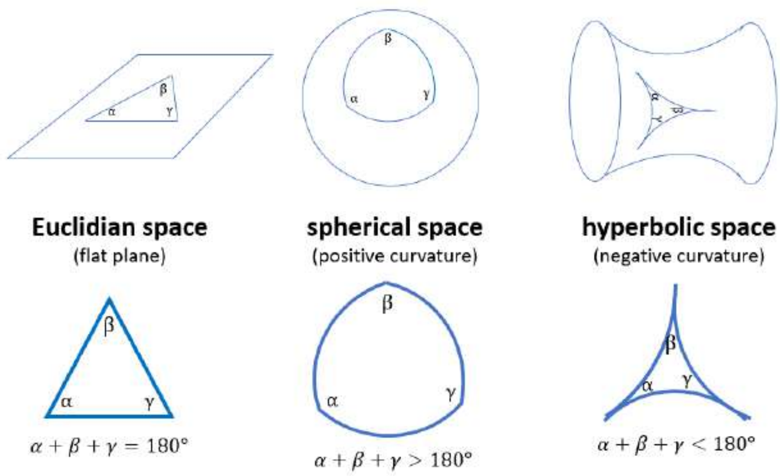

Figure 1.1: Representation of the three main types of geometry: Euclidean (flat), spherical (positive curvature), and hyperbolic (negative curvature). In each case, the sum of the angles of a triangle differs: it equals 180° in Euclidean geometry, exceeds 180° in spherical geometry, and is less than 180° in hyperbolic geometry. Source: [30].

These ideas became essential when Einstein formulated GR. In his theory, spacetime is curved by the presence of mass and energy, and the shortest paths, geodesics (see §1.2.5.4), follow the curvature of spacetime rather than straight Euclidean lines. This curved geometry, pioneered by Bernhard Riemann, provides the mathematical language to describe gravitational phenomena, replacing Newton's force with the geometry of spacetime itself.

1.2.2 Manifolds and tangent spaces

In order to describe curved spaces mathematically, we use the concept of a manifold: a space that may possess curvature or nontrivial topology, yet locally resembles Euclidean space. Formally, an N-dimensional manifold $\mathcal{M}^N$ is a space that, in a neighborhood of each point p, looks like $\mathbb{R}^N$. At every point p we can define the tangent space $T_p(\mathcal{M}^N)$, an N-dimensional vector space that provides a linear approximation to the manifold in a small region around p. Tangent spaces at different points are distinct, reflecting the fact that the manifold is not globally flat (see Fig. 1.2).

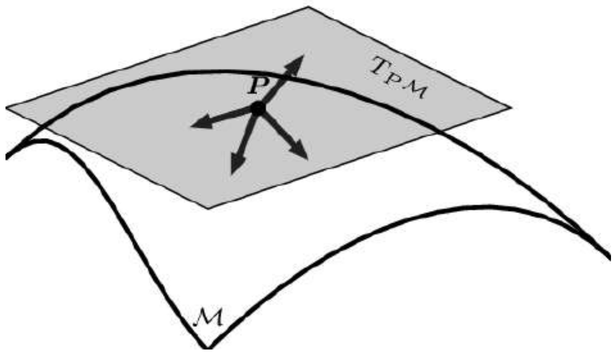

Figure 1.2: The set of all vectors defined at a point p forms the tangent space $T_p(\mathcal{M})$ at p. Assembling the tangent spaces for every point of the manifold produces the tangent bundle $T(\mathcal{M})$. Original illustration based on the source: [31].

Formally, a topological manifold satisfies three conditions: the Hausdorff property (any two distinct points can be separated by disjoint open sets), second-countability (the topology admits a countable basis), and local Euclideanness (every point has a neighborhood homeomorphic to an open subset of $\mathbb{R}^n$). The last condition ensures that, although $\mathcal{M}$ may be globally curved, every point admits a patch that looks like flat $\mathbb{R}^n$.

Manifolds appear throughout physics: in mechanics, configuration and phase spaces are manifolds; in GR, spacetime is modeled as a Lorentzian manifold (see §1.2.4), locally isometric to Minkowski space but with nontrivial global curvature determined by the distribution of matter and energy.

1.2.2.1 Mappings, charts, and atlases

A chart on $\mathcal{M}$ is a pair (U, φ) consisting of an open subset $U \subset \mathcal{M}$ and an injective map $\varphi : U \to \mathbb{R}^n$, whose components $\varphi(p) = (x^1(p), \ldots, x^n(p))$ define local coordinates on U (see Fig. 1.3). Just as no single flat map can represent the entire Earth without distortion, a single chart rarely covers an entire manifold. For instance, covering the 2-sphere S^2 requires at least two charts.

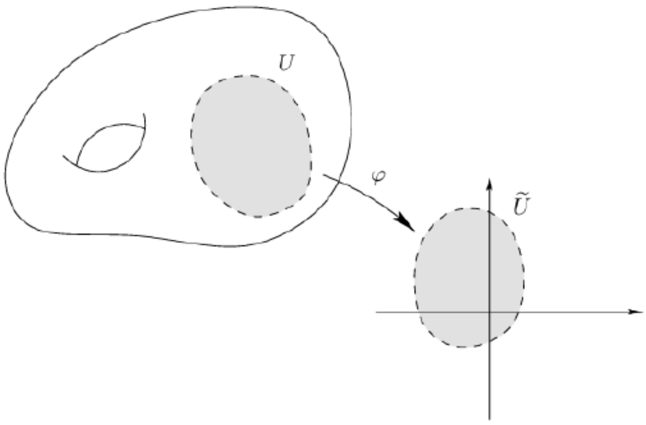

Figure 1.3: Representation of a chart (U, φ) on a manifold $\mathcal{M}$, where the subset $U \subset \mathcal{M}$ is mapped bijectively to an open set $\widetilde{U} \subset \mathbb{R}^n$ via φ. Source: [32].

An atlas is a collection of charts $\{(U_\alpha, \varphi_\alpha)\}$ whose domains cover $\mathcal{M}$, with the additional requirement that whenever $U_\alpha \cap U_\beta \neq \emptyset$, the transition map

$$\varphi_\beta \circ \varphi_\alpha^{-1} : \varphi_\alpha(U_\alpha \cap U_\beta) \to \varphi_\beta(U_\alpha \cap U_\beta) \tag{1.5}$$

is C^∞ (see Fig. 1.4). These smooth transition maps ensure that the different local coordinate patches are consistently glued together, uniquely determining the manifold's differentiable structure.

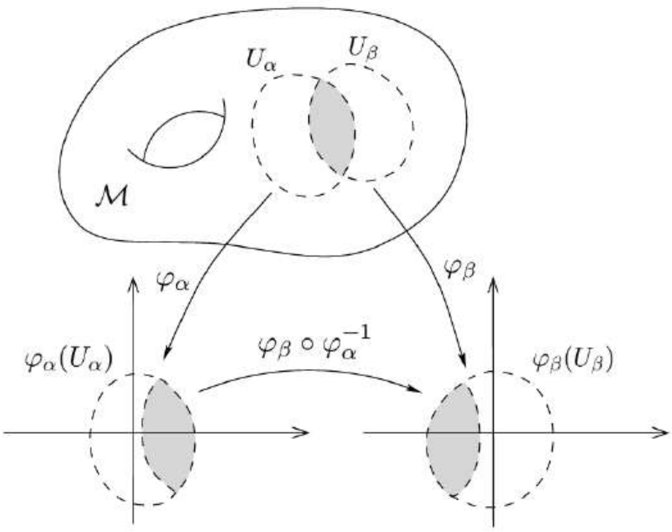

Figure 1.4: When two charts overlap, their coordinate systems are related by a smooth transition map. Original illustration based on the source: [32].

1.2.3 Tensors

Tensors are the natural language of GR: they are geometric objects whose components transform consistently under coordinate changes, ensuring that physical laws remain valid in any reference frame [33]. They generalize the familiar notions of scalars (invariant under coordinate changes), contravariant vectors (components transforming via the Jacobian of the coordinate change), and covariant vectors or one-forms (components transforming via the inverse Jacobian).

Tensors are commonly classified by their type (r, s), where r denotes the number of contravariant (upper) indices and s the number of covariant (lower) indices [34]. Some illustrative examples are:

- A scalar T carries no indices; it is a rank-0 tensor of type $(0, 0)$.
- A contravariant vector V^β is a rank-1 tensor of type $(1, 0)$, and a one-form W_α is of type $(0, 1)$.
- A mixed tensor T^α_γ is of type $(1, 1)$, and the Riemann curvature tensor $R^l{}_{ijk}$, studied in §1.2.5.5, is of type $(1, 3)$.

In general, a tensor of rank n in a d-dimensional space has d^n components, transforming under coordinate changes according to

$$T'^{i'_1 \ldots i'_r}_{j'_1 \ldots j'_s} = \left(\prod_{k=1}^{r} \frac{\partial x'^{i'_k}}{\partial x^{i_k}} \right) \left(\prod_{l=1}^{s} \frac{\partial x^{j_l}}{\partial x'^{j'_l}} \right) T^{i_1 \ldots i_r}_{j_1 \ldots j_s} \ . \tag{1.6}$$

This transformation law ensures that if a tensor equation holds in one coordinate system, it holds in all, an essential feature of any generally covariant theory such as GR. Formally, a tensor of type (r, s) can be seen as a multilinear map taking r one-forms and s vectors as inputs and returning a real number [35].

In GR, tensors are not merely mathematical tools: they encode curvature, energy, and momentum, and constitute the grammar through which the field equations are written.

1.2.4 Metric tensor

Until now, we have not introduced the most crucial element. There is a single object that, when defined on a manifold, surpasses all others in importance for understanding gravity: the metric tensor. The presence of a metric opens the door to a wide range of new concepts, which together form the framework known as Riemannian geometry [36].

A differentiable manifold equipped with a symmetric metric tensor is called a Riemannian manifold. In GR, the metric plays a fundamental role: it defines the separation between nearby points in spacetime and determines the curvature of the manifold. In flat spacetime, far from matter and energy, the metric approximates the Minkowski metric $\eta_{\mu\nu}$ of special relativity [37]. Near massive objects, such as in the Schwarzschild solution, the metric adapts to reflect spacetime curvature (see §1.4.1 for more details).

A metric tensor g on a differentiable manifold $\mathcal{M}$ is a $(0, 2)$ tensor field that allows the definition of an inner product on each tangent space $T_p(\mathcal{M})$. It satisfies:

- Symmetry: $g(X, Y) = g(Y, X)$ for all vector fields X, Y.
- Non-degeneracy: If $g(X, Y)|_p = 0$ for all $Y \in T_p(M)$, then $X_p = 0$.

In a coordinate basis $\{\partial_\mu\}$, the components of the metric are given by

$$g_{\mu\nu}(x) = g\left(\frac{\partial}{\partial x^\mu}, \frac{\partial}{\partial x^\nu}\right), \tag{1.7}$$

and by definition

$$g_{\mu\nu} = e_\mu \cdot e_\nu \ , \tag{1.8}$$

where e_μ are the coordinate basis vectors.

Given coordinates x^μ, an infinitesimal displacement in the manifold is

$$ds = dx^\mu \, e_\mu \ , \tag{1.9}$$

and its squared norm (inner product with itself) is

$$(ds)^2 = (dx^\mu e_\mu) \cdot (dx^\nu e_\nu) \overset{(i)}{=} g_{\mu\nu} \, dx^\mu dx^\nu. \tag{1.10}$$

(i) We use that $e_\mu \cdot e_\nu = g_{\mu\nu}$.

This is the line element:

$$ds^2 = g_{\mu\nu}(x) \, dx^\mu dx^\nu. \tag{1.11}$$

In an n-dimensional Riemannian manifold:

$$dl^2 = g_{ij} \, dx^i dx^j \ . \tag{1.12}$$

Similarly, the inner product of two distinct vectors is

$$V \cdot W = g_{\mu\nu} V^\mu W^\nu \ . \tag{1.13}$$

The explicit form of $g_{\mu\nu}$ depends on the coordinate system [33]. In Cartesian coordinates:

$$dl^2 = dx^2 + dy^2 + dz^2, \quad g_{ij} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}. \tag{1.14}$$

In Minkowski spacetime:

$$ds^2 = -c^2 dt^2 + dx^2 + dy^2 + dz^2, \quad \eta_{\mu\nu} = \begin{pmatrix} -1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}. \tag{1.15}$$

In spherical coordinates:

$$ds^2 = dr^2 + r^2\, d\theta^2 + r^2 \sin^2\theta\, d\varphi^2, \quad g_{ij} = \mathrm{diag}(1, r^2, r^2\sin^2\theta)^{1}. \tag{1.16}$$

In general, $g_{\mu\nu}$ is a symmetric 4×4 matrix in spacetime, with 10 independent components. Because coordinate transformations can adjust 4 of them, only 6 components contain independent physical information [39].

The metric tensor is the tool that enables the computation of lengths, angles, areas, and volumes in a manifold. In physics, it encodes the gravitational field in GR. It also acts as a bridge between contravariant and covariant vectors:

$$v_\mu = g_{\mu\nu} v^\nu, \quad v^\mu = g^{\mu\nu} v_\nu\ , \tag{1.17}$$

where $g^{\mu\nu}$ is the inverse metric, satisfying

$$g_{\mu\nu} g^{\nu\rho} = \delta^\rho_\mu\ . \tag{1.18}$$

The signature of the metric distinguishes between different geometries: Riemannian metrics are positive-definite $(+, +, +, \dots)$, while Lorentzian metrics have signature $(+, -, -, \dots, -)$ or $(-, +, +, \dots, +)$ as in spacetime. Once the metric is known, the geometry and the gravitational field (in the case of GR) are completely determined.

A particularly relevant example in the context of this work is the weak-field perturbed metric, which describes spacetime in the presence of a slowly varying Newtonian gravitational potential. This metric reduces to the Minkowski form in the absence of matter and governs the deflection of GWs as they propagate through a gravitational lens, as developed in §3.2.

1.2.4.1 Conventions

Throughout this thesis the signature $(-, +, +, +)$ is adopted for the metric, so that $\eta_{\mu\nu} = \mathrm{diag}(-1, +1, +1, +1)$ and timelike intervals satisfy $ds^2 < 0$. Newton's gravitational constant G and the speed of light c are kept explicit throughout, rather than setting $c = 1$; consequently, factors of c appear explicitly in expressions such as the Einstein field equations, $G_{\mu\nu} = (8\pi G/c^4)\, T_{\mu\nu}$, and are restored consistently in the weak-field and lensing expressions of later chapters.

1.2.5 Affine connection and curvature

We have now laid the groundwork for a deeper look into the mathematics of curved space. Ultimately, our goal is to understand:

- Geodesics (see §1.2.5.4), the "straightest possible" paths between any two points on a Riemannian manifold, and thus the natural trajectories of free particles moving through curved spacetime.

[1]The notation $\mathrm{diag}(a_1, a_2, \dots, a_n)$ denotes a diagonal matrix whose entries are zero everywhere except along the diagonal, where the components are $a_1, a_2, \dots, a_n$. In this case, it means that the metric tensor g_{ij} has components $g_{11} = 1$, $g_{22} = r^2$, $g_{33} = r^2 \sin^2\theta$, and all off-diagonal components vanish [38].

- The Riemann curvature tensor (see §1.2.5.5), which encodes how space bends. A particular version of this tensor appears on the left-hand side of the fundamental object of GR, the Einstein field equations (EFE).

Parallel transport of a vector (see §1.2.5.3) plays a key role in both of these concepts. But before we can understand how to transport vectors, we must first know how to differentiate tensor fields. To make this possible, we introduce what are known as connection coefficients or Christoffel symbols (denoted by the symbol Γ). These coefficients allow us to compare vectors located at different points on a manifold. Once the metric of the manifold is known, we can compute the connection coefficients, unlocking our ability to define derivatives and do much more. This is why the metric is central to understanding Riemannian manifolds, and therefore the geometry of spacetime itself [34].

1.2.5.1 Christoffel symbols

Our goal in GR is to express physical laws as tensor equations valid in any coordinate system. To do that, we must be able to differentiate tensors in a way that behaves well under changes of coordinates. Specifically, we want a derivative that reduces to the usual partial derivative in flat Minkowski space using Cartesian coordinates, but which still transforms as a tensor in curved spacetimes.

At first glance, differentiating a tensor might seem straightforward. For example, in Cartesian coordinates, a vector field like

$$\vec{V}(x,y,z) = (2x-y)\,\vec{e}_x + (x+3z)\,\vec{e}_y + (y-z)\,\vec{e}_z \ , \tag{1.19}$$

can be differentiated component by component:

$$\frac{\partial \vec{V}}{\partial x} = 2\,\vec{e}_x + \vec{e}_y \ , \qquad \frac{\partial \vec{V}}{\partial y} = -\,\vec{e}_x + \vec{e}_z \ , \qquad \frac{\partial \vec{V}}{\partial z} = 3\,\vec{e}_y - \vec{e}_z \ . \tag{1.20}$$

The basis vectors $\vec{e}_x, \vec{e}_y, \vec{e}_z$ are constant in space, so we do not need to worry about them changing from point to point. We just apply the partial derivative to each component function.

However, this approach fails in curved spaces or curvilinear coordinates, where the basis vectors themselves vary from point to point. In such cases, when differentiating a tensor, we must also account for the rate of change of the basis vectors. This is where the connection coefficients, or Christoffel symbols, come into play.

Let us consider a vector expressed as $\vec{V} = V^\alpha \vec{e}_\alpha$, where $\vec{e}_\alpha$ are the coordinate basis vectors. Now let us take the derivative of this with respect to some coordinate using the product rule [34]:

$$\frac{\partial \vec{V}}{\partial x^\beta} = \frac{\partial}{\partial x^\beta}\left(V^\alpha \vec{e}_\alpha\right) = \left(\frac{\partial V^\alpha}{\partial x^\beta}\right)\vec{e}_\alpha + V^\alpha \left(\frac{\partial \vec{e}_\alpha}{\partial x^\beta}\right) \ . \tag{1.21}$$

The first term reflects the change in the vector components, while the second term captures how the basis vectors change. Since $\partial_\beta \vec{e}_\alpha$ is itself a vector, it can be expressed in terms of the basis:

$$\frac{\partial \vec{e}_\alpha}{\partial x^\beta} = \Gamma^\gamma_{\alpha\beta} \vec{e}_\gamma \ . \tag{1.22}$$

Here, the Christoffel symbols $\Gamma^\gamma_{\alpha\beta}$ quantify how the basis vector $\vec{e}_\alpha$ changes in the direction of coordinate x^β, projected onto the basis $\vec{e}_\gamma$. The lower indices α and β refer to the direction and variable of differentiation, while the upper index γ indicates the output basis direction.

We can isolate the Christoffel symbol by multiplying both sides by the dual basis vector $\vec{e}^{\,\gamma}$. Relabeling the dummy index in (1.22) as λ to avoid a clash with γ, $\partial_\beta \vec{e}_\alpha = \Gamma^\lambda_{\alpha\beta} \vec{e}_\lambda$, so that

$$\vec{e}^{\,\gamma} \cdot \partial_\beta \vec{e}_\alpha = \Gamma^\lambda_{\alpha\beta}\, \vec{e}^{\,\gamma} \cdot \vec{e}_\lambda = \Gamma^\lambda_{\alpha\beta}\, \delta^\gamma_\lambda = \Gamma^\gamma_{\alpha\beta} \quad \Rightarrow \quad \Gamma^\gamma_{\alpha\beta} = \vec{e}^{\,\gamma} \cdot \partial_\beta \vec{e}_\alpha \ . \tag{1.23}$$

Here, we used the property $\vec{e}^{\,\gamma} \cdot \vec{e}_\lambda = \delta^\gamma_\lambda$, which defines the dual basis.

These coefficients do not form a tensor themselves, but they are essential for constructing tensorial derivatives.

Some key properties of Christoffel symbols are:

- They provide the rule for differentiating tensors in a curved manifold.
- They are called connection coefficients because they "connect" nearby points on the manifold, defining how vectors vary between them, crucial for defining parallel transport.
- If we know the metric $g_{\mu\nu}$, we can compute them explicitly as (see [40] for the derivation):

$$\Gamma^\rho_{\mu\nu} = \frac{1}{2} g^{\rho\lambda} \left(\partial_\mu g_{\lambda\nu} + \partial_\nu g_{\lambda\mu} - \partial_\lambda g_{\mu\nu} \right) \ . \tag{1.24}$$

- In standard GR (without torsion[2]), the Christoffel symbols are symmetric in their lower indices:

$$\Gamma^\rho_{\mu\nu} = \Gamma^\rho_{\nu\mu} \ . \tag{1.25}$$

- In an n-dimensional Riemannian manifold, there are theoretically n^3 such coefficients, though this number is reduced by the symmetry condition.

The term "Christoffel symbols" honors the mathematician Elwin Christoffel (1829–1900), who first introduced them in the context of differential geometry. Although sometimes informally dubbed "Christ-awful symbols" due to their apparent complexity, this perception tends to fade with familiarity [40]. What initially seems intricate is, in fact, a systematic way of incorporating the effects of curvature into the formulation of differential operators.

These connection coefficients are essential for extending the familiar notion of differentiation to the curved spacetimes of GR, setting the stage for the definition of the covariant derivative, which we address next.

[2] Torsion describes the amount of "asymmetry" in the way points in spacetime are connected: it indicates whether moving first in one direction and then in another ends at the same point as doing so in the reverse order. If torsion is zero, as in standard GR, both paths lead to the same point. When we study parallel transport and the Riemann curvature tensor (see §1.2.5.5), this idea will acquire a precise mathematical formulation through the torsion tensor. See [3, 41] for more details.

1.2.5.2 Covariant derivative

Having introduced the Christoffel symbols (1.22) and established how to differentiate vectors in curved spacetime using equation (1.21), we define the covariant derivative of a contravariant vector V^α:

$$\nabla_\beta V^\alpha = \frac{\partial V^\alpha}{\partial x^\beta} + V^\gamma \Gamma^\alpha_{\gamma\beta} \ . \tag{1.26}$$

Similarly, the covariant derivative of a one-form is:

$$\nabla_\beta V_\alpha = \frac{\partial V_\alpha}{\partial x^\beta} - V_\gamma \Gamma^\gamma_{\alpha\beta} \ . \tag{1.27}$$

The covariant derivative is linear and satisfies the Leibniz rule:

$$\begin{aligned} &\nabla_\mu(\alpha V^\nu + \beta W^\nu) = \alpha \nabla_\mu V^\nu + \beta \nabla_\mu W^\nu, \\ &\nabla_\mu(V^\nu W^\rho) = (\nabla_\mu V^\nu) W^\rho + V^\nu(\nabla_\mu W^\rho) \ . \end{aligned} \tag{1.28}$$

In flat spacetime with Cartesian coordinates, the Christoffel symbols vanish and the covariant derivative reduces to the partial derivative. For a scalar field ϕ, the identity $\nabla_\alpha \phi = \partial_\alpha \phi$ holds in any coordinate system, since a scalar carries no indices for the connection to act on.

For a general tensor of type (m, n), the covariant derivative is [34]:

$$\begin{aligned} \nabla_\rho T^{\mu_1 \ldots \mu_m}{}_{\nu_1 \ldots \nu_n} = {} & \partial_\rho T^{\mu_1 \ldots \mu_m}{}_{\nu_1 \ldots \nu_n} \\ & + \Gamma^{\mu_1}_{\rho\lambda} T^{\lambda \mu_2 \ldots \mu_m}{}_{\nu_1 \ldots \nu_n} + \cdots + \Gamma^{\mu_m}_{\rho\lambda} T^{\mu_1 \ldots \mu_{m-1} \lambda}{}_{\nu_1 \ldots \nu_n} \\ & - \Gamma^{\lambda}_{\rho\nu_1} T^{\mu_1 \ldots \mu_m}{}_{\lambda \nu_2 \ldots \nu_n} - \cdots - \Gamma^{\lambda}_{\rho\nu_n} T^{\mu_1 \ldots \mu_m}{}_{\nu_1 \ldots \nu_{n-1} \lambda} \ . \end{aligned} \tag{1.29}$$

The rule is: add a Christoffel term for each upper index and subtract one for each lower index.

Finally, it is postulated that the covariant derivative of the metric tensor vanishes:

$$\nabla_\beta g_{\mu\nu} = \nabla_\beta g^{\mu\nu} = 0 \ . \tag{1.30}$$

A connection satisfying this condition is called metric compatible. This ensures that inner products between vectors are preserved under parallel transport (§1.2.5.3), allowing consistent raising and lowering of indices and well-defined notions of length and angle in curved geometry [34].

1.2.5.3 Parallel transport

To understand the geometry of curved spaces, one must move beyond the familiar tools of Euclidean geometry. In flat space, vectors at different points can be naturally compared, added, subtracted, or evaluated using inner products because transporting a vector from one point to another without changing its direction is unambiguous. However, this intuition breaks down in curved manifolds, where the concept of direction is local and the comparison of vectors at different points becomes nontrivial.

This leads to a fundamental question: how can we define the idea of keeping a vector "unchanged" as we move it along a path in a curved space? The answer lies in the notion of parallel transport.

Parallel transport provides a rule for moving vectors along curves in a manifold in such a way that they remain as "constant" as possible according to the geometry of the space.

In flat space, transporting a vector around a closed loop brings it back unchanged. But in curved spaces, the result of parallel transport around a loop generally depends on the path taken, reflecting the intrinsic curvature of the manifold (see Fig. 1.5).

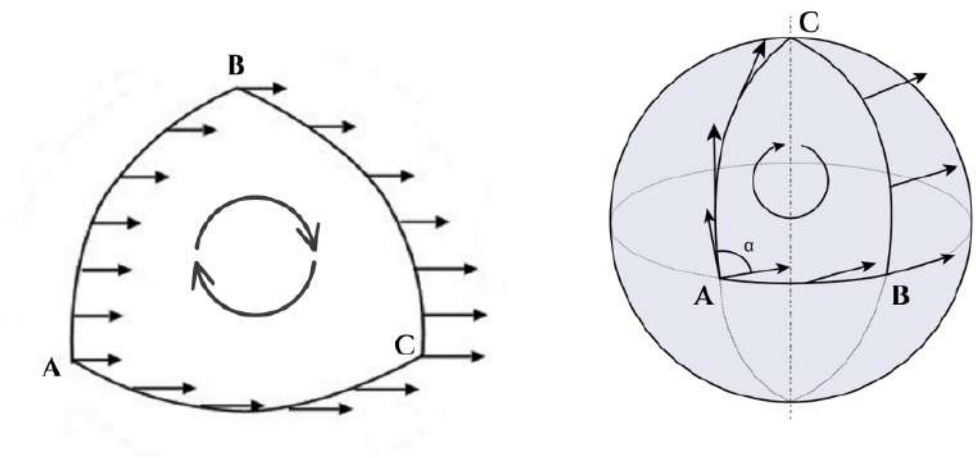

(a) Parallel transport in flat space.

(b) Parallel transport on a curved space (sphere).

Figure 1.5: Comparison between parallel transport in flat and curved spaces. In flat space, the vector returns unchanged after a closed loop; in curved space, the final result depends on the path taken, revealing the manifold's intrinsic geometry. Original illustrations based on the sources: [34, 42].

The condition for the parallel transport of the vector $\vec{V}$ along a parametrized curve is that its covariant derivative along the curve, $D\vec{V}/d\lambda \equiv (dx^\mu/d\lambda)\,\nabla_\mu \vec{V}$, must be zero, namely,

$$\frac{D\vec{V}}{d\lambda} = 0 \; . \tag{1.31}$$

To obtain the derivative, we use equation (1.21) and replace the parameter x^β with λ, obtaining:

$$\frac{d\vec{V}}{d\lambda} = \frac{dV^\alpha}{d\lambda} e_\alpha + V^\alpha \frac{de_\alpha}{d\lambda} \; . \tag{1.32}$$

The change of the basis can be expressed as

$$\frac{de_\alpha}{d\lambda} = \frac{\partial e_\alpha}{\partial x^\beta}\frac{dx^\beta}{d\lambda} = \Gamma^\gamma_{\alpha\beta}\frac{dx^\beta}{d\lambda}\, e_\gamma \; , \tag{1.33}$$

where we have used the definition of the Christoffel symbols (1.22).

Substituting this result into the previous expression and replacing the index α with γ in the first term yields

$$\frac{d\vec{V}}{d\lambda} = \left(\frac{dV^\gamma}{d\lambda} + V^\alpha\,\Gamma^\gamma_{\alpha\beta}\frac{dx^\beta}{d\lambda}\right) e_\gamma \; . \tag{1.34}$$

In component form, the absolute derivative[3] introduced above reads

$$\frac{DV^\gamma}{d\lambda} = \frac{dV^\gamma}{d\lambda} + \Gamma^\gamma_{\alpha\beta}\, V^\alpha\, \frac{dx^\beta}{d\lambda} \ . \tag{1.35}$$

Parallel transport of $\vec{V}$ along the curve is then defined by the condition

$$\frac{D\vec{V}}{d\lambda} = 0 \quad \Longleftrightarrow \quad \frac{dV^\gamma}{d\lambda} + \Gamma^\gamma_{\alpha\beta}\, V^\alpha\, \frac{dx^\beta}{d\lambda} = 0 \ . \tag{1.36}$$

This is a linear first order system of ordinary differential equations for the components $V^\gamma(\lambda)$, which, given initial data $V^\gamma(\lambda_0)$, has a unique solution along the curve.

For a tensor $T^{\mu_1 \cdots \mu_p}{}_{\nu_1 \ldots \nu_q}$ the parallel transport along a trajectory $x^\mu(\lambda)$ is defined by requiring that, for every point on the curve,

$$\left(\frac{D}{d\lambda} T\right)^{\mu_1\mu_2\ldots\mu_p}{}_{\nu_1\nu_2\ldots\nu_q} \ \equiv \ \frac{dx^\sigma}{d\lambda}\, \nabla_\sigma T^{\mu_1\mu_2\ldots\mu_p}{}_{\nu_1\nu_2\ldots\nu_q} \ \equiv \ 0 \ . \tag{1.37}$$

This expression constitutes a valid tensor equation because both the tangent vector $dx^\mu/d\lambda$ and the covariant derivative ∇_μ transform as tensors.

This equation plays a key role, as we will see in the next section that it serves as the starting point for deriving the geodesic equations. We will also find that freely moving or freely falling objects in spacetime follow the straightest possible paths, curves known as geodesics.

1.2.5.4 Geodesics

In Euclidean space, the shortest path between two points is a straight line. In a curved space, the analogous concept is a geodesic, a curve that, given the curvature of the space, is as "straight" as possible. On a sphere, for example, great circles play this role.

The idea of a geodesic can be formalized in several equivalent ways when considering two points A and B on a surface [40]:

1. The geodesic is the curve that minimizes the distance between A and B.
2. A geodesic is a curve whose length remains stationary under small variations of its shape.
3. A more local definition states that a geodesic is a curve that, at each point, is as straight as possible.

1.2.5.4.1 Geometric derivation

From a geometric perspective, a geodesic is defined as a curve whose tangent vector is parallel transported along itself. Formally, this condition requires that the covariant derivative of the tangent vector vanishes at every point of the curve [3].

[3] An absolute derivative refers to the covariant derivative taken along a curve, in this case the curve parametrized by λ.

Let $x^\mu(\lambda)$ represent a parametrized curve, where λ is an affine parameter. Its tangent vector is given by

$$U^\mu = \frac{dx^\mu}{d\lambda} \ . \tag{1.38}$$

The requirement of parallel transport along the curve is expressed as

$$\frac{DU^\mu}{d\lambda} = 0 \ , \tag{1.39}$$

which, using the explicit form of the covariant derivative, becomes

$$\frac{dU^\mu}{d\lambda} + \Gamma^\mu_{\rho\sigma} U^\rho \frac{dx^\sigma}{d\lambda} = 0 \ . \tag{1.40}$$

Substituting $U^\rho = \frac{dx^\rho}{d\lambda}$, one arrives at the standard geodesic equation [43]:

$$\frac{d^2x^\mu}{d\lambda^2} + \Gamma^\mu_{\rho\sigma} \frac{dx^\rho}{d\lambda}\frac{dx^\sigma}{d\lambda} = 0 \ . \tag{1.41}$$

Thus, the geodesic equation emerges naturally as the geometric condition ensuring that the tangent vector remains invariant under parallel transport.

The main physical significance of geodesics in GR lies in the fact that they describe the trajectories of freely falling particles, objects experiencing no proper acceleration. Within this framework, the geodesic equation serves as the relativistic analogue of Newton's second law, $\vec{F} = m\vec{a}$, specifically in the case where $\vec{F} = 0$.

1.2.5.4.2 Normalization condition

To fully characterize a geodesic, one must impose a normalization condition on the tangent vector $\dot{x}^\mu = dx^\mu/d\tau$:

$$g_{\mu\nu}\dot{x}^\mu\dot{x}^\nu = \begin{cases} -\,c^2 & \text{timelike geodesics (massive particles),} \\ \ \ 0 & \text{null geodesics (light rays or massless particles),} \\ +\,c^2 & \text{spacelike curves (not particle worldlines).} \end{cases} \tag{1.42}$$

This condition follows directly from the spacetime interval $ds^2 = g_{\mu\nu}dx^\mu dx^\nu$. For timelike geodesics, τ is the proper time of the particle; for null geodesics, proper time is undefined and τ becomes an affine parameter. In many cases, particularly when the spacetime admits symmetries, this condition alone suffices to determine the geodesic trajectory without solving the full geodesic equation [43].

1.2.5.5 Riemann curvature tensor

In §1.2.5.3 we established that parallel transport around a closed loop in a curved manifold generically returns a vector in a different orientation from its initial one (see Fig. 1.5). This path-dependence is not an artefact of the coordinate system but reflects the intrinsic geometry of the manifold. The Riemann curvature tensor provides a precise, coordinate-independent quantification of this effect.

To make this concrete, consider the parallel transport condition derived in §1.2.5.3,

$$\frac{DV^\gamma}{d\lambda} = \frac{dV^\gamma}{d\lambda} + \Gamma^\gamma_{\alpha\beta} V^\alpha \frac{dx^\beta}{d\lambda} = 0, \tag{1.43}$$

which governs how the components of a vector change as it is transported along a curve. If we transport V^λ around an infinitesimal closed loop, the net change in its components depends on the order in which we traverse the two directions spanning the loop. This order-dependence can be captured algebraically by examining the commutator of two covariant derivatives: acting first with ∇_μ and then ∇_ν, versus the reverse order. In flat space, partial derivatives commute, so this commutator vanishes identically. In a curved manifold, the Christoffel symbols $\Gamma^\gamma_{\alpha\beta}$ introduce additional terms that survive, encoding the failure of parallel transport to be path-independent.

Consider transporting a vector V^λ around an infinitesimal parallelogram with sides dx^μ and dx^ν (see Fig. 1.6).

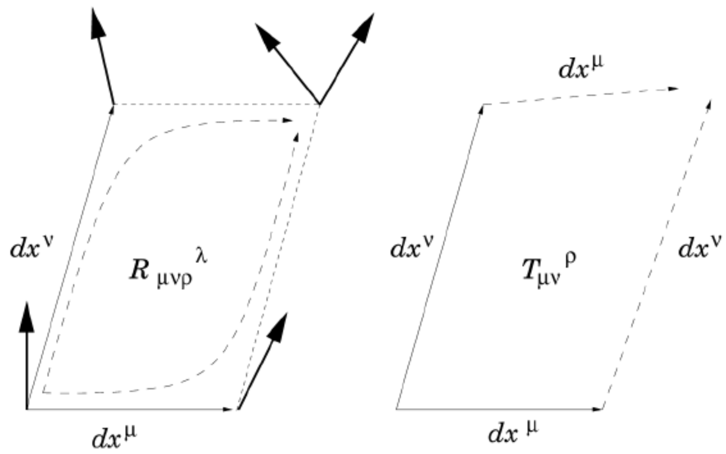

Figure 1.6: Illustration of parallel transport around an infinitesimal closed loop, showing the deviation of a vector due to curvature, characterized by the Riemann tensor $R^\lambda{}_{\mu\nu\rho}$. In the presence of torsion, described by the torsion tensor $T^\rho{}_{\mu\nu}$, the infinitesimal parallelogram fails to close, providing a geometric measure of the antisymmetric part of the connection. Source: [3].

The commutator of two covariant derivatives acting on a vector field V^ρ is defined by

$$[\nabla_\mu, \nabla_\nu]V^\rho \equiv \nabla_\mu\nabla_\nu V^\rho - \nabla_\nu\nabla_\mu V^\rho. \tag{1.44}$$

Expanding each covariant derivative in terms of the Christoffel symbols $\Gamma^\rho_{\mu\nu}$ and antisymmetrising in $(\mu \leftrightarrow \nu)$, the partial derivatives $\partial_\mu\partial_\nu V^\rho$ and the symmetric terms in $\mu\nu$ cancel, leaving

$$[\nabla_\mu, \nabla_\nu]V^\rho = \left(\partial_\mu\Gamma^\rho_{\nu\sigma} - \partial_\nu\Gamma^\rho_{\mu\sigma} + \Gamma^\rho_{\mu\lambda}\Gamma^\lambda_{\nu\sigma} - \Gamma^\rho_{\nu\lambda}\Gamma^\lambda_{\mu\sigma}\right) V^\sigma - 2\,\Gamma^\lambda_{[\mu\nu]}\,\nabla_\lambda V^\rho. \tag{1.45}$$

The first group of terms defines the Riemann curvature tensor,

$$R^\rho{}_{\sigma\mu\nu} = \partial_\mu\Gamma^\rho_{\nu\sigma} - \partial_\nu\Gamma^\rho_{\mu\sigma} + \Gamma^\rho_{\mu\lambda}\Gamma^\lambda_{\nu\sigma} - \Gamma^\rho_{\nu\lambda}\Gamma^\lambda_{\mu\sigma}\ , \tag{1.46}$$

while the second involves the antisymmetric part of the connection,

$$T^{\lambda}_{\mu\nu} = \Gamma^{\lambda}_{\mu\nu} - \Gamma^{\lambda}_{\nu\mu} \,, \tag{1.47}$$

known as the torsion tensor. Since we have already introduced the concept of torsion earlier, we can now give a more precise interpretation: torsion quantifies the failure of an infinitesimal parallelogram to close. Specifically, if the vector dx^{μ} is parallel transported along dx^{ν} to yield dx'^{μ}, and dx^{ν} is parallel transported along dx^{μ} to yield dx'^{ν}, these two resulting vectors will not, in general, meet at the same point. In other words, the sequence $dx^{\mu} \to dx^{\nu} \to dx'^{\mu} \to dx'^{\nu}$ does not necessarily form a closed parallelogram. Only for symmetric connections (such as the Levi-Civita connection) does the torsion tensor vanish, ensuring that the paths dx^{μ}-dx^{ν} and dx^{ν}-dx^{μ} yield a closed loop [3].

When torsion vanishes, only the curvature term remains, providing a direct measure of the manifold's geometric curvature. In this case, the commutator can be expressed compactly as

$$[\nabla_{\mu}, \nabla_{\nu}]V^{\rho} = R^{\rho}{}_{\sigma\mu\nu}V^{\sigma}. \tag{1.48}$$

Thus, the Riemann curvature tensor provides a local, coordinate-independent measure of how much the geometry of a manifold deviates from flatness, encapsulating all such effects within a precise mathematical framework.

1.2.5.5.1 Symmetry properties

The Riemann tensor is antisymmetric in its last two indices:

$$R^{\rho}{}_{\sigma\mu\nu} = -R^{\rho}{}_{\sigma\nu\mu} \,, \tag{1.49}$$

and possesses further symmetries that significantly reduce the number of independent components. In an n-dimensional manifold, it initially has n^4 components, but these symmetries lower the count to 1 in two dimensions, 6 in three, and 20 in four (as in spacetime geometry).

It is often convenient to express all indices in lowered form using the metric tensor, as this highlights the tensor's symmetry properties in a more transparent way:

$$R_{\alpha\beta\mu\nu} = g_{\alpha\lambda}R^{\lambda}{}_{\beta\mu\nu} \,. \tag{1.50}$$

In this form, the characteristics of $R_{\alpha\beta\mu\nu}$ can be presented as follows:

- **Antisymmetry in each index pair:** The tensor changes sign when the first two indices are swapped, and likewise for the last two:

  $$R_{\alpha\beta\mu\nu} = -R_{\beta\alpha\mu\nu}, \quad R_{\alpha\beta\mu\nu} = -R_{\alpha\beta\nu\mu} \,. \tag{1.51}$$

  This means that both (α, β) and (μ, ν) behave as antisymmetric index pairs.

- **Symmetry under exchange of index pairs:** Interchanging the first and second pair of indices leaves the tensor unchanged:

$$R_{\alpha\beta\mu\nu} = R_{\mu\nu\alpha\beta} \ . \tag{1.52}$$

This symmetry links curvature measured in one pair of directions to curvature in the orthogonal pair.

- **First Bianchi identity:** The cyclic sum over the last three indices vanishes:

$$R_{\alpha\beta\mu\nu} + R_{\alpha\nu\beta\mu} + R_{\alpha\mu\nu\beta} = 0 \ . \tag{1.53}$$

This is an algebraic identity that holds in any geometry and reflects fundamental constraints on how curvature can be arranged in a manifold.

In the special case of a flat manifold, all curvature components vanish:

$$R^{\alpha}{}_{\beta\mu\nu} = 0 \ , \tag{1.54}$$

implying that parallel transport is path-independent: any vector transported along a closed curve returns to its starting point unchanged [44].

1.2.5.6 Ricci tensor and Ricci scalar

The complete, uncontracted Riemann curvature tensor does not explicitly appear in EFE. Rather, it is contracted to yield two curvature invariants: the Ricci tensor and the Ricci scalar. These quantities are subsequently combined to construct the Einstein tensor, which forms the geometric side (left-hand side) of EFE [34].

1.2.5.6.1 Ricci tensor

By contracting the first (upper) index of the Riemann tensor with its third index, one obtains the Ricci tensor:

$$R_{\mu\nu} = R^{\rho}{}_{\mu\rho\nu} = g^{\rho\lambda} R_{\mu\rho\nu\lambda} \ . \tag{1.55}$$

Therefore, it is expressed in terms of the connection as follows:

$$R_{\mu\nu} = \partial_\rho \Gamma^{\rho}_{\mu\nu} - \partial_\nu \Gamma^{\rho}_{\rho\mu} + \Gamma^{\rho}_{\rho\lambda}\Gamma^{\lambda}_{\mu\nu} - \Gamma^{\rho}_{\nu\lambda}\Gamma^{\lambda}_{\rho\mu} \ . \tag{1.56}$$

1.2.5.6.2 Ricci scalar

By contracting both indices of the Ricci tensor, one obtains the Ricci scalar or curvature scalar:

$$R = R^{\mu}{}_{\mu} = g^{\mu\nu} R_{\mu\nu} \ . \tag{1.57}$$

As the name suggests, the Ricci scalar is a scalar and thus an invariant under general coordinate transformations.

1.3 Einstein field equations

The mathematical framework developed in the preceding sections, together with the physical insight of the equivalence principle, leads naturally to the central result of GR. The EFE relate the curvature of spacetime, encoded in the Einstein tensor

$$G_{\mu\nu} \equiv R_{\mu\nu} - \frac{1}{2} g_{\mu\nu} R \ , \tag{1.58}$$

to the energy and momentum content of matter, described by the energy-momentum tensor $T_{\mu\nu}$ (see §1.3.1):

$$G_{\mu\nu} = \frac{8\pi G}{c^4} T_{\mu\nu} \ . \tag{1.59}$$

This is a system of ten coupled, second-order nonlinear partial differential equations whose solutions determine the metric $g_{\mu\nu}$ given a distribution of matter and energy. The constant $8\pi G/c^4$ ensures consistency with Newtonian gravity in the weak-field limit. In vacuum, $T_{\mu\nu} = 0$, the equations reduce to

$$R_{\mu\nu} = \frac{1}{2} g_{\mu\nu} R \ , \tag{1.60}$$

which further simplifies, upon contraction, to $R = 0$ and therefore

$$R_{\mu\nu} = 0 \ , \tag{1.61}$$

known as the vacuum Einstein equations. Although the equations formally involve ten components, the contracted Bianchi identity

$$\nabla^{\mu} G_{\mu\nu} = 0 \ , \tag{1.62}$$

provides four constraints, leaving only six independent equations. The remaining four degrees of freedom reflect the freedom in the choice of coordinates [3].

Einstein's equations are highly nonlinear, which makes them extremely challenging to solve analytically. No general method is known, and exact solutions typically rely on assuming a high degree of symmetry. Nevertheless, they encode a remarkably rich physical content: they describe the propagation of GWs, the bending of light, the expansion of the universe, and the formation of BHs, all consequences of the same geometric framework.

1.3.1 The energy-momentum tensor

The energy-momentum tensor $T^{\mu\nu}$ appears on the right-hand side of the EFE as the source of spacetime curvature. In Newtonian physics, mass alone generates the gravitational field, but special relativity establishes that mass is just one manifestation of energy through

$$E^2 = p^2 c^2 + m^2 c^4 \ . \tag{1.63}$$

The source of gravity in GR must therefore encompass not only mass but all forms of energy and momentum. The tensor $T^{\mu\nu}$ encapsulates these contributions by describing the flux of four-momentum p^{μ} across a hypersurface of constant x^{ν}, and its components can be represented as a

4×4 matrix (see Fig. 1.7): T_{00} is the energy density, $T_{0i} = T_{i0}$ the momentum density or energy flux, and T_{ij} the momentum flux, with diagonal components giving the isotropic pressure and off-diagonal ones the shear stress.

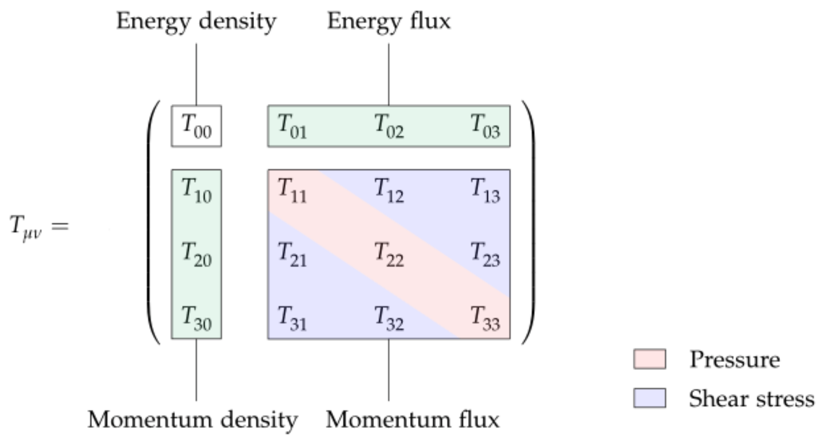

Figure 1.7: Schematic representation of the energy-momentum tensor $T_{\mu\nu}$ as a 4×4 matrix. Source: [45].

For a perfect fluid in flat spacetime, modeled as a system fully characterized by its energy density ρ and isotropic pressure P, the tensor takes the form [43]

$$T^{\mu\nu} = (\rho + P)\,\frac{U^\mu U^\nu}{c^2} + P\,\eta^{\mu\nu}\ , \tag{1.64}$$

where $U^\mu = \gamma(c, \vec{v})$ is the four-velocity of the fluid; the relative sign between the two terms is fixed by the signature adopted here, so that the spatial diagonal components reduce to the isotropic pressure P rather than $-P$. In the comoving frame, $U^\mu = (c, 0, 0, 0)$ and the tensor reduces to $T^{\mu\nu} = \mathrm{diag}(\rho, P, P, P)$.

A key property of $T^{\mu\nu}$ is its conservation law

$$\nabla_\mu T^{\mu\nu} = 0\ , \tag{1.65}$$

which ensures local conservation of energy and momentum in any physically meaningful system, both in flat and curved spacetime [34].

1.3.2 The cosmological constant

Among the possible modifications of the EFE (1.59), a particularly important one is the introduction of a cosmological constant Λ [43]. Including this term modifies the field equations to

$$R_{\alpha\beta} - \tfrac{1}{2} g_{\alpha\beta} R + g_{\alpha\beta}\Lambda = \frac{8\pi G}{c^4}\, T_{\alpha\beta}\,. \tag{1.66}$$

Historically, Einstein introduced Λ in 1917 as a way to obtain a static universe without beginning or end, effectively behaving like a repulsive pressure. However, the static solutions were later found to be unstable, and in 1929 Edwin Hubble (1889–1953) discovered that the universe is in fact expanding, a result fully consistent with the original unmodified equations. Einstein reportedly referred to the cosmological constant as "the biggest blunder" of his life, as described by George Gamow, though the remark is not documented in Einstein's own writing.

Ironically, observations of distant Type Ia supernovae in 1998, carried out by the Supernova Cosmology Project and the High-Z Supernova Search Team, revealed that the expansion is accelerating. One possible explanation invokes a pervasive dark energy effectively described by Λ, which has thus regained scientific relevance despite its troubled history [34]. The modern challenge lies in the so-called cosmological constant problem: quantum field theory predicts $\Lambda \sim m_P^4 \sim 10^{76}\,\mathrm{GeV}^4$, while observations constrain it to $\sim 10^{-47}\,\mathrm{GeV}^4$, a discrepancy of roughly 10^{123} and the largest known mismatch between theory and experiment [3].

1.4 Relevant solutions to Einstein's equations

EFE are extremely challenging to solve because of their intrinsic nonlinearity. Unlike linear systems, the superposition of two independent solutions does not in general produce a new one. The physical origin of this nonlinearity is simple: spacetime curvature is generated by the presence of matter and energy, yet the curvature itself carries energy and momentum, thereby acting as an additional source of curvature. In other words, gravity is not only coupled to matter and energy but also to itself, which makes the field equations nonlinear by nature.

Concretely, to "find a solution" means to determine a metric tensor $g_{\mu\nu}$ (and, when relevant, a compatible matter configuration) that satisfies

$$R_{\mu\nu} - \tfrac{1}{2} g_{\mu\nu} R + \Lambda g_{\mu\nu} = \frac{8\pi G}{c^4} T_{\mu\nu}, \tag{1.67}$$

subject to appropriate boundary or initial conditions.

Each solution corresponds to a possible spacetime geometry associated with a particular distribution of matter and energy. Broadly speaking, they can be grouped into three categories:

- **Exact analytical solutions**: such as the Schwarzschild, Kerr, or Friedmann-Lemaître-Robertson-Walker metrics.
- **Approximate solutions**: derived through perturbative methods, for instance weak-field expansions describing GW.
- **Numerical solutions**: obtained via large-scale simulations, such as those modelling the merger of two BHs.

The vast variety of solutions illustrates the richness of GR: each one represents a distinct possible universe governed by Einstein's equations.

Einstein himself initially believed that his equations were so complicated that no exact solution would ever be found. Remarkably, only a few months after the formulation of GR in 1915, Karl Schwarzschild (1873–1916) obtained the first exact solution, describing the spacetime of a static, spherically symmetric object. Since then, dozens, indeed hundreds, of exact solutions have been discovered [3]. In this section we will explore the Schwarzschild solution.

1.4.1 The Schwarzschild solution

In vacuum, EFE reduce to $R_{\mu\nu} = 0$. At first sight, one might be tempted to consider the trivial choice $g_{\mu\nu} = 0$. However, this is not allowed: the metric must be a non-degenerate tensor field with a well-defined inverse, a requirement already assumed in the derivation of the field equations from the variational principle[4]. This condition, $\det g_{\mu\nu} \neq 0$, immediately implies that the simplest Ricci-flat[5] solution is not the trivial one but rather Minkowski spacetime, which corresponds to flat geometry in the absence of matter or energy:

$$ds^2 = -c^2 dt^2 + dx^2 + dy^2 + dz^2 \ . \tag{1.68}$$

The first non-trivial exact solution of the vacuum equations was found by Karl Schwarzschild in 1916. The Schwarzschild metric describes the spacetime exterior to a static, spherically symmetric, non-rotating mass distribution. In Schwarzschild coordinates (t, r, θ, ϕ), the line element takes the form

$$ds^2 = -\left(1 - \frac{2GM}{c^2 r}\right) c^2 dt^2 + \left(1 - \frac{2GM}{c^2 r}\right)^{-1} dr^2 + r^2 d\theta^2 + r^2 \sin^2\theta \, d\phi^2 \, , \tag{1.69}$$

where M denotes the mass of the central object [34].

This solution has profound physical implications. It provides the theoretical framework to compute relativistic corrections to planetary orbits (e.g., the perihelion precession of Mercury), as well as to predict the bending of light by gravity. Most notably, it predicts the existence of BHs, one of the most striking consequences of GR that we will see in this section.

It is important to understand that there are two different Schwarzschild solutions. The first one is the exterior Schwarzschild solution (1.69), which applies to the region outside a spherical body of mass M and radius $R_0 > 2GM/c^2$. Since this exterior region is empty space, the solution is a vacuum one and describes the gravitational field produced by the object. According to Birkhoff's theorem, every spherically symmetric vacuum solution of the EFE is locally isometric to the Schwarzschild metric, and is therefore necessarily static in the exterior region, even if the source itself is not [3].

The second is the interior Schwarzschild solution[6], also formulated by Karl Schwarzschild in 1916, which applies inside the object ($r < R_0$) and models its matter as a perfect fluid. These two solutions are fundamentally different, but when the radius of the object satisfies $R_0 > 2GM/c^2$, they match smoothly at the boundary $r = R_0$.

An important characteristic of the Schwarzschild exterior solution is that when the radial coordinate r becomes very large, the metric tends to the Minkowski form expressed in spherical

[4] Ultimately, this constraint likely suggests that the gravitational field is not fundamental and should be replaced by another description in scenarios where $\det g_{\mu\nu}$ becomes small [36].

[5] Metrics that satisfy the vacuum equations are referred to as Ricci-flat. Although Ricci-flat solutions are generally not flat, they still represent valid vacuum solutions (assuming no cosmological constant is present) [3].

[6] The solution is given by $ds^2 = -\frac{1}{4}\left[3\sqrt{1 - \frac{2GM}{c^2 R_0}} - \sqrt{1 - \frac{2GMr^2}{c^2 R_0^3}}\right]^2 c^2 dt^2 + \left(1 - \frac{2GMr^2}{c^2 R_0^3}\right)^{-1} dr^2 + r^2 d\Omega^2$, where R_0 denotes the radius of the object [46].

coordinates, ensuring consistency with the classical limit:

$$\lim_{r\to\infty}\left(1-\frac{2GM}{c^2r}\right)=1\ . \tag{1.70}$$

This means that, far away from the source, the gravitational field is essentially negligible and spacetime is indistinguishable from flat Minkowski space[7].

1.4.1.1 The Kretschmann scalar and the Schwarzschild radius

As $r \to 0$, the metric components diverge, signaling a singularity where spacetime curvature becomes infinite and the laws of physics break down [47]. Metrics derived from EFE can present two types of singularities. Coordinate singularities are artefacts of the chosen coordinate system with no physical content: for instance, $g_{\phi\phi} = 0$ at $\theta = 0, \pi$ in spherical coordinates, which vanishes under a simple rotation. Physical singularities correspond to genuine divergences of spacetime curvature, and Penrose showed in 1965 that such singularities necessarily exist inside any closed trapped surface [48].

To distinguish between the two, one computes curvature invariants constructed from contractions of the Riemann tensor, which are independent of the coordinate choice. For the Schwarzschild solution, the Kretschmann scalar

$$K \equiv R_{\mu\nu\rho\sigma}R^{\mu\nu\rho\sigma} = \frac{48G^2M^2}{c^4r^6}, \tag{1.71}$$

diverges only at $r = 0$, confirming a physical singularity there, while remaining finite at $r = 2GM/c^2$. The apparent pathology of the metric at $r = 2GM/c^2$ is therefore a coordinate singularity. Nevertheless, this radius has profound physical significance: it defines the Schwarzschild radius and marks the event horizon of the BH, a surface beyond which standard Schwarzschild coordinates become inadequate.

1.4.1.2 Causal structure and radial null geodesics

Setting $ds^2 = 0$ and $d\theta = d\phi = 0$, the slope of radial null geodesics in the (t, r) plane is

$$\frac{dt}{dr} = \pm\frac{1}{c}\left(1-\frac{2GM}{c^2r}\right)^{-1}, \tag{1.72}$$

which recovers the Minkowski value $\pm1/c$ at large r but diverges as $r \to 2GM/c^2$, causing the light cones to close up (see Fig. 1.8). Inside the horizon, the cones flip orientation entirely, forcing all causal trajectories toward the singularity at $r = 0$. Integrating the null geodesic equation gives

$$ct = \pm\left(r + \frac{2GM}{c^2}\ln\left|\frac{c^2r}{2GM}-1\right| + C_0\right), \tag{1.73}$$

showing that an outgoing light ray requires infinite Schwarzschild coordinate time to reach $r = 2GM/c^2$. A distant observer therefore sees an infalling body freeze and fade at the horizon, while

[7] Spacetimes that exhibit this behavior are referred to as asymptotically flat and can be regarded as representing isolated systems [34].

in proper time the crossing occurs in a finite interval. The issue lies in the coordinates, not the physics.

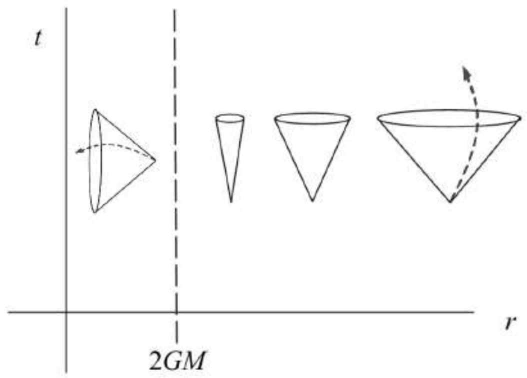

(a) Tilting of light cones in Schwarzschild coordinates.

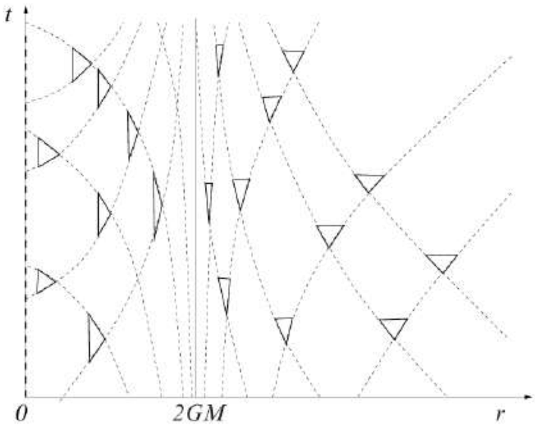

(b) Ingoing and outgoing null geodesics in Schwarzschild coordinates.

Figure 1.8: Light cones in the t-r plane of Schwarzschild spacetime. Far from the BH, cones open symmetrically at $\pm 45°$. As $r \to 2GM/c^2$ they progressively close; inside the horizon they invert, pointing inexorably toward $r = 0$. Original illustrations based on: [3, 43].

To follow geodesics continuously across the horizon, one introduces the advanced Eddington-Finkelstein coordinates [49],

$$c\tilde{t} = ct + \frac{2GM}{c^2} \log \left| \frac{c^2 r}{2GM} - 1 \right| , \tag{1.74}$$

in which the metric becomes

$$ds^2 = -\left(1 - \frac{2GM}{c^2 r}\right) c^2 d\tilde{t}^2 + \frac{4GM}{c\, r}\, d\tilde{t}\, dr + \left(1 + \frac{2GM}{c^2 r}\right) dr^2 + r^2 d\Omega^2 , \tag{1.75}$$

manifestly regular at $r = 2GM/c^2$ (see Fig. 1.9).

1.4.1.3 Black holes in general relativity

BHs are regions of spacetime where gravity is so intense that nothing can escape. They form when matter becomes sufficiently compressed, creating a boundary called the event horizon: a one-way surface beyond which all trajectories lead toward a central singularity where curvature diverges (see Fig. 1.10a).

According to the no-hair theorem, an isolated BH in GR is fully characterized by three quantities: mass M, electric charge Q, and angular momentum J [34]. Since astrophysical BHs are expected to be nearly neutral, the relevant solutions are the Schwarzschild metric (1.69) for non-rotating BHs and the Kerr metric for rotating ones. The Kerr solution describes a stationary,

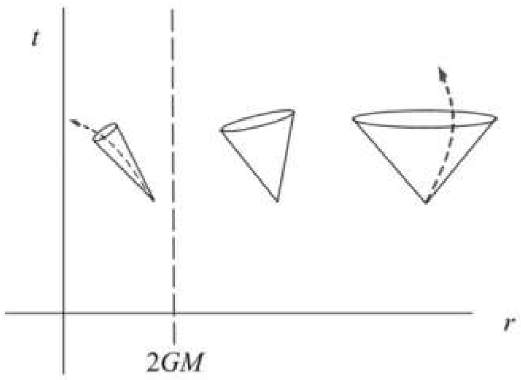

(a) Light cones in Eddington-Finkelstein coordinates.

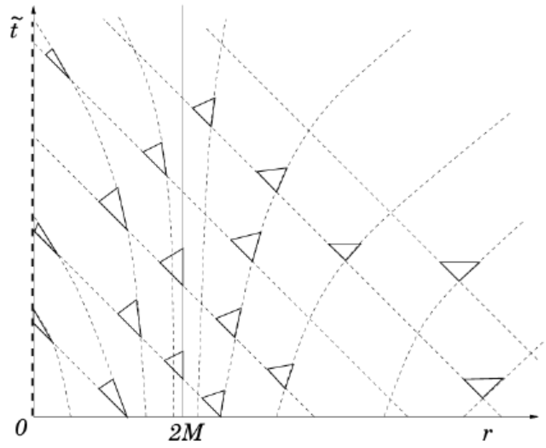

(b) Null geodesics in Eddington-Finkelstein coordinates.

Figure 1.9: In Eddington-Finkelstein coordinates the horizon is regular. Ingoing geodesics always remain at 45°; outgoing ones tilt progressively inward until at $r = 2GM/c^2$ they are trapped, and for $r < 2GM/c^2$ all light falls toward the singularity. Original illustrations based on: [3, 34].

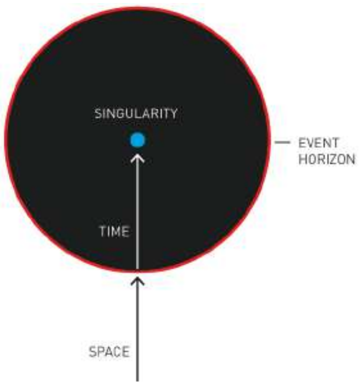

(a)

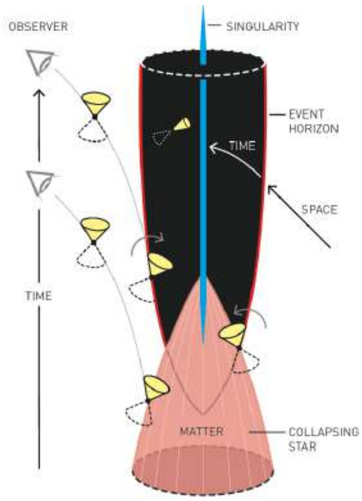

(b)

Figure 1.10: (a) Simplified depiction of a BH. The event horizon (red line) marks the boundary beyond which nothing can escape. (b) Causal structure of a Schwarzschild BH: light cones tip inward near the horizon and point entirely toward the singularity inside it. Source: [50].

axisymmetric spacetime with line element (in Boyer-Lindquist coordinates):

$$\begin{aligned} ds^2 = &-\left(1 - \frac{2GMr}{\Sigma c^2}\right) c^2 dt^2 - \frac{4GMar\sin^2\theta}{\Sigma c^2}\, c\, dt\, d\phi \\ &+ \frac{\Sigma}{\Delta} dr^2 + \Sigma\, d\theta^2 + \left(r^2 + a^2 + \frac{2GMa^2 r \sin^2\theta}{\Sigma c^2}\right) \sin^2\theta\, d\phi^2 \ , \end{aligned} \tag{1.76}$$

where $a = J/(Mc)$ is the spin parameter, $\Sigma = r^2 + a^2\cos^2\theta$, and $\Delta = r^2 - 2GMr/c^2 + a^2$ [34]. In the limit $a \to 0$, this reduces to the Schwarzschild solution. Most astrophysical BHs are expected to rotate, inheriting angular momentum from their progenitor stars or accretion processes.

BHs are classified by mass into three observationally relevant categories:

- **Stellar-mass BHs:** from a few to $\sim 100\,M_\odot$, formed by collapse of massive stars.
- **Intermediate-mass BHs (IMBHs):** 10^2–$10^5\,M_\odot$. The strongest evidence comes from GW events such as GW190521 ($142^{+28}_{-16}\,M_\odot$ remnant) [51] and GW231123 ($225^{+26}_{-43}\,M_\odot$ remnant) [52].
- **Supermassive BHs (SMBHs):** 10^6–$10^{10}\,M_\odot$, found at the centers of galaxies (e.g. Sagittarius A*, $\sim 4 \times 10^6\,M_\odot$).

Additionally, primordial BHs may have formed from density fluctuations in the early universe, spanning a wide mass range from $\sim 10^{-5}$ g to $\sim 10^5\,M_\odot$ [53].

In summary, this chapter has laid the mathematical and physical foundations of GR, progressing from the geometry of curved spacetime and the tools of differential geometry (parallel transport, geodesics, and the Riemann tensor) to the EFE and their most astrophysically relevant solutions: BHs. These concepts constitute the theoretical foundation upon which the remainder of this thesis is built. In the next chapter we turn to the central subject of this work: GWs, which emerge naturally as propagating solutions of the linearized Einstein equations and represent one of the most profound predictions of GR.

Chapter 2

Gravitational wave theory

We have now reached the central topic of this thesis and one of the most remarkable predictions of Einstein's GR: GWs. In this chapter, we introduce their theoretical foundations, beginning with their natural emergence from the linearized EFE (§2.1), and characterize their polarization states and physical action on test masses via the geodesic deviation equation. We then survey the main astrophysical source classes (§2.2) and discuss the detection principle underlying current interferometric observatories, covering antenna pattern functions, noise sources, and the global detector network (§2.3). The chapter closes with an overview of the Bayesian parameter estimation (PE) framework used to extract source properties from GW signals (§2.4), establishing the statistical foundation for the data analysis presented in subsequent chapters.

2.1 Gravitational waves in linearized gravity

In 1916, shortly after formulating the final version of his general theory of relativity, Einstein predicted the existence of GWs [54].[1] He showed that, within the weak-field or linearized approximation of his equations, spacetime admits wave-like solutions: transverse oscillations of the metric that propagate at the speed of light. These waves are produced by time-dependent variations of the mass quadrupole moment of the source. Einstein also realized that their amplitudes would be extremely small, which led many to doubt their physical reality for decades[2] [5].

GWs can be understood as ripples in the curvature of spacetime that carry energy and information about their astrophysical sources. They are generated by asymmetric acceleration of mass or energy, typically in compact binary systems involving BHs (BH-BH), NSs (NS-NS), or mixed BH-NS mergers. The most direct and elegant derivation of their existence arises from the linearized formulation of GR, obtained by expanding the EFE around the flat Minkowski metric $\eta_{\mu\nu}$ [57].

The first indirect evidence for GWs came from the discovery of the binary pulsar system PSR B1913+16 by Hulse and Taylor in 1974 [58]. Subsequent measurements of its orbital energy loss provided compelling proof of their existence. This milestone, together with advances in astrophysics, established the foundation for direct GW astronomy [5].

Nearly a century after Einstein's prediction, on September 14, 2015, the LIGO detectors achieved the first direct detection of GWs, GW150914, announced jointly by the LIGO Scientific Collaboration and the Virgo Collaboration [5], a signal originating from a binary black hole merger

[1] Einstein's original 1916 derivation of the quadrupole formula contained a numerical error in its coefficient. A 1918 follow-up paper corrected most of the 1916 mistakes, but a remaining factor of two error in the coefficient was only identified and fixed by Eddington in 1922 [55].

[2] Indeed, until the Chapel Hill conference in 1957 [56], the question of whether GWs were genuine physical phenomena remained under debate.

(BBH). This event not only confirmed Einstein's theoretical insight but also inaugurated a new era of observational astrophysics, allowing us to probe relativistic systems in the dynamic, strong-field regime of gravity.

2.1.1 Linearized gravity

GR describes gravity as the manifestation of spacetime curvature. The dynamics of the gravitational field are governed by EFE [57],

$$G_{\mu\nu} \equiv R_{\mu\nu} - \frac{1}{2} g_{\mu\nu} R = \frac{8\pi G}{c^4} T_{\mu\nu}, \tag{2.1}$$

where $R_{\mu\nu}$ is the Ricci tensor, R the Ricci scalar, and $T_{\mu\nu}$ the energy-momentum tensor.

The curvature of spacetime is encoded in the Riemann tensor, which is constructed from the Christoffel symbols. For a torsion free connection, the Christoffel symbols are defined as [3]

$$\Gamma^{\rho}_{\mu\nu} = \frac{1}{2} g^{\rho\lambda} \left(\partial_\mu g_{\lambda\nu} + \partial_\nu g_{\lambda\mu} - \partial_\lambda g_{\mu\nu} \right). \tag{2.2}$$

In terms of these quantities, the components of the Riemann curvature tensor are given by

$$R^{\lambda}{}_{\rho\mu\nu} = \partial_\mu \Gamma^{\lambda}_{\rho\nu} - \partial_\nu \Gamma^{\lambda}_{\rho\mu} + \Gamma^{\lambda}_{\sigma\mu} \Gamma^{\sigma}_{\rho\nu} - \Gamma^{\lambda}_{\sigma\nu} \Gamma^{\sigma}_{\rho\mu}. \tag{2.3}$$

The Riemann tensor fully characterizes the intrinsic curvature of spacetime and vanishes identically in flat manifolds. By contracting its indices, one obtains the Ricci tensor [34],

$$R_{\mu\nu} = R^{\rho}{}_{\mu\rho\nu}, \tag{2.4}$$

and the Ricci scalar,

$$R = g^{\mu\nu} R_{\mu\nu}, \tag{2.5}$$

which together determine the Einstein tensor appearing in the field equations. These geometric quantities provide the starting point for the linearized description of gravity developed in the following sections.

2.1.2 The weak field limit

As a first step toward the understanding of GWs, we wish to study the expansion of the Einstein equations around the flat space metric in the so called weak field limit. Mathematically, we can write this as

$$g_{\mu\nu} \approx \eta_{\mu\nu} + h_{\mu\nu}, \qquad |h_{\mu\nu}| \ll 1, \tag{2.6}$$

where $\eta_{\mu\nu}$ is the flat space Minkowski metric $\eta_{\mu\nu} = \mathrm{diag}(-1, +1, +1, +1)$ and $h_{\mu\nu}$ represents a small perturbation [57].

The assumption that $h_{\mu\nu}$ is small allows us to neglect all terms beyond first order in this quantity. The resulting approximation is known as the linearized theory of gravity, and it leads to

$$g^{\mu\nu} \approx \eta^{\mu\nu} - h^{\mu\nu} + \mathcal{O}(h^2). \tag{2.7}$$

All index operations from this point onward will be performed using the Minkowski metric $\eta_{\mu\nu}$ instead of the full metric $g_{\mu\nu}$ [43]. For instance,

$$h^{\mu\nu} = \eta^{\mu\rho}\eta^{\nu\sigma}h_{\rho\sigma}. \tag{2.8}$$

Our goal is to determine the equation of motion satisfied by the perturbations $h_{\mu\nu}$, which can be obtained by expanding Einstein's equations to first order. We begin by considering the Christoffel symbols (2.2), given by

$$\Gamma^{\rho}_{\mu\nu} = \frac{1}{2}g^{\rho\lambda}\left(\partial_\mu g_{\nu\lambda} + \partial_\nu g_{\mu\lambda} - \partial_\lambda g_{\mu\nu}\right) \approx \frac{1}{2}\eta^{\rho\lambda}\left(\partial_\mu h_{\nu\lambda} + \partial_\nu h_{\mu\lambda} - \partial_\lambda h_{\mu\nu}\right). \tag{2.9}$$

Here, $g^{\rho\lambda}$ has been replaced by $\eta^{\rho\lambda}$, since all terms of order higher than $h_{\mu\nu}$ are neglected. Using these connection coefficients, we can compute the Riemann tensor (2.3) as

$$\begin{aligned} R^{\rho}{}_{\mu\nu\sigma} &= \partial_\nu\Gamma^{\rho}_{\mu\sigma} - \partial_\sigma\Gamma^{\rho}_{\mu\nu} + \Gamma^{\delta}_{\mu\sigma}\Gamma^{\rho}_{\delta\nu} - \Gamma^{\delta}_{\mu\nu}\Gamma^{\rho}_{\delta\sigma} \\ &\approx \partial_\nu\Gamma^{\rho}_{\mu\sigma} - \partial_\sigma\Gamma^{\rho}_{\mu\nu} + \mathcal{O}(h^2) \\ &\approx \frac{1}{2}\eta^{\rho\lambda}\left(\partial_\mu\partial_\nu h_{\lambda\sigma} - \partial_\nu\partial_\lambda h_{\mu\sigma} - \partial_\mu\partial_\sigma h_{\lambda\nu} + \partial_\lambda\partial_\sigma h_{\mu\nu}\right) + \mathcal{O}(h^2). \end{aligned} \tag{2.10}$$

Since the connection coefficients $\Gamma^{\rho}_{\mu\nu}$ are already first order quantities, the only contributions to the Riemann tensor in the linear approximation come from their derivatives, while the quadratic terms in Γ can be safely discarded [43].

From the Riemann tensor derived above, we can obtain the linearized Ricci tensor by contracting the indices ρ and ν:

$$\begin{aligned} R_{\mu\sigma} = R^{\rho}{}_{\mu\rho\sigma} &\approx \frac{1}{2}\eta^{\rho\lambda}(\partial_\mu\partial_\rho h_{\lambda\sigma} - \partial_\rho\partial_\lambda h_{\mu\sigma} - \partial_\mu\partial_\sigma h_{\lambda\rho} + \partial_\lambda\partial_\sigma h_{\mu\rho}) + \mathcal{O}(h^2) \\ &= \frac{1}{2}\Big(\partial_\mu\partial^\lambda h_{\lambda\sigma} - \partial_\rho\partial^\rho h_{\mu\sigma} - \partial_\mu\partial_\sigma h^{\rho}{}_{\rho} + \partial^\rho\partial_\sigma h_{\mu\rho}\Big) + \mathcal{O}(h^2) \\ &= \frac{1}{2}\Big(\partial_\mu\partial^\lambda h_{\lambda\sigma} + \partial^\rho\partial_\sigma h_{\mu\rho} - \Box h_{\mu\sigma} - \partial_\mu\partial_\sigma h\Big) + \mathcal{O}(h^2), \end{aligned} \tag{2.11}$$

where $h \equiv h^{\rho}{}_{\rho} \equiv \eta^{\rho\sigma}h_{\rho\sigma} + \mathcal{O}(h^2)$ is the trace of $h_{\mu\nu}$ and $\Box \equiv \partial_\rho\partial^\rho \equiv \eta^{\rho\sigma}\partial_\rho\partial_\sigma$ is the D'Alembertian operator [34].

The Ricci scalar is obtained by contracting the Ricci tensor with $\eta^{\mu\sigma}$:

$$\begin{aligned} R \approx \eta^{\mu\sigma}R_{\mu\sigma} &\approx \frac{1}{2}\Big(\partial_\mu\partial^\lambda h^{\mu}_{\lambda} + \partial^\rho\partial_\sigma h^{\sigma}_{\rho} - \Box h - \partial_\mu\partial^\mu h\Big) \\ &= \partial_\sigma\partial^\rho h^{\sigma}_{\rho} - \Box h = \partial^\nu\partial^\rho h_{\nu\rho} - \Box h, \end{aligned} \tag{2.12}$$

where we have interchanged the dummy indices ($\lambda \leftrightarrow \rho$) and ($\mu \leftrightarrow \sigma$) in the first term, and used $h^{\sigma}_{\rho} = \eta^{\sigma\nu}h_{\rho\nu}$ in the final equality.

Once we have these, using the expressions (2.11) and (2.12) we can compute the linearization of the Einstein tensor $G_{\mu\nu} = R_{\mu\nu} - (1/2)\eta_{\mu\nu}R$ for a given source $T_{\mu\sigma}$ and obtain the linearized form of the EFE (2.1)

$$\frac{1}{2}\Big(\partial_\mu\partial^\lambda h_{\lambda\sigma} + \partial^\rho\partial_\sigma h_{\mu\rho} - \Box h_{\mu\sigma} - \partial_\mu\partial_\sigma h\Big) - \frac{1}{2}\eta_{\mu\sigma}\left(\partial^\nu\partial^\rho h_{\nu\rho} - \Box h\right) = \frac{8\pi G}{c^4}T_{\mu\sigma}. \tag{2.13}$$

We first study the propagation of the resulting perturbation, appropriate far from the source where $T_{\mu\sigma} = 0$; the generation of GWs by a matter source is treated separately in §2.1.5. Setting $T_{\mu\sigma} = 0$ and multiplying through by 2, we obtain

$$2G_{\mu\sigma} \equiv \partial_\mu \partial^\lambda h_{\lambda\sigma} + \partial^\rho \partial_\sigma h_{\mu\rho} - \Box h_{\mu\sigma} - \partial_\mu \partial_\sigma h - \eta_{\mu\sigma} \left(\partial^\nu \partial^\rho h_{\nu\rho} - \Box h\right) = 0. \tag{2.14}$$

To simplify this expression, we introduce the trace reversed perturbation [57], defined as

$$\bar{h}_{\mu\nu} \equiv h_{\mu\nu} - \frac{1}{2}\eta_{\mu\nu} h. \tag{2.15}$$

Observe that the trace of the perturbation satisfies $\bar{h} \equiv \eta^{\mu\nu}\bar{h}_{\mu\nu} = h - 2h = -h$, which allows us to invert the definition in (2.15) to obtain

$$h_{\mu\nu} = \bar{h}_{\mu\nu} - \frac{1}{2}\eta_{\mu\nu}\bar{h}. \tag{2.16}$$

Using this relation, we can now express $2G_{\mu\sigma}$ entirely in terms of the trace reversed perturbation:

$$\begin{aligned} 2G_{\mu\sigma} \equiv{}& \partial_\mu \partial^\lambda \left(\bar{h}_{\lambda\sigma} - \tfrac{1}{2}\eta_{\lambda\sigma}\bar{h}\right) + \partial^\rho \partial_\sigma \left(\bar{h}_{\mu\rho} - \tfrac{1}{2}\eta_{\mu\rho}\bar{h}\right) - \Box\left(\bar{h}_{\mu\sigma} - \tfrac{1}{2}\eta_{\mu\sigma}\bar{h}\right) - \partial_\mu\partial_\sigma(-\bar{h}) \\ & - \eta_{\mu\sigma}\left[\partial^\nu\partial^\rho\left(\bar{h}_{\nu\rho} - \tfrac{1}{2}\eta_{\nu\rho}\bar{h}\right) - \Box(-\bar{h})\right] = \partial_\mu\partial^\lambda\bar{h}_{\lambda\sigma} - \tfrac{1}{2}\partial_\mu\partial_\sigma\bar{h} + \partial^\rho\partial_\sigma\bar{h}_{\mu\rho} \\ & - \tfrac{1}{2}\partial_\sigma\partial_\mu\bar{h} - \Box\bar{h}_{\mu\sigma} + \tfrac{1}{2}\eta_{\mu\sigma}\Box\bar{h} + \partial_\mu\partial_\sigma\bar{h} - \eta_{\mu\sigma}\left(\partial^\nu\partial^\rho\bar{h}_{\nu\rho} - \tfrac{1}{2}\Box\bar{h} + \Box\bar{h}\right). \end{aligned} \tag{2.17}$$

Finally, by collecting and simplifying all terms, we obtain the compact linearized form of the EFE in vacuum, written in terms of the trace reversed perturbation $\bar{h}_{\mu\nu}$:

$$2G_{\mu\sigma} \equiv \partial_\mu\partial^\lambda\bar{h}_{\lambda\sigma} + \partial^\rho\partial_\sigma\bar{h}_{\mu\rho} - \Box\bar{h}_{\mu\sigma} - \eta_{\mu\sigma}\partial^\nu\partial^\rho\bar{h}_{\nu\rho} = 0. \tag{2.18}$$

2.1.2.1 Gauge freedom and the Lorenz condition

With the linearized field equations at hand, the next step is to address the issue of gauge freedom. In GR, the theory remains covariant under arbitrary coordinate transformations. This means that even when we decompose the metric as a small perturbation around flat spacetime $g_{\mu\nu} = \eta_{\mu\nu} + h_{\mu\nu}$, the choice of coordinates is not unique. There can exist other coordinate systems in which the metric can still be expressed as the Minkowski background plus a small perturbation, but with a different $h_{\mu\nu}$. Hence, this decomposition is not unique [43].

In the linear approximation, an infinitesimal coordinate transformation

$$x^\mu \to x'^\mu = x^\mu + \xi^\mu(x), \tag{2.19}$$

induces a transformation of the perturbation tensor given by

$$h'_{\mu\nu} = h_{\mu\nu} - \partial_\mu\xi_\nu - \partial_\nu\xi_\mu, \tag{2.20}$$

where the derivatives $|\partial_\mu\xi_\nu|$ are taken to be of the same perturbative order as the metric perturbations $h_{\mu\nu}$.

To simplify the equations, we use the trace reversed perturbation (2.15) [57]:

$$\bar{h}'_{\mu\nu} = \bar{h}_{\mu\nu} - \partial_\mu \xi_\nu - \partial_\nu \xi_\mu + \eta_{\mu\nu} \partial_\rho \xi^\rho. \tag{2.21}$$

Taking the divergence of this expression leads to

$$\partial^\nu \bar{h}'_{\mu\nu} = \partial^\nu \bar{h}_{\mu\nu} - \Box \xi_\mu. \tag{2.22}$$

We can now exploit the coordinate freedom by choosing the vector ξ_μ to satisfy

$$\Box \xi_\mu = \partial^\nu \bar{h}_{\mu\nu}. \tag{2.23}$$

With this choice, the transformed perturbation fulfills

$$\partial^\nu \bar{h}'_{\mu\nu} = 0, \tag{2.24}$$

which defines the Lorenz gauge condition [57] (also called the harmonic gauge or the De Donder gauge):

$$\partial^\nu \bar{h}_{\mu\nu} = 0. \tag{2.25}$$

Imposing the Lorenz condition simplifies the linearized Einstein equations (2.18) to

$$\Box \bar{h}_{\mu\nu} = 0, \tag{2.26}$$

which is the standard wave equation in flat spacetime.

2.1.2.2 Solving the wave equation and the Transverse-Traceless gauge

We are now interested in solving the linearized equation (2.26) to find the form of GWs. The general solution of the linearized wave equation can be written as a superposition of plane waves:

$$\bar{h}_{\mu\nu}(x) = \mathrm{Re}\left\{ A_{\mu\nu} e^{ik_\alpha x^\alpha} \right\}, \tag{2.27}$$

where $A_{\mu\nu}$ is a complex, symmetric polarization tensor and k_α is the wave four-vector [59].

Substituting this ansatz into the wave equation (2.26), we obtain

$$k^\sigma k_\sigma = 0. \tag{2.28}$$

This implies that, for non-trivial solutions $\bar{h}_{\mu\nu} \neq 0$, the wave four-vector must be null. That is, the perturbations propagate at the speed of light[3] (GWs in vacuum) [43].

Furthermore, applying the Lorenz gauge condition (2.25) to the wave equation (2.27), we can prove that:

$$A_{\mu\nu} k^\mu = 0, \tag{2.29}$$

[3] A null wave four-vector satisfies $k^\mu k_\mu = 0$, which in Minkowski spacetime with $k^\mu = (\omega/c, \mathbf{k})$ reads $\omega^2/c^2 - |\mathbf{k}|^2 = 0$. This relation implies $\omega = c\,|\mathbf{k}|$, so the phase velocity of the perturbation is $v = \omega/|\mathbf{k}| = c$. Therefore, GWs correspond to massless excitations of the gravitational field that propagate at the speed of light in vacuum [36].

which shows that GWs are transverse with respect to their direction of propagation.

After applying the Lorenz gauge condition (2.24), there still remains a residual gauge freedom that allows further simplification. This residual freedom is a consequence of the fact that the Lorenz gauge does not completely eliminate coordinate freedom. As follows from (2.22), under the transformation (2.19) the quantity $\partial^\nu \bar{h}_{\mu\nu}$ transforms homogeneously. As a result, the gauge condition $\partial^\nu \bar{h}_{\mu\nu} = 0$ remains invariant under additional coordinate transformations of the form $x^\mu \to x^\mu + \xi^\mu$, provided that the vector field ξ^μ satisfies

$$\Box \xi^\mu = 0. \tag{2.30}$$

If $\Box \xi^\mu = 0$, then the symmetric tensor

$$\xi_{\mu\nu} = \partial_\mu \xi_\nu + \partial_\nu \xi_\mu - \eta_{\mu\nu} \partial_\rho \xi^\rho, \tag{2.31}$$

also fulfills $\Box \xi_{\mu\nu} = 0$, since the flat d'Alembertian operator commutes with partial derivatives. Consequently, solutions of the wave equation for $\bar{h}_{\mu\nu}$ remain invariant under the subtraction of such terms.

This residual gauge freedom allows one to impose four additional conditions on $h_{\mu\nu}$. In particular, one may choose the gauge such that the trace vanishes, $h = 0$, implying $\bar{h}_{\mu\nu} = h_{\mu\nu}$. Furthermore, the remaining gauge freedom can be used to set $h_{0\mu} = 0$. The Lorenz condition then reduces to

$$\partial^j h_{ij} = 0, \tag{2.32}$$

while the traceless condition becomes $h^i{}_i = 0$.

In summary, we can impose the conditions

$$h_{0\mu} = 0, \qquad h^i{}_i = 0, \qquad \partial^j h_{ij} = 0, \tag{2.33}$$

which define the transverse-traceless (TT) gauge [57]. The first condition indicates that GWs possess only spatial components (since the symmetry enforces $h_{\mu 0} = h_{0\mu} = 0$). The second condition implies that the perturbation is traceless, meaning the sum of the diagonal elements of the perturbation matrix vanishes. Finally, the last relation reveals that the perturbation is transverse, such that the amplitude tensor is orthogonal to the direction of propagation.

As a result, the degrees of freedom of $A_{\mu\nu}$ are reduced from ten to only two independent components, indicating that a GW has only two physical degrees of freedom, corresponding to its two polarization states [60].

An important feature of working in the TT gauge is that the vanishing trace condition implies

$$\bar{h}_{\mu\nu} = h_{\mu\nu}. \tag{2.34}$$

Therefore, the plane wave solution can be expressed as

$$h^{\rm TT}_{\mu\nu} = A_{\mu\nu} e^{ik_\alpha x^\alpha}, \tag{2.35}$$

where $h^{\rm TT}_{\mu\nu}$ represents the metric perturbation in the TT gauge. Following the convention of

Eq. (2.27), here and in the remainder of this section the physical perturbation is understood as the real part of the complex expression, which is left implicit to lighten the notation.

2.1.3 Polarization states

Considering a wave propagating in the z-direction, we can choose the wave four-vector as $k^\mu = (\omega/c, 0, 0, \omega/c)$. The Lorenz gauge (2.25) implies $A_{\mu 0} + A_{\mu 3} = 0$[4] and the additional TT gauge conditions lead to $A_{\mu 0} = 0$. Thus, the only non-zero components of the amplitude tensor are A_{11}, A_{22}, and $A_{12} = A_{21}$. The traceless condition $A^i{}_i = 0$ further implies that $A_{11} + A_{22} = 0$. Defining $h_+^{TT}(t, z) \equiv A_+ \, e^{-i\omega(t-z/c)}$ and $h_\times^{TT}(t, z) \equiv A_\times \, e^{-i\omega(t-z/c)}$, the TT perturbation can be written as

$$h_{\mu\nu}^{TT}\,(t, z) = \begin{pmatrix} 0 & 0 & 0 & 0 \\ 0 & h_+^{TT} & h_\times^{TT} & 0 \\ 0 & h_\times^{TT} & -h_+^{TT} & 0 \\ 0 & 0 & 0 & 0 \end{pmatrix}_{\mu\nu} , \tag{2.36}$$

where h_+^{TT} and $h_\times^{TT}$ represent the two independent and orthogonal polarization states of the GW, commonly referred to as the plus and cross modes [36]. Returning from the complex to the real domain, the strain can finally be expressed as

$$h_{\mu\nu}^{TT}\,(t, z) = \begin{pmatrix} 0 & 0 & 0 & 0 \\ 0 & A_+ & A_\times & 0 \\ 0 & A_\times & -A_+ & 0 \\ 0 & 0 & 0 & 0 \end{pmatrix}_{\mu\nu} \cos\left[\,\omega\,(t - z/c)\right], \tag{2.37}$$

where A_+ and $A_\times$ denote the amplitudes of the two polarization modes.

2.1.4 Effect of a gravitational wave on test masses

The physical influence of a GW on a system of test particles is captured, coordinate independently, by the relative motion of nearby particles. This behavior is governed by the geodesic deviation equation [43], which describes how a GW modifies the separation between neighboring geodesics

$$a^\mu = \frac{D^2 S^\mu}{d\tau^2} = R^\mu{}_{\nu\rho\sigma}\, U^\nu U^\rho S^\sigma, \tag{2.38}$$

where S^μ denotes the separation between two neighboring geodesics, $U^\mu = dx^\mu/d\tau$ is the tangent vector to the reference geodesic, a^μ represents the relative acceleration between the two test particles, and $R^\mu{}_{\nu\rho\sigma}$ is the Riemann curvature tensor.

It is instructive to first examine the effect of a GW on a single particle before turning to an extended system of test masses. We consider a particle initially at rest in flat spacetime, prior to the arrival of the wave. An inertial reference frame is attached to this particle, with the z-axis chosen along the direction of propagation of an incoming TT GW. The particle subsequently follows a geodesic of the perturbed spacetime generated by the wave.

[4] If the gauge condition requires $h_{\mu 0} = 0$ for all spacetime points and the metric perturbation is written as $h_{\mu\nu} = A_{\mu\nu} e^{ik_\rho x^\rho}$, the exponential factor never vanishes. Therefore, the only way for $h_{\mu 0}$ to be identically zero is that the corresponding amplitude components vanish $A_{\mu 0} = 0$ (and, by symmetry, $A_{0\mu} = 0$) [36].

The motion of the particle is governed by the geodesic equation,

$$\frac{d^2x^\alpha}{d\tau^2} + \Gamma^\alpha_{\mu\nu}\frac{dx^\mu}{d\tau}\frac{dx^\nu}{d\tau} = \frac{dU^\alpha}{d\tau} + \Gamma^\alpha_{\mu\nu}U^\mu U^\nu = 0. \tag{2.39}$$

At $\tau = 0$, the particle is at rest in this frame, so that $U^\alpha = (c, 0, 0, 0) + \mathcal{O}(h)$. The acceleration induced by the GW at this instant is therefore given by

$$\left.\frac{dU^\alpha}{d\tau}\right|_{\tau=0} = -c^2\,\Gamma^\alpha_{00} = -\frac{c^2}{2}\eta^{\alpha\beta}\left(h_{\beta 0,0} + h_{0\beta,0} - h_{00,\beta}\right). \tag{2.40}$$

However, in the TT gauge all components $h_{0\mu}$ vanish. As a consequence,

$$\left.\frac{dU^\alpha}{d\tau}\right|_{\tau=0} = 0. \tag{2.41}$$

The four-velocity thus remains constant at all times, implying that the particle experiences no coordinate acceleration due to the passage of the GW. Its coordinate position therefore remains unchanged. This result shows that the motion of a single freely falling particle is insufficient to reveal the presence of a GW. Detectable effects arise only when considering the relative motion between multiple nearby test masses [60].

We now turn to the physically relevant case of a system of test particles, whose relative motion is invariant under coordinate transformations.

We consider a congruence of neighboring timelike geodesics describing freely falling test particles. In the absence of gravitational radiation, these particles are initially at rest with respect to a common inertial frame, so that their four-velocity is

$$U^\mu = (c, 0, 0, 0) + \mathcal{O}(h). \tag{2.42}$$

At some reference time $t = 0$, a plane GW propagating along the z-direction passes through the system. We work throughout in the TT gauge, where the only non vanishing components of the metric perturbation lie in the plane orthogonal to the direction of propagation.

The relative motion of neighboring particles is governed by the geodesic deviation equation (2.38). In the linearized regime of gravity, the Riemann tensor is already first order in the metric perturbation $h_{\mu\nu}$. Consequently, it is sufficient to evaluate the equation using the unperturbed four-velocity $U^\mu = (c, 0, 0, 0)$, and to identify the coordinate time t with the proper time along the geodesics, up to corrections of order $\mathcal{O}(h^2)$.

Using the already computed expression for the linearized Riemann tensor in the TT gauge,

$$R^\rho{}_{\mu\nu\sigma} = \frac{1}{2}\eta^{\rho\lambda}\left(\partial_\mu\partial_\nu h_{\lambda\sigma} - \partial_\nu\partial_\lambda h_{\mu\sigma} - \partial_\mu\partial_\sigma h_{\lambda\nu} + \partial_\lambda\partial_\sigma h_{\mu\nu}\right) + \mathcal{O}(h^2), \tag{2.43}$$

one finds that the only non vanishing components relevant for the geodesic deviation equation are

$$R^i{}_{00j} = \frac{1}{2}\eta^{i\lambda}\left(\partial_0\partial_0 h_{\lambda j} - \partial_0\partial_\lambda h_{0j} - \partial_0\partial_j h_{\lambda 0} + \partial_\lambda\partial_j h_{00}\right) + \mathcal{O}(h^2) = \frac{1}{2c^2}\,\partial_t^2 h_{ij} + \mathcal{O}(h^2), \tag{2.44}$$

where $i, j = 1, 2, 3$ label spatial coordinates and $\partial_0 = c^{-1}\partial_t$. Contracting with $U^0U^0 = c^2$, the equation of motion for the spatial components of the separation vector then reduces to

$$\frac{d^2 S^i}{dt^2} = \frac{1}{2}\,\partial_t^2 h_{ij}^{TT}\, S^j. \tag{2.45}$$

This result shows explicitly that the GW induces tidal forces only in the plane transverse to the direction of propagation. For a wave travelling along the z-axis, the longitudinal component S^3 remains unaffected, while the transverse components (S^1, S^2) encode the physical deformation of the system.

Integrating the equation we find

$$S^i(t) = S^i(0) + \frac{1}{2}\, h_{ij}^{TT}(t)\, S^j(0). \tag{2.46}$$

For a general superposition of the two polarizations, the metric perturbation in the TT gauge can be written as

$$h_{ij}^{TT}(t,z) = \begin{pmatrix} A_+ & A_\times & 0 \\ A_\times & -A_+ & 0 \\ 0 & 0 & 0 \end{pmatrix}_{ij} e^{-i\omega(t-z/c)}. \tag{2.47}$$

Restricting to the spatial components and inserting this expression into the solution of the geodesic deviation equation (2.46), one obtains

$$\begin{aligned} S^x(t) &= \left(1 + \frac{1}{2}A_+ e^{-i\omega(t-z/c)}\right) S^x(0) + \frac{1}{2}A_\times e^{-i\omega(t-z/c)}\, S^y(0), \\ S^y(t) &= \left(1 - \frac{1}{2}A_+ e^{-i\omega(t-z/c)}\right) S^y(0) + \frac{1}{2}A_\times e^{-i\omega(t-z/c)}\, S^x(0), \\ S^z(t) &= S^z(0), \end{aligned} \tag{2.48}$$

which are coupled equations for the transverse components of the separation vector.

Since the geodesic deviation equation is linear in the metric perturbation $h_{\mu\nu}^{TT}$, the response of the system to a GW is also linear, so the effects of a general superposition of the two polarizations are simply the sum of their individual responses. It is therefore sufficient to analyze purely $+$ and purely $\times$ polarized waves separately, without loss of generality [36].

If the cross polarization vanishes, $A_\times = 0$, the metric perturbation reduces to

$$h_{ij}^{TT}(t,z) = \begin{pmatrix} A_+ & 0 & 0 \\ 0 & -A_+ & 0 \\ 0 & 0 & 0 \end{pmatrix}_{ij} e^{-i\omega(t-z/c)}. \tag{2.49}$$

The corresponding evolution of the separation vector is then

$$\begin{aligned} S^x(t) &= \left(1 + \frac{1}{2}A_+ e^{-i\omega(t-z/c)}\right) S^x(0), \\ S^y(t) &= \left(1 - \frac{1}{2}A_+ e^{-i\omega(t-z/c)}\right) S^y(0), \\ S^z(t) &= S^z(0), \end{aligned} \tag{2.50}$$

which stretches and squeezes the particle distribution along the transverse axes (see Fig. 2.1).

If the plus polarization vanishes, $A_+ = 0$, the metric perturbation becomes

$$h_{ij}^{TT}(t,z) = \begin{pmatrix} 0 & A_\times & 0 \\ A_\times & 0 & 0 \\ 0 & 0 & 0 \end{pmatrix}_{ij} e^{-i\omega(t-z/c)}. \tag{2.51}$$

In this case, the solution of the geodesic deviation equation reads

$$\begin{aligned} S^x(t) &= S^x(0) + \frac{1}{2}A_\times e^{-i\omega(t-z/c)}\, S^y(0), \\ S^y(t) &= S^y(0) + \frac{1}{2}A_\times e^{-i\omega(t-z/c)}\, S^x(0), \\ S^z(t) &= S^z(0). \end{aligned} \tag{2.52}$$

This polarization induces a shear deformation in the transverse plane, with principal axes rotated by 45° with respect to the x and y directions (see Fig. 2.1).

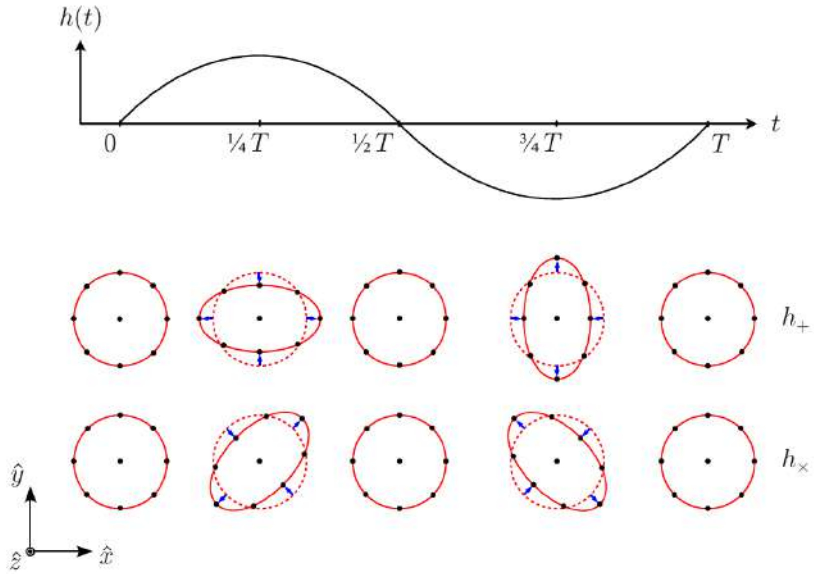

Figure 2.1: A monochromatic GW with angular frequency $\omega = 2\pi/T$ propagates in the $\hat{z}$ direction. The lower panel illustrates the deformation produced by the + and × polarization modes on a ring of freely falling test particles in a local inertial frame. Source: [35].

2.1.5 Generation of gravitational waves

In the previous sections we described GWs as propagating perturbations of the spacetime metric and analyzed their effect on matter. We now turn to the inverse question: how dynamical matter sources generate gravitational radiation.

A conceptual subtlety immediately appears. GW emission can take place in strongly non-linear regimes, such as the late stages of CBCs, where the full Einstein equations are required to describe the dynamics near the source. However, the radiation detected far from the source is typically very weak, so that the metric can be written as a small perturbation of flat spacetime. This justifies studying GW generation in the local wave zone within the framework of linearized gravity (see Fig. 2.2).

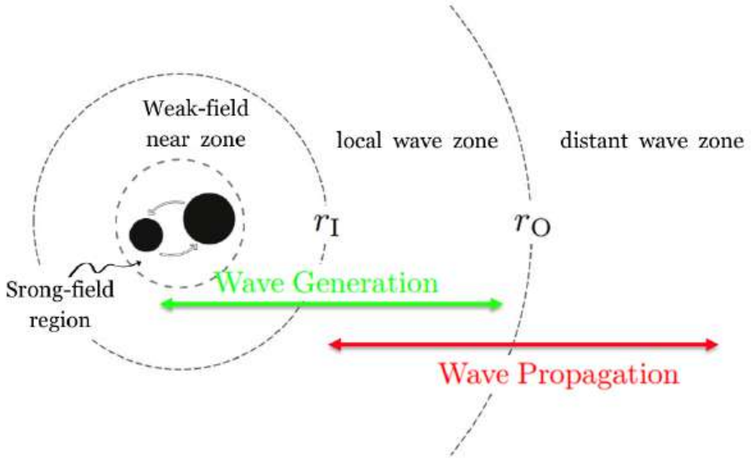

Figure 2.2: Spatial regions associated with a compact binary source of GWs. The strong-field region, weak-field near zone, and the local and distant wave zones are indicated, separated by the characteristic scales r_I and r_O. GW generation predominantly occurs in the near zone (green arrow), while wave propagation is described in the distant wave zone (red arrow). Original illustration based on the sources: [61, 62].

We therefore consider the linearized Einstein equations in the presence of matter, assuming weak gravitational fields and slowly moving sources, so that the background metric can be approximated as Minkowski and internal velocities are small compared to the speed of light. Under these conditions the problem becomes analytically tractable and closely analogous to electromagnetic radiation from accelerated charges [35].

In contrast with electromagnetism, conservation laws strongly constrain gravitational radiation. Mass conservation excludes monopole radiation, while conservation of linear and angular momentum eliminates dipole contributions. Consequently, the leading radiative term arises from the time variation of the mass quadrupole moment of the source.

2.1.5.1 Einstein's quadrupole formula

The starting point is the linearized Einstein equation (2.26) in the presence of matter, written in harmonic gauge,

$$\Box \bar{h}_{\mu\nu} = -\frac{16\pi G}{c^4} T_{\mu\nu}, \tag{2.53}$$

where $\bar{h}_{\mu\nu} = h_{\mu\nu} - \frac{1}{2}\eta_{\mu\nu} h$ denotes the trace-reversed metric perturbation and $T_{\mu\nu}$ is the energy-momentum tensor of the source (see §1.3.1). The Lorenz gauge condition $\partial^\mu \bar{h}_{\mu\nu} = 0$ is consistent provided that energy-momentum is conserved, $\partial^\mu T_{\mu\nu} = 0$.

As in electromagnetism, this set of decoupled wave equations can be solved using Green's function techniques [57]. For an isolated source localized within a finite spatial region Σ, the retarded solution is

$$\bar{h}_{\mu\nu}(t, \mathbf{x}) = \frac{4G}{c^4} \int_{\Sigma} \frac{T_{\mu\nu}(t_{\text{ret}}, \mathbf{x}')}{|\mathbf{x} - \mathbf{x}'|} \, d^3x', \tag{2.54}$$

where the retarded time is

$$t_{\text{ret}} = t - \frac{|\mathbf{x} - \mathbf{x}'|}{c}. \tag{2.55}$$

This expression explicitly enforces causality: the metric perturbation at $(t, \mathbf{x})$ depends on the configuration of the source at the earlier time t_{ret} (see Fig. 2.3) [36].

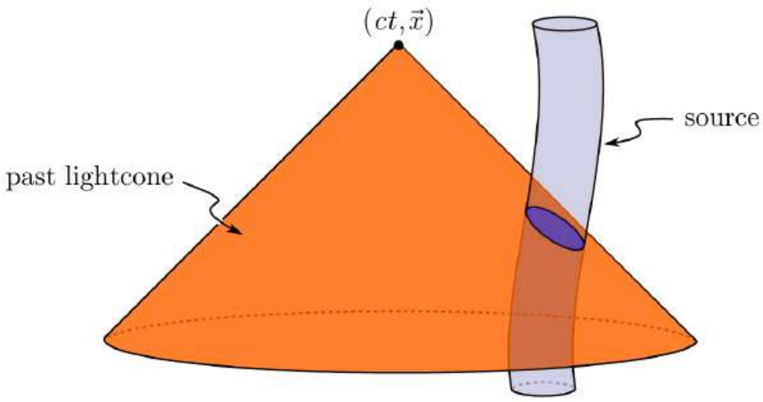

Figure 2.3: Representation of the past light cone of the observation event $(ct, \mathbf{x})$. GWs propagate at the finite speed c, so the metric perturbation measured at the detector is determined by the configuration of the source at the retarded time $t' = t - |\mathbf{x} - \mathbf{y}|/c$. Only the portion of the source worldtube intersecting the past light cone contributes to the observed gravitational field, ensuring causal propagation. Source: [35].

We are interested in the radiation observed far from the source, in the wave zone, where $r = |\mathbf{x}|$ is much larger than both the source size R and the gravitational wavelength λ. In this regime one expands

$$|\mathbf{x} - \mathbf{x}'| \simeq r - \mathbf{n} \cdot \mathbf{x}', \qquad \mathbf{n} = \frac{\mathbf{x}}{r}, \tag{2.56}$$

and keeps only the leading terms in $1/r$ and in v/c.

Conservation laws severely restrict the possible radiative multipoles. By Birkhoff's theorem, any spherically symmetric mass distribution produces a static exterior solution, so a time-varying monopole term cannot radiate. Furthermore, the first time derivative of the mass dipole moment is proportional to the total linear momentum of the system, which is conserved for an isolated source. As a result, both monopole and dipole radiation are absent in GR.

The leading contribution to gravitational radiation therefore arises at quadrupole order. Keeping the lowest non-vanishing term in the multipole expansion and restricting to non-relativistic sources, one finds that the dominant spatial components of the metric perturbation are

$$\bar{h}_{ij}(t, \mathbf{x}) = \frac{2G}{c^4 r} \, \ddot{I}_{ij}(t - r/c), \tag{2.57}$$

where

$$I_{ij}(t) = \int_{\mathcal{S}} \rho(t, \mathbf{x}) \, x_i x_j \, d^3x, \tag{2.58}$$

is the (Newtonian) mass moment of inertia tensor of the source $\mathcal{S}$ and ρ is the (Newtonian) mass density.

At this stage the solution is still written in harmonic gauge and is not yet transverse nor traceless. The radiative degrees of freedom are obtained by extracting the TT part of the spatial metric perturbation. Outside the source, where $T_{\mu\nu} = 0$, this can be done by applying the TT projection operator,

$$h_{ij}^{\rm TT} = \Lambda_{ij,kl}\, \bar{h}_{kl}, \tag{2.59}$$

with

$$\Lambda_{ij,kl} = P_{ik}P_{jl} - \frac{1}{2}P_{ij}P_{kl}, \qquad P_{ij} = \delta_{ij} - n_i n_j. \tag{2.60}$$

It is convenient to introduce the traceless mass quadrupole moment,

$$\mathcal{I}_{ij}(t) = \int \rho(t, \mathbf{x}) \left(x_i x_j - \frac{1}{3}|\mathbf{x}|^2 \delta_{ij} \right) d^3x, \tag{2.61}$$

which differs from I_{ij} only by a trace term. Since the TT projection automatically removes traces, one may equivalently replace I_{ij} by $\mathcal{I}_{ij}$ in the radiative formula.

Combining the far-zone expansion with the TT projection, the metric perturbation takes the form

$$h_{ij}^{\rm TT}(t, \mathbf{x}) = \frac{2G}{c^4 r} \Lambda_{ij,kl}\, \ddot{\mathcal{I}}_{kl}(t - r/c). \tag{2.62}$$

This is Einstein's quadrupole formula: the gravitational radiation measured far from an isolated source is proportional to the second time derivative of its traceless mass quadrupole moment, evaluated at the retarded time [35]. In particular, for compact binary systems the symmetry of the quadrupole tensor implies that the dominant gravitational radiation oscillates at twice the orbital frequency.

The quadrupole moment formalism provides the leading-order description of GW emission from isolated systems. Although originally derived by Einstein for non-gravitationally driven sources, it was later shown to hold also for self-gravitating systems in the weak-field, Newtonian regime. In this limit, the quadrupole moment is computed from the Newtonian mass density, and the resulting metric perturbation in the TT gauge gives the dominant radiative contribution, with higher-order multipoles and relativistic corrections appearing at post-Newtonian order [63].

2.1.5.2 Radiated energy

A system that radiates GWs necessarily undergoes an energy loss. It is therefore important to quantify the amount of energy carried away by the radiation. In other words, we aim to determine the energy flux transported by GWs away from the source.

The formula for the energy flux of GWs can be derived by considering the effective stress-energy tensor of the gravitational field in the wave zone. In the TT gauge, the energy flux per unit area is given by

$$\frac{dE}{dA dt} = \frac{c^3}{32\pi G} \langle \dot{h}_{ij}^{\rm TT} \dot{h}_{ij}^{\rm TT} \rangle, \tag{2.63}$$

where the angle brackets denote an average over several wavelengths or periods of the wave. This expression shows that the energy flux is proportional to the square of the time derivative of the metric perturbation, reflecting the fact that GWs carry energy away from their sources. Integrating the energy flux over a sphere of radius r centered on the source, one obtains the total power radiated in GWs,

$$P = \frac{c^3 r^2}{32\pi G} \int \langle \dot{h}_{ij}^{\mathrm{TT}} \dot{h}_{ij}^{\mathrm{TT}} \rangle d\Omega, \tag{2.64}$$

where the integral is taken over the solid angle Ω. Substituting the quadrupole formula for h_{ij}^{TT}, one finds the well-known expression for the power radiated by a source with mass quadrupole moment $\mathcal{I}_{ij}$,

$$P = \frac{G}{5c^5} \langle \dddot{\mathcal{I}}_{ij} \dddot{\mathcal{I}}_{ij} \rangle, \tag{2.65}$$

where the triple dot denotes the third time derivative. This formula quantifies the energy loss due to GW emission in terms of the third time derivative of the mass quadrupole moment of the source. For compact binary systems, this energy loss leads to a gradual inspiral of the orbit, which can be observed as a characteristic chirp signal in the GW waveform [57].

2.2 Astrophysical sources

GW astronomy identifies four broad categories of astrophysical signals, each characterized by distinct time-scales, amplitudes, and expected sources. These categories are commonly referred to as transient (compact binary coalescences (CBCs) and bursts), continuous waves (CWs) and stochastic gravitational wave background (SGWB) (see Fig. 2.4) [64].

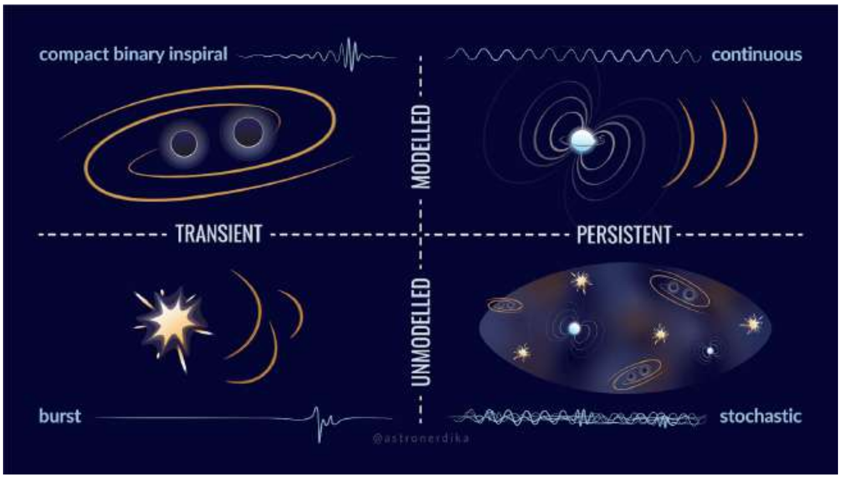

Figure 2.4: Classification of GW sources according to their temporal character (transient vs. persistent) and the availability of accurate waveform models (modeled vs. unmodeled), encompassing CBCs, burst events, CWs, and the stochastic background. Credit: Shanika Galaudage (@astronerdika).

2.2.1 Compact binary coalescences

CBCs are systems formed by two compact objects such as BHs, NSs, or a combination of both, which inspiral and merge due to the emission of GWs (see Fig. 2.5).

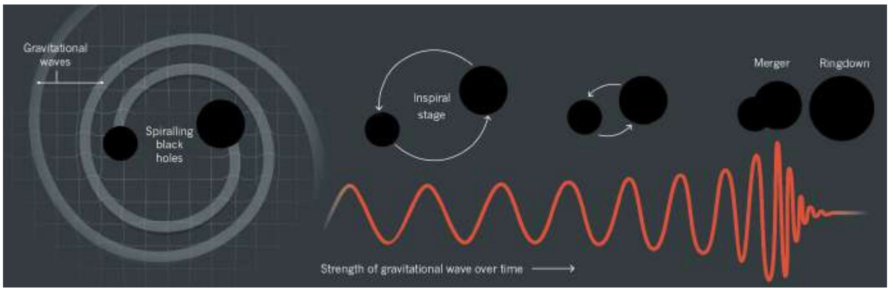

Figure 2.5: The three phases of a CBC: the inspiral, during which the two compact objects lose energy through GW emission and progressively approach each other; the merger, when the objects collide and coalesce; and the ringdown, in which the remnant settles into a stationary configuration by radiating the remaining perturbations. The corresponding gravitational waveform, showing the characteristic chirp of increasing frequency and amplitude, is displayed beneath each phase. Credit: Illustration by Lucy Reading-Ikkanda.

These systems are classified as BBHs, binary neutron stars (BNS), or NS and BH binaries (NSBH). The emitted signal is driven by the time varying quadrupole moment and exhibits a characteristic chirp pattern, with increasing frequency and amplitude as the orbit shrinks. The duration of the signal within the detector's sensitive frequency band (§2.3.3) depends on the total mass of the system: the rate at which the orbital frequency sweeps through this band scales with the chirp mass, so more massive systems evolve rapidly and merge at lower frequencies, leaving only fractions of a second of observable signal for BBH mergers, whereas the lower-mass BNS systems sweep slowly through the band and can remain visible for tens of seconds to minutes before merging [57].

To date, all GW detections made by ground based interferometers, including the LIGO Scientific Collaboration and the Virgo Collaboration, are consistent with CBC sources. During O1 and O2 [65], 11 events were identified, including the first BNS detection, GW170817 [66]. The number of detections increased significantly during O3 [67, 68] to around 90 events. During O4a [69], 128 new CBC candidates were identified, bringing the total to 218 when combined with previous catalogs. The most recent release, GWTC-5.0, extends the catalog through the end of O4b and adds 104 further observations, bringing the cumulative total to over 300 significant candidates ($p_{\mathrm{astro}} \geq 0.5$) [70], with BBH mergers continuing to dominate the observed population.

The success in detecting CBCs is largely due to accurate waveform modelling, which enables matched filtering techniques to identify signals in noisy data [71]. Future space based detectors such as the Laser Interferometer Space Antenna (LISA) (see §2.3.5.2) will extend observations to other sources, including white dwarf binaries and extreme mass ratio inspirals (EMRIs), providing further insight into compact object populations and their formation [72].

2.2.2 Burst events

GW bursts are short duration signals produced by highly energetic and often poorly modeled astrophysical events. Typical sources include core collapse supernovae [73], gravitational collapses, highly eccentric compact binary mergers, and more speculative scenarios such as cosmic strings [74]. In addition, burst like emission can be associated with transient phenomena such as gamma ray bursts linked to NS or NS-BH mergers.

These signals typically last from a fraction of a second to a few tens of seconds and are characterized by their unpredictable morphology. Unlike CBCs or CWs, burst signals do not have accurate waveform models, which makes their detection more challenging. Instead of matched filtering techniques, searches rely on identifying excess power or coherent transients across multiple detectors [75].

2.2.3 Continuous waves

CWs are expected to be emitted by rapidly rotating compact objects, particularly NSs with small deviations from perfect axial symmetry. These asymmetries produce a time varying quadrupole moment, enabling the continuous emission of gravitational radiation. The resulting signal is nearly monochromatic, typically at a frequency related to the star's rotation, and evolves very slowly due to spindown effects [76].

In contrast to CBCs, CW signals are extremely long lived but much weaker in amplitude, as they originate from small deformations in the source. As a result, their detection requires long observation times, often spanning months or years, in order to accumulate sufficient signal-to-noise ratio (SNR). Despite their low amplitude, once detected, these signals can be tracked over extended periods, allowing for precise studies of the source properties.

The strength of the emission depends on factors such as the rotation rate and internal structure of the NS, including its equation of state and crust properties [76]. In addition to isolated NSs, other potential sources of CWs have been proposed, such as boson clouds [77] around rapidly spinning BHs or compact binaries in very early inspiral stages. Although no CW signal has been confirmed so far, current searches have placed increasingly stringent upper limits, improving the prospects for future detections [78].

2.2.4 Stochastic backgrounds

The SGWB is a diffuse signal produced by the superposition of many unresolved sources, which are either too weak or too numerous to be individually detected. This background can have both astrophysical and cosmological origins [79]. Astrophysical contributions arise from the cumulative effect of distant CBCs or populations of compact binaries such as WD systems, while cosmological contributions may be generated in the early Universe through processes such as inflation or phase transitions.

Although no definitive detection has yet been achieved by ground based detectors, current analyses have placed increasingly stringent upper limits on the energy density of the background [80]. Detecting the SGWB would provide unique information about the population of unresolved sources and could offer a direct probe of physical processes in the early Universe, complementing observations of the cosmic microwave background.

2.3 Detectors

The first experimental attempts to detect GWs began in the 1960s with resonant-mass detectors, which aimed to measure the mechanical response of massive bars to passing waves [81]. However, their limited sensitivity and narrow frequency range prevented successful detections. This motivated the development of interferometric detectors, capable of measuring extremely small changes in distance between test masses. After decades of technological progress, this approach led to the first direct detection of GWs in 2015 [5].

Today, GW observations are routinely performed by a global network of ground-based interferometers, marking the beginning of GW astronomy.

2.3.1 Interferometric detection principle

Current GW detectors are based on laser interferometry, using a modified Michelson configuration, as illustrated in Fig. 2.6. A coherent laser beam is split into two perpendicular arms by a beam splitter. The light travels along each arm, is reflected by mirrors, and recombines at the beam splitter, producing an interference pattern that is measured by a photodetector [82].

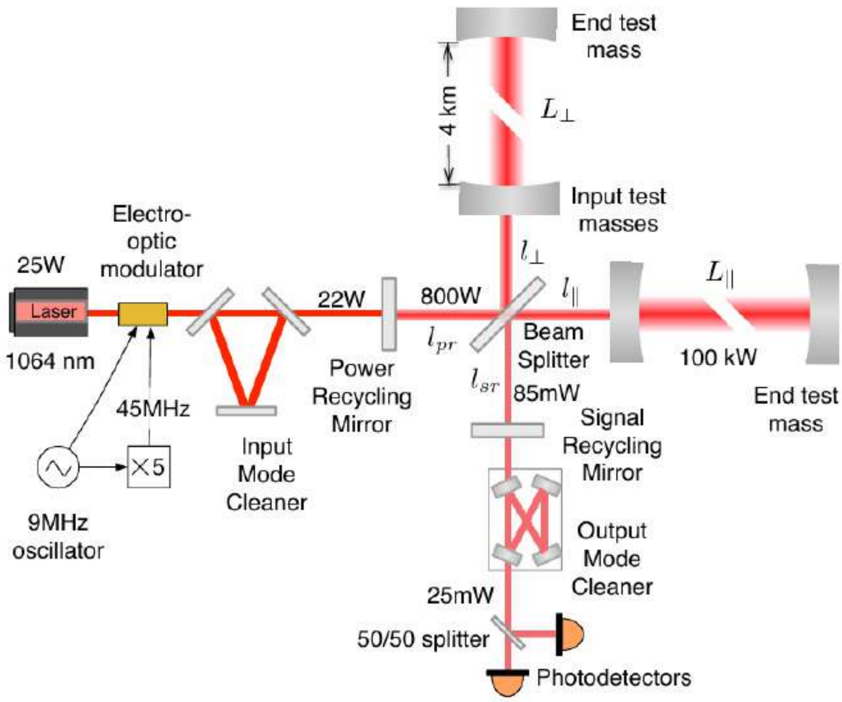

Figure 2.6: Optical layout of the advanced LIGO interferometer, illustrating the Fabry-Pérot arm cavities, the power recycling mirror (PRM) at the input port, and the signal recycling mirror (SRM) at the output port. Source: [82].

In the absence of a GW, the lengths of the two arms are adjusted so that the interference at the output is destructive, resulting in minimal detected light. When a GW passes through the detector, it induces a differential stretching and compression of the arms, analogous to the tidal deformation described in §2.1.4 by the geodesic deviation equation (2.38). This changes the optical path lengths and produces a phase difference between the two beams, leading to a measurable signal at the photodetector.

The quantity measured is the strain $h(t)$, defined as the relative change in the arm lengths, $h \sim \Delta L/L$ (see §2.3.2.1). This is precisely the quantity appearing in the TT metric perturbation (2.37): the two polarization amplitudes A_+ and $A_\times$ directly set the scale of the fractional length change experienced by the detector arms. For typical detectors with arm lengths of a few kilometers, the expected variations are extremely small, of order 10^{-21}, requiring highly sensitive instrumentation.

Modern interferometers incorporate several enhancements to increase their sensitivity. These include Fabry-Pérot cavities in the arms, which effectively increase the optical path length, and power and signal recycling mirrors, which enhance the circulating power and the detector response. Despite these improvements, the measurement remains extremely challenging due to the tiny magnitude of the effect, set ultimately by the quadrupole formula discussed in §2.1.5.1 and the astrophysical distances involved.

2.3.2 Detector response

The response of an interferometric detector depends on the relative orientation between the incoming GW and the detector arms. This is described by antenna pattern functions, which determine the sensitivity to different sky locations and polarization states. As a result, a single detector has limited directional sensitivity, making a network of detectors essential for accurate sky localization and PE.

2.3.2.1 Antenna pattern of a ground-based interferometric detector

Here we derive the response of a ground-based Michelson interferometer to a plane GW in the long-wavelength approximation, where the detector arm length is much smaller than the reduced wavelength of the wave.

A complete description of the detection problem requires three distinct Cartesian frames [83]. The wave frame (x', y', z') has its z'-axis aligned with the propagation direction; in this frame the TT metric perturbation takes the standard form

$$H = \begin{pmatrix} h_+ & h_\times & 0 \\ h_\times & -h_+ & 0 \\ 0 & 0 & 0 \end{pmatrix}, \tag{2.66}$$

where h_+ and $h_\times$ denote the plus and cross polarizations. The detector frame $(\bar{x}, \bar{y}, \bar{z})$ has its $\bar{z}$-axis pointing toward the local zenith and the two arms lying in the $(\bar{x}, \bar{y})$ plane along unit vectors $\hat{\mathbf{n}}_1$ and $\hat{\mathbf{n}}_2$, separated by an opening angle ζ. Finally, an ecliptic frame (x, y, z) centered at the Solar System Barycenter serves as a common inertial reference for source localization. The relative orientation of the wave and detector frames is parametrized by the polar angles $(\bar{\theta}_N, \bar{\phi}_N)$ specifying the source direction in the detector frame, and the polarization angle ψ (see Fig. 2.7).

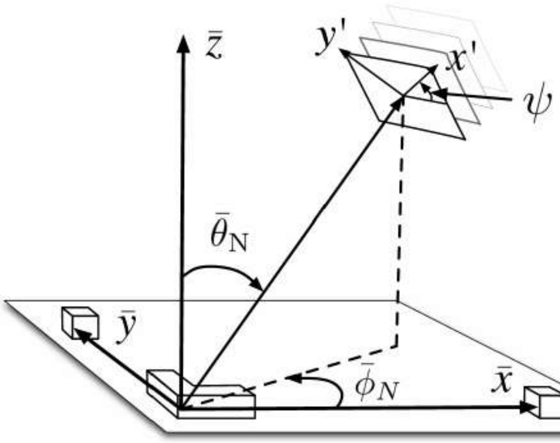

Figure 2.7: Geometric relationship between the detector frame $(\bar{x}, \bar{y}, \bar{z})$ and the wave frame (x', y', z'). The interferometer arms lie in the $(\bar{x}, \bar{y})$ plane, with $\bar{z}$ pointing toward the local zenith. The angles $\bar{\theta}_N$ and $\bar{\phi}_N$ give the polar and azimuthal coordinates of the incoming wave direction in the detector frame, and ψ denotes the rotation of the polarization axes about the propagation direction. Original illustration based on the source: [84].

The observable quantity is the differential arm-length variation induced by the passing wave. Working to first order in h and applying the geodesic deviation equation (2.45) in the detector frame, the fractional length change of a single arm oriented along $\hat{\mathbf{n}}$ is

$$\frac{\delta L(t)}{L} = \frac{1}{2}\,\hat{\mathbf{n}}^T \tilde{H}(t)\,\hat{\mathbf{n}}, \tag{2.67}$$

where $\tilde{H}(t) = MH(t)M^T$ is the wave-frame strain tensor (2.66) rotated into the detector frame by the orthogonal matrix M. The scalar strain measured by the interferometer is then

$$h(t) = \frac{\Delta L_{\text{arms}}(t)}{L} = \frac{L_1(t) - L_2(t)}{L} = \frac{\delta L_1(t)}{L} - \frac{\delta L_2(t)}{L} = \frac{1}{2}\,\hat{\mathbf{n}}_1^T \cdot \tilde{H}(t)\,\hat{\mathbf{n}}_1 - \frac{1}{2}\,\hat{\mathbf{n}}_2^T \cdot \tilde{H}(t)\,\hat{\mathbf{n}}_2, \tag{2.68}$$

where the sign convention follows from choosing $\hat{\mathbf{n}}_1 \times \hat{\mathbf{n}}_2$ to point outward from the Earth's surface [85].

The frame transformation is given by the sequence of Euler rotations

$$M = R_z(-\bar{\phi}_N)\,R_y(-\bar{\theta}_N + \pi)\,R_z(\psi), \tag{2.69}$$

which encodes both the sky location of the source and the orientation of the polarization axes relative to the detector.

Substituting the explicit form of M and the arm vectors into (2.68), and simplifying with standard trigonometric identities, the detector response takes the form

$$h(t) = F_+(\psi, \zeta, \bar{\theta}_N, \bar{\phi}_N)\,h_+(t) + F_\times(\psi, \zeta, \bar{\theta}_N, \bar{\phi}_N)\,h_\times(t), \tag{2.70}$$

where the antenna pattern functions are [83, 85]

$$F_+ = \sin\zeta\left[\frac{1}{2}(1+\cos^2\bar\theta_N)\cos 2\bar\phi_N\cos 2\psi - \cos\bar\theta_N \sin 2\bar\phi_N \sin 2\psi\right], \tag{2.71}$$

$$F_\times = -\sin\zeta\left[\frac{1}{2}(1+\cos^2\bar\theta_N)\cos 2\bar\phi_N\sin 2\psi + \cos\bar\theta_N \sin 2\bar\phi_N \cos 2\psi\right]. \tag{2.72}$$

These expressions satisfy the symmetry relations

$$F_+(\psi + \tfrac{\pi}{4}) = F_\times(\psi), \qquad F_{+,\times}(\psi+\tfrac{\pi}{2}) = -F_{+,\times}(\psi), \tag{2.73}$$

reflecting the spin-2 character of GWs: a $\pi/4$ rotation of the polarization angle exchanges the two modes, while a $\pi/2$ rotation reverses their sign.

A polarization-independent measure of the directional sensitivity is provided by the quadrature sum

$$|F| = \sqrt{F_+^2 + F_\times^2}, \tag{2.74}$$

shown in Fig. 2.8 for $\zeta = \pi/2$. The characteristic quadrupolar structure reflects the tensorial nature of the GW strain and the geometric projection onto the detector arms. The individual patterns F_+ and $F_\times$ are displayed in Fig. 2.9 for several values of ψ; each successive column corresponds to a $\pi/8$ increment in polarization angle, illustrating the interchange symmetry above.

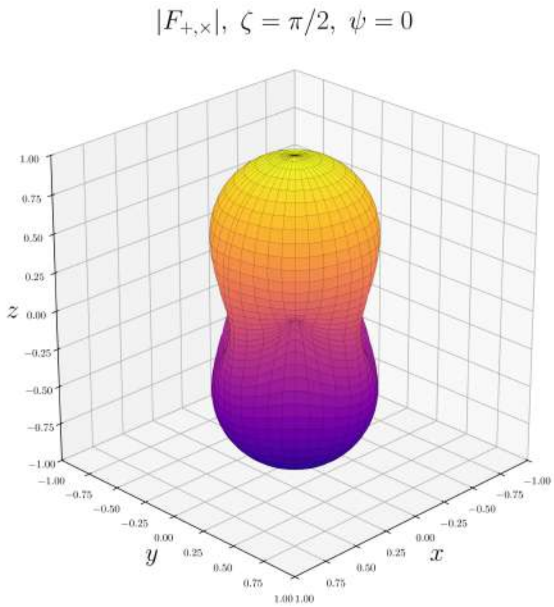

Figure 2.8: Polarization-averaged antenna pattern $|F|$ of a ground-based interferometer with orthogonal arms ($\zeta = \pi/2$), as a function of sky angles $(\bar\theta_N, \bar\phi_N)$. Original result.

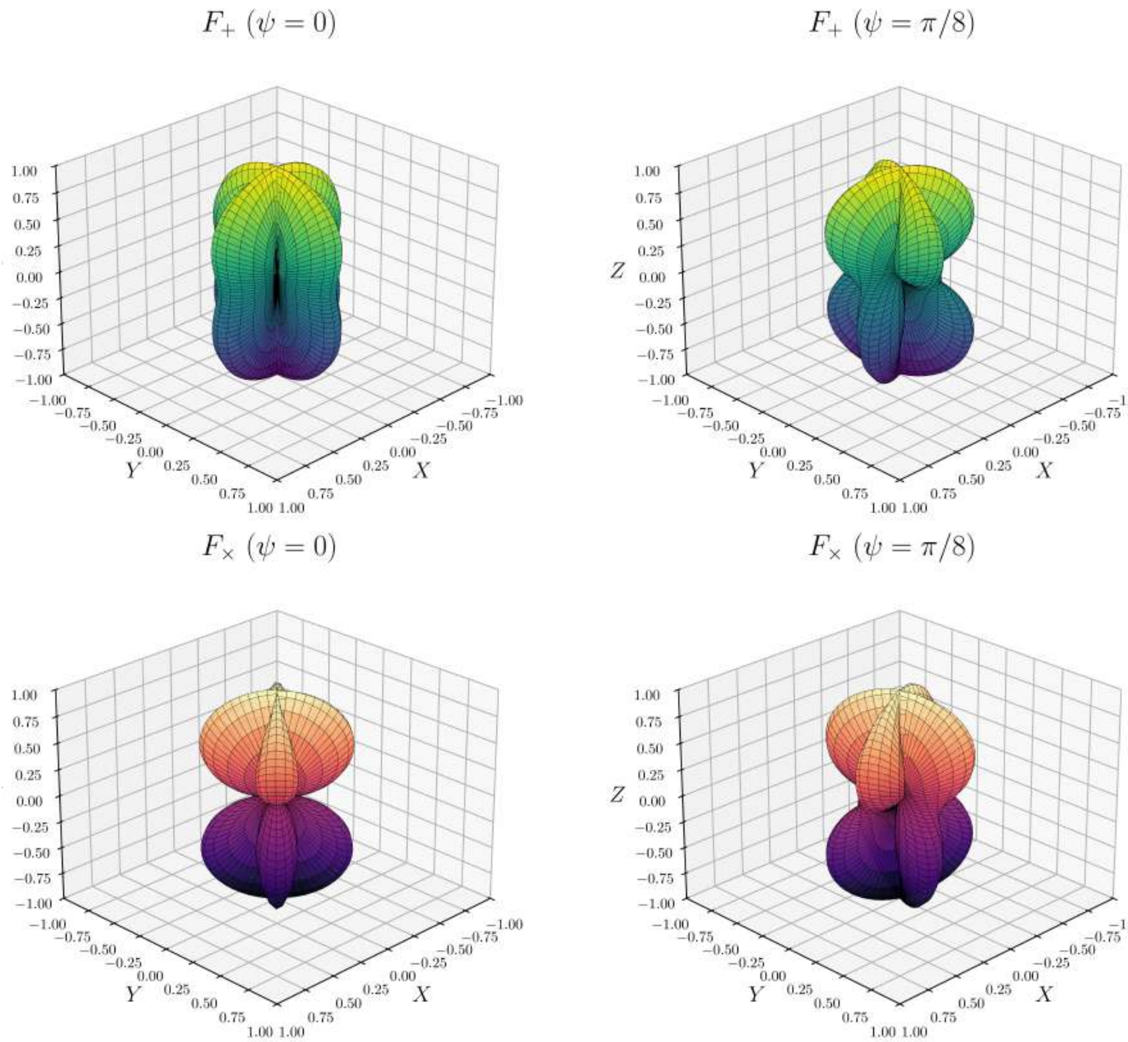

Figure 2.9: Antenna patterns for the plus (F_+, top) and cross ($F_\times$, bottom) polarizations for $\zeta = \pi/2$ and representative values of ψ. Original result.

2.3.3 Noise and sensitivity

A GW with strain $h \sim 10^{-21}$ produces a differential arm-length change $\delta L = hL/2 \sim 10^{-18}\,\mathrm{m}$ for $L = 4\,\mathrm{km}$, three orders of magnitude smaller than the diameter of a proton. Achieving sensitivity at this level requires identifying and suppressing a broad set of noise sources that would otherwise overwhelm the signal. These are conveniently characterized by the one-sided power spectral density (PSD) $S_n(f)$ [86], defined as the Fourier transform of the autocorrelation function[5] of the detector output:

$$S_n(f) = 2\int_{-\infty}^{\infty} C(\tau)\, e^{-i2\pi f\tau}\, d\tau, \qquad f \geq 0, \tag{2.75}$$

[5]The autocorrelation function measures how strongly a signal is correlated with a delayed version of itself as a function of the time lag τ. In the context of detector noise, it quantifies the degree to which noise fluctuations at different times remain statistically related [86].

where $C(\tau) = \langle n(t)n(t+\tau)\rangle$ is the autocorrelation function of the noise process $n(t)$, and $\langle\cdot\rangle$ denotes an ensemble average. The factor of 2 converts the two-sided power spectrum given by the Wiener-Khinchin theorem into the one-sided convention used throughout this thesis, for which all the noise power at a given frequency is folded into $f \geq 0$. Its square root $\sqrt{S_n(f)}$, expressed in units of $\mathrm{Hz}^{-1/2}$, defines the sensitivity curve of the instrument and sets the threshold below which signals are inaccessible.

Noise sources are broadly divided into two categories. Fundamental noise arises directly from the measurement principle and cannot be removed without redesigning the detector. Technical noise is not intrinsic to the design but still limits performance. The main contributions are the following [87, 88]:

- **Shot noise.** The Poissonian arrival statistics of photons at the output photodetector produce a phase uncertainty scaling as $1/\sqrt{N_\gamma}$, where N_γ is the number of detected photons per unit time. This contribution decreases with increasing circulating power and dominates the sensitivity above $\sim$500 Hz. It can be partially mitigated by injecting squeezed vacuum at the output port, a technique already incorporated in the Advanced LIGO and Advanced Virgo detectors.

- **Radiation-pressure noise.** Amplitude fluctuations of the intracavity field exert a stochastic force on the mirrors, inducing position noise that grows with laser power and dominates at low frequencies. Shot noise and radiation-pressure noise are related by the Heisenberg uncertainty principle and together define the standard quantum limit for a free test mass. At a fixed squeezing angle, squeezed-light injection only redistributes the quantum uncertainty between quadratures, improving one at the expense of the other. Reducing both below the standard quantum limit requires frequency-dependent squeezing, where the squeezed vacuum is reflected off a detuned filter cavity so the squeezing angle rotates with frequency; this was implemented in Advanced LIGO for O4 [89].

- **Thermal noise.** Any component at finite temperature undergoes stochastic mechanical motion. Via the fluctuation-dissipation theorem[6], the fluctuation amplitude is proportional to the material's mechanical loss angle. It affects the test masses, their coatings, and the suspension fibers, dominating the sensitivity around 40-200 Hz. Low-loss fused-silica substrates and fibers minimize this contribution; cryogenic operation, as adopted by KAGRA, offers an alternative.

- **Seismic noise.** Ground motion couples mechanically to the test masses through the suspension chain. The quadruple-pendulum isolation system attenuates seismic vibrations by more than ten orders of magnitude above $\sim$10 Hz, but residual motion from microseisms, wind, and anthropogenic activity still limits the sensitivity below $\sim$40 Hz.

- **Newtonian noise.** Density fluctuations in the surrounding ground and atmosphere produce direct gravitational forces on the test masses that cannot be shielded mechanically. This contribution, also known as gravity-gradient noise, sets a fundamental floor at low frequencies for surface-based detectors.

[6] The fluctuation-dissipation theorem states that any dissipative mechanism in a system in thermal equilibrium is necessarily accompanied by a fluctuating force of the same origin. In interferometric detectors, mechanical losses in the test masses, coatings, and suspension fibers therefore set a fundamental limit on their positional stability [90].

- **Control noise.** Servo loops are required to maintain the arm cavities on resonance and keep the optical components aligned. The sensors and actuators in these feedback systems introduce noise that couples into the GW channel, representing the dominant technical limitation at low frequencies.

- **Residual gas noise.** Despite the ultra-high vacuum environment ($< 1\,\mu$Pa in the beam tubes), residual molecules scatter laser photons and cause fluctuations in the optical path length. This contribution is generally subdominant but must be controlled to avoid limiting the sensitivity in specific frequency bands.

- **Other technical noise sources.** Laser frequency and intensity fluctuations, electromagnetic coupling, scattered light, and photodetector dark noise contribute across various frequency ranges. These are monitored and subtracted using a network of $\gtrsim 2 \times 10^5$ auxiliary channels (including magnetometers, accelerometers, microphones and photodetectors) that record the state of the instrument and its environment continuously.

The noise process is approximately Gaussian but not stationary. Two classes of artefact further degrade the data quality. Lines are persistent narrow-band features visible as peaks above the noise floor in the PSD, arising from power-line harmonics (50 or 60 Hz and their multiples), mirror suspension violin modes, calibration injections, and environmental electromagnetic coupling (see Fig. 2.10) [91].

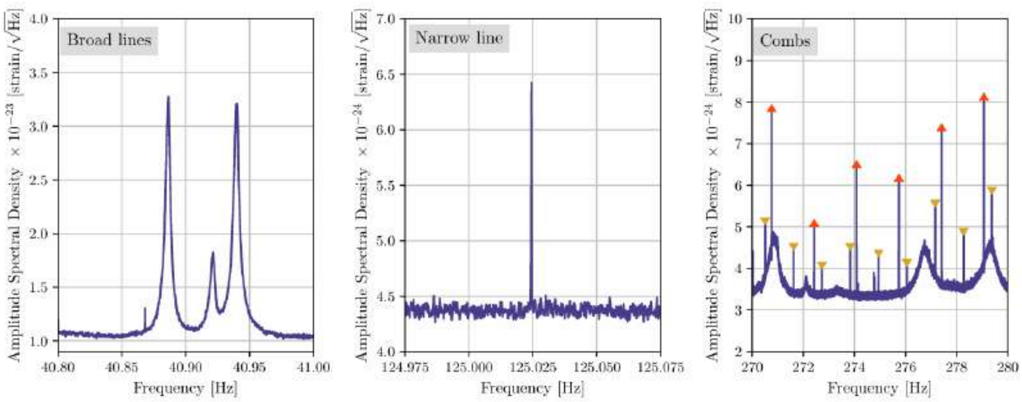

Figure 2.10: High-resolution amplitude spectral density ($\Delta f = 1/1800\,$Hz) of LIGO Hanford strain data illustrating three representative line artefacts. *Left:* a broad spectral feature associated with mirror suspension roll-mode resonances. *Center:* an isolated narrow line of unknown origin. *Right:* two frequency combs with distinct spacings, marked by upright and inverted triangles, alongside broader noise structures. Source: [92].

Glitches are short-duration transient bursts of excess power whose morphology can mimic astrophysical signals, reducing the efficiency of detection pipelines (see Fig. 2.11) [93].

The interplay of all these contributions shapes the detector sensitivity curve shown in Fig. 2.12, where three distinct regimes are clearly visible: control and seismic noise below $\sim 50\,$Hz, a combination of thermal and quantum noise in the intermediate band (50-500 Hz), and shot noise above $\sim 500\,$Hz. Together they define the accessible frequency range of $\sim$10-5000 Hz for current ground-based detectors. The sharp spectral peaks superimposed on the broadband noise floor correspond to the lines discussed above, most prominently the 60 Hz power-line harmonic and the mirror suspension violin modes.

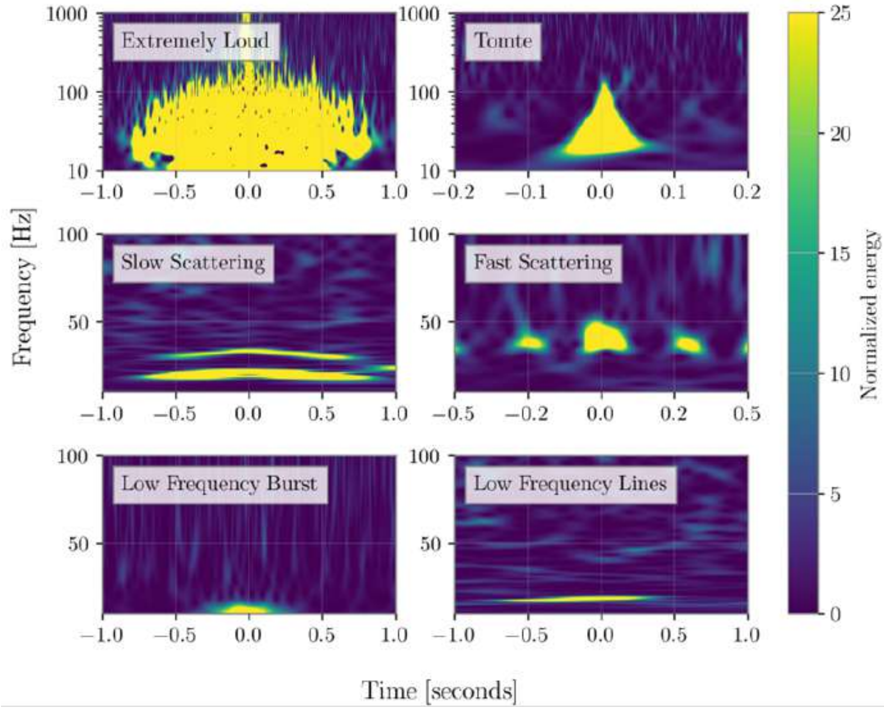

Figure 2.11: Time-frequency spectrograms of representative glitch classes recorded at LHO and LLO. Slow and Fast Scattering artefacts are attributed to light scattered by optics under increased ground motion, while the origin of Extremely Loud, Tomte, Low Frequency Burst, and Low Frequency Lines glitches remains unidentified. Source: [92].

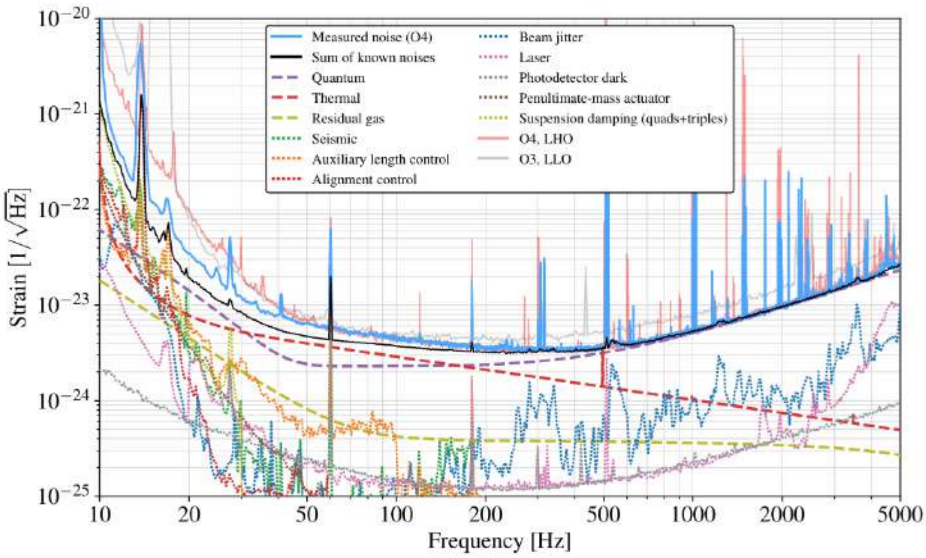

Figure 2.12: Strain noise amplitude spectral density of the LLO detector during O4, with the O3 LLO and O4 LHO curves for comparison. Dashed lines are noise contributions from analytical models or independent component measurements; dotted lines are noise projected from auxiliary witness channels. The black curve is the quadrature sum of all known sources. Three regimes stand out: control noise below $\sim 50\,\mathrm{Hz}$, quantum and thermal noise in the intermediate band, and shot noise above $\sim 500\,\mathrm{Hz}$. Narrow lines from power-line harmonics, violin modes, and calibration injections appear throughout. Source: [88].

2.3.4 Current detector network

A single interferometer provides limited directional information: the antenna pattern functions derived in §2.3.2.1 show that, at any given time, a ground-based detector is nearly blind along certain sky directions, an effect that continuously shifts as the Earth's rotation changes the detector's orientation relative to the sky. A global network of detectors is therefore essential, both to confirm detections through coincident observations and to localize sources on the sky via time delay triangulation. Sky localization precision improves with the number of detectors, their mutual separation, and their individual sensitivities. The current network consists of four facilities operating in the frequency band $\sim$10-5000 Hz (see Fig. 2.13):

- **Advanced LIGO** [87] comprises two 4-km detectors: LIGO Hanford (H1) in Washington State and LIGO Livingston (L1) in Louisiana, separated by $\sim$3000 km. Operated by the LIGO Scientific Collaboration (LSC), they have been the primary instruments for GW detections since the first observation of a BBH merger in September 2015 [5].

- **Advanced Virgo (V1)** [94] is a 3-km interferometer located in Cascina, Italy, operated by the Virgo Collaboration. It joined the network in the final month of O2, contributing to the first three-detector observation of a BBH merger (GW170814 [95]) and to the precise sky localization of the BNS merger GW170817 [66]. Although its sensitivity is somewhat below that of the LIGO instruments, its geographic separation from the US sites significantly improves source localization.

- **KAGRA (K1)** [96] is a 3-km cryogenic underground detector situated 200 m beneath the Kamioka mountain in Gifu Prefecture, Japan. Its subterranean location provides natural seismic and atmospheric shielding, and its test masses are cooled to cryogenic temperatures ($\sim$20 K) to suppress thermal noise, making it the first second-generation detector to adopt this approach. KAGRA participated in a short joint observing run with GEO600 at the end of O3, and has been part of the network since O4.

- **GEO600** [97] is a 600-m detector near Hannover, Germany, operated jointly by British and German institutions. Its shorter arms make it significantly less sensitive than the kilometre-scale instruments, so it does not contribute to standard astrophysical searches. Instead, it operates in astrowatch mode during maintenance periods of the other detectors and serves as a testbed for advanced instrumentation techniques (including squeezed-light injection) before their deployment in larger facilities. GEO600 is scheduled to conclude operations in 2026, marking the end of its role as the longest-running interferometric GW detector in Europe.

The evolution of the network across observing runs is illustrated in Fig. 2.14. During O1 and the majority of O2, only H1 and L1 were operational; Virgo joined for the final weeks of O2 and participated throughout O3. A brief joint run (O3GK) at the close of O3 included GEO600 and KAGRA. Since O4, the network comprises all four facilities, with the BNS inspiral range (the standard figure of merit for detector sensitivity) improving progressively across runs as a result of the upgrades discussed in §2.3.3.

Figure 2.13: Aerial and underground views of the current LVK detector network. Top row, left to right: LIGO Hanford (Washington, USA), LIGO Livingston (Louisiana, USA), and Advanced Virgo (Cascina, Italy). Bottom row, left to right: KAGRA underground tunnel (Kamioka, Japan) and GEO600 (Hannover, Germany). Sources: [98, 99, 100, 101, 102].

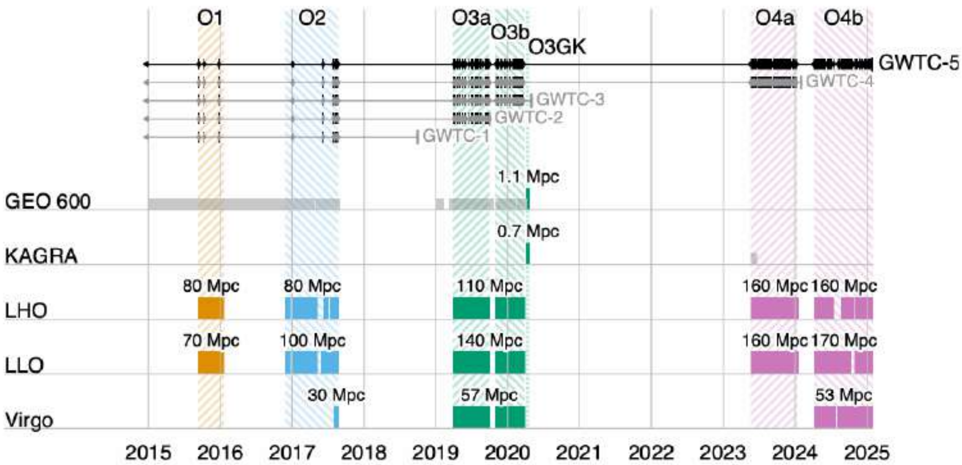

Figure 2.14: Timeline of LVK observing runs from the first detection in 2015 through the start of O4c in January 2025. Colored bars indicate the periods during which each detector was operating, together with the typical BNS inspiral range achieved in each run. GEO600 astrowatch and KAGRA O4a periods (light gray) were not included in the main GW analyses. In O1 and O4a only LHO and LLO participated; Virgo joined for the final month of O2 and operated throughout O3a, O3b and O4b. A short joint O3GK run also included GEO600 and KAGRA. Tick marks along the top indicate GW candidates from GWTC-1 through GWTC-5 with astrophysical probability $\geq 50\%$. Source: [70].

2.3.5 Future detectors

The current LVK network has demonstrated the transformative potential of GW astronomy, but its sensitivity and frequency coverage remain limited by fundamental and technical noise sources discussed in §2.3.3. The next generation of detectors, both ground-based and space-based, aims to extend the accessible parameter space by one to two orders of magnitude in strain sensitivity and to cover a much broader frequency band.

2.3.5.1 Ground-based third-generation detectors.

- **Einstein Telescope** (ET) [103] is the European third-generation concept. Its final geometry has not yet been decided and remains linked to the ongoing site-selection process, expected to conclude in 2027. The originally proposed reference design is an equilateral triangle with 10-km arms hosted underground at a single site to suppress seismic and Newtonian noise, where each side hosts a xylophone configuration: two nested interferometers, one optimised for high frequencies at room temperature and one for low frequencies operating cryogenically, jointly covering the band $\sim$2-10^4 Hz, providing full sky coverage and simultaneous measurement of both GW polarizations from a single facility. The alternative design, denoted 2L, instead places two separate 15-km L-shaped detectors at widely separated European sites; a detailed comparison of the scientific performance of both geometries is presented in [104].
- **Cosmic Explorer** (CE) [105] is the US counterpart, conceived as an L-shaped interferometer with 40-km arms (roughly ten times the length of Advanced LIGO) retaining the proven Michelson topology while achieving a strain sensitivity improvement of approximately one order of magnitude. Together, ET and CE would form a third-generation network capable of detecting CBC mergers throughout the majority of the observable Universe.
- **LIGO India** [106] is a planned Advanced LIGO-class detector to be built in Maharashtra, India. Although it will not exceed the sensitivity of the current LIGO instruments, its geographic location will substantially improve sky localization and detection confidence for the global network.

2.3.5.2 Space-based detectors.

Ground-based detectors are fundamentally limited below $\sim$10 Hz by seismic and Newtonian noise. Space-based missions overcome this barrier by operating in a quiet gravitational environment, opening the mHz frequency window inaccessible from the ground.

- **LISA** [107] is the flagship ESA mission, scheduled for launch in the mid-2030s. It consists of three spacecraft in heliocentric orbit forming a near-equilateral triangle with 2.5×10^6 km arms. Because the arm lengths are unequal and variable, standard Michelson cancellation of laser frequency noise is not possible; instead, Time Delay Interferometry (TDI) is employed to suppress this noise in post-processing. LISA will cover the band $\sim 10^{-4}$–10^{-1} Hz, giving access to sources including massive BH mergers, EMRIs, and the early inspiral of compact binaries centuries before they enter the ground-based band. The feasibility of the key technologies was demonstrated by the LISA Pathfinder mission in 2015-2017.

Further space-based concepts under study include the Japanese DECIGO [108], targeting the decihertz band ($\sim$ 0.1-10 Hz) that bridges LISA and ground-based detectors, and the Chinese missions TianQin and Taiji [109], which share a similar frequency range to LISA with different orbital configurations.

The combination of third-generation ground-based detectors and space-based observatories is expected to provide continuous frequency coverage from $\sim 10^{-4}$ Hz to several kHz by the late 2030s, enabling a new era of multi-band GW astronomy.

2.4 Gravitational wave data analysis

The preceding sections established the theoretical framework for GW generation and propagation, and the experimental techniques employed to detect them. The remaining challenge is to extract astrophysical information from the noisy output of the detectors, a task that combines signal processing, physical modelling, and statistical inference. This relies on three crucial ingredients: matched filtering, which searches for candidate events buried in the noise by correlating the data against a bank of templates (§2.4.1); accurate waveform models, which provide these templates and are required both for the search and for the PE that follows (§2.4.2); and Bayesian PE, which builds directly on the matched-filtering framework introduced above to infer the full posterior distribution over the source parameters once a candidate has been identified (§2.4.3).

2.4.1 Matched filtering and signal-to-noise ratio

2.4.1.1 Detector output and noise model

The output of a GW detector can be written as the sum of a GW signal and instrumental noise,

$$s(t) = h(t; \boldsymbol{\theta}) + n(t), \tag{2.76}$$

where $h(t; \boldsymbol{\theta})$ is the GW strain parametrized by the source parameters $\boldsymbol{\theta}$, in practice provided by a theoretical waveform model (§2.4.2), and $n(t)$ is the noise contribution. With the sensitivity of current detectors, the noise amplitude is typically several orders of magnitude larger than the signal, $|h(t; \boldsymbol{\theta})| \ll |n(t)|$, so in most cases a detection cannot be claimed by direct inspection of $s(t)$ [57]. A systematic method is therefore required to reliably extract the vast majority of buried signals.

Although the true noise is non-stationary and non-Gaussian (exhibiting glitches and spectral lines as discussed in §2.3.3) a tractable analytic framework is obtained by modelling it as a zero-mean, stationary Gaussian random process fully characterized by its one-sided power spectral density (PSD) $S_n(f)$ (2.75),

$$\langle \tilde{n}^*(f)\, \tilde{n}(f') \rangle = \frac{1}{2} S_n(f)\, \delta(f - f'), \tag{2.77}$$

where tildes denote Fourier transforms, $\delta(f - f')$ is the Dirac delta encoding the statistical independence of different frequency components, and $\langle \cdot \rangle$ denotes an ensemble average, which for a single detector is identified with a time average via the ergodic theorem[7]. Under this Gaussian

[7]The ergodic theorem states that, for a stationary random process, time averages computed from a single sufficiently long realization are equivalent to ensemble averages over all possible realizations. In practice, this means that the PSD $S_n(f)$ can be estimated directly from a single stretch of detector data, for example via the Welch method [110], without requiring multiple independent noise realizations.

noise model, the probability density of a noise realization $n(t)$ is

$$p(n) \propto \exp\left(-\frac{1}{2}(n|n)\right), \tag{2.78}$$

where we have introduced the noise-weighted inner product (Wiener scalar product) between two real time series $a(t)$ and $b(t)$,

$$(a|b) \equiv 4\,\mathrm{Re}\int_0^\infty \frac{\tilde{a}^*(f)\,\tilde{b}(f)}{S_n(f)}\,df. \tag{2.79}$$

This inner product weights frequency bins inversely by the noise power, so that frequency bands where the detector is more sensitive contribute more to the overlap between two waveforms.

2.4.1.2 The matched filter and optimal SNR

Given a template waveform $h_m(t;\boldsymbol{\theta})$ representing our best theoretical model of the expected signal (see §2.4.2), the matched filter is the linear filter $K(t)$ that maximizes the SNR. The SNR is defined as the ratio of the expected filtered output when a signal is present to the root-mean-square of the filtered output when only noise is present. Concretely, one filters the detector output by computing the scalar product

$$\hat{s} = \int_{-\infty}^{\infty} s(t)\,K(t)\,dt, \tag{2.80}$$

and seeks the $K(t)$ that maximizes the ratio of the expected value of $\hat{s}$ when a signal is present to the root-mean-square of $\hat{s}$ when only noise is present. Applying the Cauchy-Schwarz inequality to this optimization in the frequency domain, one finds that the optimal filter is proportional to the signal template weighted by the inverse PSD,

$$\tilde{K}(f) \propto \frac{\tilde{h}_m(f)}{S_n(f)}, \tag{2.81}$$

which down-weights frequency bands where the detector is noisier. This is the key intuition behind matched filtering: the optimal filter is shaped by the signal itself, so that correlating the data against $K(t)$ effectively searches for a signal whose morphology matches the template $h_m(t;\boldsymbol{\theta})$. The additional weighting by $1/S_n(f)$ refines this correlation for realistic, colored detector noise, suppressing frequency bands where the detector is noisier and amplifying those where it is most sensitive.

With this optimal filter, the SNR is

$$\mathrm{SNR} \equiv \rho = \frac{(h_m|s)}{\sqrt{\langle (h_m|n)(n|h_m)\rangle}} = \frac{(h_m|s)}{\sqrt{(h_m|h_m)}} = (\hat{h}_m|s), \tag{2.82}$$

where we have used the property $\langle (a|n)(n|b)\rangle = (a|b)$ of the Wiener inner product, and defined the normalized template

$$\hat{h}_m \equiv \frac{h_m}{(h_m|h_m)^{1/2}}. \tag{2.83}$$

Substituting (2.76) and assuming the noise is uncorrelated with the template, $\langle (h_m|n)\rangle = 0$, the expectation value of ρ reduces to $\langle\rho\rangle = (\hat{h}_m|h)$, which is maximized when the template exactly matches the signal, $h_m = h$. In that ideal case one recovers the optimal SNR,

$$\rho^2_{\rm opt} = (h|h) = 4\int_0^\infty \frac{|\tilde{h}(f)|^2}{S_n(f)}\,df, \tag{2.84}$$

which quantifies the intrinsic detectability of a waveform given the detector noise [111].

For a network of $N_{\rm det}$ detectors, assuming their noise realizations are independent, the individual matched-filter outputs combine in quadrature to give the network SNR,

$$\rho_{\rm net} = \left(\sum_{d=1}^{N_{\rm det}} \rho_d^2\right)^{1/2}, \tag{2.85}$$

where ρ_d is the SNR recovered in detector d [71]. This is the quantity referred to throughout this thesis as the "network SNR", e.g. $\rho_{\rm net}$ for the H1+L1 network of §4.1. In practice, ρ is computed as a time series by sliding the template across the data; a detection candidate is declared when ρ exceeds a predefined threshold, typically $\rho \gtrsim 8$ for a single detector and requiring coincident triggers across the network to suppress noise transients [71].

The matched filter is optimal only under the Gaussian noise assumption of (2.77). In practice, non-Gaussianities and glitches can produce large values of ρ in the absence of a real signal. To mitigate this, matched-filter searches are implemented through dedicated pipelines such as PyCBC [112] and GstLAL [113], which supplement the SNR with a χ^2 consistency test (verifying that the signal power across different frequency bands is consistent with the template) and require coincident triggers across multiple detectors to suppress noise transients. The data is further classified by data quality vetoes that exclude periods of poor detector performance. A candidate event is then ranked by a combined detection statistic and assigned a false alarm rate; events passing a significance threshold are passed to the PE stage described in §2.4.3. Finally, since ρ depends directly on the match between the template and the true signal, the effectiveness of the entire pipeline is tied to the accuracy of the waveform models discussed in the following section.

2.4.2 Waveform models for CBCs

The matched filter derived in §2.4.1 requires a template $h_m(t;\boldsymbol{\theta})$ that accurately represents the expected GW signal as a function of the source parameters. The quality of both the search and the subsequent PE depends directly on how faithfully these templates reproduce the true signal: a mismatch between template and signal reduces the recovered SNR (2.82) and biases the inferred parameters. In practice, matched-filter searches correlate the data against banks of $\sim 10^5$ templates covering the CBC parameter space, while PE requires evaluating $\sim 10^7$–10^9 waveforms [114]. Each template evaluation must therefore be both accurate and computationally cheap.

For a CBC, the GW signal evolves through the three phases described in §2.2 (inspiral, merger, and ringdown) and no single analytic method covers all three. During the inspiral, Post-Newtonian (PN) theory provides accurate closed-form waveforms by expanding the EFE in powers of v^2/c^2 [115], but breaks down near merger where velocities approach c. The merger phase requires full numerical relativity (NR) simulations of the EFE [116], which are highly accurate but computationally prohibitive and serve primarily to calibrate analytic models. The ringdown is described analytically by the quasi-normal modes (QNMs) of the remnant compact object [117].

To cover the full frequency evolution, these approaches are combined into inspiral-merger-ringdown (IMR) waveform families. Effective One Body (EOB) models reformulate the two-body problem as an effective single-body problem, resumming the PN series and matching to NR near merger; current implementations include `SEOBNRv5HM` and `SEOBNRv5PHM` [118, 119], which incorporate higher multipole modes and spin precession. Phenomenological models (IMRPhenom) fit closed-form frequency-domain expressions to hybrid PN-NR waveforms across the three phases; a widely used example, `IMRPhenomXPHM` [120], is calibrated to NR simulations at comparable mass ratios up to $m_1/m_2 \approx 18$ and dimensionless spins $|\chi_i| \lesssim 1$, with coverage extended to extreme mass ratios up to $m_1/m_2 \sim 1000$ through BH perturbation theory waveforms rather than additional NR calibration. NR surrogate models (NRSur) interpolate directly over NR simulations without an underlying analytic ansatz, achieving the highest fidelity at the cost of a restricted parameter space; `NRHybSur3dq8` [121] combines a PN-EOB inspiral with NR near merger, but is limited to non-precessing BBHs with $q \leq 8$. All three families have been progressively extended to include precessing spins, higher multipole modes, and tidal effects for NS components. Fast frequency-domain approximants such as `TaylorF2` or `IMRPhenomD` are favored for matched-filter searches, while more sophisticated models are employed in PE, where waveform accuracy directly impacts posterior reliability.

2.4.3 Bayesian parameter estimation

Once a GW candidate has been identified, the goal is to determine what kind of source produced it and to measure its physical properties. Concretely, we seek to answer questions such as: what were the masses and spins of the merging objects? From which direction on the sky did the signal arrive? How far away was the source? The answers are encoded in the shape of the detected waveform, since the parameters of the binary uniquely determine the GW signal it emits. PE is the systematic procedure by which we compare the space of possible waveforms to the observed data and assign credible ranges to each source parameter [122].

The parameters of a GW signal fall into two categories. Intrinsic parameters describe the binary itself, and extrinsic parameters encode how the merger is observed from Earth. Together these constitute the parameter vector $\boldsymbol{\theta}$ over which inference is performed.

The intrinsic parameters of a quasi-circular CBC include the component masses $m_1 \geq m_2$ and the spin angular momenta $\mathbf{S}_1, \mathbf{S}_2$ of each object. The spins are conveniently decomposed into their dimensionless magnitudes $\chi_{1,2} = c|\mathbf{S}_{1,2}|/(Gm_{1,2}^2) \leq 1$, the tilt angles $\theta_{1,2}$ between each spin and the orbital angular momentum $\mathbf{L}$, and the azimuthal angle ϕ_{12} between the two spin components in the orbital plane. For BNS or NSBH systems, the tidal deformabilities $\Lambda_{1,2}$ are additional intrinsic parameters encoding the NS equation of state. In practice, several derived combinations of the intrinsic parameters appear naturally in the waveform and are therefore better constrained by the data: the chirp mass $\mathcal{M} = (m_1 m_2)^{3/5}/(m_1+m_2)^{1/5}$, which governs the leading-order frequency evolution; the effective inspiral spin $\chi_{\text{eff}} = (m_1\chi_{1z} + m_2\chi_{2z})/(m_1 + m_2)$

[123], which captures the spin-orbit coupling along $\mathbf{L}$; and the effective precession spin χ_p [124], which parametrizes the in-plane spin components responsible for orbital precession.

The extrinsic parameters describe the location and orientation of the binary relative to the detector network: the sky position (α, δ) (right ascension and declination), the luminosity distance D_L[8], the inclination angle θ_{JN} between the total angular momentum $\mathbf{J}$ and the line of sight, the polarization angle ψ, the coalescence phase ϕ_c, and the geocentric coalescence time t_c. The angle ϕ_{JL} between $\mathbf{J}$ and $\mathbf{L}$ projected onto the plane perpendicular to the line of sight is also required when spins are precessing [125]. A complete description of all parameters, including their definitions and dimensions, is provided in Appendix §A.

Not all parameters are equally well constrained by a given observation. The chirp mass is typically the best-measured intrinsic quantity, as it directly controls the rate of frequency evolution visible in the detector band. Individual masses and mass ratio $q = m_2/m_1$ are less well constrained, and spin parameters are generally poorly measured unless the signal has high SNR or a long inspiral in band. Among the extrinsic parameters, the sky position and distance are generally poorly constrained by a single detector, improving significantly with a multi-detector network via time-delay triangulation [126].

In the present work, the injection parameters used to simulate the CBC signal are $\{m_1, m_2, D_L, \chi_1, \chi_2, \theta_1, \theta_2, \phi_{12}, \phi_{JL}, \theta_{JN}, \psi, \phi_c, t_c, \alpha, \delta\}$, corresponding to the full 15-dimensional parameter space of a precessing BBH system.

2.4.3.1 Bayes' theorem and posterior inference

The natural framework for GW PE is Bayesian inference, which provides a systematic way to update prior knowledge about the source parameters using the observed data. Unlike frequentist statistics, which interprets probability as the limiting frequency of repeated experiments, Bayesian inference treats probability as a degree of belief that is updated as new information becomes available. This distinction is particularly relevant in GW astronomy, where astrophysical events are unique and non-repeatable, making frequentist interpretations difficult to apply. The direct link between data and model provided by GR (where a BBH is fully described by just 15 parameters) makes Bayesian inference especially well-suited to this field [122].

Given detector strain data d and a waveform model $\mathcal{M}$, Bayes' theorem relates the posterior distribution $p(\boldsymbol{\theta}|d, \mathcal{M})$ to the likelihood $\mathcal{L}(d|\boldsymbol{\theta}, \mathcal{M})$ and the prior $p(\boldsymbol{\theta}|\mathcal{M})$ through

$$p(\boldsymbol{\theta}|d, \mathcal{M}) = \frac{\mathcal{L}(d|\boldsymbol{\theta}, \mathcal{M})\, p(\boldsymbol{\theta}|\mathcal{M})}{\mathcal{Z}(\mathcal{M})}, \tag{2.86}$$

where the evidence $\mathcal{Z}(\mathcal{M}) = \int \mathcal{L}(d|\boldsymbol{\theta}, \mathcal{M})\, p(\boldsymbol{\theta}|\mathcal{M})\, d\boldsymbol{\theta}$ normalizes the posterior. Each term has a distinct physical interpretation. The likelihood $\mathcal{L}(d|\boldsymbol{\theta}, \mathcal{M})$ quantifies how probable the observed data are given a particular set of source parameters (it encodes the noise model of the detector). The prior $p(\boldsymbol{\theta}|\mathcal{M})$ encodes our knowledge of the parameters before the measurement; for instance, an isotropic prior is assigned to the sky location, reflecting the absence of any preferred direction.

[8] Throughout this work D_L denotes the luminosity distance to the source, except in §3.2 where the same symbol is used for the angular diameter distance to the lens, following standard lensing conventions; it should be clear from context which one is meant.

The posterior is the central quantity of interest: it encodes the full probability distribution of the source parameters given the data, and its one-dimensional marginals,

$$p(\theta_i|d, \mathcal{M}) = \int \prod_{j\neq i} d\theta_j \; p(\boldsymbol{\theta}|d, \mathcal{M}), \tag{2.87}$$

yield credible intervals for individual parameters after integrating out all others.

For Gaussian, stationary noise characterized by the PSD $S_n(f)$ introduced in §2.3.3, the likelihood takes the form of a Gaussian in the residual between the data and the waveform template,

$$\mathcal{L}(d|\boldsymbol{\theta}) \propto \exp\left(-\frac{1}{2}(d - h(\boldsymbol{\theta}) \,|\, d - h(\boldsymbol{\theta}))\right), \tag{2.88}$$

where $(\cdot|\cdot)$ is the noise-weighted inner product defined in (2.79) and $h(\boldsymbol{\theta})$ is the waveform template evaluated at parameters $\boldsymbol{\theta}$. Taking the logarithm of (2.88) gives the log-likelihood form used in practice. This expression makes explicit the connection between PE and matched filtering: maximizing $\ln\mathcal{L}$ over $\boldsymbol{\theta}$ is equivalent to finding the template that best matches the data in the noise-weighted sense of §2.4.1, with the prior regularizing the solution across the full parameter space.

Because the posterior (2.86) is a 15-dimensional distribution for a precessing BBH, it cannot be evaluated analytically. Stochastic sampling methods are therefore required to explore it efficiently, as discussed in §2.4.3.3.

2.4.3.2 Hypothesis testing and the Bayes factor

Beyond PE within a fixed model, Bayesian inference also provides a rigorous framework for comparing competing hypotheses. The natural quantity for this is the odds ratio between two models $\mathcal{M}_1$ and $\mathcal{M}_2$,

$$\mathcal{O}_{12} = \frac{p(\mathcal{M}_1|d)}{p(\mathcal{M}_2|d)} = \frac{\mathcal{Z}(\mathcal{M}_1)}{\mathcal{Z}(\mathcal{M}_2)} \cdot \frac{p(\mathcal{M}_1)}{p(\mathcal{M}_2)}, \tag{2.89}$$

where the second factor is the prior odds, encoding our relative belief in the two models before observing the data. When both models are considered equally plausible a priori, the prior odds is unity and the odds ratio reduces to the Bayes factor,

$$\mathcal{B}_{12} = \frac{\mathcal{Z}(\mathcal{M}_1)}{\mathcal{Z}(\mathcal{M}_2)}, \tag{2.90}$$

which measures the factor by which the data update our relative belief between the two models. A value $\mathcal{B}_{12} \gg 1$ favors $\mathcal{M}_1$, while $\mathcal{B}_{12} \ll 1$ favors $\mathcal{M}_2$; by convention, $|\ln\mathcal{B}_{12}| \geq 8$ is often taken as the threshold for strong evidence in favor of one hypothesis over the other [127, 122].

The most fundamental application in GW astronomy is the signal-vs-noise Bayes factor, where $\mathcal{M}_1$ is the hypothesis that the data contain a GW signal and $\mathcal{M}_2$ is the noise-only hypothesis. In this case,

$$\mathcal{Z}_S = \int \mathcal{L}(d|\boldsymbol{\theta})\, p(\boldsymbol{\theta})\, d\boldsymbol{\theta}, \qquad \mathcal{Z}_N = \mathcal{L}(d|h=0), \tag{2.91}$$

where $h = 0$ denotes the absence of any waveform in the data, and $\mathcal{B}_{SN} = \mathcal{Z}_S/\mathcal{Z}_N$ quantifies how much more probable the data are under the signal hypothesis than under pure noise.

Beyond detection, Bayes factors are used throughout GW data analysis to address a variety of model selection questions: whether a signal requires spin precession or tidal deformability, which waveform approximant better describes the data, or whether the observations are consistent with GR or favor an alternative theory of gravity [125]. In each case the Bayes factor naturally penalizes more complex models through an Occam factor: a model with a large prior volume that is only marginally better supported by the data will yield a lower evidence than a simpler model with a tighter parameter space, reflecting the Bayesian preference for parsimony.

It is important to note that computing $\mathcal{Z}$ requires integrating the likelihood over the entire prior volume, which is precisely what nested sampling is designed to do efficiently, as discussed in §2.4.3.3.2.

2.4.3.3 Stochastic sampling methods

As we mentioned in this section, the posterior distribution (2.86) lives in a high-dimensional parameter space (15 dimensions for a precessing BBH) and cannot be evaluated analytically. A brute-force grid approach is equally intractable: with just 10 bins per dimension, evaluating the likelihood on the full grid would require 10^{15} waveform computations, far beyond any feasible computational budget. Stochastic sampling algorithms are therefore employed to explore the posterior efficiently, concentrating evaluations in the regions of parameter space where the posterior has significant support. The two dominant families used in GW astronomy are Markov Chain Monte Carlo (MCMC) [128, 129] and nested sampling [130], whose complementary strategies are illustrated in Fig. 2.15.

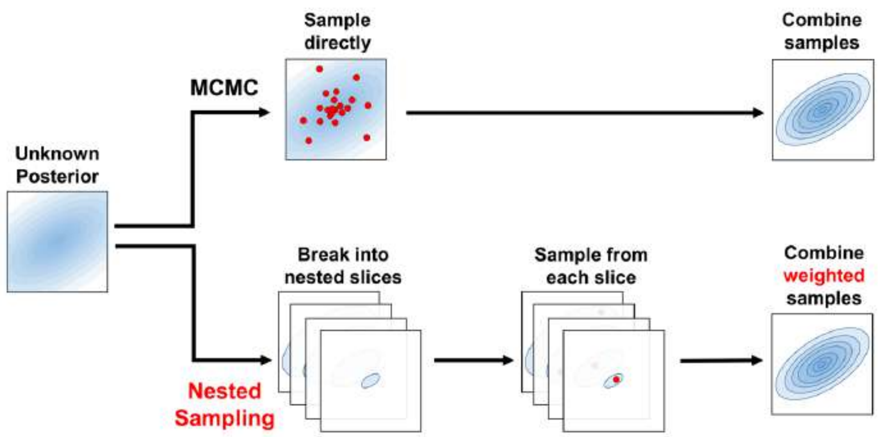

Figure 2.15: Schematic comparison of the two main stochastic sampling strategies. MCMC draws samples directly from the posterior via a random walk, combining them to reconstruct the distribution. Nested sampling instead decomposes the posterior into nested likelihood slices, samples from each slice independently, and combines the weighted samples to recover the posterior and compute the evidence $\mathcal{Z}$. Source: [131].

The output of both methods is a set of posterior samples $\{\boldsymbol{\theta}_k\}$ drawn from $p(\boldsymbol{\theta}|d, \mathcal{M})$, such that the density of samples in any region of parameter space is proportional to the posterior probability there. These samples can then be used to compute expectation values of any quantity of interest, to construct marginalized one- and two-dimensional posterior distributions for individual parameters, and to report credible intervals (the Bayesian analog of confidence intervals).

2.4.3.3.1 Markov Chain Monte Carlo.

MCMC methods construct a random walk through parameter space whose stationary distribution is the target posterior, as shown in Fig. 2.16. The most widely used variant is the Metropolis-Hastings algorithm, which proceeds as follows: starting from an initial point $\boldsymbol{\theta}_0$ (often drawn from the prior), at each iteration t a new point $\boldsymbol{\theta}'$ is proposed from a proposal distribution $q(\boldsymbol{\theta}'|\boldsymbol{\theta}_t)$ (a common choice for q is a Gaussian distribution with mean $\boldsymbol{\theta}_t$ and some fixed covariance matrix), and accepted with probability

$$A = \min\left(1, \frac{p(\boldsymbol{\theta}'|d)\, q(\boldsymbol{\theta}_t|\boldsymbol{\theta}')}{p(\boldsymbol{\theta}_t|d)\, q(\boldsymbol{\theta}'|\boldsymbol{\theta}_t)}\right). \tag{2.92}$$

If the ratio inside the minimum exceeds unity the proposal is always accepted; otherwise it is accepted with probability A. When the proposal distribution q is symmetric, i.e. $q(\boldsymbol{\theta}'|\boldsymbol{\theta}_t) = q(\boldsymbol{\theta}_t|\boldsymbol{\theta}')$, the ratio of proposals cancels and A reduces to the simple posterior ratio $p(\boldsymbol{\theta}'|d)/p(\boldsymbol{\theta}_t|d)$. This acceptance rule ensures detailed balance and guarantees that the chain converges to the correct stationary distribution. By recording the position of the chain at each accepted step, one accumulates draws from the posterior.

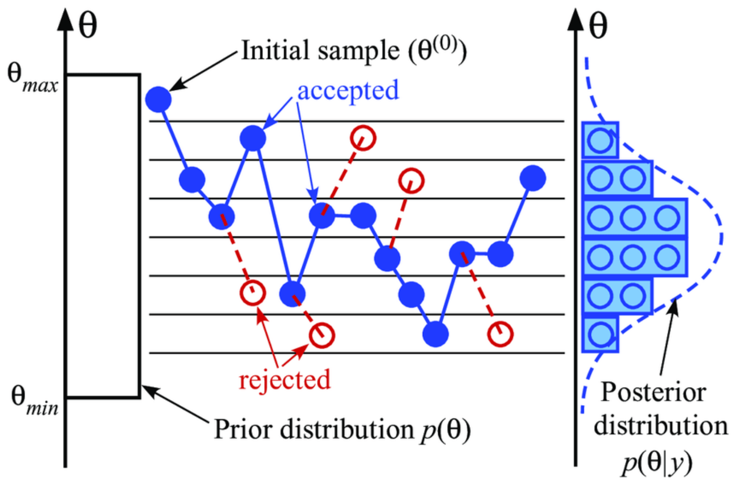

Figure 2.16: Illustration of the Metropolis-Hastings random walk in parameter space. Starting from initial samples drawn from the prior $\pi(\boldsymbol{\theta})$, walkers explore the parameter space following the acceptance criterion (2.92), progressively converging toward the high-probability region of the posterior $p(\boldsymbol{\theta}|y)$. Source: [122].

Several practical considerations arise. The early phase of the chain (the burn-in) is strongly dependent on the starting point and must be discarded before the chain has reached its stationary distribution. Successive samples within a chain are autocorrelated, so the chain must be thinned by retaining only samples separated by the autocorrelation length to obtain approximately independent draws. Finally, MCMC walkers can become trapped in a single mode of a multimodal posterior, a common situation in GW PE, where the sky position degeneracy produces two nearly symmetric modes. Parallel tempering mitigates this by running multiple chains at different temperatures (i.e., with artificially flattened posteriors) that periodically exchange states, allowing the cold chain to escape local maxima by borrowing proposals from hotter, more diffuse chains [125].

2.4.3.3.2 Nested sampling.

Nested sampling is an algorithm designed primarily to compute the evidence $\mathcal{Z}$ directly, with posterior samples produced as a byproduct. It is particularly well-suited to GW PE because it handles multimodal posteriors naturally and provides the evidence required for model comparison via the Bayes factor (2.90).

The algorithm proceeds as follows. A fixed number N_{live} of live points are drawn uniformly from the prior. At each iteration, the live point with the lowest likelihood $\mathcal{L}_{\text{min}}$ is removed and replaced by a new point drawn from the prior subject to the constraint $\mathcal{L} > \mathcal{L}_{\text{min}}$, progressively shrinking the prior volume explored. The discarded point contributes to the evidence integral via the trapezoid rule,

$$\mathcal{Z} \approx \sum_i \mathcal{L}_i \, w_i, \qquad w_i = X_{i-1} - X_i, \tag{2.93}$$

where X_i is the fraction of the prior volume enclosed by the i-th likelihood contour, estimated statistically from the number of live points. The algorithm terminates when the fractional contribution of the remaining prior volume to $\mathcal{Z}$ falls below a predefined threshold. Posterior samples are obtained by reweighting the discarded points by $\mathcal{L}_i w_i / \mathcal{Z}$. The geometric interpretation of the algorithm is shown in Fig. 2.17.

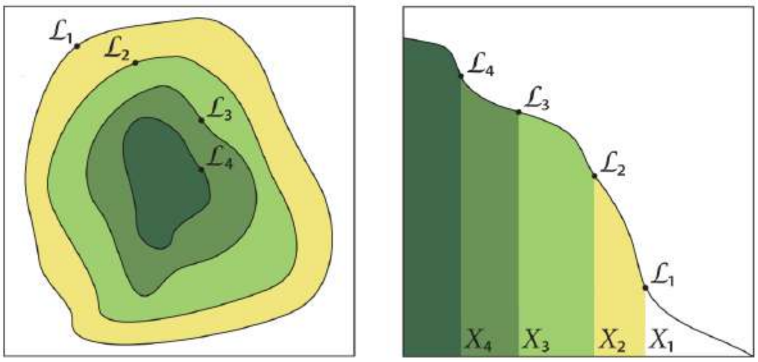

Figure 2.17: Geometric interpretation of nested sampling. *Left:* iso-likelihood contours in parameter space, with the discarded live points $\mathcal{L}_1 > \mathcal{L}_2 > \mathcal{L}_3 > \mathcal{L}_4$ marking successive likelihood shells. *Right:* the corresponding one-dimensional mapping $\mathcal{L}(X)$ as a function of the enclosed prior volume X, where the evidence $\mathcal{Z}$ is approximated as the area under the curve via the trapezoid rule. Source: [132].

The main algorithmic challenge is drawing new live points efficiently from the likelihood-constrained prior. The most widely used implementations in GW astronomy are MultiNest [133], which approximates the iso-likelihood contour with a set of ellipsoids that can be sampled exactly, and Dynesty [131], which uses dynamic live-point allocation to concentrate computational effort where the posterior has the most support. Both are accessible through the Bilby [134] inference library, which provides a modular Python interface for CBC PE and is used in this work.

In summary, this chapter has established the theoretical and experimental foundations of GW astronomy. Beginning from the linearized EFE, we derived the existence of GWs as propagating spacetime perturbations, characterized their two polarization states, and quantified their effect on test masses through the geodesic deviation equation. We then surveyed the main astrophysical source classes and described the detection principle underlying current interferometric observatories, from the antenna pattern functions and noise budget to the global detector network. The chapter closed with the Bayesian PE framework (matched filtering, waveform models, and stochastic sampling) that allows us to extract source properties from detected signals. These tools constitute the methodological foundation of the analysis presented in the following chapters.

Chapter 3

Gravitational lensing of gravitational waves

Having established the theoretical and observational foundations of GW astronomy in the preceding chapter, we now turn to a phenomenon that lies at the heart of this thesis: the gravitational lensing of GWs. We begin with the basics of gravitational lensing and its particular features in the GW context (§3.1), before developing the geometric optics formalism that governs strong lensing, including the lens equation, image formation, magnification, time delay, and the main lens models used in the literature (§3.2). We then characterize the signatures that strong lensing imprints on GW signals, including the amplification factor, image classification, the Morse phase, and their effect on the observed strain and inferred source parameters (§3.3). The chapter closes with a discussion of how the joint analysis of lensed image pairs can be exploited to improve sky localization (§3.4), motivating the methodology developed in the subsequent chapters.

3.1 Gravitational lensing

Gravitational lensing is a direct consequence of GR: massive objects curve spacetime, and any particle following a geodesic (including massless ones such as photons or GWs) will have its trajectory deflected [135, 136, 137]. The phenomenon was confirmed observationally by Eddington during the solar eclipse of 1919, when the apparent positions of stars near the Sun were displaced by exactly the amount predicted by GR [9]. Since then, gravitational lensing has become one of the most powerful tools in observational astrophysics, providing information about the mass distribution of the Universe independently of its luminous content [138].

When a GW propagates from its source to a detector, massive structures along the line of sight act as gravitational lenses (galaxies, galaxy clusters, and other compact objects). Depending on the geometry and lens mass, three regimes are distinguished [139]: strong lensing, caused by galaxies or galaxy clusters, produces multiple images of the same source arriving at different times and with different amplitudes; microlensing, caused by compact objects, produces overlapping images whose interference can modulate the observed strain; and weak lensing, due to large-scale structure, induces small coherent distortions. Figure 3.1 illustrates the general picture of GW lensing, including both strong-lensing and microlensing effects. This work focuses exclusively on the strong lensing regime.

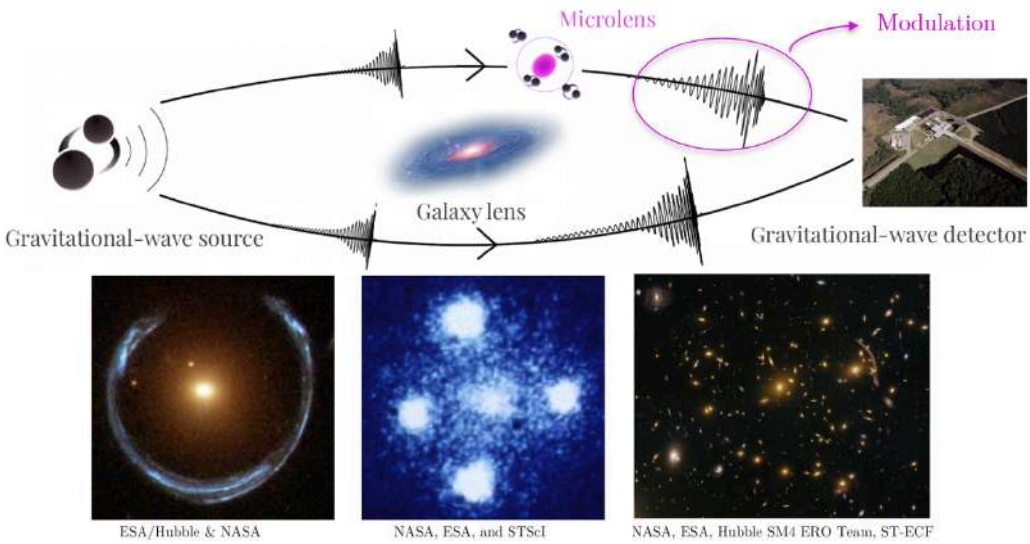

Figure 3.1: Schematic representation of GW lensing. A galaxy acting as a macrolens can produce multiple images of the same GW signal, arriving at the detector with different magnifications and time delays. Compact objects embedded within the lens galaxy (microlenses) can further perturb these images, introducing additional amplification, interference, and diffraction effects. The images shown at the bottom illustrate examples of strong gravitational lensing observed in electromagnetic astronomy. Original illustration based on sources: [140, 141].

3.1.1 Particularities of gravitational wave lensing

The deflection of GWs by a gravitational field follows the same geometric principles as for light, but several properties of GWs lead to qualitatively different lensing phenomenology.

The most fundamental difference is the wavelength regime. GWs detectable by ground-based interferometers, spanning the $\sim$10-5000 Hz band (§2.3.3), have wavelengths ranging from roughly 60 km to 30 000 km, many orders of magnitude larger than EM wavelengths. The key parameter governing the lensing regime is the ratio $\lambda_{\rm GW}/r_E$, where $r_E = D_L\theta_E$ is the physical Einstein radius of the lens (see §3.2.3). When $\lambda_{\rm GW} \ll r_E$, each image travels along a distinct geometric path and the standard geometric optics formalism applies. When $\lambda_{\rm GW} \gtrsim r_E$, diffraction becomes important and wave optics must be used, producing frequency-dependent modulations of the waveform amplitude and phase with no EM analog [142]. For galaxy-scale lenses ($r_E \sim \mathcal{O}(\text{kpc})$), ground-based GW detectors always operate in the geometric optics regime. Wave optics effects become relevant for lower mass lenses such as stars or IMBHs, which fall outside the scope of this work. Fig. 3.2 summarizes the lens-mass scales associated with each regime.

A second key difference is coherence. GWs are coherent waves with a well-defined amplitude and phase, so when two lensed images arrive within a time shorter than the signal duration they can interfere, producing observable beating patterns in the waveform. This effect has no counterpart in standard EM lensing, where sources are typically incoherent.

Finally, unlike EM radiation, GWs propagate through the intergalactic medium and through the lens galaxy itself without absorption or dispersion. The lensed GW signal therefore encodes clean information about the source and the lensing potential, uncontaminated by dust extinction or plasma effects that complicate EM observations.

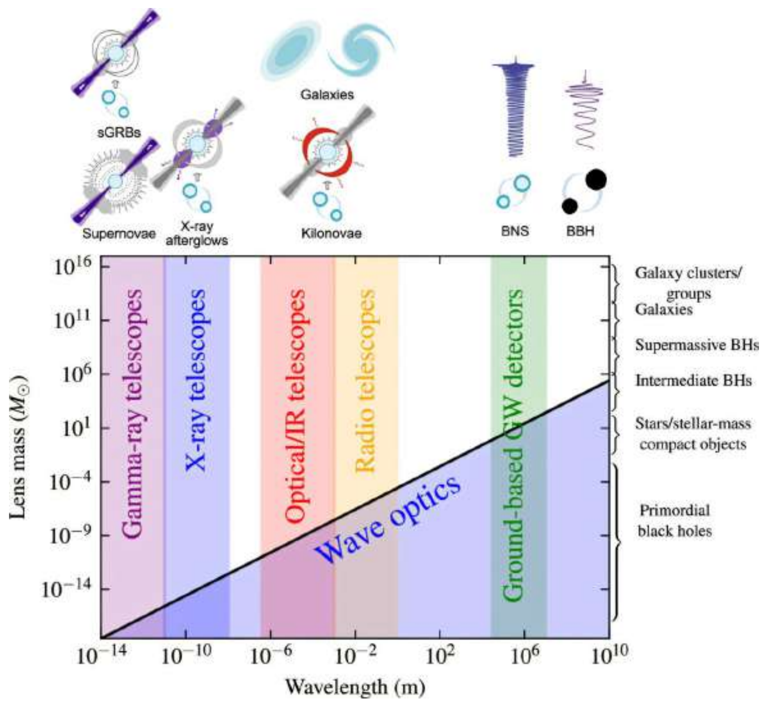

Figure 3.2: Illustration of the lens-mass scales relevant for gravitationally lensed signals. Depending on the relation between the radiation wavelength and the characteristic scale of the lensing potential, lensing occurs in either the geometric optics or wave optics regime. While geometric optics describes distinct propagation paths, wave optics accounts for diffraction and interference effects. Due to their long wavelengths, GW can exhibit wave-optics effects for relatively low lens masses. Source: [143].

3.1.2 Relevance for gravitational wave astronomy

No confirmed detection of gravitational lensing of GWs has been reported despite dedicated searches in GWTC-1 through GWTC-4 [144]. As detector sensitivities improve, detection is expected to become feasible within the next observing runs, with predictions suggesting more than one strongly lensed event per year at design sensitivity [145].

Such detections would open a rich set of applications: measuring the Hubble constant via lensing time delays combined with GW standard sirens [146, 147], probing dark matter distributions in lens galaxies [148], and testing GR by comparing the propagation speed of GWs and light through the same lensing potential [149]. Lensing also biases the apparent properties of detected sources (a magnified image appears closer and more massive than it truly is) making the identification of lensed events important for controlling systematics in population studies [150].

In addition, multiple images of the same source provide repeated observations that can be combined to improve PE and sky localization [147, 151]. This last application is the central motivation of the present work: by performing a joint Bayesian analysis of two strongly lensed images of the same BBH merger, we investigate how the sky localization of the source can be dramatically improved compared to the analysis of a single image.

3.2 Gravitational lensing formalism

3.2.1 Deflection of light by a gravitational field

The deflection of light by a gravitational field can be derived by studying null geodesics in a perturbed spacetime. Our Universe is well described by a spatially flat Friedmann-Lemaître-Robertson-Walker (FLRW) metric, but for the scales relevant to lensing we can work locally with Minkowski spacetime as a starting point and incorporate cosmological corrections later through angular diameter distances.

The presence of a massive lens introduces inhomogeneities in the spacetime metric. Treating the lens in the weak-field limit (slow-moving, $v \ll c$, and small curvature), the energy-momentum tensor is dominated by the energy density component $T^{00} = \rho$, with all other components negligible. Linearizing the EFE to first order in the metric perturbation $\gamma_{\mu\nu}$, and neglecting time derivatives since the lens evolves slowly, the spatial components satisfy a Poisson equation whose solution yields the perturbed metric [152]

$$ds^2 = -\left(1 + \frac{2U}{c^2}\right) c^2\, dt^2 + \left(1 - \frac{2U}{c^2}\right) \delta_{ij}\, dx^i\, dx^j, \tag{3.1}$$

where $U(\mathbf{x})$ is the Newtonian gravitational potential of the lens, satisfying $\nabla^2 U = 4\pi G\rho$. This is the same linearization procedure applied in §2.1 to derive the GW wave equation, now applied to a static mass distribution rather than a radiating source.

A photon travelling along a null geodesic of (3.1) at closest approach distance ξ from a mass M (the impact parameter) accumulates a total deflection angle [153]

$$\hat{\alpha} = \frac{4GM}{c^2\xi}, \tag{3.2}$$

twice the value predicted by Newtonian gravity, and confirmed by Eddington's 1919 measurement [9]. Since the deflection is linear in the perturbation, the contributions of multiple masses add linearly, so any mass distribution can be treated as a superposition of point-mass deflectors.

3.2.2 The lens equation

The standard framework for gravitational lensing adopts the thin lens approximation: all deflection is assumed to occur at a single plane perpendicular to the line of sight, the lens plane, while the source lies on a parallel source plane [154]. This approximation is valid when the physical extent of the lens is much smaller than the distances between observer, lens, and source. The geometry of the problem is illustrated in Fig. 3.3, which defines the angular positions $\boldsymbol{\theta}$ (image) and $\boldsymbol{\beta}$ (source), the angular diameter distances D_L, D_S, and D_{LS}, and the impact parameter $\boldsymbol{\xi} = D_L\boldsymbol{\theta}$.

Projecting the geometry onto the source plane and applying the small-angle approximation, one obtains the lens equation

$$\boldsymbol{\beta} = \boldsymbol{\theta} - \boldsymbol{\alpha}(\boldsymbol{\theta}), \tag{3.3}$$

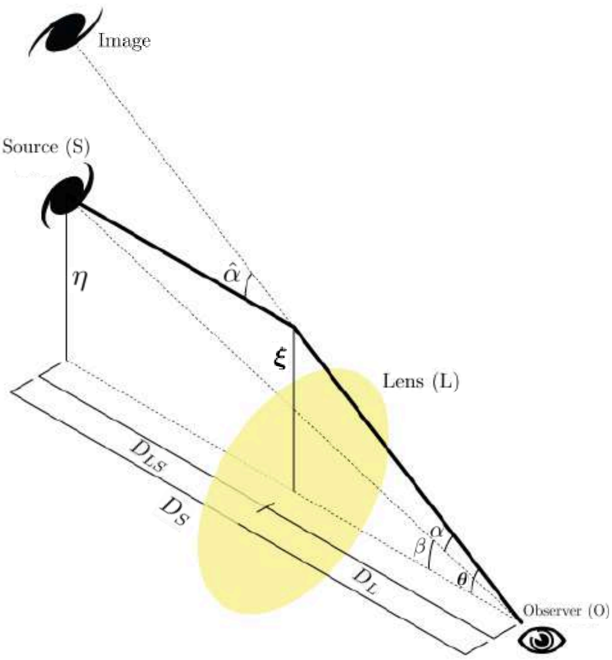

Figure 3.3: Geometry of gravitational lensing in the thin lens approximation. The observer (O), lens (L), and source (S) define the angular diameter distances D_L, D_S, and D_{LS}. The true source position $\boldsymbol{\beta}$ and the apparent image position $\boldsymbol{\theta}$ are related by the lens equation through the reduced deflection angle $\boldsymbol{\alpha}$. Adapted from the source: [155].

where the reduced deflection angle is defined as

$$\boldsymbol{\alpha}(\boldsymbol{\theta}) \equiv \frac{D_{LS}}{D_S}\,\hat{\boldsymbol{\alpha}}(\boldsymbol{\theta}). \tag{3.4}$$

The lens equation maps each image position $\boldsymbol{\theta}$ to a unique source position $\boldsymbol{\beta}$, but the inverse mapping can be multi-valued: a single source can produce multiple images. Equivalently, by Fermat's principle, images form at the stationary points of the Fermat potential (or scaled time delay surface)

$$\phi(\boldsymbol{\theta}, \boldsymbol{\beta}) = \frac{1}{2}|\boldsymbol{\theta} - \boldsymbol{\beta}|^2 - \psi(\boldsymbol{\theta}), \tag{3.5}$$

where $\psi(\boldsymbol{\theta})$ is the projected gravitational potential of the lens, obtained from the Newtonian potential $U(\boldsymbol{\xi}, Z)$ by integration along the line of sight,

$$\psi(\boldsymbol{\theta}) = \frac{2D_{LS}}{c^2 D_S D_L} \int U(\boldsymbol{\xi}, Z)\, dZ. \tag{3.6}$$

The deflection angle is then $\boldsymbol{\alpha} = \nabla_{\boldsymbol{\theta}}\psi$, and the lens equation (3.3) is recovered from $\nabla_{\boldsymbol{\theta}}\phi = 0$.

Since stationary points of ϕ always appear and disappear in pairs (a minimum together with a saddle point, or a maximum together with a saddle point) as the source position varies, the total number of images is always odd [156]: there is always at least one image corresponding to the global minimum of ϕ, and any additional images appear in pairs of two.[1] The wavefront

[1]This holds for a smooth and finite lensing potential. The point-mass and singular isothermal sphere models introduced in §3.2.5 have a divergent surface density at the origin, which removes the additional odd, typically

interpretation of this result is illustrated in Fig. 3.4.

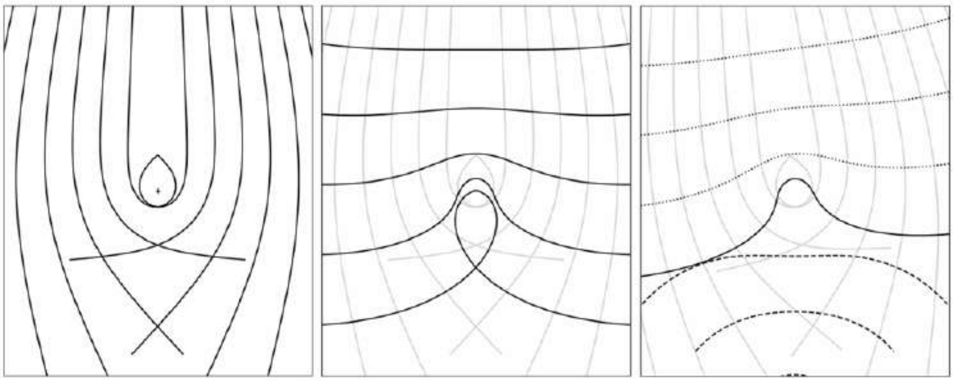

Figure 3.4: Wavefront picture of gravitational lensing. *Left*: light rays from a distant source are deflected near the center, where the gravitational field is strongest. Center: the corresponding wavefronts at a fixed time. The small closed loop arises from rays undergoing very large deflections. *Right*: a second wavefront traced backwards from the observer (dashed) is superimposed on the frozen source wavefront (solid). Lensed images form where the two wavefronts share the same local direction of propagation, which corresponds to the stationary points of the Fermat potential (3.5). Source: [154].

3.2.3 Einstein radius and multiple image formation

For a point mass M located at the origin, the lens equation (3.3) admits an analytic solution. When the source lies exactly behind the lens ($\boldsymbol{\beta} = 0$), perfect circular symmetry produces a ring image at the Einstein radius

$$\theta_E = \sqrt{\frac{4GM}{c^2}\frac{D_{LS}}{D_L D_S}}, \tag{3.7}$$

which sets the characteristic angular scale of lensing [154]. For a general source position, the lens equation reduces to a quadratic in $|\boldsymbol{\theta}|$, always producing exactly two images at positions

$$\theta_{1,2} = \frac{1}{2}\left(\beta \pm \sqrt{\beta^2 + 4\theta_E^2}\right), \tag{3.8}$$

where $\beta = |\boldsymbol{\beta}|$. The image at $\theta_1 > \theta_E$ lies outside the Einstein ring, while $\theta_2 < \theta_E$ lies inside. More generally, multiple images form when the projected surface mass density of the lens,

$$\Sigma(\boldsymbol{\xi}) = \int \rho(\boldsymbol{\xi}, Z)\, dZ, \tag{3.9}$$

obtained by integrating the three-dimensional mass density ρ of the lens along the line of sight coordinate Z, exceeds the critical surface density

$$\Sigma_{\rm cr} = \frac{c^2}{4\pi G}\frac{D_S}{D_L D_{LS}}, \tag{3.10}$$

strongly demagnified, image predicted by the theorem; both therefore produce an even number of images without contradicting it.

a combination of distances that sets the lensing strength for a given geometry. It is convenient to introduce the dimensionless convergence

$$\kappa(\boldsymbol{\xi}) \equiv \frac{\Sigma(\boldsymbol{\xi})}{\Sigma_{\rm cr}}, \tag{3.11}$$

which measures the projected mass density in units of the critical value. When $\kappa > 1$ in some region of the lens plane, the lensing effect is strong enough to produce multiple images. Equivalently, κ satisfies the two-dimensional Poisson equation $\nabla^2\psi = 2\kappa$, linking the lensing potential to the projected mass distribution of the lens.

3.2.4 Magnification and time delay

Since gravitational lensing preserves surface brightness but distorts the solid angle subtended by a source, it produces a net magnification. The magnification of each image is given by the inverse determinant of the Jacobian matrix of the lens mapping,

$$\mu = \frac{1}{\det \mathbf{A}}, \qquad A_{ij} = \delta_{ij} - \frac{\partial^2\psi}{\partial\theta_i\,\partial\theta_j}. \tag{3.12}$$

The sign of μ encodes the parity of the image: images at minima and maxima of ϕ have positive magnification, while images at saddle points have negative magnification (i.e. they are mirror-inverted). Explicit expressions for the magnifications of each image type are given in §3.2.5 for the specific lens models considered in this work.

The arrival time of each image is delayed with respect to an unlensed ray by two contributions. The geometric delay arises from the longer path length of the deflected ray, while the Shapiro delay [157, 158] results from the slowing of the signal in the gravitational potential of the lens. Together, they yield the time delay between two images i and j

$$\Delta t_{ij} = \frac{(1+z_L)}{c}\frac{D_L D_S}{D_{LS}}\left[\phi(\boldsymbol{\theta}_i, \boldsymbol{\beta}) - \phi(\boldsymbol{\theta}_j, \boldsymbol{\beta})\right], \tag{3.13}$$

where the $(1+z_L)$ factor accounts for the cosmological redshifting of the delay at the lens. For galaxy-scale lenses, Δt_{ij} ranges from minutes to months [145], and constitutes one of the key observables in GW lensing searches.

3.2.5 Lens models

The choice of lens model determines the deflection angle $\hat{\boldsymbol{\alpha}}$ and hence all lensing observables. While real lenses are complex, extended, and often non-circular, a small set of idealized models captures the essential physics and is widely used in both analytical work and population studies [154].

3.2.5.1 Point mass lens

The simplest model treats the entire lens mass M as concentrated at a single point, so that the projected surface mass density is $\Sigma(\boldsymbol{\xi}) = M\delta^{(2)}(\boldsymbol{\xi})$ [142]. The deflection angle is then exactly (3.2), and the lens equation reduces to a quadratic whose two solutions are the image positions (3.8). The corresponding magnifications are

$$\mu_{\pm} = \frac{1}{2} \pm \frac{u^2+2}{2u\sqrt{u^2+4}}, \qquad u \equiv \frac{\beta}{\theta_E}, \tag{3.14}$$

where u is the normalized source-lens separation [154]. The positive image (μ_+) is always magnified, while the negative image (μ_-) can be magnified or demagnified. The total magnification diverges as $u \to 0$ (Einstein ring) and approaches unity for $u \gg 1$ (no lensing). The time delay between the two images is

$$\Delta t = \frac{4GM_{Lz}}{c^3}\left[\frac{u\sqrt{u^2+4}}{2} + \ln\frac{\sqrt{u^2+4}+u}{\sqrt{u^2+4}-u}\right], \tag{3.15}$$

where $M_{Lz} = M_L(1+z_L)$ is the redshifted lens mass [142].

3.2.5.2 Singular isothermal sphere

A more realistic description of galaxy-scale lenses is the singular isothermal sphere (SIS), which models the lens as a self-gravitating gas of stars with a Maxwellian velocity distribution characterized by a constant velocity dispersion σ_v. The three-dimensional density profile is

$$\rho(r) = \frac{\sigma_v^2}{2\pi G r^2}, \tag{3.16}$$

which yields a flat rotation curve $v_c = \sqrt{2}\,\sigma_v$, consistent with the observed flatness of galaxy rotation curves and with the kinematics of early-type galaxies [139]. Projecting onto the lens plane gives the surface mass density

$$\Sigma(\xi) = \frac{\sigma_v^2}{2G\xi}, \tag{3.17}$$

which depends only on the radial distance $\xi = |\boldsymbol{\xi}|$, confirming that the SIS is an axisymmetric lens [159]. Expressing this in terms of the convergence gives

$$\kappa(\theta) = \frac{\theta_E}{2\theta}, \qquad \theta_E = 4\pi\frac{\sigma_v^2}{c^2}\frac{D_{LS}}{D_S}, \tag{3.18}$$

and the deflection angle has the particularly simple form $\boldsymbol{\alpha} = \theta_E\,\hat{\boldsymbol{\theta}}$, independent of the image position. When $u \leq 1$ (source inside the Einstein radius) two images form with magnifications $\mu_+ = 1+1/u$ and $|\mu_-| = 1/u-1$ (the second image sits at a saddle point of the Fermat potential, so its signed magnification is negative), and time delay $\Delta t = (8GM_{Lz}/c^3)\,u$; for $u > 1$ only a single image is produced [142].

The SIS is the simplest axisymmetric model for strong lensing by individual galaxies. Its elliptical generalization, the singular isothermal ellipsoid (SIE), extends the convergence to elliptical isocontours parametrized by the projected axis ratio q_l, and has traditionally been the model of choice in electromagnetic surveys of galaxy-galaxy lensing [160]. External shear can be added to either model to account for the tidal field of surrounding mass. For galaxy cluster lenses, more complex profiles are required, typically a Navarro-Frenk-White (NFW) dark matter halo combined with a SIS profile for the brightest central galaxy [139].

3.2.5.3 Extended power-law lens models

Observations of early-type galaxies show that their density slopes are close to isothermal but not exactly so, with a distribution of slopes $\gamma \sim \mathcal{N}(2.0, 0.2^2)$ [161]. A more flexible family is the elliptical power law (EPL), which generalizes the SIE by allowing the density slope to vary freely. The three-dimensional density profile is

$$\rho(r) \propto r^{-\gamma}, \tag{3.19}$$

where $\gamma = 2$ recovers the isothermal case [162]. The projected convergence for an elliptical mass distribution with axis ratio q_l takes the form

$$\kappa(\boldsymbol{\theta}) = \frac{3-\gamma}{2}\left(\frac{\theta_E}{\sqrt{q_l\theta_1^2 + \theta_2^2/q_l}}\right)^{\gamma-1}, \tag{3.20}$$

where θ_1 and θ_2 are coordinates aligned with the major and minor axes of the lens. For $\gamma = 2$ and $q_l = 1$ this reduces to the SIS convergence (3.18).

The EPL model has become the preferred choice for population studies of galaxy-scale GW lensing, as its variable slope and ellipticity naturally produce the range of image configurations observed in strongly lensed systems [163]. In practice, external shear is added to account for the tidal contribution of the lens environment, yielding the EPL + shear model that constitutes the current standard in strong lensing analyses [164] and the model adopted in this work for the generation of the lensed GW population.

3.3 Strong lensing of gravitational waves

Strong gravitational lensing occurs when a GW propagates through the gravitational potential of a sufficiently massive foreground object, producing multiple images of the same source. In GW astronomy these images manifest as repeated signals arriving at the detector with identical intrinsic parameters but different amplitudes, arrival times, and phases [137]. This section develops the theoretical framework that describes these effects within the geometric optics limit relevant for ground-based detectors and galaxy-scale lenses.

3.3.1 From geometric optics to the amplification factor

The effect of lensing on a GW signal can be modeled by an amplification factor $F(f;\boldsymbol{\phi})$ modifying the unlensed waveform in the frequency domain [142]

$$\tilde{h}_L(f;\,\boldsymbol{\theta},\boldsymbol{\phi}) = F(f;\,\boldsymbol{\phi})\,\tilde{h}_U(f;\,\boldsymbol{\theta}), \tag{3.21}$$

where $\tilde{h}_U(f;\boldsymbol{\theta})$ is the frequency-domain unlensed strain with binary parameters $\boldsymbol{\theta}$, $\tilde{h}_L(f;\boldsymbol{\theta},\boldsymbol{\phi})$ is the corresponding lensed strain, and $\boldsymbol{\phi}$ collects the lensing parameters. The simplicity of this relation arises because F depends solely on the mass density profile of the lens and the source position relative to it, decoupling cleanly from the intrinsic source parameters.

In the general wave optics case, $F(f;\boldsymbol{\phi})$ is given by a Fresnel-Kirchhoff diffraction integral over the lens plane [165]

$$F(w) = \frac{w}{2\pi i}\int d^2x\, e^{iw\,\phi(x,y)}, \tag{3.22}$$

where x and y are dimensionless coordinates in the lens and source planes, $\phi(x,y)$ is the Fermat potential (3.5), and

$$w \equiv \frac{8\pi G M_{Lz} f}{c^3} = \frac{4\pi R_S f}{c} = \frac{2R_S}{\bar{\lambda}}, \tag{3.23}$$

is a dimensionless frequency proportional to the product of the lens mass and the GW frequency, with $M_{Lz} = M_L(1+z_L)$ the redshifted effective lens mass, $R_S = 2GM_{Lz}/c^2$ the corresponding Schwarzschild radius, and $\bar{\lambda} \equiv \lambda/(2\pi) = c/(2\pi f)$ the reduced GW wavelength. The value of w controls the lensing regime: for $w \ll 1$ (equivalently $\bar{\lambda} \gg R_S$) lensing effects are negligible (perturbative regime); for $w \sim 1$ diffraction and interference are important (wave optics regime); and for $w \gg 1$ (equivalently $\bar{\lambda} \ll R_S$) the geometric optics limit applies. For a fixed GW frequency f, the transition between regimes is set by the lens mass, as illustrated in Fig. 3.5.

In the geometric optics limit $w \to \infty$, the integral is dominated by the stationary points of ϕ, and each image j contributes one term, which we write following the convention of [166]:

$$F(f;\boldsymbol{\phi}) = \sum_j |\mu_j|^{1/2}\, e^{-2\pi i f t_j + i\pi n_j\,\mathrm{sign}(f)}, \tag{3.24}$$

where μ_j is the magnification (3.12) evaluated at image j, t_j is the arrival time of image j relative to a reference, and n_j is the Morse index introduced in the following section.

3.3.2 Image classification and the Morse phase

The Morse index n_j classifies each image according to the topology of the Fermat potential at its location [167]

$$n_j = \begin{cases} 0 & \text{Type I (minimum of } \phi) \\ 1/2 & \text{Type II (saddle point of } \phi) \\ 1 & \text{Type III (maximum of } \phi) \end{cases} \tag{3.25}$$

This classification is directly related to the parity of the image: Type I and III images have $\det \mathbf{A} > 0$ (positive parity), while Type II images have $\det \mathbf{A} < 0$ (negative parity, mirror-inverted), consistent with the magnification matrix defined in (3.12). The structure of the arrival-

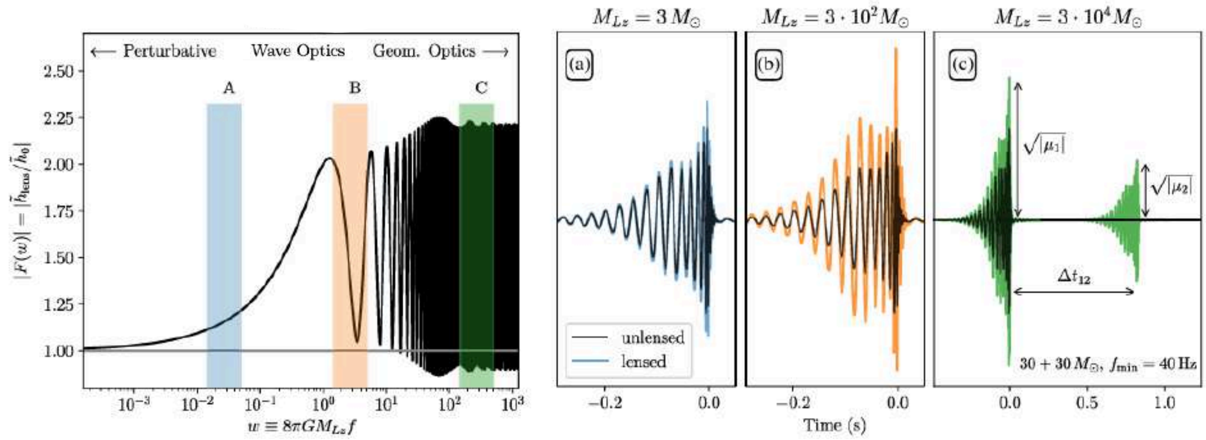

Figure 3.5: *Left:* modulus of the dimensionless amplification factor $|F(w)|$ as a function of dimensionless frequency $w = 8\pi G M_{Lz} f/c^3$ for a SIS lens with impact parameter $y = 0.7$. Three regimes are distinguished: perturbative (low w, negligible lensing), wave optics (intermediate w, diffraction and interference), and geometric optics (high w, $|F|$ converges to the sum of image magnifications). *Right:* time-domain waveforms of a $30+30\, M_\odot$ BBH system (unlensed in black, lensed in color) for three effective lens masses representative of each regime (panels a, b, c). In the geometric optics limit (panel c), the two images are clearly separated in time by Δt_{12} and have distinct amplitudes $\sqrt{|\mu_1|}$ and $\sqrt{|\mu_2|}$. Source: [159].

time surface and the resulting image positions are illustrated in Fig. 3.6 for three different source positions with the same lens potential.

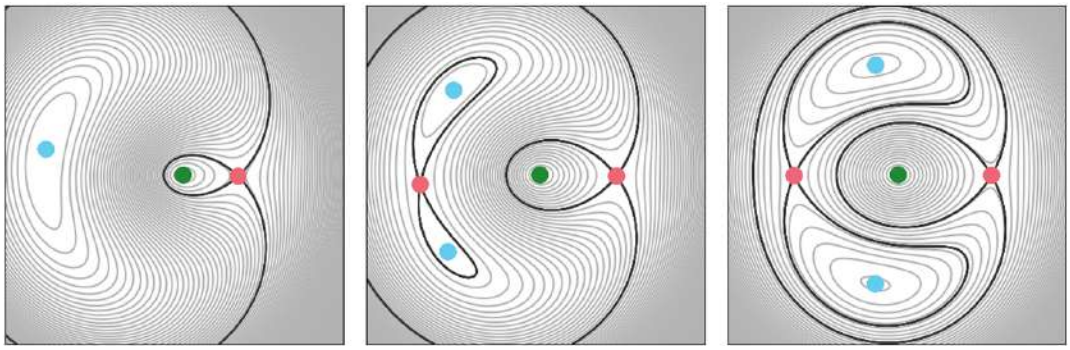

Figure 3.6: Contour maps of the Fermat potential $\phi(\boldsymbol{\theta})$ for three different source positions with the same lens potential, shown qualitatively in the image plane $\boldsymbol{\theta}$. Type I images (minima) are marked in blue, Type II (saddle points) in red, and Type III (maxima) in green. The thick black contours pass through the saddle points. The figure is intended to illustrate the topological structure of the arrival-time surface rather than a specific lens model or parameter space. Source: [154].

Galaxy-scale lenses typically produce two bright images of Type I and II. A faint central image of Type III can in principle exist, but is usually demagnified below detectability by the central mass concentration of the lens.

The factor $e^{i\pi n_j\, \mathrm{sign}(f)}$ in (3.24) is the Morse phase. A Type I image ($n = 0$) leaves the waveform morphology unchanged. A Type III image ($n = 1$) introduces an overall sign flip. A Type II image ($n = 1/2$) applies a $\pi/2$ phase shift equivalent to a Hilbert transform of the time-domain signal. For non-precessing binaries with only the dominant quadrupole mode, this phase shift is degenerate with the coalescence phase and cannot be detected. For signals with spin precession or higher-order multipole moments, such as those considered in this work, the Morse phase produces measurable imprints on the waveform [137].

3.3.3 Lensed gravitational wave strain and apparent source parameters

Combining the geometric optics amplification factor (3.24) with the unlensed waveform, the frequency-domain strain of the j-th lensed image is

$$\tilde{h}_L^j(f;\,\boldsymbol{\theta},\boldsymbol{\phi}_j) = |\mu_j|^{1/2}\, e^{-2\pi i f t_j + i\pi n_j\,\mathrm{sign}(f)}\, \tilde{h}_U(f;\,\boldsymbol{\theta}), \tag{3.26}$$

where $\boldsymbol{\phi}_j = \{\mu_j, t_j, n_j\}$ collects the lensing parameters of image j, consistent with the general form (3.21) [166]. The three multiplicative factors encode amplitude rescaling, time delay, and Morse phase respectively. All images share the same intrinsic parameters $\boldsymbol{\theta}$, which is the fundamental property exploited in joint PE analyses of lensed pairs.

Because (3.26) enters the detector identically to an unlensed signal, a PE pipeline that does not account for lensing will infer biased source parameters. The factor $|\mu_j|^{1/2}$ is degenerate with the luminosity distance, since $h \propto D_L^{-1}$, yielding an effective distance

$$D_L^{\mathrm{eff},j} = \frac{D_L}{\sqrt{|\mu_j|}}, \tag{3.27}$$

so a magnified image appears closer than the true source. Since redshift is inferred from D_L, the source-frame masses are also biased: the inferred chirp mass $\mathcal{M}_c^{\mathrm{obs}} = (1+\tilde{z})\mathcal{M}_c^{\mathrm{src}}$ corresponds to an incorrect redshift $\tilde{z}$, and lensed events at high redshift may be misidentified as nearby events with anomalously massive BHs [168]. Similarly, the effective coalescence time is

$$t_c^{\mathrm{eff},j} = t_c + t_j, \tag{3.28}$$

so the two images of the same event appear to arrive at different times.

3.3.4 Lensing observables

In a joint analysis of two lensed images the absolute lensing parameters $\boldsymbol{\phi}_j$ are not individually measurable; only their relative values between the two images are observable. The primary lensing observables for an image pair are the relative magnification

$$\mu_{\mathrm{rel}} \equiv \frac{\mu_2}{\mu_1}, \tag{3.29}$$

the time delay between images $\Delta t = t_2 - t_1$ (3.13), and the difference in Morse indices $\Delta n = n_2 - n_1$ [169]. The relative magnification can span several orders of magnitude depending on the source-lens alignment, while Δn is only detectable when the signal carries precession or higher-mode content. Together, $\{\mu_{\mathrm{rel}}, \Delta t, \Delta n\}$ constitute the primary lensing observables exploited in a joint analysis of lensed image pairs; how these are parameterized in the specific joint PE analysis carried out in this work is described in §3.4.2.

3.4 Sky localization with lensed image pairs

Strong gravitational lensing offers a natural opportunity to improve the sky localization of GW sources beyond what is achievable from a single image [2]. This section describes the limitations of single-image sky localization, explains how the joint analysis of lensed image pairs overcomes them, and introduces the statistical framework used to quantify the evidence for the lensed hypothesis.

3.4.1 Limitations of single-image sky localization

The sky localization of a GW source relies primarily on triangulation: the difference in signal arrival times across the detector network constrains the direction of propagation, while the antenna pattern functions (see §2.3.2.1) and strain amplitudes (see §2.3.3) provide complementary information [139]. The achievable localization depends critically on the number and geometry of the detectors: a single detector provides no directional information for short-duration transients; two detectors constrain the source to two arc-shaped regions on the celestial sphere, lying along the ring of constant time delay around the axis connecting them, since the antenna pattern response depends on direction and splits this ring into two disconnected high probability arcs; a third non-collinear detector intersects this ring with a second independent one, reducing the arcs to a small bimodal pair of regions; and a fourth non-coplanar detector resolves the remaining reflection degeneracy to a single sky area [126].

With only two detectors, a single GW image therefore produces a sky region of typically hundreds of square degrees, containing millions of potential host galaxies and tens of thousands of galaxies within the 90% credible volume [147], making electromagnetic follow-up or lens identification essentially impossible in the absence of an independent counterpart. The analyses presented in this work use only H1 and L1, reflecting the most conservative scenario for strong lensing searches, since it corresponds to the lower bound on the localization improvement achievable once additional detectors are included in the network.

Virgo, currently reaching 50-60 Mpc in O4, is not expected to improve further during IR1 (the intermediate observing run between O4 and O5), remaining at 50-60 Mpc, before reaching a projected 90-130 Mpc by O5, while KAGRA is expected to improve from a few Mpc in O4 to 25-90 Mpc by O5 [170]. Extending the joint analysis presented in this work to include these additional detectors, once their sensitivity and duty cycle make them reliably available, is a natural direction for future improvements in sky localization beyond the H1+L1 network considered here.

3.4.2 Joint parameter estimation of lensed pairs

When a GW source is strongly lensed, the multiple images arrive at the detector network at different times. Because the Earth rotates during the delay between arrivals (typically minutes to months for galaxy-scale lenses, see §3.2.4), each image is observed with a different detector orientation, introducing varying antenna pattern functions and effectively enlarging the triangulation baseline [2]. Each lensed image therefore contributes independent constraints on the source position through different noise realizations and detector responses, and their combination reduces the overall sky localization uncertainty compared to a single image [147].

This is exploited through a coherent joint PE of the two images, testing between two competing hypotheses: the lensed hypothesis $\mathcal{H}_L$, under which the two signals are lensed images of the same source, and the unlensed hypothesis $\mathcal{H}_U$, under which they originate from two independent sources. Under $\mathcal{H}_L$, both data streams d_1 and d_2 share the same intrinsic binary parameters $\boldsymbol{\theta}$, as established in (3.26), so the joint likelihood is a product of the individual likelihoods evaluated at the same $\boldsymbol{\theta}$, with separate lensing parameters $\boldsymbol{\phi}_j = \{\mu_j, t_j, n_j\}$ for each image:

$$p(d_1, d_2 \mid \mathcal{H}_L) \equiv \mathcal{Z}_L = \int p(d_1 \mid \boldsymbol{\theta}, \boldsymbol{\phi}_1)\, p(d_2 \mid \boldsymbol{\theta}, \boldsymbol{\phi}_2)\, p(\boldsymbol{\theta}, \boldsymbol{\phi}_1, \boldsymbol{\phi}_2)\, d\boldsymbol{\theta}\, d\boldsymbol{\phi}_1\, d\boldsymbol{\phi}_2. \tag{3.30}$$

In this work, this joint analysis is carried out with HANABI [169], a Python package that extends BILBY to the coherent joint PE of strongly lensed image pairs, implementing the shared-parameter likelihood of Eq. (3.30) together with the lensing prior described above.

The shared parameters include the masses, spins, and crucially the sky location (α, δ), while the lensing parameters per image are the image-type label IT_j, sampled independently for each image from a discrete uniform prior over $\{1, 2, 3\}$ (the integer labelling convention used by HANABI for Type I, II, and III respectively, distinct from the Morse index $n_j \in \{0, 1/2, 1\}$ of Eq. (3.25)), and the relative magnification μ_{rel}, which in the implementation used in this work is not sampled independently but is instead already captured by the two independently fitted luminosity distances (§5.3). Under $\mathcal{H}_U$, the two signals are independent and the joint evidence reduces to the product of the individual evidences,

$$p(d_1, d_2 \mid \mathcal{H}_U) = \mathcal{Z}_1 \cdot \mathcal{Z}_2, \tag{3.31}$$

obtained from the single-image analyses. The improvement in sky localization is quantified through the 90% credible area of the posterior on (α, δ), defined as the minimum sky region $\mathcal{R}_{90}$ satisfying

$$\int_{\mathcal{R}_{90}} p(\alpha, \delta \mid d)\, d\Omega = 0.90, \tag{3.32}$$

where $d\Omega = \cos\delta\, d\alpha\, d\delta$ is the solid angle element and $p(\alpha, \delta \mid d)$ is obtained by marginalizing the full PE posterior over all remaining parameters.

3.4.3 Identifying lensed events: the lensing Bayes factor

The Bayesian hypothesis testing framework introduced in §2.4.3.2 applies directly to the identification of lensed GW events, comparing the lensed and unlensed hypotheses $\mathcal{H}_L$ and $\mathcal{H}_U$ introduced in §3.4.2. The coherence ratio

$$\mathcal{C}^{\mathrm{L}}_{\mathrm{U}} \equiv \frac{\mathcal{Z}_L}{\mathcal{Z}_1 \cdot \mathcal{Z}_2}, \tag{3.33}$$

quantifies how much more probable the data are under $\mathcal{H}_L$ than under $\mathcal{H}_U$, without accounting for selection effects or population priors [139]. In practice, since each PE run outputs a signal-vs-noise Bayes factor $\mathcal{B}_{SN}$ and the noise evidences cancel in the ratio, it can be computed as

$$\ln \mathcal{C}_{\mathrm{U}}^{\mathrm{L}} = \ln \mathcal{B}_{SN}^{\mathrm{joint}} - \left(\ln \mathcal{B}_{SN}^{(1)} + \ln \mathcal{B}_{SN}^{(2)} \right). \tag{3.34}$$

A more complete statistic is the full lensing Bayes factor $\mathcal{B}_{\mathrm{U}}^{\mathrm{L}}$, which additionally incorporates selection effects and astrophysical population priors [169]. Computing $\mathcal{B}_{\mathrm{U}}^{\mathrm{L}}$ for the injected population is beyond the scope of this work, whose primary focus is the improvement in sky localization achieved by the joint analysis; the simpler coherence ratio $\mathcal{C}_{\mathrm{U}}^{\mathrm{L}}$, which requires no additional ingredients beyond the Bayes factors already computed for each run, is instead evaluated for the full sample as a validation check in §5.5.5.

Part II

Original results

With the theoretical framework for gravitational lensing of GWs now established, including the geometric optics formalism, the signatures of strong lensing on GW signals, and the joint PE approach that exploits lensed image pairs to improve sky localization, we are in a position to present the original results of this thesis. Chapter 4 describes the simulation setup and the population of strongly lensed BBH events generated with LeR. Chapter 5 presents the PE pipeline. Each image is first analyzed individually with Bilby, and the two images of each event are then analyzed jointly with Hanabi, which combines them through a coherent likelihood that shares their parameters (§5.3), on top of the Bilby inference backend. The chapter also describes the automation infrastructure deployed on the Picasso HPC cluster, and validates the recovered parameters against the injected values, including a correction for the magnification-distance degeneracy and a consistency check of the lensed hypothesis via the coherence ratio. Finally, Chapter 6 presents the sky localization results, comparing the 90% credible sky area obtained from single-image and joint analyses across the final set of analyzed events.

Chapter 4

Simulation setup and implementation

This chapter describes the construction of the lensed BBH population used as input to the PE pipeline. The simulation proceeds in two stages: first, a large Monte Carlo (MC) population of strongly lensed BBH events is generated with LeR v0.4.2 [171], and then a realistic detectability filter is applied to select the subset of events that would be observed by the H1 and L1 detectors under O5 sensitivity conditions.

4.1 Lensed BBH population with LeR

LeR (LVK Event Rate calculator) is a Python framework for the joint simulation and rate computation of unlensed and strongly lensed GW populations [171]. It integrates source and lens population sampling, lens equation solving via Lenstronomy [172], and detector selection into a single reproducible pipeline. All lensing calculations are performed in the geometric optics regime and under the thin lens approximation, both of which are valid for galaxy-scale lenses in the LIGO frequency band (see §3.1).

The simulation is initialized with the following configuration, specifying the detector network, noise curves, and waveform approximant:

```
ler = LeR(
    npool=6,
    event_type="BBH",
    ifos=['L1', 'H1'],
    psds={'L1': 'T2500310-v2_O5a_strain_psd.txt',
          'H1': 'T2500310-v2_O5a_strain_psd.txt'},
    waveform_approximant="IMRPhenomXPHM",
    spin_zero=False,
    spin_precession=True,
)
```

Here, `npool` sets the number of parallel worker processes used for the MC sampling and SNR calculations. The `event_type="BBH"` argument selects the BBH source population model described in Table 4.1, as opposed to the BNS or NS-BH options also supported by LeR. The `ifos` argument specifies the detector network over which the network SNR ρ_{net} (§2.4.1) is computed, here restricted to the two LIGO detectors, H1 and L1. Throughout this thesis, every quoted SNR, whether computed by LeR for the injected population or recovered from the Bilby/Hanabi posteriors (§5.5), is the optimal SNR ρ_{opt} of Eq. (2.84), evaluated directly from a template and the detector PSD, rather than the noise-dependent matched-filter SNR of Eq. (2.82) that would be obtained by filtering real, noisy data.

The `psds` argument sets the noise curve assigned to each detector, here the O5a design sensitivity PSD [173]. This choice is motivated by the goal of simulating a realistic future observing scenario. The default PSD in LeR approximates O4 sensitivity, but no strongly lensed events were detected in O4. Simulating at O5a sensitivity instead, the first epoch of the O5 run with a BNS range of 225 Mpc [173], provides a more informative and forward-looking setting for this study. Figure 4.1 shows the projected O5a strain sensitivity curve used in this work, compared against the L1 O4 sensitivity achieved in September 2024 and the later O5b and O5c projections.

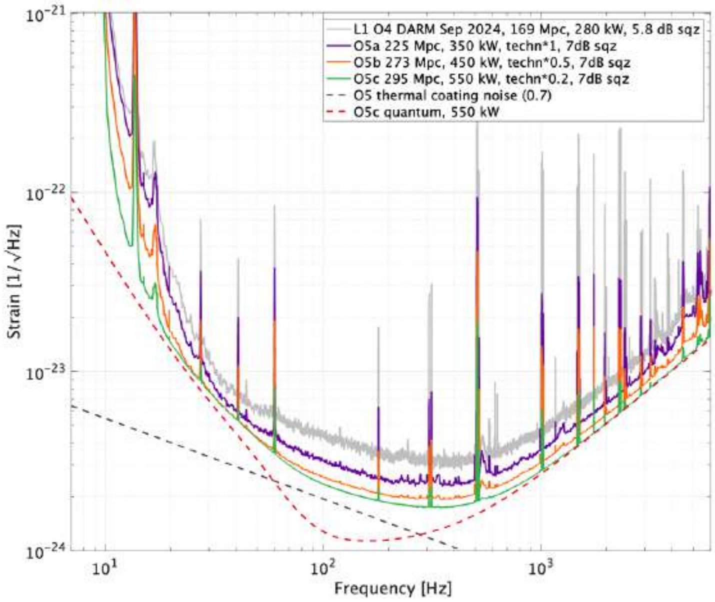

Figure 4.1: Projected strain sensitivity curves for the O5 observing run (O5a, O5b, and O5c), compared against the L1 O4 DARM sensitivity achieved in September 2024 (grey). The O5a curve (purple), adopted in this work, corresponds to a BNS range of 225 Mpc. Dashed lines show the O5 thermal coating noise and O5c quantum noise budgets. Source: [173].

The `waveform_approximant` argument selects `IMRPhenomXPHM` [120] for the SNR computation, which incorporates spin precession and higher-order multipole moments. The population presented here uses the original release of this model, with no additional `waveform-arguments-dict` options passed to LeR. A more recent variant [174] refines the treatment of precession and the final-spin prescription through additional arguments,

```
waveform-approximant = IMRPhenomXPHM
waveform-arguments-dict = {"PhenomXHMReleaseVersion": 122022,
                           "PhenomXPFinalSpinMod": 2,
                           "PhenomXPrecVersion": 320}
```

but this improved configuration was not used for the population generation in this work. Finally, `spin_zero=False` enables non-zero component spins in the source population, and `spin_precession=True` allows for spin tilts misaligned with the orbital angular momentum, so that the resulting waveforms include precession effects. Both are necessary for the spin parameters listed in Table 4.1 to be sampled rather than fixed to zero.

Parameter	Distribution / value
Cosmology	
H_0, Ω_m, Ω_Λ	$70\,\mathrm{km\,s^{-1}\,Mpc^{-1}}$, 0.3, 0.7 (flat ΛCDM) [171]
BBH source parameters	
Component masses m_1, m_2	Power Law + Gaussian (LIGO O3) [175]
Spin magnitudes a_1, a_2	Uniform [0, 1]
Spin tilt angles θ_1, θ_2	Sine prior on $[0, \pi]$
Azimuthal spin angles ϕ_{12}, ϕ_{JL}	Uniform $[0, 2\pi]$
Inclination θ_{JN}	Sine prior on $[0, \pi]$
Polarization ψ	Uniform $[0, \pi]$
Coalescence phase ϕ_c	Uniform $[0, 2\pi]$
Sky location (α, δ)	Isotropic
Source redshift z_s	Merger rate density [176]
Lens galaxy parameters	
Optical depth $\tau(z_s)$	$(D_c[\mathrm{Gpc}]/62.2)^3$ (SIS analytic) [177]
Velocity dispersion σ_v	$161 \cdot \mathrm{GenGamma}(0.869, 2.67)\ \mathrm{km\,s^{-1}}$ [178]
Lens redshift z_L	Conditioned on z_s (SDSS catalogue)
Density slope γ	$\mathcal{N}(2.0, 0.2^2)$ [179, 180]
Axis ratio q	$\mathrm{Rayleigh}(s(\sigma_v))$, $q \in [0.2, 1]$ [181]
External shear γ_1, γ_2	$\mathcal{N}(0, 0.05^2)$ [171]
Lens model	EPL + external shear [182]
Detectability	
Network SNR threshold	$\rho_{\mathrm{net}} \geq 10$
Minimum detectable images	2

Table 4.1: Source and lens population parameters used by LeR v0.4.2 [171] to generate the strongly lensed BBH population. All models correspond to its default configuration.

The population is generated by drawing events in batches of 20 000 until 500 strongly lensed detectable events are collected, each requiring that at least two images independently satisfy $\rho_{\mathrm{net}} \geq 10$:

```
lensed_params = ler.selecting_n_lensed_detectable_events(
    size=500,
    batch_size=20000,
    snr_threshold=10,
    num_img=2,
    resume=False,
    output_jsonfile="n_lensed_params_bbh.txt",
)
```

The SNR threshold of $\rho_{\mathrm{net}} \geq 10$ and the requirement of at least two detectable images follow standard detection criteria adopted in GW searches, so that the resulting population reflects events that a real detector network could plausibly identify and confidently claim as lensed, rather than a theoretical population extending down to marginal or sub-threshold significance.

A total of approximately 152 000 candidate events were drawn to collect the 500 detectable ones, yielding a total lensed event rate of $\mathcal{R}_L \approx 0.75\,\mathrm{yr}^{-1}$ for this detector configuration and sensitivity. This value is consistent with early forecasts for second-generation detector networks at design-like sensitivity, which span roughly 0.1-10 yr^{-1} depending on the assumed high-redshift merger rate and detector network [181], and is somewhat lower than forecasts for the full three- or four-detector network, reflecting the more limited H1+L1 configuration adopted here.

4.2 Injected parameter distributions

The 500 events produced by LeR represent the strongly lensed population detectable under ideal observing conditions. To obtain a more realistic injection set, a duty cycle filter is applied: for each image of each event, both H1 and L1 are independently activated with probability $p_{\mathrm{det}} = 0.70$, reflecting typical O4 single-detector duty cycles [70].

An image is retained only if both detectors are simultaneously active and $\rho_{\mathrm{net}} \geq 10$. Events with fewer than two surviving images are discarded; when more than two images survive, only the two with the highest network SNR are kept. After applying this filter once to the fixed 500-event population, 157 events are selected as the final injection set analyzed with Bilby. Table 4.2 summarizes the successive stages of this selection process, from the initial MC draws performed by LeR to this final injection set.

Selection stage	Number of events	Retained (%)
Candidate events drawn by LeR	$\sim 152\,000$	—
Detectable lensed events (≥ 2 images with $\rho_{\mathrm{net}} \geq 10$)	500	100
Duty cycle filter ($p_{\mathrm{det}} = 0.70$ per detector)	309	61.8
Duty cycle + SNR cut (≥ 2 images with $\rho_{\mathrm{net}} \geq 10$)	157	31.4

Table 4.2: Number of events surviving each stage of the selection process described in the text above, from the initial LeR population to the final 157-event injection set. All percentages are relative to the 500 detectable events.

Figures 4.2–4.5 show the distributions of the key parameters for the full detectable population ($N = 500$, dashed) and the final injection set ($N = 157$, solid). Figure 4.2 compares the optimal SNR distributions for L1, H1, and the network, $\rho_{\mathrm{opt}}^{\mathrm{L1}}$, $\rho_{\mathrm{opt}}^{\mathrm{H1}}$, and $\rho_{\mathrm{opt}}^{\mathrm{net}}$ (Eqs. (2.84)-(2.85)), abbreviated as ρ_{net} throughout the rest of this chapter. The network SNR distribution for the filtered population is truncated at $\rho_{\mathrm{net}} = 10$ by construction, and its peak shifts to higher SNR relative to the full population, reflecting the removal of images that fall below the threshold once only the active detectors are considered; the individual L1 and H1 distributions can extend below $\rho = 10$ even after filtering, since the selection criterion applies to the network SNR rather than to each detector separately.

Figure 4.3 shows the primary source-frame mass m_1^{src} (left) and the number of lensed images per event (right). The mass distribution reflects the underlying LIGO O3 Power Law + Gaussian population model adopted by LeR; the duty cycle and SNR selection preferentially retains events

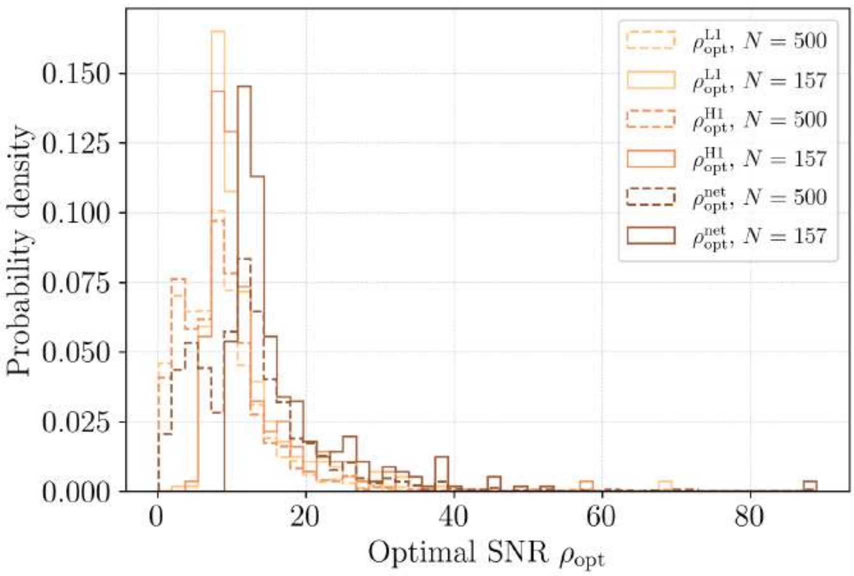

Figure 4.2: Optimal SNR for L1 ($\rho_{\rm opt}^{\rm L1}$, light orange), H1 ($\rho_{\rm opt}^{\rm H1}$, orange), and the network ($\rho_{\rm opt}^{\rm net}$, brown), for the full detectable population ($N = 500$, dashed) and the final injection set ($N = 157$, solid).

with $m_1^{\rm src}$ in the 25-35 $M_\odot$ range, consistent with higher-mass systems producing intrinsically louder signals and thus being more robust to the loss of detector uptime, while the high-mass tail above $\sim 40\,M_\odot$ is largely absent from the filtered population. While the underlying population is dominated by four-image (quad) systems ($\sim 75\%$), the joint requirement of both detectors active and $\rho_{\rm net} \geq 10$ typically leaves only two images per event above the threshold, and almost never leaves all four. This reflects the combined effect of the two filters: the fainter pair rarely reaches the SNR threshold even with both detectors active, and the duty cycle filter can remove any of the four images, including the brighter pair, whenever a detector happens to be inactive.

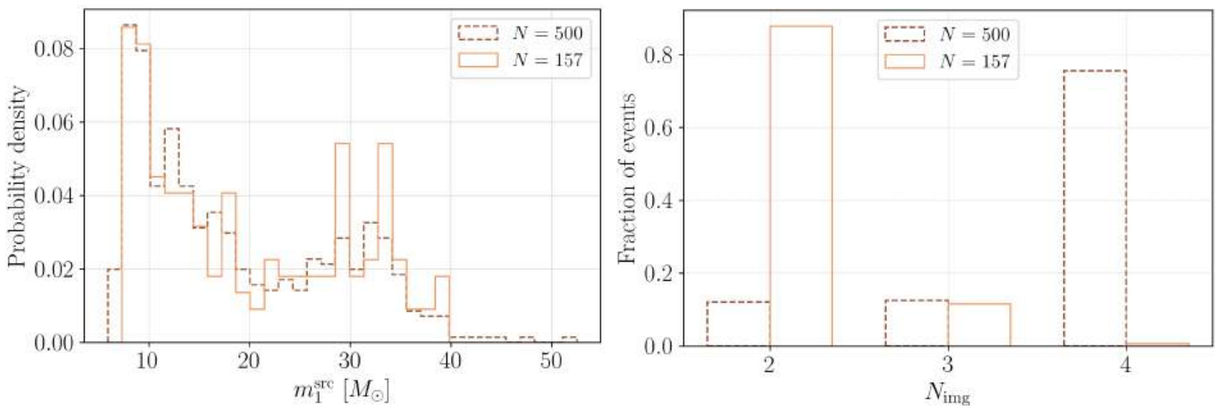

Figure 4.3: *Left*: primary source-frame mass $m_1^{\rm src}$. *Right*: number of lensed images per event, comparing the true image multiplicity ($N = 500$, dashed) against the number surviving the duty cycle and SNR filters before the top-2 cut ($N = 157$, solid).

Figure 4.4 compares the true luminosity distance D_L against the effective luminosity distance $D_L^{\rm eff}$. Because magnification boosts the observed signal amplitude, $D_L^{\rm eff} = D_L/\sqrt{|\mu|}$ is systematically smaller than the true distance, shifting the effective-distance distribution toward lower values relative to D_L.

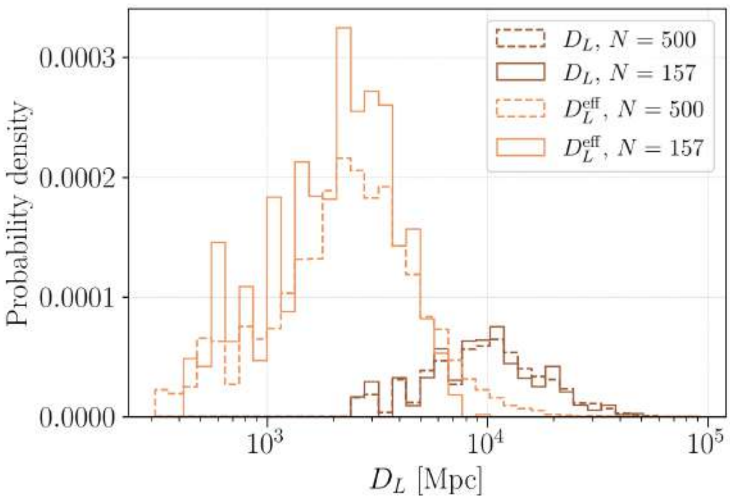

Figure 4.4: True luminosity distance D_L (brown) versus effective luminosity distance $D_L^{\text{eff}} = D_L/\sqrt{|\mu|}$ (orange), for the full detectable population ($N = 500$, dashed) and the final injection set ($N = 157$, solid).

Figure 4.5 shows the corresponding image magnification μ. On a linear scale (left), the distribution is strongly peaked near $\mu \sim 0$ with extended tails; negative values correspond to Type II images (saddle points of the Fermat potential), which have mirror-inverted parity and acquire a $\pi/2$ Morse phase shift (see §3.3). Its horizontal axis is clipped to the 0.5-99.5 percentile range of the combined data to make the shape of the peak and tails visible, excluding a small number of extreme outliers reaching $|\mu| \sim 10^3$. The right panel shows the same distribution for $|\mu|$ on a logarithmic axis, which accommodates this full dynamic range without clipping and shows the tail out to $|\mu| \sim 10^3$ directly. In this panel, the injection set (solid) is visibly truncated below $|\mu| \sim 1$ relative to the full population (dashed), which still extends into $|\mu| < 1$: the SNR threshold underlying the top-2 selection preferentially retains magnified images, since magnification boosts the observed amplitude at fixed distance. This selection effect is picked up again quantitatively in §5.5, where every image among the 150 analyzed events has $|\mu_{\text{true}}| > 1$.

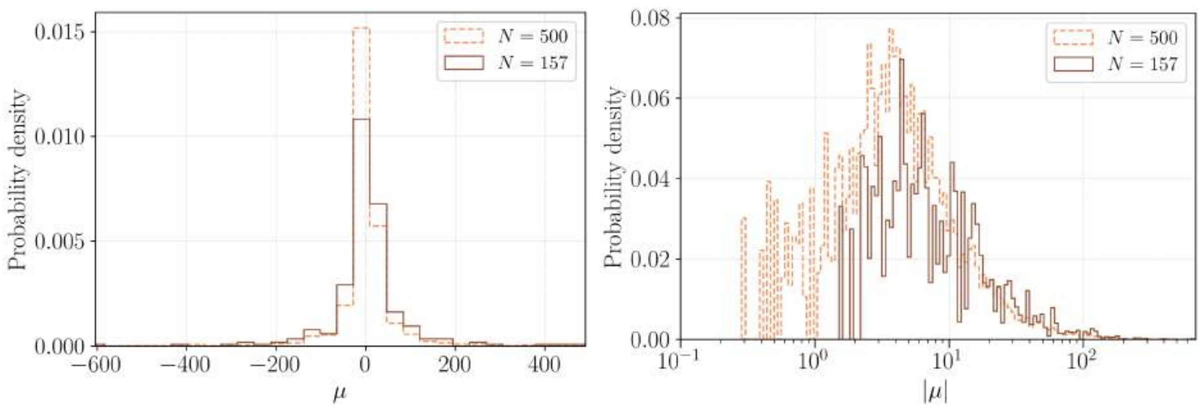

Figure 4.5: *Left*: signed magnification μ, clipped to the 0.5-99.5 percentile range. *Right*: absolute magnification $|\mu|$ on a logarithmic axis, showing the full tail without clipping.

Figure 4.6 shows the joint m_1^{src}-m_2^{src} distribution for the final injection set of 157 events, color-coded by the injected network SNR $\rho_{\text{net}} = \sqrt{\rho_{\text{net},1}^2 + \rho_{\text{net},2}^2}$, the quadrature sum of the per-image network SNRs (§5.5); this pair-combined quantity is denoted ρ_{net} here and in Figure 5.2 below, matching the axis labels of both figures, and should not be confused with the per-image $\rho_{\rm net}$ of Eq. (2.85) used elsewhere in this chapter. The population is concentrated in the region $10 \lesssim m_1^{\rm src} \lesssim 40\,M_\odot$ with most events lying close to the equal-mass line $q = 1$, reflecting the shape of the underlying LIGO O3 Power Law + Gaussian mass model (§4.1). The event density falls sharply for $m_2^{\rm src} < 10\,M_\odot$ and for $q \lesssim 0.5$, consistent with the SNR-driven detection bias: unequal-mass and low-mass binaries produce weaker signals and are less likely to survive the $\rho \geq 10$ threshold applied in both images. The highest-SNR events are all found near $q \simeq 1$ and around $M \sim 40$-$70\,M_\odot$, the sweet spot of the O5a network sensitivity for stellar-mass BBH mergers.

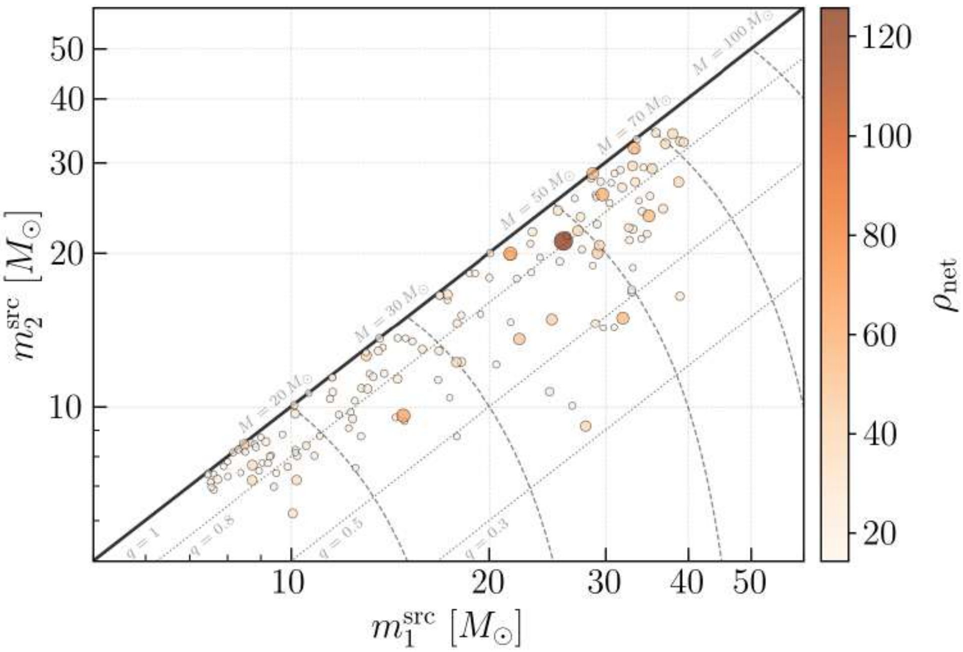

Figure 4.6: Source-frame component masses $m_1^{\rm src}$ and $m_2^{\rm src}$ for the 157 injected lensed BBH events, colored by the injected network SNR $\rho_{\rm net}$ (pair-combined, see text). Dashed lines mark constant total mass M and dotted lines mark constant mass ratio q. The distribution clusters near $q = 1$ and $M \sim 40$-$70\,M_\odot$, where the network sensitivity is highest.

The time delay Δt between lensed images is the second key ingredient of the joint analysis, since it sets the rotation angle of the Earth (and hence the change in the detector antenna pattern) between successive image arrivals. Figure 4.7 shows the distribution of Δt_{1j} ($j = 2, 3, 4$) with respect to the first-arriving image, obtained from the full LeR detectable population before applying the duty cycle and SNR cuts. The Δt_{12} and Δt_{13} distributions peak between days and weeks and span from minutes to years, consistent with the characteristic delays produced by galaxy-scale lenses in the geometric optics regime. The Δt_{14} distribution, which corresponds to quadruply lensed events only, is shifted toward longer values because the last-arriving image typically emerges from a saddle point of the Fermat potential further from the central caustic[1]. Delays below roughly one hour are rare, which works in favor of the joint analysis: the Earth usually has time to rotate enough between the two images to noticeably reorient the H1+L1 baseline relative to the source, giving the extra angular information exploited in §3.4.2.

[1] A caustic is a curve in the source plane where the lens magnification formally diverges; sources close to it are strongly magnified and produce closely spaced image pairs, while the last-arriving image of a quadruple system typically corresponds to a source located further from it, and hence more weakly magnified and more delayed.

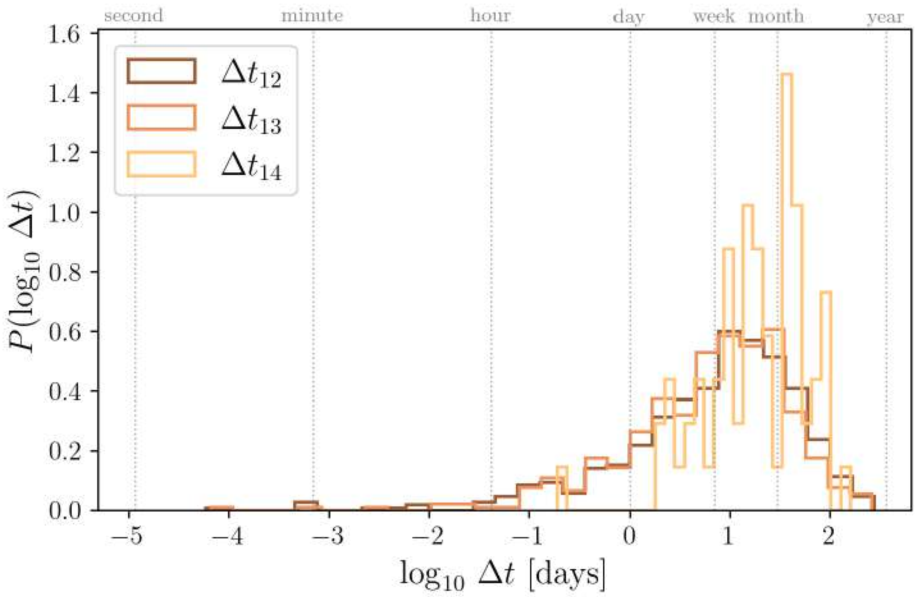

Figure 4.7: Distribution of the time delays Δt_{1j} (with $j = 2, 3, 4$) between the first-arriving image and each subsequent image, on a logarithmic scale, for the full LeR detectable population. The top axis marks characteristic timescales from seconds to years. The detector antenna pattern is periodic with the Earth's rotation, so it is only delays much shorter than this period, of order minutes, that are guaranteed to leave it essentially unchanged; the bulk of the delays lie well above this scale, in the hours-to-months range where the antenna pattern has generically reoriented enough to enable the sky localization improvement exploited by the joint analysis (§3.4.2).

Chapter 5

Parameter estimation pipeline and results

5.1 Pipeline overview

This chapter presents the PE pipeline built on the injection set described in Chapter 4, together with the automation infrastructure used to run it for all 157 events. Figure 5.1 places this pipeline within the full analysis developed in this thesis, from the population generation of Chapter 4 (Stage 1) to the sky localization results of Chapter 6 (Stage 4). The rest of this section focuses on Stages 2 and 3: the automated generation of BILBY/HANABI `.ini` files, and the single-image plus joint PE itself.

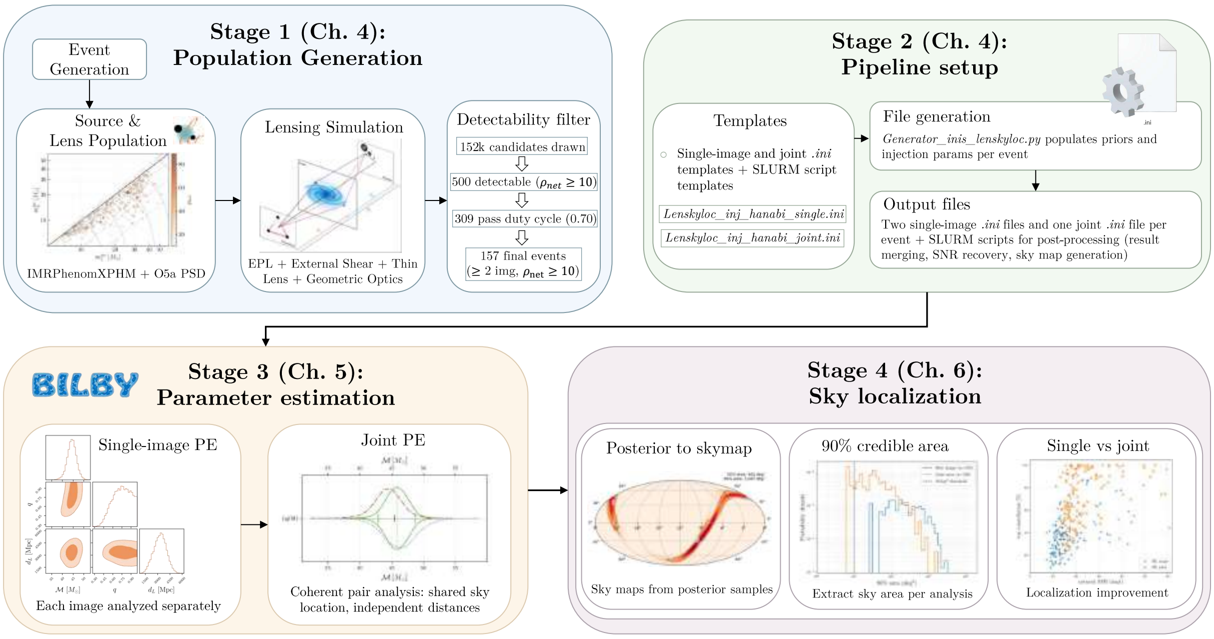

Figure 5.1: Overview of the full analysis pipeline developed in this thesis. Stage 1 (§4.1-§4.2): the lensed BBH population is generated with LER and reduced from 500 detectable events to the final injection set of 157 through the duty cycle and SNR selection process (Table 4.2). Stage 2 (§5.4): `Generator_inis_lenskyloc.py` converts each event's injected parameters into single-image and joint `.ini` configurations and the corresponding post-processing scripts. Stage 3 (§5.2-§5.3): each image is analyzed independently with BILBY, then jointly with HANABI using the shared-parameter likelihood described in §5.3. Stage 4 (§6.1-§6.2): sky maps are generated from the posteriors and the recovered 90% credible areas are compared between single-image and joint analyses. A higher-resolution version of this diagram is available at `https://github.com/XxavierGarcia/TFM_reproducibility_material/blob/main/diagrama_pipeline_reduced.pdf`.

Starting from the filtered population `n_lensed_params_bbh_filtered.txt` produced in §4.2 (Stage 1), a single script, `Generator_inis_lenskyloc.py`, generates for every event the two single-image `.ini` configurations, the joint `.ini` configuration, and the SLURM[1] (Simple Linux Utility for Resource Management) scripts needed for the steps described below, populating priors and injection parameters directly from the corresponding LeR output (Stage 2, §5.4).

Each single-image `.ini` file is analyzed independently with `bilby_pipe` [183], using the lensed waveform source model provided by Hanabi (§5.2). This yields, for every event, two independent posterior distributions, one per image. Once both single-image jobs have completed, the joint `.ini` file is submitted with `hanabi_joint_pipe` [169], which jointly analyzes the image pair through the shared-parameter likelihood described in §5.3, on top of the same Bilby inference framework used for the single-image runs.

From this point the pipeline splits into two independent post-processing tracks. In the first, since `hanabi_joint_pipe` does not compute the recovered SNR on its own, the joint posterior, together with the single-image data realizations, is passed to `hanabi_postprocess_result` (included in the Hanabi software package) with the `--generate-snrs` option, which recomputes the single-image and network SNRs consistently across the single and joint analyses. A dedicated script then extracts these recovered SNRs into a summary table for each event, later combined across all events for the SNR consistency checks of §5.5. In the second track, the single-image and joint posteriors are combined with PESummary [184] into a single metafile, from which sky maps are generated for each image and for the joint analysis using `ligo-skymap-from-samples`, `ligo-skymap-stats`, and `ligo-skymap-constellations` [185, 186], providing the 90% credible sky areas and the constellation localization compared between single-image and joint analyses in Chapter 6.

The entire pipeline, from `.ini` generation to job submission and the SNR post-processing track, is automated on the Picasso HPC cluster[2] (SLURM scheduler); the sky-map track is run separately, as described in §5.4.

5.2 Single-image analysis with Bilby

As introduced in §5.1, each image is analyzed with `bilby_pipe` as an isolated signal, characterized by its own data segment, noise realization, and set of extrinsic parameters. The `.ini` configuration is generated automatically for each image from a shared template (§5.4). Table 5.1 summarizes the detector and sampler settings common to all 157 events. The waveform generation frequency listed there, 13.33 Hz, is lower than the $f_{\min} = 20\,\mathrm{Hz}$ of the analysis band: this is because the waveform generation frequency refers to the dominant $(\ell, m) = (2, 2)$ mode, while a subdominant mode's instantaneous frequency scales as $f_{\ell m} = (m/2) f_{22}$, so starting the $(2, 2)$ mode at $f_{22} = 20/1.5 = 13.33\,\mathrm{Hz}$ ensures the $(3, 3)$ mode, the most significant subdominant harmonic in this band, already has full content by $f_{\min} = 20\,\mathrm{Hz}$ [187].

The three parallel runs (`n-parallel`) are independent full Dynesty analyses combined into a single posterior to mitigate sampler-to-sampler variance, while `npool` parallelizes the likelihood evaluation within each individual run and matches the number of CPUs requested per SLURM

[1] https://slurm.schedmd.com/overview.html

[2] https://www.scbi.uma.es/web/es/supercomputacion/#documentation

Setting	Value
Detector network	H1, L1
PSD	O5a strain PSD (same as LER, §4.1)
Sampling frequency	2048 Hz
Reference frequency	$f_{\mathrm{ref}} = 20\,\mathrm{Hz}$
Minimum and maximum frequency	$f_{\mathrm{min}} = 20\,\mathrm{Hz}$, $13.33\,\mathrm{Hz}$ (waveform) and $f_{\mathrm{max}} = 896\,\mathrm{Hz}$
Segment duration	Computed per event (see below)
Waveform approximant	`IMRPhenomXPNR` [188]
Sampler and sampling method	DYNESTY [131] and acceptance-walk
Live points and `naccept`	1000 and 60
Parallel runs per image	`n-parallel` = 3
CPUs requested per job (SLURM)	128

Table 5.1: Detector and sampler settings for the single-image analysis, common to all 157 events except the segment duration, which is computed individually per event (see below).

job.

The same O5a PSD used to generate the population with LER (§4.1) is adopted for the analysis, so that the noise assumptions are consistent between simulation and recovery. Unlike the data duration, which is fixed for the whole population in LER, the segment duration adopted for each image is computed individually from its injected component masses and spins, using LALSIMULATION [189] bounds on the merger and ringdown time together with the expected waveform length at $f_{\mathrm{min}} = 20\,\mathrm{Hz}$, rounded up to the next power of two. This keeps each analysis as short as the signal allows without truncating the inspiral, rather than using a single conservative duration for the whole population.

The injected and recovered waveforms both use the `IMRPhenomXPNR` approximant [188], rather than the `IMRPhenomXPHM` approximant used by LER to generate the population and compute the detection SNRs that define the selection process of §4.2. This choice is intentional: `IMRPhenomXPHM` is computationally cheap and sufficiently accurate for the population-level detectability calculation performed by LER, whereas `IMRPhenomXPNR` incorporates precession effects calibrated against NR simulations and is the more accurate of the two, better suited for PE, following standard practice in recent LVK analyses. Indeed, the most recent O4 lensing search still uses `IMRPhenomXPHM` [166], while the standard (unlensed) LVK PE pipeline has already moved to `IMRPhenomXPNR` in GWTC-5 [71]. For the parameter ranges considered here, the two models are expected to agree closely, so this choice between the population-selection and PE stages is not expected to introduce a significant waveform-systematics effect.

The waveform is generated through HANABI's lensed source model, `hanabi.lensing.waveform.strongly_lensed_BBH_waveform`[3], which implements the lensed strain expression of §3.3, combining the intrinsic `IMRPhenomXPNR` waveform with an image-dependent amplitude and Morse phase factor. In the single-image configuration this factor is fixed to `image_type=1` (Type I, a minimum of the Fermat potential) rather than treated as a free parameter. Type I images carry no Morse phase shift ($n = 0$, §3.3), so their waveform is identical in phase to an ordinary

[3] `https://github.com/ricokaloklo/hanabi/blob/master/hanabi/lensing/waveform.py`

unlensed signal; the only lensing imprint, an amplitude rescaling and a shifted arrival time, is already absorbed into the standard `luminosity_distance` and `geocent_time` parameters. The single-image prior fixes `image_type` to this value regardless of the true injected type, so each image is analyzed as an ordinary CBC signal; identifying the true image type of each image, together with the relative magnification between them, is instead the specific role of the joint analysis (§5.3), which samples `image_type` independently for each image from a discrete prior.

The remaining priors follow standard choices for BBH PE: a chirp mass prior adapted individually for each event to a ±30% range around its injected value (rather than a single wide range for the whole population, which would waste live points on regions of parameter space far from every event's true chirp mass), mass ratio and component-mass constraints, uniform spin magnitudes and precession angles, an isotropic sky location, and a 50-20 000 Mpc uniform luminosity distance prior; the full `prior-dict` is reproduced in Appendix §C.

This chirp mass prior can be centered on the injected value only because the true value is already known here; doing so also keeps each run computationally cheap by concentrating live points where the posterior actually has support, a non-negligible consideration given the total computational cost of the campaign (§5.4). A dedicated check across all 300 single-image chirp mass posteriors finds that only 4 (1.3%) show noticeably more density right next to one prior edge than near the center of the range, and in every case the affected posterior belongs to the lower-SNR image of its pair, with single-image SNR of 7.2 (`inj030`), 10.7 (`inj104`), 15.8 (`inj085`), and 10.3 (`inj151`), among the lowest in the sample; for the remaining 296 posteriors the credible interval is unaffected by the prior boundary. For a real candidate, where the true value is not known, a wide prior would be required instead.

The spin magnitude priors are bounded to $a_{1,2} \in [0, 0.99]$, whereas LeR draws the injected spins from Uniform$[0, 1]$ (Table 4.1); an injected value above 0.99 would therefore fall outside the recovery prior. Across all 157 injected events, the maximum injected spin magnitude is 0.78 (a_1) and 0.80 (a_2), so no event is affected by this truncation in practice.

5.3 Joint analysis with Hanabi

As introduced in §5.1, the joint analysis is performed with `hanabi_joint_pipe`, treating the two triggers of an event as two images of the same underlying source rather than as independent signals, with each trigger pointing to its corresponding single-image `.ini` file and otherwise inheriting its own data segment, noise realization, and PSD. The parameters shared between the two images are the intrinsic binary parameters (masses and spins), the sky location, the inclination, the polarization angle, and the coalescence phase, reflecting the fact that both images are gravitationally lensed copies of the same source. The luminosity distance, by contrast, is fitted independently for each image.

Formally, Hanabi's lensing model requires specifying a relative magnification, an image type, and a time delay for each image pair. Table 5.2 summarizes these lensing-specific priors: the relative magnification is fixed to 1 for both images (not sampled), since its physical content is already captured by the two independently fitted luminosity distances discussed above, while the image type is the only genuinely free lensing-specific parameter, sampled independently for each image from a discrete uniform prior over the three Morse-index classes introduced in §3.3. The time delay is not assigned a dedicated lensing prior either: it is simply the difference between the two images' arrival times, each sampled with BILBY's default ± 0.1 s uniform `geocent_time` prior around its own trigger time, the same as an ordinary, unlensed single-image analysis.

Lensing parameter	Prior
Relative magnification $\mu_{\text{rel}}^{(1)}$, $\mu_{\text{rel}}^{(2)}$	Fixed to 1 (not sampled)
Image type $\text{IT}^{(1)}$, $\text{IT}^{(2)}$	Discrete uniform over Type I/II/III
Time delay $\Delta t = t_c^{(2)} - t_c^{(1)}$	Same as single-image analysis: $t_c^{(i)} \sim \text{Unif}(t_{\text{trig}}^{(i)} \pm 0.1\,\text{s})$

Table 5.2: Lensing priors of the joint analysis (`lensing-prior-dict`); all other priors are shared between images as described in the text.

Sharing the sky location across two images observed at different times, typically separated by minutes to months (Figure 4.7), is the key mechanism behind the improvement in sky localization, as described in §3.4.2.

The joint likelihood is evaluated over both images simultaneously, using the same sampler settings as the single-image runs (Table 5.1), but with substantially higher resource requirements, reflecting the added cost of a coherent likelihood evaluated across two detectors and two images: 180 GB of memory and a 71.5 h wall-time limit, compared to 110 GB and 48 h for each single-image run.

5.4 Automation and infrastructure on Picasso

Running this pipeline for all 157 events relies on two components: a script that generates every configuration and job file automatically, and a two-stage submission mechanism that enforces the job dependencies of §5.1. Appendix §B complements this section with a practical, day-to-day guide to working on these two HPC clusters (SSH access, job monitoring, troubleshooting), aimed at future students rather than at reproducing this specific pipeline. Jobs execute directly inside a dedicated `conda`/`miniforge` environment (`hanabi_bilby260`[4]) on Picasso, whose exact specification is provided for reproducibility.

Summed over the sampling time reported in the logs of every parallel DYNESTY run (`n-parallel =3` per image or per joint analysis, each requesting 128 CPUs), the full PE campaign consumed approximately 830,000 CPU-hours (∼6,500 node-hours, or ∼270 node-days on 128-core nodes), split roughly evenly between the 942 single-image parallel runs and the 450 completed joint parallel runs. This scale of computational cost motivates several of the efficiency choices described below and in §5.2, such as the per-event chirp-mass prior and the segment duration computed individually from each event's parameters, rather than adopting a single conservative configuration for the whole population.

[4] `https://github.com/XxavierGarcia/TFM_reproducibility_material/blob/main/environment_full.yml`

5.4.1 Configuration and script generation

Beyond the duration and chirp-mass prior already computed for each event (§5.2), `Generator_inis_lenskyloc.py` populates the `injection-dict` of both single-image templates directly from the corresponding LeR parameters, mapping, e.g., `effective_luminosity_distance` (the magnification-rescaled apparent distance of Eq. (3.27)) to `luminosity_distance`, and `effective_geocent_time` (the lensed arrival time of Eq. (3.28)) to `geocent_time`, so that each image is injected with its own apparent extrinsic parameters rather than the source's true values; the same per-image mapping applies to `image_type`, so each image is injected with its own true Morse type from LeR rather than a fixed value. Appendix §C reproduces a representative single-image and joint `.ini` pair.

5.4.2 Submission and job dependencies

Job submission is split into two stages. The first, `submit_all_part1.sh`, submits the two single-image `bilby_pipe` jobs for each event, followed by the joint `hanabi_joint_pipe` job with an explicit SLURM dependency on both single-image job IDs. The second stage, `submit_all_part2.sh`, runs after the first stage completes and launches the SNR-recovery track described in §5.1. Figure C.1 (Appendix §C) lays out this Picasso (SLURM) branch of the pipeline in full, script by script.

The sky-map track, generating PESummary pages and sky maps for every event, is instead run on LIGO's computing cluster at Caltech (CIT)[5] using HTCondor [190], within the IGWN Conda Distribution[6] (the `igwn-py311` series distributed via CVMFS, a rolling release updated periodically; the post-processing was run with the snapshot current at execution time): `Pesummaries_automated_condor_all.py` and `Skymap_condor.py` are each launched as an HTCondor job array of 157 processes, generating the summary pages and sky maps for all events in parallel. The PE itself, both single-image and joint, was always run on Picasso and never on CIT; only this post-processing track uses HTCondor, since the version of PESummary available in the Picasso environment produced errors, while the `igwn` environment on CIT ran it correctly. Because this track only produces summary pages and sky maps from posteriors already sampled on Picasso, the exact snapshot of the rolling `igwn` distribution used does not affect the scientific results; the relevant environment is the Picasso environment `hanabi_bilby260` in which all sampling was performed. Figure C.2 (Appendix §C) lays out this CIT (HTCondor) branch in the same level of detail. The complete pipeline code, including both the Picasso and CIT branches, is publicly available at `https://github.com/XxavierGarcia/TFM_reproducibility_material`.

5.4.3 Incomplete events

The joint analysis is the most expensive step of the pipeline, requiring up to 71.5 h of wall time per event (§5.3). Partway through the production run, the project's total compute allocation on Picasso was exhausted. New jobs could still be submitted afterward, but the project lost its scheduling priority and the maximum wall-time allowed per job dropped to 24 h, well below

[5] `https://computing.docs.ligo.org/guide/computing-centres/cit/`
[6] `https://computing.docs.ligo.org/conda/`

the 71.5 h a joint analysis needs to complete in a single run. In practice this meant that each remaining joint job would have had to be split across multiple 24 h segments, resumed from its checkpoint file each time (Appendix §B), while also queuing far longer than before due to the reduced priority; finishing the remaining joint analyses this way was no longer practical within the time available for this thesis. The last seven joint analyses[7] were left interrupted mid-run as a result. The next allocation period was not due to begin for approximately two months, and rather than delay the analysis by that much to recover seven events, it was finalized with the 150 events that had completed successfully. This loss of 7 out of 157 events is a consequence of the compute allocation available at the time, not a scientific selection criterion, but it is not entirely random with respect to the injected parameters either: all seven incomplete events fall in the low end of the sample's chirp mass range (11.3-20.0 $M_{\odot}$, against a full range of 10.9-79.7 $M_{\odot}$ and a median of 44.8 $M_{\odot}$), while their network SNR, time delay, and magnification are unremarkable compared to the rest of the sample. This is consistent with the origin of the cutoff: the data segment duration is computed individually per event from its masses (§5.2), and lower-mass systems require longer segments, so their joint analyses take longer and were more likely to still be running when the allocation was exhausted. So the final sample of 150 events has a small gap at low chirp mass compared to the full 157, caused by the analysis taking longer for these events, not by anything astrophysical.

5.5 Parameter recovery and validation

This section evaluates how well the pipeline recovers the injected parameters of the 150 successfully analyzed events: first through a qualitative view of the recovered joint mass posteriors across the sample's SNR range, then through a direct comparison of the full joint and single-image posteriors for a few representative events (shown for the complete sample in Appendix §D), followed by systematic residual checks against the injected values for every recovered parameter, and finally through a quantitative look at the mass bias introduced by gravitational lensing, together with an attempt to correct it.

5.5.1 Masses

Figure 5.2 shows the recovered (m_1, m_2) 90% credible contours for six representative events spanning the pair-combined network SNR range of the sample ($\rho_{\rm net}$ from 15.0 to 65.3, see Fig. 4.6 above), overlaid on the contours of the remaining events in grey. The two examples with the lowest SNR, `inj084` and `inj052`, show visibly elongated, curved contours: the leading-order phase evolution constrains mainly the chirp mass $\mathcal{M}$, so individual masses remain only weakly determined along lines of near-constant $\mathcal{M}$, and this degeneracy widens as the SNR decreases. The shape of the contour, however, is not set by SNR alone: it also depends on how much of the signal is dominated by the inspiral. For lower mass systems, many PN cycles are observed in-band, and a high SNR is then enough to also break the mass-ratio degeneracy, recovering both individual masses precisely, as for `inj117` ($\rho_{\rm net} = 65.3$). Higher-mass systems, by contrast, spend comparatively little time in the inspiral before merging within the band; there, a high SNR still pins down the total mass tightly through the merger and ringdown, but does not by

[7] Events `inj110`, `inj115`, `inj118`, `inj137`, `inj138`, `inj142`, and `inj150`.

itself supply the extra inspiral cycles needed to separate m_1 from m_2, so the mass ratio can remain comparatively broad even when the SNR is high. Masses are shown here in the detector frame $(1+z)m_i$, rather than the source frame reported in LVK catalogs [71]. In this analysis, each image's luminosity distance is sampled independently rather than sharing a common source distance across the pair, so what the posterior actually constrains for each image is the effective distance, $D_L^{\rm eff} = D_L^{\rm true}/\sqrt{|\mu|}$, consistent with how it was injected (§5.4). With only the two images' apparent distances, the true source distance $D_L^{\rm true}$ and the magnification μ of each image cannot be separated: two observed quantities constrain three unknowns, and additional information (an electromagnetic counterpart or an independent lens model) would be needed to break this degeneracy [191]. Converting the apparent distance directly to a redshift, and hence to a source-frame mass, would therefore absorb the unknown magnification with the true distance and bias the result; masses are reported in the detector frame instead to avoid this.

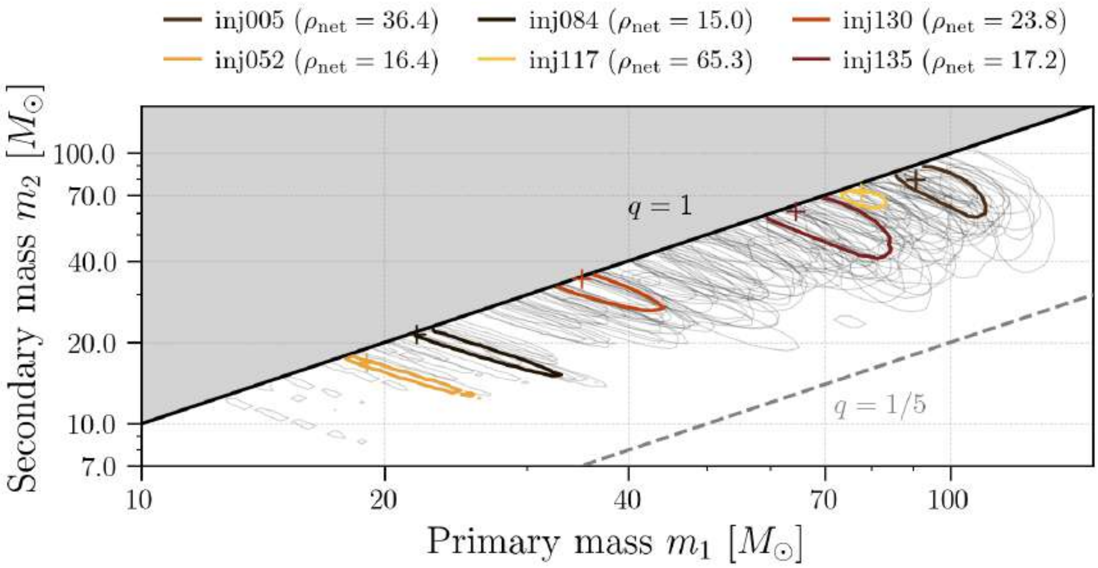

Figure 5.2: Recovered 90% credible (m_1, m_2) contours for six events spanning the sample's network SNR range (colored), overlaid on the contours of all other events (grey). Crosses mark the injected (m_1, m_2) value for each highlighted event. The region above $q = 1$ is excluded by the $m_1 \geq m_2$ convention; the region below $q = 1/5$ is a visual reference, not the mass-ratio prior actually used (§5.2, $q_{\rm min} = 0.05$).

Before turning to the population-level residuals, it is instructive to look at the full recovered posteriors of individual events, and in particular at what the joint analysis adds over analyzing each image separately. Because the intrinsic binary parameters are shared between the two images in the joint analysis (§5.3), the joint posterior draws on both data streams at once and can be narrower than either single-image posterior. Figure 5.3 illustrates this for three representative events, showing the joint posterior (green) together with the two single-image posteriors (image 1 in red, image 2 in blue) and the injected value (black tick), for the mass ratio q, the effective spin $\chi_{\rm eff}$, and the effective luminosity distance $D_L^{\rm eff}$. For `inj152` the joint 90% credible interval on q is a factor of 1.7 narrower than the better of the two single-image intervals, and for `inj113` the joint analysis tightens $\chi_{\rm eff}$ by a factor of 2.2 and $D_L^{\rm eff}$ by a factor of 3.5, the largest single improvement in the whole sample. The gain is not universal: for the high-SNR event `inj117`, already well measured from a single image, the joint and single-image posteriors nearly coincide across all three parameters. The improvement is therefore real but strongly event-dependent, largest where a single image constrains a parameter only weakly; the residual checks below quantify the recovery for the population as a whole. The equivalent figures for all 150 events and every recovered parameter, together with a quantitative ranking of these event-by-event improvements, are collected in Appendix §D, where a summary table lists the events with the largest gains alongside a few with little or no improvement.

That same ranking also includes a handful of events with $\mathrm{IF}_\theta < 1$ for the chirp mass $\mathcal{M}$ or the mass ratio q, meaning the joint posterior is broader than the better single-image posterior rather than narrower. This is not a problem with the joint analysis: it happens when a single high-SNR image happens to produce an anomalously narrow posterior for that parameter, and the joint analysis, required to stay consistent with both images at once, correctly returns a wider, more honest interval instead of inheriting that overconfidence. The clearest case is `inj085` ($\mathrm{IF}_q = 0.80$): image 2's mass ratio posterior alone is narrower than the joint one, but the joint interval still recovers the injected value at least as reliably, so the lower IF_θ here reflects a correction of an overconfident single-image result rather than a genuine loss of information.

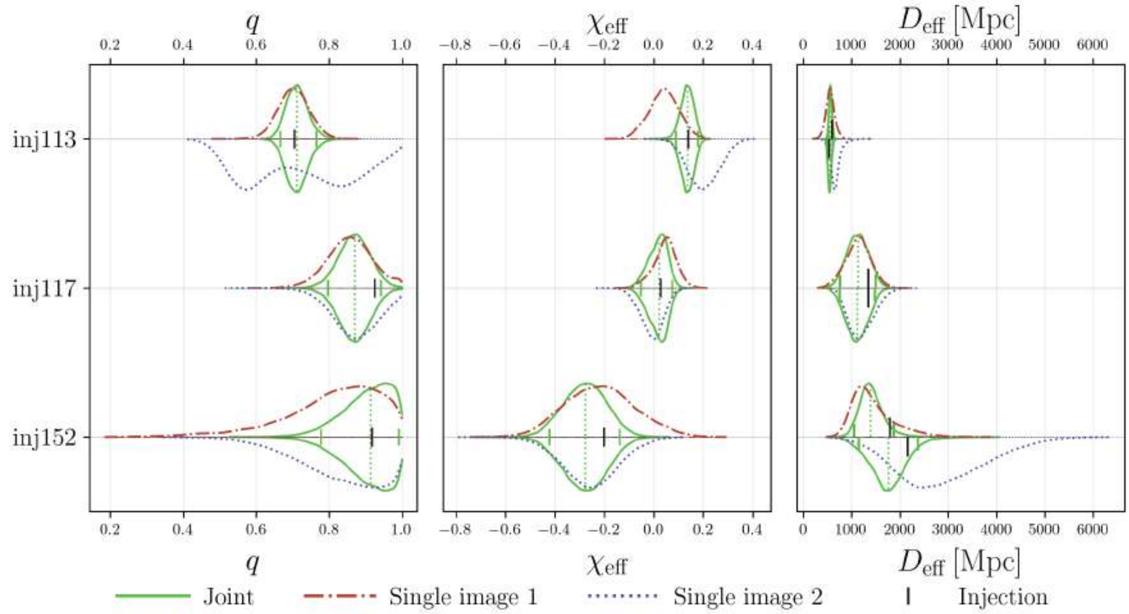

Figure 5.3: Recovered posteriors of the mass ratio q, effective spin χ_{eff}, and effective luminosity distance D_L^{eff} for three representative events. Joint posterior in green (solid), single-image posteriors in red (dash-dot, image 1) and blue (dotted, image 2); the black tick marks the injected value, and the green dotted vertical lines mark the 90% credible interval of the joint posterior. For D_L^{eff} each row is shown as a split violin, with the upper half corresponding to image 1 and the lower half to image 2, since this parameter is fitted independently for each image. The joint analysis tightens q for `inj152` and χ_{eff} and D_L^{eff} for `inj113`, while for the high-SNR event `inj117` it adds little beyond the single-image result.

For the remaining parameters, each event's joint posterior is summarized by its median ($\tilde{y}$) and 90% credible interval (5th to 95th percentile), compared against the corresponding injected value from LeR. Figure 5.4 shows, across all 150 events, the residual between the posterior median and the injected value as a function of the injected value, with error bars spanning the 90% credible interval, for the chirp mass $\mathcal{M}$ (left) and mass ratio q (right). The chirp mass residuals are centered on zero across the full mass range ($\tilde{y} = -0.002\, M_\odot$), consistent with $\mathcal{M}$ being the best determined parameter in the waveform phase. The mass ratio shows substantially more scatter ($\tilde{y} = -0.078$, $\sigma = 0.114$), increasing with $q_{\rm inj}$ and biased toward lower recovered values for injected mass ratios near unity. This is a systematic trend, not random noise: because $\mathcal{M}$ is tightly measured but q is not, different (m_1, m_2) combinations that keep $\mathcal{M}$ fixed remain compatible with the data, so the recovered q can drift from the injected value without moving $\mathcal{M}$. The consistent downward bias specifically for $q_{\rm inj}$ close to unity is a prior-boundary effect: since $q \leq 1$ by the $m_1 \geq m_2$ convention, posterior support that would otherwise extend above $q = 1$ is reflected back below it, pulling the recovered median down whenever the true value sits close to this edge.

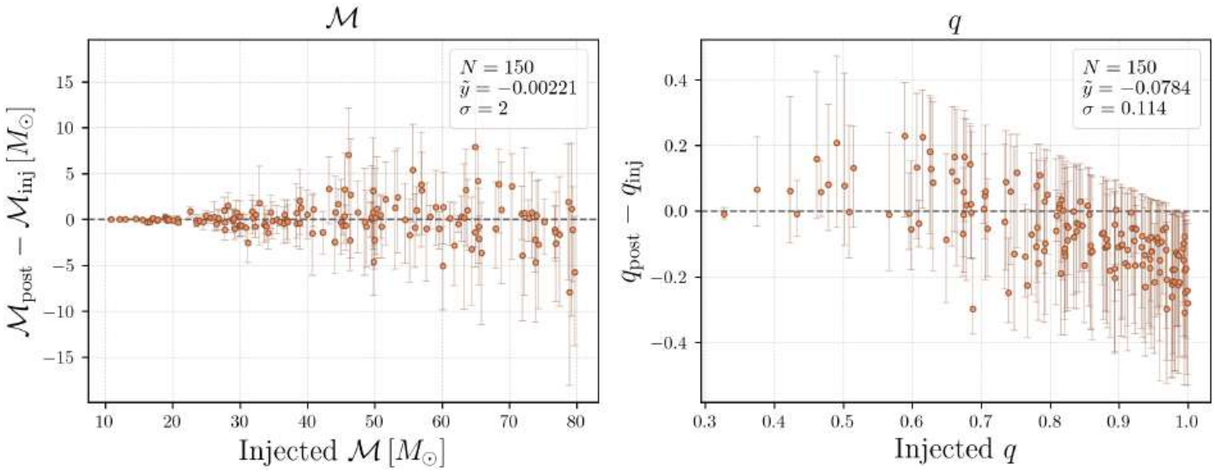

Figure 5.4: *Left*: chirp mass residual $\mathcal{M}_{\rm post} - \mathcal{M}_{\rm inj}$. *Right*: mass ratio residual $q_{\rm post} - q_{\rm inj}$. Both as a function of the injected value, for all 150 events.

Recall from Fig. 5.2 that the true source distance and the magnification of each image cannot be separated from the two apparent distances alone, which is why component masses were reported in the detector frame in the residual checks above.

The size of this bias can now be quantified directly: for each image, a naive source-frame chirp mass is computed by treating its effective luminosity distance as if it were the true source distance: converting it to a redshift under the same cosmology adopted by LeR ($H_0 = 70\,\mathrm{km\,s^{-1}\,Mpc^{-1}}$, $\Omega_m = 0.3$, flat ΛCDM, Table 4.1), and dividing the (detector-frame) recovered component masses by $(1 + z_{\rm naive})$.

Figure 5.5 shows the resulting bias, $\mathcal{M}^{\rm source}_{\rm naive} - \mathcal{M}^{\rm source}_{\rm inj}$, against the injected source-frame chirp mass, for image 1 (left) and image 2 (right). The bias is systematically positive and an order of magnitude larger than the chirp mass measurement noise seen in Fig. 5.4 ($\sigma \approx 2\,M_\odot$ there, versus a median of $11.6\,M_\odot$ and $11.4\,M_\odot$ here for images 1 and 2), so this is not simply propagated distance measurement noise.

To confirm that this bias is driven by magnification rather than by measurement noise, Figure 5.6 shows the same residual against the true injected magnification $\mu_{\rm true}$ of each image, plotted on a logarithmic axis to cover the several orders of magnitude spanned by $|\mu_{\rm true}|$ across the sample as in Fig. 4.5. The bias grows with $\log_{10}|\mu_{\rm true}|$ over most of the range, in the direction expected from the underlying physics: a magnified image ($|\mu| > 1$) has an effective distance smaller than the true source distance, so treating it as the true distance underestimates the redshift

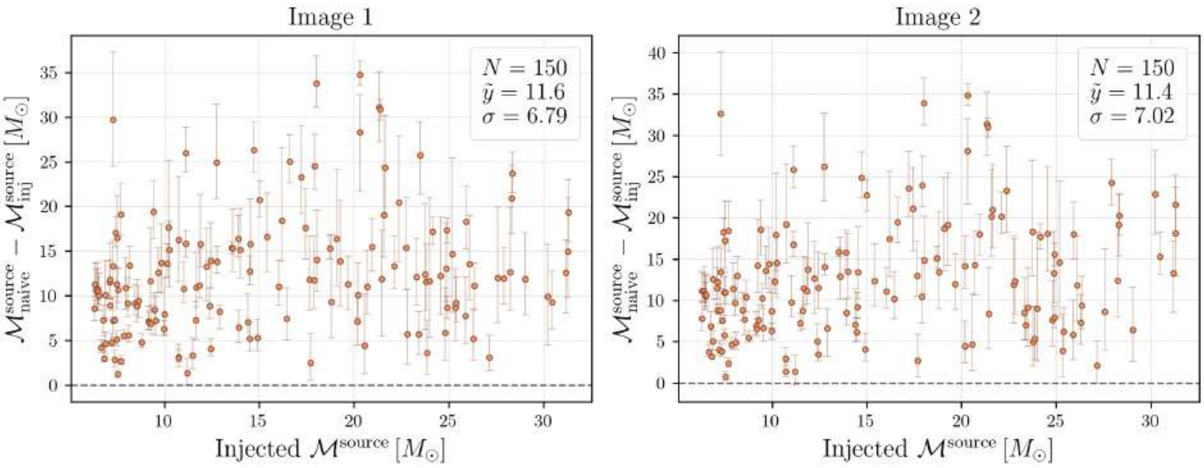

Figure 5.5: Bias of the naive source-frame chirp mass (obtained by treating each image's effective luminosity distance as the true source distance) against the injected source-frame chirp mass, for image 1 (*left*) and image 2 (*right*), for all 150 events.

and overestimates the recovered source-frame mass. This trend, rather than a random scatter centered on zero, confirms that magnification is the dominant driver of the bias seen in Fig. 5.5.

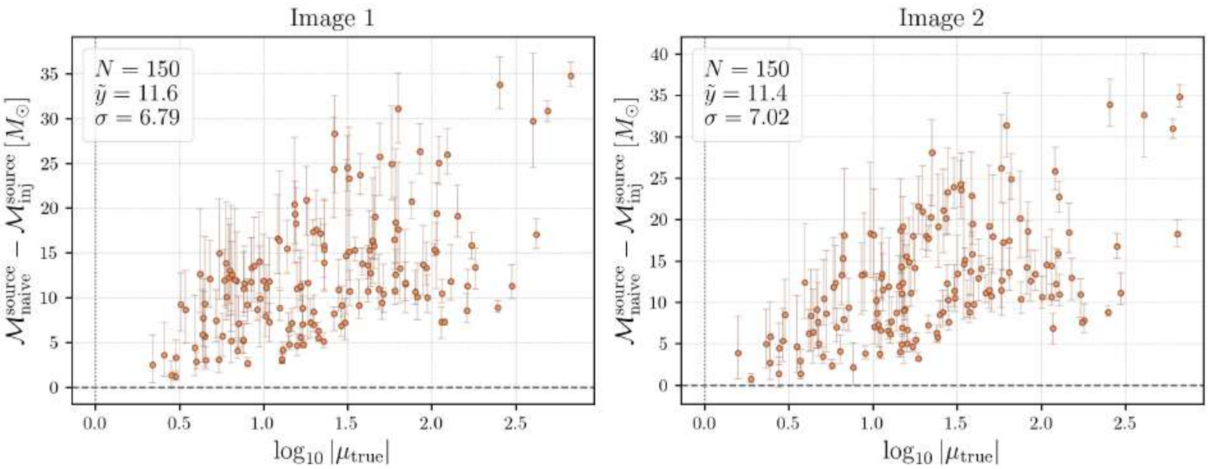

Figure 5.6: Same bias as in Fig. 5.5, plotted against the true injected magnification of each image, for image 1 (*left*) and image 2 (*right*).

5.5.2 Spins

Figure 5.7 shows the same comparison for the aligned effective spin $\chi_{\rm eff}$ (left) and the precessing spin parameter $\chi_{\rm p}$ (right). $\chi_{\rm eff}$ is recovered tightly and without systematic offset ($\tilde{y} = 0.009$, $\sigma = 0.082$), as expected since it enters the phase through the spin-orbit coupling at 1.5PN order, close to the leading mass-ratio dependence at 1PN, both comparatively early terms in the PN expansion. $\chi_{\rm p}$ is much less well constrained ($\tilde{y} = 0.040$, $\sigma = 0.165$), with the residual growing increasingly negative at higher injected values. As with the individual spin magnitudes below, this is a prior effect rather than a measurement one: $\chi_{\rm p}$ is not sampled directly but derived from a_1, a_2, and the tilt angles, whose priors induce an effective prior on $\chi_{\rm p}$ favoring low-to-moderate values; when the data constrain the underlying spin components only weakly, the recovered χ_{p} reverts toward this prior-typical range rather than toward a higher true value.

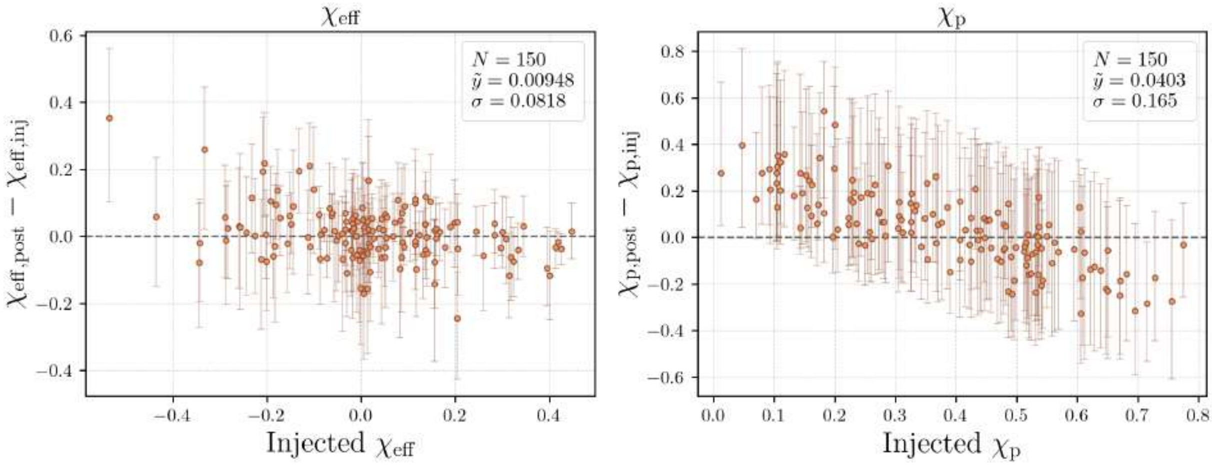

Figure 5.7: *Left*: effective spin residual $\chi_{\mathrm{eff,post}} - \chi_{\mathrm{eff,inj}}$. *Right*: precessing spin residual $\chi_{\mathrm{p,post}} - \chi_{\mathrm{p,inj}}$. Both as a function of the injected value, for all 150 events.

Figure 5.8 shows the residuals for the individual spin magnitudes a_1 (left) and a_2 (right). Both display a clear downward trend, with recovered values above the injected value at low a_{inj} and below it at high a_{inj} ($\tilde{y} = 0.014$ and 0.028, $\sigma = 0.196$ and 0.229, respectively). This compression toward intermediate spin values is the expected signature of a weakly measured parameter combined with a bounded prior: with limited information from the data, the posterior partially reverts toward the prior's typical range rather than to the true injected value.

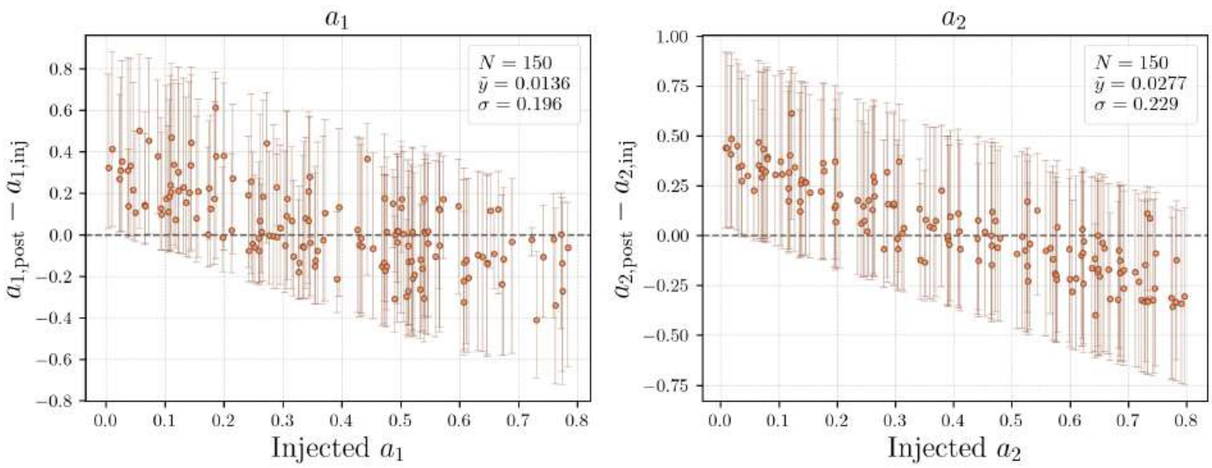

Figure 5.8: *Left*: primary spin residual $a_{1,\mathrm{post}} - a_{1,\mathrm{inj}}$. *Right*: secondary spin residual $a_{2,\mathrm{post}} - a_{2,\mathrm{inj}}$. Both as a function of the injected value, for all 150 events.

5.5.3 Distance and SNR consistency

Figure 5.9 shows the relative residual $(D^{\mathrm{eff}}_{L,\mathrm{post}} - D^{\mathrm{eff}}_{L,\mathrm{inj}})/D^{\mathrm{eff}}_{L,\mathrm{inj}}$ for the effective luminosity distance of image 1 (left) and image 2 (right). Both show a relative scatter of about 30% ($\tilde{y} = 0.027$, $\sigma = 0.279$ for image 1; $\tilde{y} = 0.018$, $\sigma = 0.301$ for image 2) without a strong systematic trend, consistent with luminosity distance being one of the least well determined parameters for a two-detector network due to its degeneracy with the inclination angle.

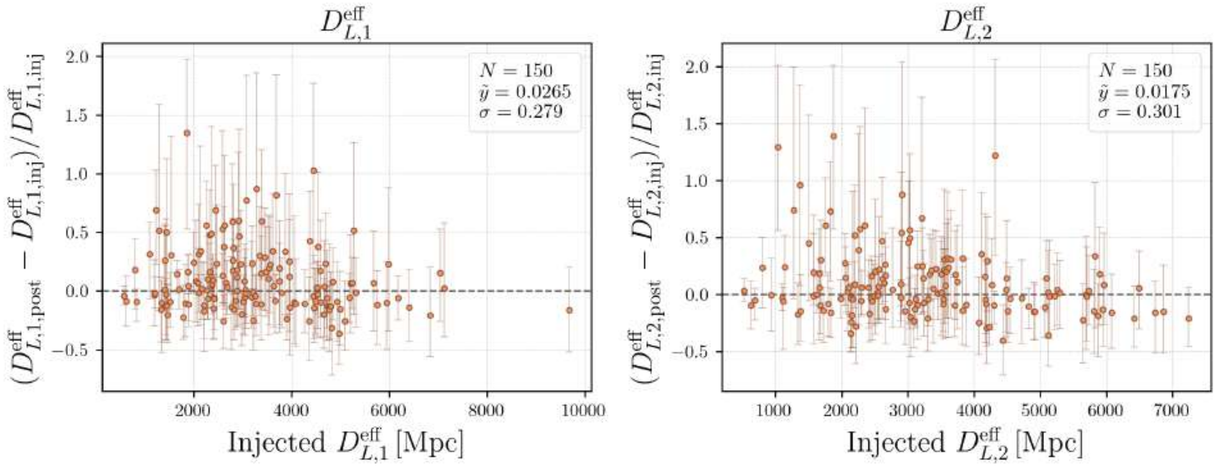

Figure 5.9: *Left*: relative residual for image 1, $(D^{\mathrm{eff}}_{L,1,\mathrm{post}} - D^{\mathrm{eff}}_{L,1,\mathrm{inj}})/D^{\mathrm{eff}}_{L,1,\mathrm{inj}}$. *Right*: same for image 2. Both as a function of the injected value, for all 150 events.

Figure 5.10 closes this residual comparison with a quantity that is not itself a source parameter but a consistency check on the pipeline: the recovered network SNR ρ_{post} against the injected network SNR ρ_{inj}. Because the joint analysis combines both lensed images, ρ_{post} is itself a two-image quantity, the median of $\sqrt{\rho^2_{\mathrm{net},1} + \rho^2_{\mathrm{net},2}}$ recomputed from the joint posterior samples (§5.4); like ρ_{inj}, this is the optimal SNR evaluated for the recovered template rather than a matched-filter SNR against the noisy data, so no random noise-realization scatter is expected between the two, only the systematic effects discussed below. For a consistent comparison, ρ_{inj} is defined here as the quadrature sum of the two per-image network SNRs reported by LeR for that event, $\rho_{\mathrm{inj}} = \sqrt{\rho^2_{\mathrm{net},1,\mathrm{inj}} + \rho^2_{\mathrm{net},2,\mathrm{inj}}}$, rather than the SNR of either image alone. With this definition the residual is centered close to zero and shows no systematic trend with ρ_{inj} ($\tilde{y} = -0.335$). Its scatter, however, is substantial ($\sigma = 2.89$), well above the few-percent level expected between two optimal-SNR estimates of the same signals, and is dominated by a single pronounced outlier, `inj113`, at $\rho_{\mathrm{post}} - \rho_{\mathrm{inj}} = +18.3$. This clear discrepancy led us to look more closely at how the injected SNR is computed by LeR.

The outlier, `inj113`, sits in an extreme corner of the parameter space: an extreme magnification ($|\mu_1| \approx 491$ and $|\mu_2| \approx 604$, the third-highest in the sample and within the $|\mu| \sim 10^3$ tail of Fig. 4.5) that boosts its effective SNR, combined with a high total mass ($M_{\mathrm{det}} \approx 136\, M_\odot$, 72nd percentile). `gwsnr`'s interpolation estimator reproduces the exact inner-product SNR to better than 0.5% on average across randomly sampled parameters [192], but it is built on the inspiral-based FINDCHIRP amplitude scaling, $\rho_{\mathrm{opt}} \propto \mathcal{M}^{5/6}/D_{\mathrm{eff}}$ [192], and is expected to be least accurate for high-total-mass, merger-dominated systems such as this one, where much of the SNR accumulates in the merger and ringdown rather than the inspiral. This suggests the injected SNR, rather than the recovery, as the origin of the discrepancy.

To confirm this, ρ_{inj} was recomputed for the full sample with `gwsnr`'s exact estimator (`snr_type="inner_product"`), which evaluates the noise-weighted frequency-domain inner product directly instead of interpolating a precomputed grid. For `inj113` this raises the injected network SNR from 38.9 to 59.7, an underestimate of about 35% by the interpolation, far larger than its sub-

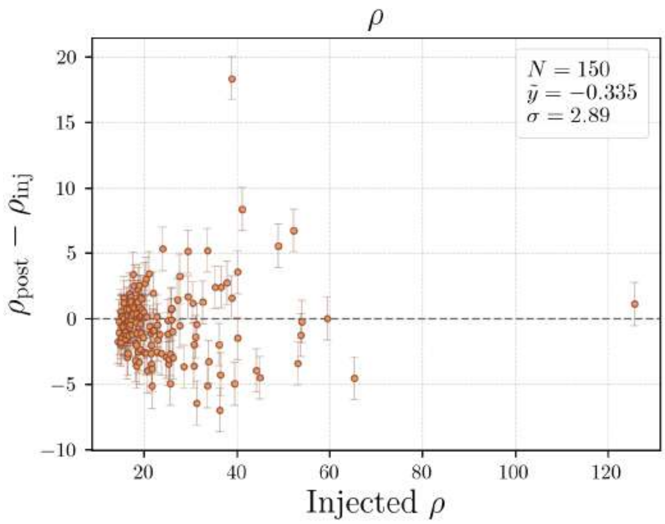

Figure 5.10: Difference between the recovered posterior network SNR $\rho_{\rm post}$ (median of the combined network SNR recomputed from the joint posterior samples, §5.4) and the injected network SNR $\rho_{\rm inj}$ (quadrature sum of the two per-image network SNRs computed by LeR, §4.1), as a function of $\rho_{\rm inj}$, for all 150 events. The residual is centered close to zero with no systematic trend ($\tilde{y} = -0.335$, $\sigma = 2.89$).

percent average accuracy and confirming that this event is an extreme corner of the method rather than a generic failure. The exact value is in close agreement with the recovered $\rho_{\rm post} = 57.2$ and reduces the residual from $+18.3$ to -2.5, turning a 6σ outlier into an ordinary event. Across the whole sample, the exact recomputation removes the outlier entirely (the largest absolute residual drops from 18.3 to 6.6), roughly halves the scatter (σ from 2.89 to 1.43), and tightens the median toward zero (-0.34 to -0.06), as shown in Fig. 5.11. This confirms that the large scatter and the `inj113` outlier in Fig. 5.10 originate in the accuracy limitation of LeR's fast interpolation SNR estimator, not in the joint PE recovery. The residual that remains after the exact recomputation ($\sigma = 1.43$) is at the few-percent level and is consistent with two genuine, smaller differences between the two calculations: `gwsnr` integrates the SNR up to a default 2048 Hz, wider than the 896 Hz upper cutoff of the Bilby analysis (Table 5.1), and the population is injected with the `IMRPhenomXPHM` waveform while the recovery uses `IMRPhenomXPNR` (§5.2), two models expected to agree closely but not exactly for these systems.

Importantly, this SNR mismatch does not affect any of the results of this thesis. The injected SNR computed by LeR enters the analysis only through the $\rho_{\rm net} \geq 10$ detectability selection (§4.2) and as the reference for this consistency check; it is never itself an input to the PE, which operates directly on the injected waveform and detector noise rather than on any scalar SNR value. The interpolation inaccuracy, moreover, appears only at high SNR, far above the $\rho_{\rm net} = 10$ threshold, so it does not change which events pass the selection. The recovered parameters and the sky localization results of Chapter 6 therefore depend only on the full PE of the injected data and are unaffected by the accuracy of LeR's fast SNR estimate. Figure 5.10 is retained with the interpolation-based $\rho_{\rm inj}$ precisely because it is this quantity, not the exact one, that defined the injected population.

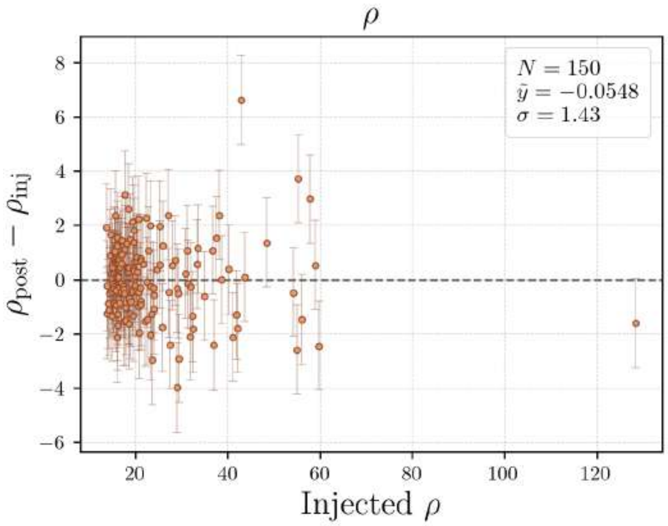

Figure 5.11: Same comparison as Fig. 5.10, but with the injected network SNR $\rho_{\rm inj}$ recomputed using `gwsnr`'s exact inner-product estimator instead of the default interpolation. The `inj113` outlier disappears and the scatter is roughly halved ($\tilde{y} = -0.06$, $\sigma = 1.43$), confirming that the discrepancy in Fig. 5.10 originates in the interpolation-based SNR estimate rather than in the joint PE recovery.

5.5.4 Recovering the true source-frame mass with an assumed lens model

Recovering the true source distance, and hence the true source-frame mass, requires resolving the magnification-distance degeneracy identified above. For a general lens model this is not possible from the two apparent distances alone: even a full reconstruction of the lens from all its lensed-image observables leaves residual degeneracies, such as an overall similarity transformation, that can only be broken with electromagnetic follow-up of the identified lens galaxy (its redshift and Einstein radius, plus an independent mass estimate from its velocity dispersion) [191], or in the presence of wave-optics lensing effects, which imprint interference patterns on the waveform that break this degeneracy without requiring an electromagnetic counterpart [193]. Since this is a simulation-only study, no such follow-up exists, and substituting LeR's known lens truth in its place would be circular, recovering the injected value by construction rather than testing a method.

Instead, the degeneracy is broken by assuming a simple, generic lens model [194], reducing the unknowns to a single free parameter: the normalized source-lens separation $u \equiv \beta/\theta_E$ introduced in Eq. (3.14). In the geometric optics limit, the magnifications of the two images of a point-mass (PM) or singular isothermal sphere (SIS) lens are

$$\begin{aligned} \text{PM}: \quad & \mu_+(u) = \frac{1}{2}\left(\frac{u^2+2}{u\sqrt{u^2+4}}+1\right), \quad |\mu_-(u)| = \frac{1}{2}\left(\frac{u^2+2}{u\sqrt{u^2+4}}-1\right), \\ \text{SIS}: \quad & \mu_+(u) = 1+\frac{1}{u}, \quad |\mu_-(u)| = \frac{1}{u}-1 \quad (u<1). \end{aligned} \tag{5.1}$$

Both expressions are written in terms of $|\mu_-|$ rather than the signed value of Eq. (3.14), since only the magnitude of the magnification enters the apparent distance and hence the observable relative magnification below. The relative magnification $\mu_{\rm rel} = |\mu_-|/\mu_+$ is measurable directly from the two apparent distances already in the joint posterior, without any external information,

$$\mu_{\rm rel} = \left(\frac{D^{\rm eff}_{L,+}}{D^{\rm eff}_{L,-}}\right)^2, \tag{5.2}$$

where the "+" image is defined, per event, as the one with the smaller apparent distance. Equations (5.1) and (5.2) are inverted numerically to solve for u, which is then substituted back into the corresponding expression for $\mu_+(u)$ to obtain a corrected true distance $D_L = \sqrt{\mu_+}\, D^{\rm eff}_{L,+}$, from which the source-frame chirp mass is recomputed as in the naive case above.

Figure 5.12 shows the resulting bias for both models, evaluated with only the joint posterior's own apparent distances, never the injected truth. Unlike the naive bias, which is positive for every event, the corrected bias takes both signs, so the median alone can partly reflect cancellation between events corrected too much and events corrected too little, rather than a genuine reduction. The median absolute bias, $\widetilde{|y|}$, avoids this and is used here as the main summary statistic: the naive bias corresponds to $\widetilde{|y|} \approx 11.5\, M_\odot$ (Figure 5.5, where every residual has the same sign), the point-mass correction reduces this to $\widetilde{|y|} = 6.18\, M_\odot$ ($\tilde{y} = 5.86$, $\sigma = 6.33$), and the SIS correction reduces it further to $\widetilde{|y|} = 4.34\, M_\odot$ ($\tilde{y} = 3.52$, $\sigma = 5.62$).

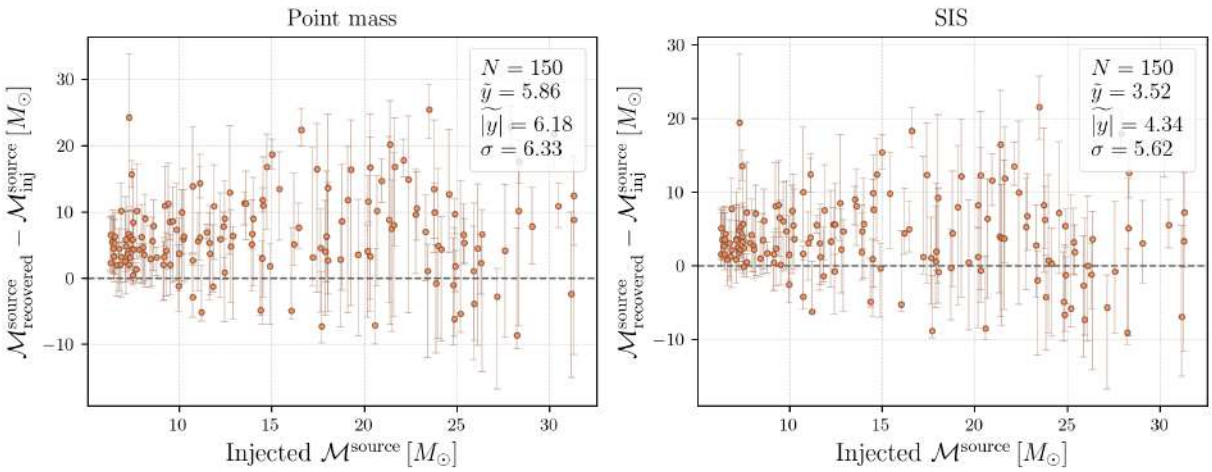

Figure 5.12: Bias of the source-frame chirp mass after correcting for magnification with an assumed point-mass (*left*) or SIS (*right*) lens model, against the injected source-frame chirp mass, for all 150 events.

The SIS correction outperforms the point-mass correction on both statistics, consistent with LeR's true lens model being an EPL profile with external shear (Table 4.1, $\gamma \sim \mathcal{N}(2.0, 0.2^2)$), structurally closer, on average, to an isothermal profile than to a point mass. Neither model brings the bias down to the noise level seen in the detector-frame chirp mass residual ($\sigma \approx 2\, M_\odot$, Figure 5.4); what remains comes from the difference between the simple, round lens assumed here and the true elliptical, sheared lens of each event.

5.5.5 Confirming the lensed nature of the injected pairs

The coherence ratio $\mathcal{C}^{\rm L}_{\rm U}$ (Eq. (3.33), §3.4.3) tests whether the joint analysis is favored over treating the two images as independent sources, and can be evaluated directly from the signal-vs-noise

Bayes factor already reported by each BILBY and HANABI run, without any additional computation, using Eq. (3.34). For each of the 150 successfully analyzed pairs (§5.4.3), $\ln \mathcal{C}^{\mathrm{L}}_{\mathrm{U}}$ is computed from the $\ln \mathcal{B}_{SN}$ values of the two single-image runs and the joint run. Table 5.3 summarizes the resulting distribution, and Figure 5.13 shows it in full. Every one of the 150 pairs yields $\ln \mathcal{C}^{\mathrm{L}}_{\mathrm{U}} > 0$, including the weakest event at $\ln \mathcal{C}^{\mathrm{L}}_{\mathrm{U}} = 0.76$, with a median of 6.92 corresponding to odds of about 1000:1 in favor of the lensed hypothesis. This is the expected outcome for a population of pairs that are lensed by construction, and serves as an internal consistency check of the joint analysis: the coherent likelihood correctly favors $\mathcal{H}_L$ over $\mathcal{H}_U$ across the full range of SNRs and time delays present in the sample, without incorporating selection effects or population priors, which would be required for the full lensing Bayes factor $\mathcal{B}^{\mathrm{L}}_{\mathrm{U}}$ of §3.4.3 but are beyond the primary scope of this work.

Statistic	$\ln \mathcal{C}^{\mathrm{L}}_{\mathrm{U}}$
Median	6.92
Mean	7.70
Standard deviation	3.82
Minimum	0.76
Maximum	26.66
Fraction with $\ln \mathcal{C}^{\mathrm{L}}_{\mathrm{U}} > 0$	150/150 (100%)

Table 5.3: Summary statistics of the coherence ratio $\ln \mathcal{C}^{\mathrm{L}}_{\mathrm{U}}$ (Eq. (3.34)) across the 150 successfully analyzed lensed pairs, computed from the $\ln \mathcal{B}_{SN}$ values already reported by the single-image and joint BILBY/HANABI runs.

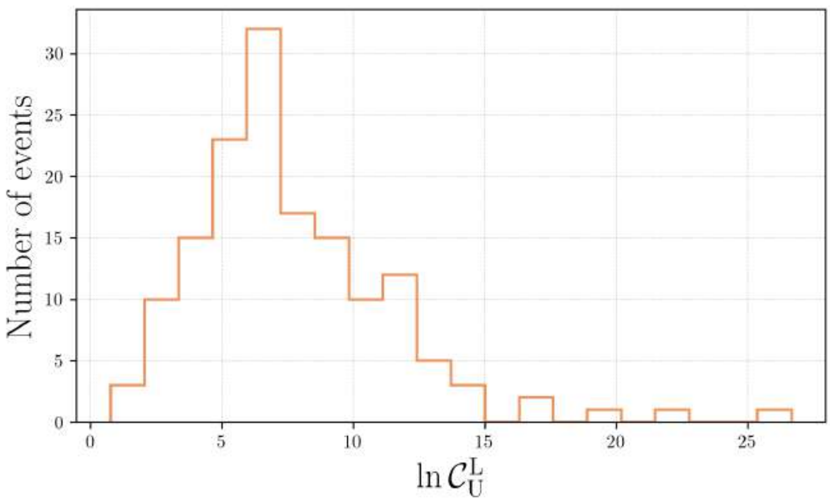

Figure 5.13: Distribution of the coherence ratio $\ln \mathcal{C}^{\mathrm{L}}_{\mathrm{U}}$ across the 150 successfully analyzed pairs. All events lie above $\ln \mathcal{C}^{\mathrm{L}}_{\mathrm{U}} = 0$, the threshold above which the lensed hypothesis $\mathcal{H}_L$ is favored over the unlensed hypothesis $\mathcal{H}_U$.

Chapter 6

Sky localization

6.1 From posterior samples to sky maps

The sky position posteriors obtained from each PE run are converted into all-sky probability maps using `ligo-skymap-from-samples`, which performs an adaptive kernel density estimate on the right ascension and declination samples and produces a multi-resolution HEALPix [195] map. For each of the 150 successfully analyzed lensed pairs (§5.4.3), three sky maps are generated: one for each single-image analysis (image 1 and image 2 independently) and one for the joint analysis. From each map, the 50% and 90% credible areas are extracted with `ligo-skymap-stats`, following the definition of the minimum credible region $\mathcal{R}_{90}$ given in Eq. (3.32). The dominant constellation and its contained probability fraction are obtained with `ligo-skymap-constellations` (§6.2.2).

Figure 6.1 illustrates the effect of the joint analysis on sky localization for a representative event (injection 002, $\Delta t \approx 5.6\,\mathrm{h}$ between images). The first single-image analysis yields a broad arc spanning $A_{90} \approx 970\,\mathrm{deg}^2$, characteristic of a two-detector network where the source direction is constrained to an arc of constant time delay between H1 and L1 (§3.4.1). For this particular event, the second image happens to be more strongly magnified than the first ($|\mu_2| \approx 129$ versus $|\mu_1| \approx 80$) and therefore has higher SNR; its 90% area is correspondingly smaller, $\approx 416\,\mathrm{deg}^2$, and concentrated in a single region rather than an extended arc. This is a property of this specific pair rather than a general rule, since either image could equally be the more magnified one. The joint posterior, which requires both images to share the same sky position (Eq. (3.30)), collapses the credible region to $\approx 27\,\mathrm{deg}^2$: a factor of $\sim$36 improvement over the broader single image and ~ 15 over the narrower one. This large reduction, well beyond what either image achieves on its own regardless of its individual SNR, confirms the mechanism described in §3.4.2: the 5.6 h time delay is long enough to significantly reorient the H1-L1 antenna pattern between the two images, so combining them breaks a degeneracy that improving either image alone could not.

6.2 Sky localization: single-image versus joint analysis

6.2.1 90% credible area

The 90% credible sky area A_{90} is computed for each of the 150 events, separately for the best single-image analysis (defined as $\min(A_{90}^{\mathrm{img1}}, A_{90}^{\mathrm{img2}})$) and for the joint analysis. The improvement factor per event is defined as $\mathrm{IF}_{A_{90}} \equiv A_{90}^{\text{best single}}/A_{90}^{\text{joint}}$.

Figure 6.2 shows the cumulative and probability density distributions of A_{90} for the best single-

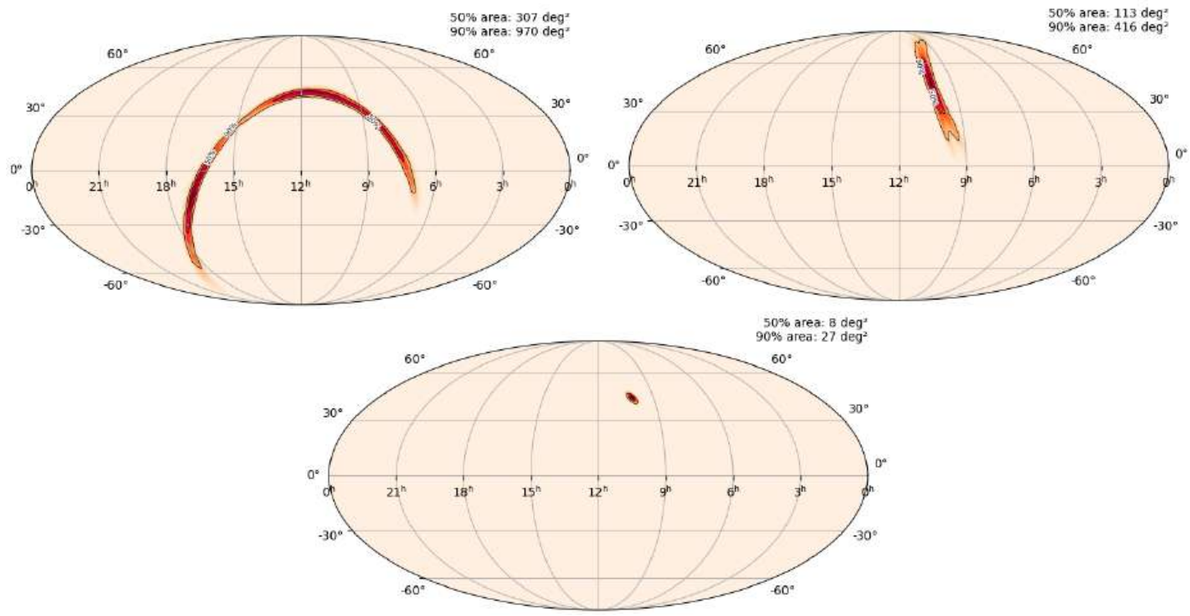

Figure 6.1: Sky localization for injection 002 ($\Delta t \approx 5.6\,\mathrm{h}$). *Top left*: image 1 ($A_{90} \approx 970\,\mathrm{deg}^2$). *Top right*: image 2 ($A_{90} \approx 416\,\mathrm{deg}^2$). *Bottom*: joint analysis ($A_{90} \approx 27\,\mathrm{deg}^2$).

image and joint analyses. The joint distribution is shifted to smaller areas across the entire range. The median 90% area drops from $571\,\mathrm{deg}^2$ (best single image) to $99.9\,\mathrm{deg}^2$ (joint), a ratio of medians of 5.7. The median per-event improvement factor is 5.1, confirming that the improvement is not driven by a small number of outliers but is systematic across the population. The joint analysis improves the sky localization for all 150 events without exception ($\mathrm{IF}_{A_{90}} > 1$ in every case), with individual improvement factors ranging from 1.1 to 28.0. Five events (3.3%) reach a joint $A_{90} < 10\,\mathrm{deg}^2$, while half of the sample (75/150) falls below $100\,\mathrm{deg}^2$. The $10\,\mathrm{deg}^2$ threshold is commonly adopted as a benchmark for host galaxy identification and multimessenger follow-up with wide-field survey telescopes [196]: sky areas at or below this scale contain a manageable number of candidate host galaxies for spectroscopic follow-up, and can be covered by a small number of pointings with instruments such as DECam, ZTF, or the Rubin Observatory.

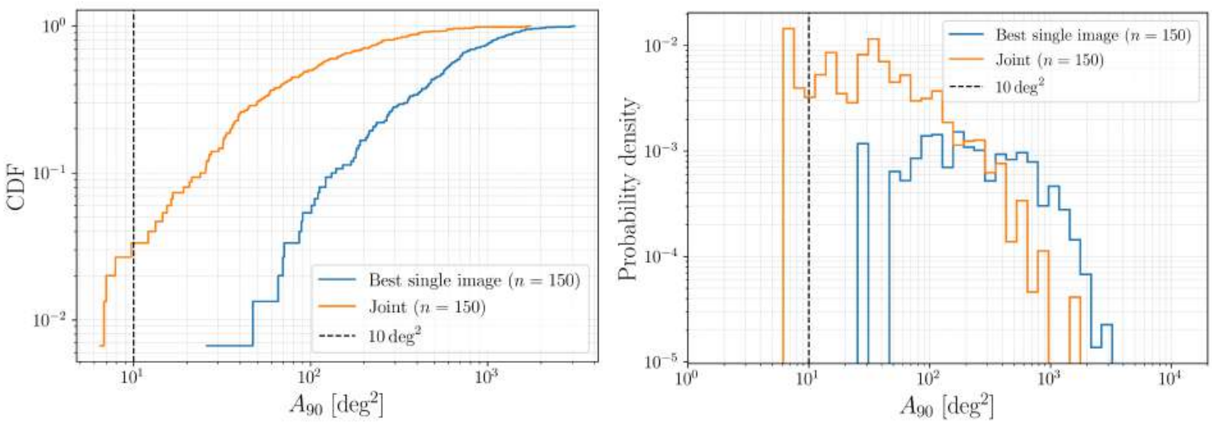

Figure 6.2: CDF (*left*) and probability density (*right*) of the 90% credible sky area for the best single-image analysis and the joint analysis across the 150 events. The vertical dashed line marks $10\,\mathrm{deg}^2$.

In addition to the 90% credible area, the tighter 50% credible area A_{50} provides a complementary summary of the core localization precision, less sensitive to the extended low-probability tails of the posterior visible in Fig. 6.1. Table 6.1 compares both statistics for the best single-image and joint analyses. The median A_{50} drops from $145.4\,\mathrm{deg}^2$ (best single image) to $27.2\,\mathrm{deg}^2$ (joint), a median improvement factor $\mathrm{IF}_{A_{50}} = 4.7$, consistent with, though slightly smaller than, the $\mathrm{IF}_{A_{90}} = 5.1$ found above. Only one event has $\mathrm{IF}_{A_{50}}$ marginally below unity (0.97), the sole case among either metric where the joint analysis does not improve the localization; even there, the joint A_{50} is left essentially unchanged rather than substantially worse.

	A_{50} [deg^2]		A_{90} [deg^2]	
	Best single	Joint	Best single	Joint
Median	145.4	27.2	571.0	99.9
Mean	172.4	48.2	699.3	185.7

Table 6.1: Median and mean 50% and 90% credible sky areas for the best single-image and joint analyses across the 150 events. The corresponding median improvement factors are $\mathrm{IF}_{A_{50}} = 4.7$ and $\mathrm{IF}_{A_{90}} = 5.1$.

Figure 6.3 shows the 90% credible area as a function of network SNR for each image, comparing the single-image and joint results. For single-image analyses, A_{90} shows a broad scatter at all SNR values, reflecting the dominant role of sky position and antenna pattern orientation relative to SNR in determining the localization precision of a two-detector network. The joint analysis consistently reduces A_{90}, with the improvement being most pronounced for events whose single-image areas are largest.

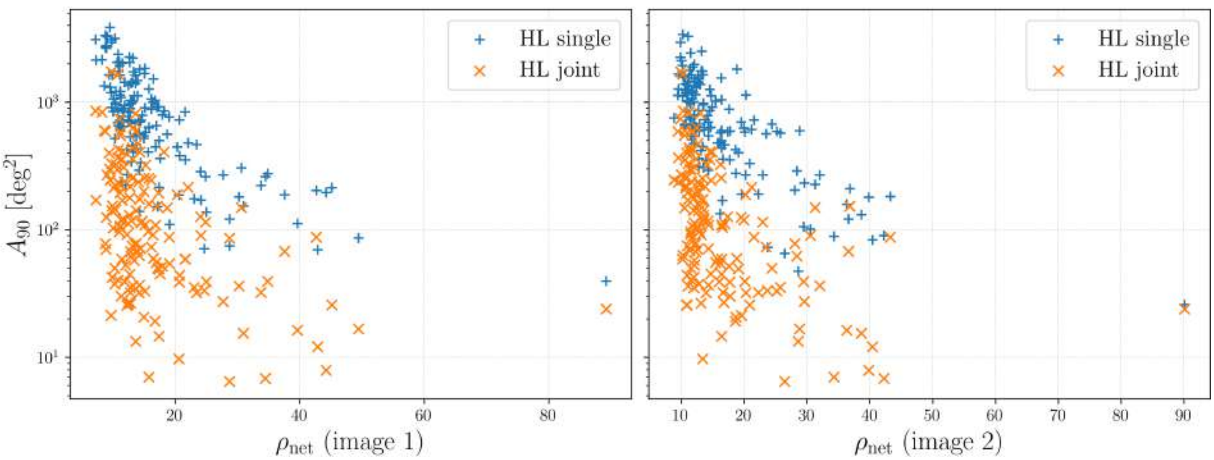

Figure 6.3: 90% credible sky area versus network SNR for image 1 (*left*) and image 2 (*right*). Blue crosses: single-image analysis; orange crosses: joint analysis for the same event.

6.2.2 Constellation localization

The `ligo-skymap-constellations` tool [185] identifies the IAU constellation that contains the largest fraction of the posterior probability from each sky map. Unlike the credible area of §6.2.1, this metric maps the localization onto a recognizable region of the sky rather than a raw solid angle, which is useful for visualization and outreach, and, as discussed below, offers a possible naming convention for lensed pairs. A caveat of this metric is that the 88 IAU constellations span a wide range of angular sizes, from $\sim$70 to $\sim$1300 deg^2, so the constellation fraction mixes together how well localized a source is with where on the sky it happens to fall; the comparison with real GWTC-5 events below, which is subject to the same effect, provides a control for this dependence.

Figure 6.4 shows the probability fraction contained in the dominant constellation as a function of network SNR. For single-image analyses, the median fraction in the top constellation is $\sim$41% (best of the two images), with only 50/150 events exceeding 50%. The joint analysis raises the median to 68.6%, with 115/150 events (76.7%) above 50% and 47/150 (31.3%) above 90%. The clustering of joint results toward high constellation fractions, visible in the upper part of the scatter plot, reflects the same mechanism as the sky area reduction: the joint posterior concentrates into a compact region that fits within a single constellation boundary.

A high constellation fraction also has a practical use beyond visualization: naming lensed pairs. Lensed image pairs are currently identified by concatenating the two individual event names (e.g. GW200105-GW200115), a notation that gets harder to read as more pairs are found, particularly when their events share similar dates. Constellation-based names are already the default convention in several branches of astronomy, and a comparable scheme could in principle be adopted for confidently localized lensed pairs. The 47/150 joint analyses reaching above 90% found here suggest this is not a marginal possibility: for a non-negligible fraction of events, the joint posterior is confidently confined to a single constellation. This fraction is expected to improve substantially if the analysis is extended to a three-detector network (H1, L1, and Virgo), which would further tighten the joint localization.

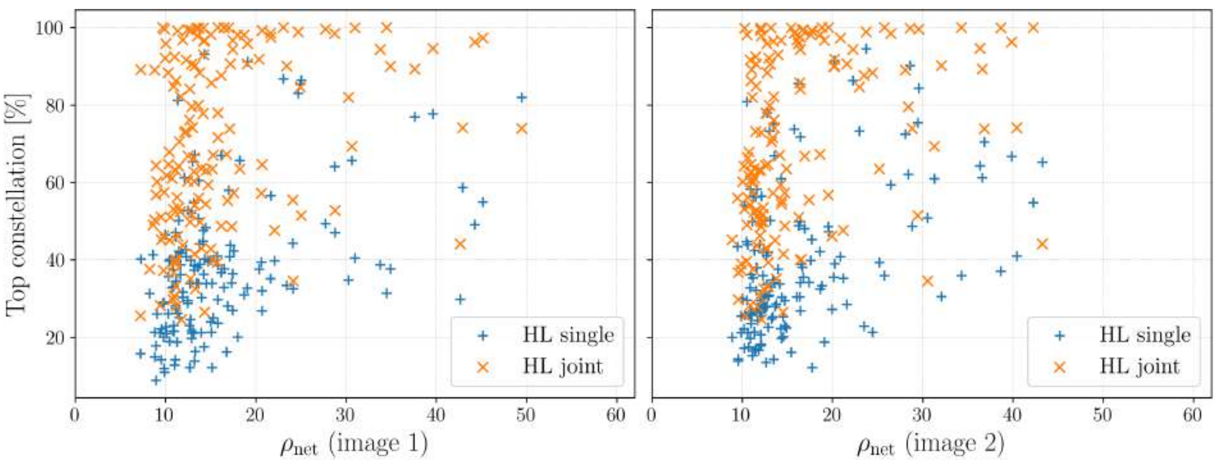

Figure 6.4: Probability fraction in the dominant constellation versus network SNR for image 1 (*left*) and image 2 (*right*). Blue crosses: single-image analysis; orange crosses: joint analysis.

To place these results in the context of real observations, Figure 6.5 shows the same metric computed from the published sky maps of 104 GWTC-5 events [70], obtained from the Gravitational Wave Open Science Center [1], analyzed with the `IMRPhenomXPNR` waveform model. Each event is color-coded by the detector network that observed it. Events detected by the full three-detector network (HLV, 75 events) span the entire range of constellation fractions from $\sim$15% to 100%, with the highest values reached at moderate to high SNR. In contrast, the two-detector subsets (HL, HV, LV) and the single-detector event (H) are confined to low constellation fractions, typically below 50%. The HL events, which use the same detector pair as this work, cluster between 10% and 50%, consistent with the single-image results of the simulated population (median $\sim 41\%$). This comparison highlights the localization advantage that a third detector provides in standard analyses, and underscores the value of the joint lensed-image approach: even with only two detectors, the joint analysis achieves constellation fractions (median 68.6%) that are comparable to those of the three-detector network on real data.

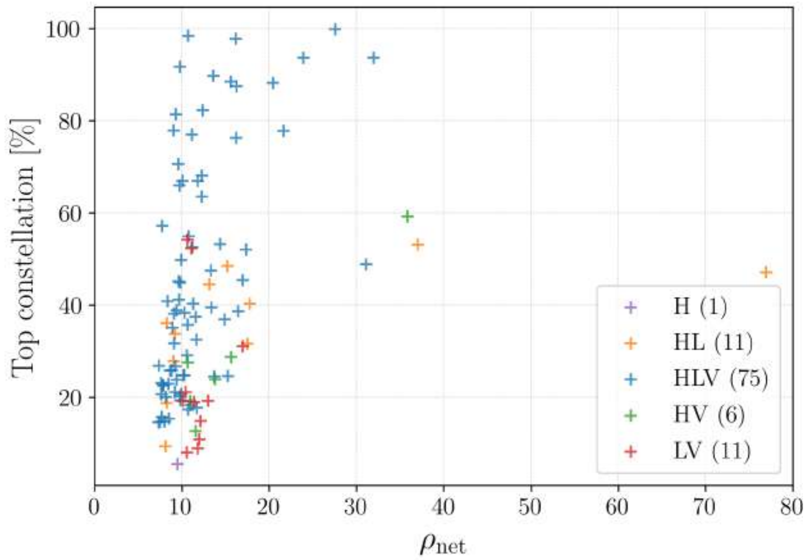

Figure 6.5: Probability fraction in the dominant constellation versus network SNR for 104 GWTC-5 events, colored by detector network. The HL subset (orange), which matches the detector configuration used in this work, is confined below $\sim 50\%$.

6.2.3 Comparison with the literature

Two recent studies have investigated sky localization improvements from strongly lensed GW images, using complementary approaches.

Hannuksela et al. [147] demonstrated that quadruply lensed GW events observed by the full LIGO/Virgo/KAGRA network can be localized to one or a few candidate host galaxies by combining GW sky maps with EM constraints on the lens galaxy properties, ultimately achieving sub-arcsecond precision within the identified host. Their approach relies on the availability of four images and three or more detectors, and targets a different regime from the present work: galaxy-scale identification rather than population-level sky area statistics.

Li and Hannuksela [2] performed the first systematic study of repeated-signal localization using Bayestar [185] sky maps on a simulated lensed population observed with the three-detector HLV network (H1, L1, and Virgo), adopting Bilby's default detector sensitivity curves (the Advanced LIGO O4 and Advanced Virgo design PSDs). They found that combining two images reduces the 90% credible area by roughly an order of magnitude, somewhat larger than the median improvement factor of 5.1 found in this work. The key differences with respect to the present analysis are: (i) this work uses a two-detector network (H1+L1) with the projected O5a PSD, rather than the three-detector HLV network with Bilby's default design curves, (ii) the localization is based on full Bayesian PE with Bilby and Hanabi rather than the rapid

BAYESTAR algorithm, and (iii) the population is generated with LeR including a realistic duty cycle filter. The smaller improvement factor obtained here is expected: with only two detectors, the single-image localization is already less constrained by triangulation and relies more heavily on antenna pattern information, so the additional gain from combining a second image, which acts through the same antenna pattern channel, is comparatively smaller than in a three-detector network where the two images can each contribute additional triangulation baselines. Despite this difference in magnitude, both studies converge on the same qualitative conclusion: combining two lensed images produces a substantial, consistent improvement in sky localization across detector configurations and analysis methods.

6.3 Summary of sky localization results

The results presented in this chapter demonstrate that the joint PE analysis of strongly lensed BBH image pairs produces a systematic and substantial improvement in sky localization compared to the analysis of individual images, even with a two-detector network. The median 90% credible area drops from $571\,\mathrm{deg}^2$ (best single image) to $99.9\,\mathrm{deg}^2$ (joint), corresponding to a median per-event improvement factor of 5.1. The improvement is universal across the sample: all 150 events are better localized by the joint analysis, with improvement factors ranging from 1.1 to 28.0. Five events (3.3%) reach a joint sky area below the $10\,\mathrm{deg}^2$ threshold commonly adopted for host galaxy identification and multimessenger follow-up [196], placing them within reach of targeted EM searches with wide-field survey telescopes. Comparison with real GWTC-5 events shows that the joint lensed analysis with two detectors achieves constellation localization fractions comparable to those of the three-detector network, and comparison with prior theoretical work confirms that combining two lensed images produces a substantial and consistent sky localization gain across different detector configurations and analysis methods, albeit with a magnitude that depends on the detector network considered.

Chapter 7

Conclusions

In this thesis we have investigated whether the joint PE improves sky localization compared to the analysis of individual images. To this end, we have developed and executed a complete simulation and analysis pipeline, from population generation to sky map comparison, targeting the H1+L1 LIGO network under projected O5a sensitivity. The main original results of this work are:

1. The construction of a realistic strongly lensed BBH injection set of 157 events (150 successfully analyzed), generated with LeR under O5a sensitivity and filtered through a detector duty cycle, and the development of a fully automated PE pipeline that processes each event through single-image analysis with Bilby and joint analysis with Hanabi on the Picasso HPC cluster.

2. A systematic validation of the recovered parameters across the 150 analyzed events, demonstrating unbiased recovery of the chirp mass and effective spin, quantifying the magnification-distance degeneracy and its impact on source-frame mass estimation, and confirming the lensed nature of all 150 injected pairs through the coherence ratio.

3. The demonstration that the joint analysis of two lensed images produces a universal improvement in sky localization: the median 90% credible area drops from $571\,\mathrm{deg}^2$ (best single image) to $99.9\,\mathrm{deg}^2$ (joint), corresponding to a median per-event improvement factor of 5.1, with all 150 events showing improvement and five reaching below the $10\,\mathrm{deg}^2$ threshold relevant for multimessenger follow-up.

Table 7.1 collects the key figures behind these three results.

These results are presented in Chapters 4–6. Chapter 4 described the construction of the population with LeR v0.4.2, drawing approximately 152 000 candidate events to collect 500 detectable ones ($\mathcal{R}_L \approx 0.75\,\mathrm{yr}^{-1}$), reduced to the final 157-event injection set by a realistic duty cycle filter. Chapter 5 presented the PE pipeline, consuming approximately 830,000 CPU-hours, and its validation: the chirp mass and effective spin are recovered without significant bias, the individual spin magnitudes and χ_p remain prior-dominated as expected from their weak imprint on the signal, and the magnification-distance degeneracy introduces a source-frame mass bias (median $\sim 11.5\,M_\odot$) that an assumed SIS lens model partially corrects (down to $4.34\,M_\odot$). All 150 pairs yield a positive coherence ratio (median $\ln \mathcal{C}^\mathrm{L}_\mathrm{U} = 6.92$), confirming that the joint likelihood correctly favors the lensed hypothesis throughout the sample. Chapter 6 presented the sky localization results themselves: individual improvement factors range from 1.1 to 28.0, the constellation localization fraction rises from a median of 41% to 68.6%, and the joint analysis

Quantity	Value
Injection set (successfully analyzed)	157 (150)
Compute cost (Picasso HPC)	$\sim$ 830,000 CPU-hours
Chirp mass recovery bias	$\tilde{y} = -0.002\,M_\odot$
Naive source-frame mass bias (median)	$\sim 11.5\,M_\odot$
Source-frame mass bias, SIS-corrected (median)	$4.34\,M_\odot$
Coherence ratio, all pairs	$\ln \mathcal{C}^{\mathrm{L}}_{\mathrm{U}} > 0$ (median 6.92)
Median A_{90}: best single image $\rightarrow$ joint	$571 \rightarrow 99.9\,\mathrm{deg}^2$
Median per-event improvement factor $\mathrm{IF}_{A_{90}}$	5.1 (range 1.1-28.0)
Events with joint $A_{90} < 10\,\mathrm{deg}^2$	5/150 (3.3%)
Median constellation fraction: single $\rightarrow$ joint	41% $\rightarrow$ 68.6%

Table 7.1: Summary of the key quantitative results of this thesis, drawn from Chapters 4–6.

with two detectors reaches constellation fractions comparable to those of the three-detector HLV network on real GWTC-5 events.

These results extend the first systematic study of lensed-pair sky localization by Li and Hannuksela [2], who found an order-of-magnitude improvement using Bayestar sky maps on a three-detector (HLV) simulated population. The smaller median improvement factor found here (5.1) is expected from the use of only two detectors, where single-image localization already relies primarily on antenna pattern rather than triangulation, leaving comparatively less room for the joint analysis to improve upon it. Both studies nonetheless agree qualitatively: combining two lensed images produces a substantial, consistent sky localization gain regardless of detector network or analysis method, and the present work extends this conclusion to full Bayesian PE, a more realistic O5a population with a duty cycle filter, and a direct comparison with real GWTC-5 data.

Several idealizations should be kept in mind when interpreting these results: the noise is stationary and Gaussian, with none of the glitches present in real LIGO data; the same O5a PSD is assumed for both detectors, whereas real detectors have distinct, time-varying sensitivities; detector calibration is treated as perfect; and the waveform model used for recovery (`IMRPhenomXPNR`) differs from the one used for injection (`IMRPhenomXPHM`). None of these idealizations are expected to change the qualitative conclusion, but they mark natural directions for follow-up work, together with extending the analysis to a three-detector network (H1, L1, Virgo), which should further tighten both the absolute localization and the improvement factor; computing the full population-informed lensing Bayes factor $\mathcal{B}^{\mathrm{L}}_{\mathrm{U}}$ rather than the coherence ratio alone; refining the source-frame mass correction with more sophisticated lens reconstruction, particularly where electromagnetic identification of the lens is available; and simulating targeted electromagnetic follow-up campaigns for the events reaching joint sky areas below $10\,\mathrm{deg}^2$, to assess the practical feasibility of host galaxy identification.

Taken together, the results of this thesis demonstrate that the joint analysis of strongly lensed GW image pairs is a powerful tool for improving sky localization, even with a two-detector network. As detector sensitivity improves and the rate of detections grows, lensed pairs will become an increasingly common feature of the GW landscape, and the pipeline and methodology developed here provide a foundation for analyzing them systematically.

Bibliography

[1] LIGO Scientific Collaboration and Virgo Collaboration and KAGRA Collaboration. *Open Data from the Second Part of the Fourth Observing Run of Advanced LIGO, Advanced Virgo, and KAGRA*. Data obtained from the Gravitational Wave Open Science Center (https://gwosc.org). 2026. arXiv: 2605.27090 [gr-qc] (cit. on pp. IV, 109).

[2] Alvin K. Y. Li and Otto A. Hannuksela. *A First Investigation of Repeated-Signal Localization of Strongly Lensed Gravitational Waves for Multimessenger Astronomy*. 2026. arXiv: 2604.16561 [astro-ph.IM] (cit. on pp. V, 77, 78, 110, 113).

[3] Bert Janssen. *General Relativity Theory*. Lecture notes, University of Granada, 4th year Physics course. Department of Theoretical Physics and Cosmology, Mecenas Building, Fuente Nueva Campus, 18071 Granada, Spain, 2013 (cit. on pp. VI, 3, 4, 12, 15, 17, 18, 20, 22, 23, 25, 26, 29).

[4] J.A. Wheeler and K. Ford. *Geons, Black Holes, and Quantum Foam: A Life in Physics*. W. W. Norton, 2010. ISBN: 9780393079487. URL: https://books.google.es/books?id=zGFkK2tTXPsC (cit. on p. VI).

[5] B. P. Abbott et al. "Observation of Gravitational Waves from a Binary Black Hole Merger". In: *Phys. Rev. Lett.* 116.6 (2016), p. 061102. DOI: 10.1103/PhysRevLett.116.061102 (cit. on pp. VI, 3, 28, 44, 52).

[6] Max Planck Society. *100 Years of General Relativity: A Chronology*. Accessed: 2025-07-17. 2015. URL: https://www.mpg.de/9700434/chronology (cit. on p. 2).

[7] Albert Einstein. "Zur Elektrodynamik bewegter Körper". In: *Annalen der Physik* 322.10 (1905). English translation: *On the Electrodynamics of Moving Bodies*, pp. 891–921. DOI: 10.1002/andp.19053221004 (cit. on p. 2).

[8] Karl Schwarzschild. *On the Gravitational Field of a Point-Mass According to Einstein's Theory*. Tech. rep. Translated in 2008 by Larissa Borissova and Dmitri Rabounski. Sitzungsberichte der Königlich Preußischen Akademie der Wissenschaften, Berlin (Mathematisch-Physikalische Klasse), 1916, pp. 189–196. URL: https://www.ipcms.fr/uploads/2023/11/On-the-Gravitational-Field-of-a-Point-Mass.pdf (cit. on p. 2).

[9] A. S. Eddington. "The total eclipse of 1919 May 29 and the influence of gravitation on light". In: *The Observatory* 42 (Mar. 1919), pp. 119–122 (cit. on pp. 2, 65, 68).

[10] Michael Way and Harry Nussbaumer. "Lemaître's Hubble relationship". In: *Physics Today* 64.8 (Aug. 2011), pp. 8–8. ISSN: 0031-9228. DOI: 10.1063/PT.3.1194 (cit. on p. 2).

[11] Edwin Hubble. “A relation between distance and radial velocity among extra-galactic nebulae”. In: *Proceedings of the National Academy of Sciences* 15.3 (1929), pp. 168–173. DOI: 10.1073/pnas.15.3.168 (cit. on p. 2).

[12] Leah Poffenberger. “February 8, 1967: Joseph Weber submits first gravitational wave paper to Physical Review Letters”. In: *APS News* (Feb. 2020). URL: https://www.aps.org/apsnews/2020/02/joseph-weber-first-gravitational-wave (cit. on p. 2).

[13] Harald Lück. *The upgrade of GEO 600. arXiv preprint* arXiv:1004.0338 [gr-qc]. 2010. URL: https://arxiv.org/abs/1004.0338 (cit. on p. 2).

[14] Albert Einstein. “Zur allgemeinen Relativitätstheorie”. In: *Sitzungsberichte der Preußischen Akademie der Wissenschaften zu Berlin* (1915). Plenary session of November 4, 1915, pp. 778–786. URL: https://einsteinpapers.press.princeton.edu/vol6-trans/110 (cit. on p. 3).

[15] Albert Einstein. “Zur allgemeinen Relativitätstheorie (Nachtrag)”. In: *Sitzungsberichte der Königlich Preußischen Akademie der Wissenschaften zu Berlin* (1915). Plenary session of November 11, 1915, pp. 799–801. URL: https://einsteinpapers.press.princeton.edu/vol6-trans/120 (cit. on p. 3).

[16] Albert Einstein. “Erklärung der Perihelbewegung des Merkur aus der allgemeinen Relativitätstheorie”. In: *Sitzungsberichte der Königlich Preußischen Akademie der Wissenschaften zu Berlin* (1915). Plenary session of November 18, 1915, pp. 831–839. URL: https://einsteinpapers.press.princeton.edu/vol6-trans/124 (cit. on p. 3).

[17] Albert Einstein. “Die Feldgleichungen der Gravitation”. In: *Sitzungsberichte der Königlich Preußischen Akademie der Wissenschaften zu Berlin* (1915). Plenary session of November 25, 1915; published December 2, 1915, pp. 844–847. URL: https://einsteinpapers.press.princeton.edu/vol6-trans/129 (cit. on p. 3).

[18] Max Planck Institute for Radio Astronomy. *Confirming Einstein’s Most Fortunate Thought: Radio astronomers use a dance of three exotic stars to test the universality of free fall.* Press release. Accessed: 2025-07-23. June 2020. URL: https://www.mpifr-bonn.mpg.de/pressreleases/2020/5 (cit. on p. 3).

[19] Abraham Pais. *‘Subtle is the Lord . . . ’: The Science and the Life of Albert Einstein.* Oxford University Press, Aug. 2005. ISBN: 9780192806727. DOI: 10.1093/oso/9780192806727.001.0001 (cit. on p. 3).

[20] Stefano Salvia. “From Archimedean Hydrostatics to Post-Aristotelian Mechanics: Galileo’s Early Manuscripts ’De motu antiquiora’ (ca. 1590)”. In: *Physics in Perspective* 19.2 (2017), pp. 105–150. DOI: 10.1007/S00016-017-0202-Y (cit. on p. 3).

[21] G. Galilei. *Discorsi e dimostrazioni matematiche intorno a due nuove scienze attenenti alla mecanica ed i movimenti locali.* Gli Elsevirii, 1638. URL: https://books.google.es/books?id=zB5aAAAAcAAJ (cit. on p. 3).

[22] R. V. Eötvös, D. Pekár, and E. Fekete. “Beiträge zum Gesetze der Proportionalität von Trägheit und Gravität”. In: *Annalen der Physik* 373.9 (1922), pp. 11–66. DOI: 10.1002/andp.19223730903 (cit. on p. 4).

[23] P. G. Roll, R. Krotkov, and R. H. Dicke. “The equivalence of inertial and passive gravitational mass”. In: *Annals of Physics* 26.3 (1964), pp. 442–517. DOI: 10.1016/0003-4916(64)90259-3 (cit. on p. 4).

[24] Pierre Touboul et al. "MICROSCOPE Mission: Final Results of the Test of the Equivalence Principle". In: *Phys. Rev. Lett.* 129.12 (Sept. 2022), p. 121102. DOI: 10.1103/PhysRevLett.129.121102 (cit. on p. 4).

[25] Peter Galison. "Minkowski's Space-Time: From Visual Thinking to the Absolute World". In: *Historical Studies in the Physical and Biological Sciences* 30.1 (1999), pp. 95–130. URL: https://projects.iq.harvard.edu/sites/projects.iq.harvard.edu/files/andrewhsmith/files/galison_-_minkowski.pdf (cit. on p. 5).

[26] Euclid. *The Elements, Book I.* Ed. by David E. Joyce. http://aleph0.clarku.edu/~djoyce/java/elements/bookI/bookI.html. Accessed: 2025-07-23. 1996 (cit. on p. 5).

[27] Roberto Bonola. *Non-Euclidean Geometry: A Critical and Historical Study of Its Development.* Courier Corporation, 1955 (cit. on p. 5).

[28] Nikolai I. Lobachevsky. *Geometrical Researches on the Theory of Parallels.* University of Texas, 1891 (cit. on p. 5).

[29] János Bolyai. *The Science Absolute of Space.* Vol. 3. The Neomon, 1896 (cit. on p. 5).

[30] Olgica Milenković and Coordinated Science Laboratory. *Milenkovic gaining biological insights by analyzing data embedded in non-Euclidean spaces.* News, Coordinated Science Laboratory, University of Illinois. Aug. 2023. URL: https://csl.illinois.edu/news-and-media/58879 (cit. on p. 5).

[31] Masoud Faraki, Maziar Palhang, and Conrad Sanderson. "Log-Euclidean Bag of Words for Human Action Recognition". In: *IET Computer Vision* 9 (June 2014). DOI: 10.1049/iet-cvi.2014.0018 (cit. on p. 6).

[32] John M. Lee. *Introduction to Smooth Manifolds.* 2nd ed. Vol. 218. Graduate Texts in Mathematics. Springer, 2013 (cit. on p. 7).

[33] Dwight E. Neuenschwander. *Tensor Calculus for Physics: A Concise Guide.* Baltimore, MD: Johns Hopkins University Press, 2014. ISBN: 978-1421415659 (cit. on pp. 8, 9).

[34] Peter Collier. *A Most Incomprehensible Thing: Notes Towards a Very Gentle Introduction to the Mathematics of Relativity.* 3rd. Kindle Edition, version 108. Incomprehensible Books, 2017 (cit. on pp. 8, 11, 13, 14, 19, 21–26, 29, 30).

[35] Alexandre Le Tiec and Jérôme Novak. "Theory of Gravitational Waves". In: *An Overview of Gravitational Waves.* WORLD SCIENTIFIC, Feb. 2017, pp. 1–41. ISBN: 9789813141766. DOI: 10.1142/9789813141766_0001 (cit. on pp. 8, 37–40).

[36] David Tong. *Lectures on General Relativity.* http://www.damtp.cam.ac.uk/user/tong/gr.html. 2012 (cit. on pp. 8, 23, 32, 34, 36, 39).

[37] Bernard F. Schutz. *A First Course in General Relativity.* 2nd. Cambridge University Press, 2009. ISBN: 9780521887052 (cit. on p. 8).

[38] Sadri Hassani. *Mathematical Physics: A Modern Introduction to Its Foundations.* 2nd. Springer, 2009. ISBN: 9780387985794 (cit. on p. 10).

[39] James B. Hartle. *Gravity: An Introduction to Einstein's General Relativity.* San Francisco: Addison-Wesley, 2003. ISBN: 978-0-8053-8662-2 (cit. on p. 10).

[40] Leonard Susskind and André Cabannes. *General Relativity: The Theoretical Minimum.* Penguin Books, 2023. ISBN: 9780241562581 (cit. on pp. 12, 15).

[41] Tomoki Watanabe and Mitsuo J. Hayashi. *General Relativity with Torsion*. 2004. arXiv: gr-qc/0409029 [gr-qc]. URL: https://arxiv.org/abs/gr-qc/0409029 (cit. on p. 12).

[42] Fred the Oyster. *Parallel Transport*. Wikimedia Commons. Sept. 2014. URL: https://ca.m.wikipedia.org/wiki/Fitxer:Parallel_Transport.svg (cit. on p. 14).

[43] Sean M. Carroll. *Spacetime and Geometry: An Introduction to General Relativity*. Lecture notes version available at https://arxiv.org/abs/gr-qc/9712019. San Francisco: Addison Wesley, 2004. ISBN: 978-0-8053-8732-2 (cit. on pp. 16, 21, 25, 30–32, 34).

[44] OMEI, Department of Mathematics, University of Cape Town. *Properties of the Riemann curvature tensor (Chapter 6, Section 10)*. http://math_research.uct.ac.za/omei/gr/chap6/node10.html. Accessed: 2025-08-10. (cit. on p. 19).

[45] Johann Colombano-Rut. *General Relativity Loud and Clear*. Nature Loves Math (blog). Jan. 2017. URL: https://www.naturelovesmath.com/en/mathematical-physics/general-relativity-loud-and-clear/ (cit. on p. 21).

[46] Philip Beltracchi and Paolo Gondolo. "An exact time-dependent interior Schwarzschild solution". In: *Phys. Rev. D* 99.8 (2019), p. 084021. DOI: 10.1103/PhysRevD.99.084021 (cit. on p. 23).

[47] J. R. Mureika. *Schwarzschild Solution – In-Class Notes for PHYS 471: General Relativity and Cosmology*. https://jmureika.lmu.build/PHYS471/InClass/SchwarzschildSolution.pdf. Lecture notes, accessed 2025-08-29. 2019 (cit. on p. 24).

[48] Roger Penrose. "Gravitational Collapse and Space-Time Singularities". In: *Phys. Rev. Lett.* 14.3 (1965), pp. 57–59. DOI: 10.1103/PhysRevLett.14.57 (cit. on p. 24).

[49] Fay Dowker. *Black Holes*. Course notes. Accessed: 2025-08-29. London, United Kingdom, 2015. URL: https://www.sns.ias.edu/~bkocsis/Assets/Teaching/BH2015/dowker.pdf (cit. on p. 25).

[50] The Nobel Committee for Physics. *Figure 2: Cross-section of a black hole (from Nobel Prize in Physics 2020)*. Accessed: 2025-09-29. 2020. URL: https://www.nobelprize.org/uploads/2020/10/fig2-phy-en-cross-section.pdf (cit. on p. 26).

[51] R. Abbott et al. "GW190521: A Binary Black Hole Merger with a Total Mass of 150 solar masses". In: *Phys. Rev. Lett.* 125.10 (Sept. 2020), p. 101102. ISSN: 1079-7114. DOI: 10.1103/PhysRevLett.125.101102 (cit. on p. 27).

[52] R. Abbott et al. *GW231123: a Binary Black Hole Merger with Total Mass of 190–265 solar masses*. 2025. arXiv: 2507.08219 [astro-ph.HE]. URL: https://arxiv.org/abs/2507.08219 (cit. on p. 27).

[53] Bernard Carr et al. "Constraints on primordial black holes". In: *Reports on Progress in Physics* 84.11 (2021), p. 116902. DOI: 10.1088/1361-6633/ac1e31 (cit. on p. 27).

[54] Albert Einstein. "Approximative Integration of the Field Equations of Gravitation". In: *Sitzungsberichte der Königlich Preußischen Akademie der Wissenschaften (Berlin)* vol. 1916 (1916), pp. 688–696 (cit. on p. 28).

[55] Galina Weinstein. *Einstein's Discovery of Gravitational Waves 1916-1918*. 2016. arXiv: 1602.04040 [physics.hist-ph] (cit. on p. 28).

[56] Dean Rickles. "Behind the scenes of the 1957 Chapel Hill Conference on the role of gravitation in physics". In: *Journal & Proceedings of the Royal Society of New South Wales* 154.Part 2 (2021), pp. 224–228. ISSN: 0035-9173/21/020224-05. URL: https://www.royalsoc.org.au/wp-content/uploads/2021/12/154-2-Rickles.pdf (cit. on p. 28).

[57] Michele Maggiore. *Gravitational Waves: Volume 1: Theory and Experiments*. Oxford University Press, Oct. 2007. ISBN: 9780198570745. DOI: 10.1093/acprof:oso/9780198570745.001.0001 (cit. on pp. 28, 29, 31–33, 38, 41, 42, 55).

[58] R. A. Hulse and J. H. Taylor. "Discovery of a pulsar in a binary system." In: *The Astrophysical Journal Letters* 195 (Jan. 1975), pp. L51–L53. DOI: 10.1086/181708 (cit. on p. 28).

[59] Saúl Ramos-Sánchez. *Relatividad: para futuros físicos*. Mexico City, Mexico: CopIt-arXives, 2018. URL: https://copitarxives.fisica.unam.mx/LT0001ES/LT0001ES.pdf (cit. on p. 32).

[60] Valeria Ferrari, Leonardo Gualtieri, and Paolo Pani. *General Relativity and its Applications: Black Holes, Compact Stars and Gravitational Waves*. 1st. CRC Press, 2020. DOI: 10.1201/9780429491405 (cit. on pp. 33, 35).

[61] Tan Liu, Wen-Fan Feng, and Zong-Kuan Guo. "Influence of cosmic expansion on gravitational waveforms". In: *Phys. Rev. D* 112.2 (July 2025). ISSN: 2470-0029. DOI: 10.1103/wg63-t8bp (cit. on p. 38).

[62] Kip S. Thorne. *Gravitational Radiation: Astrophysical and Cosmological Sources of Gravitational Waves, Lorentz Lectures*. Lecture notes, Lorentz Chair, University of Leiden. 2009. URL: https://www.lorentz.leidenuniv.nl/lorentzchair/thorne/Thorne2.pdf (cit. on p. 38).

[63] Luc Blanchet. "Gravitational Radiation from Post-Newtonian Sources and Inspiralling Compact Binaries". In: *Living Reviews in Relativity* 17.1 (2014), p. 2. DOI: 10.12942/lrr-2014-2 (cit. on p. 40).

[64] Virgo Collaboration. *Astrophysical Sources of Gravitational Waves*. Accessed: 2025-07-17. 2024. URL: https://www.virgo-gw.eu/science/gw-universe/astrophysical-sources-of-gw/ (cit. on p. 41).

[65] B. P. Abbott et al. "GWTC-1: A Gravitational-Wave Transient Catalog of Compact Binary Mergers Observed by LIGO and Virgo during the First and Second Observing Runs". In: *Phys. Rev. X* 9 (3 Sept. 2019), p. 031040. DOI: 10.1103/PhysRevX.9.031040 (cit. on p. 42).

[66] B. P. Abbott et al. "GW170817: Observation of Gravitational Waves from a Binary Neutron Star Inspiral". In: *Phys. Rev. Lett.* 119 (16 Oct. 2017), p. 161101. DOI: 10.1103/PhysRevLett.119.161101 (cit. on pp. 42, 52).

[67] R. Abbott et al. "GWTC-2: Compact Binary Coalescences Observed by LIGO and Virgo during the First Half of the Third Observing Run". In: *Phys. Rev. X* 11 (2 June 2021), p. 021053. DOI: 10.1103/PhysRevX.11.021053 (cit. on p. 42).

[68] R. Abbott et al. "GWTC-3: Compact Binary Coalescences Observed by LIGO and Virgo during the Second Part of the Third Observing Run". In: *Phys. Rev. X* 13 (4 Dec. 2023), p. 041039. DOI: 10.1103/PhysRevX.13.041039 (cit. on p. 42).

[69] The LIGO Scientific Collaboration et al. *GWTC-4.0: Updating the Gravitational-Wave Transient Catalog with Observations from the First Part of the Fourth LIGO-Virgo-KAGRA Observing Run*. 2025. arXiv: 2508.18082 [gr-qc] (cit. on p. 42).

[70] The LIGO Scientific Collaboration et al. *GWTC-5.0: An Introduction to Version 5.0 of the Gravitational-Wave Transient Catalog*. 2026. arXiv: 2605.27223 [gr-qc] (cit. on pp. 42, 53, 84, 109).

[71] The LIGO Scientific Collaboration et al. *GWTC-5.0: Methods for Identifying and Characterizing Gravitational-wave Transients*. 2026. arXiv: 2605.27224 [gr-qc] (cit. on pp. 42, 57, 91, 96).

[72] Rahul Srinivasan et al. "Simulation-based population inference of LISA's Galactic binaries: Bypassing the global fit". In: *Phys. Rev. D* 112 (10 Nov. 2025), p. 103043. DOI: 10.1103/shym-w46f (cit. on p. 42).

[73] Anthony Mezzacappa and Michele Zanolin. *Gravitational Waves from Neutrino-Driven Core Collapse Supernovae: Predictions, Detection, and Parameter Estimation*. 2024. arXiv: 2401.11635 [astro-ph.HE] (cit. on p. 43).

[74] M B Hindmarsh and T W B Kibble. "Cosmic strings". In: *Reports on Progress in Physics* 58.5 (May 1995), p. 477. DOI: 10.1088/0034-4885/58/5/001 (cit. on p. 43).

[75] Marco Drago et al. "coherent WaveBurst, a pipeline for unmodeled gravitational-wave data analysis". In: *SoftwareX* 14 (2021), p. 100678. ISSN: 2352-7110. DOI: https://doi.org/10.1016/j.softx.2021.100678 (cit. on p. 43).

[76] Magdalena Sieniawska and Michał Bejger. "Continuous Gravitational Waves from Neutron Stars: Current Status and Prospects". In: *Universe* 5.11 (2019). ISSN: 2218-1997. DOI: 10.3390/universe5110217 (cit. on p. 43).

[77] Sylvia J. Zhu et al. "Characterizing the continuous gravitational-wave signal from boson clouds around Galactic isolated black holes". In: *Phys. Rev. D* 102 (6 Sept. 2020), p. 063020. DOI: 10.1103/PhysRevD.102.063020 (cit. on p. 43).

[78] Karl Wette. "Searches for continuous gravitational waves from neutron stars: A twenty-year retrospective". In: *Astroparticle Physics* 153 (2023), p. 102880. ISSN: 0927-6505. DOI: https://doi.org/10.1016/j.astropartphys.2023.102880 (cit. on p. 43).

[79] Nick van Remortel, Kamiel Janssens, and Kevin Turbang. "Stochastic gravitational wave background: Methods and implications". In: *Progress in Particle and Nuclear Physics* 128 (2023), p. 104003. ISSN: 0146-6410. DOI: https://doi.org/10.1016/j.ppnp.2022.104003 (cit. on p. 43).

[80] The LIGO Scientific Collaboration et al. *Upper Limits on the Isotropic Gravitational-Wave Background from the first part of LIGO, Virgo, and KAGRA's fourth Observing Run*. 2025. arXiv: 2508.20721 [gr-qc] (cit. on p. 43).

[81] Odylio Denys Aguiar. "Past, present and future of the Resonant-Mass gravitational wave detectors". In: *Research in Astronomy and Astrophysics* 11.1 (Jan. 2011), p. 1. DOI: 10.1088/1674-4527/11/1/001 (cit. on p. 44).

[82] D. V. Martynov et al. "Sensitivity of the Advanced LIGO detectors at the beginning of gravitational wave astronomy". In: *Phys. Rev. D* 93 (11 June 2016), p. 112004. DOI: 10.1103/PhysRevD.93.112004 (cit. on p. 44).

[83] Miquel Trias Cornellana. “Gravitational wave observation of compact binaries: detection, parameter estimation and template accuracy”. PhD thesis. Universitat de les Illes Balears, 2011 (cit. on pp. 45, 47).

[84] J. Abadie et al. “Calibration of the LIGO gravitational wave detectors in the fifth science run”. In: *Nuclear Instruments and Methods in Physics Research Section A: Accelerators, Spectrometers, Detectors and Associated Equipment* 624.1 (2010), pp. 223–240. ISSN: 0168-9002. DOI: https://doi.org/10.1016/j.nima.2010.07.089 (cit. on p. 46).

[85] Piotr Jaranowski, Andrzej Królak, and Bernard F. Schutz. “Data analysis of gravitational-wave signals from spinning neutron stars: The signal and its detection”. In: *Phys. Rev. D* 58 (6 Aug. 1998), p. 063001. DOI: 10.1103/PhysRevD.58.063001 (cit. on pp. 46, 47).

[86] J.D.E. Creighton and W.G. Anderson. *Gravitational-Wave Physics and Astronomy: An Introduction to Theory, Experiment and Data Analysis*. Wiley Series in Cosmology. Wiley, 2011. ISBN: 9783527408863. URL: https://books.google.es/books?id=jLRLVaR97AUC (cit. on p. 48).

[87] The LIGO Scientific Collaboration et al. “Advanced LIGO”. In: *Classical and Quantum Gravity* 32.7 (Mar. 2015), p. 074001. DOI: 10.1088/0264-9381/32/7/074001 (cit. on pp. 49, 52).

[88] E. Capote et al. “Advanced LIGO detector performance in the fourth observing run”. In: *Phys. Rev. D* 111 (6 Mar. 2025), p. 062002. DOI: 10.1103/PhysRevD.111.062002 (cit. on pp. 49, 51).

[89] LIGO O4 Detector Collaboration. “Broadband Quantum Enhancement of the LIGO Detectors with Frequency-Dependent Squeezing”. In: *Phys. Rev. X* 13.4 (2023), p. 041021. DOI: 10.1103/PhysRevX.13.041021 (cit. on p. 49).

[90] R Kubo. “The fluctuation-dissipation theorem”. In: *Reports on Progress in Physics* 29.1 (Jan. 1966), p. 255. DOI: 10.1088/0034-4885/29/1/306 (cit. on p. 49).

[91] P. B. Covas et al. “Identification and mitigation of narrow spectral artifacts that degrade searches for persistent gravitational waves in the first two observing runs of Advanced LIGO”. In: *Phys. Rev. D* 97 (8 Apr. 2018), p. 082002. DOI: 10.1103/PhysRevD.97.082002 (cit. on p. 50).

[92] S Soni et al. “LIGO Detector Characterization in the first half of the fourth Observing run”. In: *Classical and Quantum Gravity* 42.8 (Apr. 2025), p. 085016. DOI: 10.1088/1361-6382/adc4b6 (cit. on pp. 50, 51).

[93] J Glanzer et al. “Data quality up to the third observing run of advanced LIGO: Gravity Spy glitch classifications”. In: *Classical and Quantum Gravity* 40.6 (Feb. 2023), p. 065004. DOI: 10.1088/1361-6382/acb633 (cit. on p. 50).

[94] F Acernese et al. “Advanced Virgo: a second-generation interferometric gravitational wave detector”. In: *Classical and Quantum Gravity* 32.2 (Dec. 2014), p. 024001. DOI: 10.1088/0264-9381/32/2/024001 (cit. on p. 52).

[95] B. P. Abbott et al. “GW170814: A Three-Detector Observation of Gravitational Waves from a Binary Black Hole Coalescence”. In: *Phys. Rev. Lett.* 119 (14 Oct. 2017), p. 141101. DOI: 10.1103/PhysRevLett.119.141101 (cit. on p. 52).

[96] T. Akutsu et al. "KAGRA: 2.5 generation interferometric gravitational wave detector". In: *Nature Astronomy* 3.1 (2019), pp. 35–40. DOI: 10.1038/s41550-018-0658-y (cit. on p. 52).

[97] B Willke et al. "The GEO 600 gravitational wave detector". In: *Classical and Quantum Gravity* 19.7 (Mar. 2002), p. 1377. DOI: 10.1088/0264-9381/19/7/321 (cit. on p. 52).

[98] Yakima Herald. *Hanford Observatory's place in scientific history recognized*. Accessed: 2026-05-12. 2018. URL: https://www.yakimaherald.com/news/local/hanford-observatory-s-place-in-scientific-history-recognized/article_0c37c614-6a60-11e8-9255-1fb7cfdab9fd.html (cit. on p. 53).

[99] LIGO Lab, Caltech. *LIGO Livingston*. Accessed: 2025-05-12. 2015. URL: https://www.ligo.caltech.edu/image/ligo20150731b (cit. on p. 53).

[100] eoPortal. *Virgo Interferometer*. Accessed: 2025-05-12. 2024. URL: https://www.eoportal.org/other-space-activities/virgo-interferometer#virgo-interferometer (cit. on p. 53).

[101] Academia Sinica Institute of Physics, Gravitational Wave Lab. *KAGRA*. Accessed: 2025-05-12. URL: https://www.phys.sinica.edu.tw/~gwlab/r_kagra.htm (cit. on p. 53).

[102] GEO600 Collaboration. *GEO600 Gravitational-Wave Detector*. Accessed: 2025-05-12. URL: https://www.geo600.org/ (cit. on p. 53).

[103] Adrian Abac et al. "The Science of the Einstein Telescope". In: *Journal of Cosmology and Astroparticle Physics* 2026.03 (Mar. 2026), p. 081. DOI: 10.1088/1475-7516/2026/03/081 (cit. on p. 54).

[104] Marica Branchesi et al. "Science with the Einstein Telescope: a comparison of different designs". In: *Journal of Cosmology and Astroparticle Physics* 2023.07 (July 2023), p. 068. DOI: 10.1088/1475-7516/2023/07/068 (cit. on p. 54).

[105] Matthew Evans et al. *A Horizon Study for Cosmic Explorer: Science, Observatories, and Community*. 2021. arXiv: 2109.09882 [astro-ph.IM] (cit. on p. 54).

[106] C. S. Unnikrishnan. *LIGO-India: A Decadal Assessment on Its Scope, Relevance, Progress, and Future*. 2023. arXiv: 2301.07522 [astro-ph.IM] (cit. on p. 54).

[107] Monica Colpi et al. *LISA Definition Study Report*. 2024. arXiv: 2402.07571 [astro-ph.CO] (cit. on p. 54).

[108] Seiji Kawamura et al. "Current status of space gravitational wave antenna DECIGO and B-DECIGO". In: *Progress of Theoretical and Experimental Physics* 2021.5 (May 2021), 05A105. ISSN: 2050-3911. DOI: 10.1093/ptep/ptab019 (cit. on p. 54).

[109] Yungui Gong, Jun Luo, and Bin Wang. "Concepts and status of Chinese space gravitational wave detection projects". In: *Nature Astronomy* 5.9 (2021), pp. 881–889. DOI: 10.1038/s41550-021-01480-3 (cit. on p. 54).

[110] P. Welch. "The use of fast Fourier transform for the estimation of power spectra: A method based on time averaging over short, modified periodograms". In: *IEEE Transactions on Audio and Electroacoustics* 15.2 (1967), pp. 70–73. DOI: 10.1109/TAU.1967.1161901 (cit. on p. 55).

[111] Éanna É. Flanagan and Scott A. Hughes. "Measuring gravitational waves from binary black hole coalescences. I. Signal to noise for inspiral, merger, and ringdown". In: *Phys. Rev. D* 57 (8 Apr. 1998), pp. 4535–4565. DOI: 10.1103/PhysRevD.57.4535 (cit. on p. 57).

[112] Samantha A Usman et al. "The PyCBC search for gravitational waves from compact binary coalescence". In: *Classical and Quantum Gravity* 33.21 (Oct. 2016), p. 215004. DOI: 10.1088/0264-9381/33/21/215004 (cit. on p. 57).

[113] Cody Messick et al. "Analysis framework for the prompt discovery of compact binary mergers in gravitational-wave data". In: *Phys. Rev. D* 95 (4 Feb. 2017), p. 042001. DOI: 10.1103/PhysRevD.95.042001 (cit. on p. 57).

[114] Javier Roulet et al. "Template bank for compact binary coalescence searches in gravitational wave data: A general geometric placement algorithm". In: *Phys. Rev. D* 99 (12 June 2019), p. 123022. DOI: 10.1103/PhysRevD.99.123022 (cit. on p. 57).

[115] Luc Blanchet. "Post-Newtonian Expansion of Gravitational Radiation". In: *Progress of Theoretical Physics Supplement* 136 (Apr. 1999), pp. 146–157. ISSN: 0375-9687. DOI: 10.1143/PTPS.136.146 (cit. on p. 58).

[116] Carlos Palenzuela. "Introduction to Numerical Relativity". In: *Frontiers in Astronomy and Space Sciences* Volume 7 - 2020 (2020). ISSN: 2296-987X. DOI: 10.3389/fspas.2020.00058 (cit. on p. 58).

[117] Kostas D. Kokkotas and Bernd G. Schmidt. "Quasi-Normal Modes of Stars and Black Holes". In: *Living Reviews in Relativity* 2.1 (1999), p. 2. ISSN: 1433-8351. DOI: 10.12942/lrr-1999-2 (cit. on p. 58).

[118] Lorenzo Pompili et al. "Laying the foundation of the effective-one-body waveform models SEOBNRv5: Improved accuracy and efficiency for spinning nonprecessing binary black holes". In: *Phys. Rev. D* 108 (12 Dec. 2023), p. 124035. DOI: 10.1103/PhysRevD.108.124035 (cit. on p. 58).

[119] Antoni Ramos-Buades et al. *SEOBNRv5PHM: Next generation of accurate and efficient multipolar precessing-spin effective-one-body waveforms for binary black holes.* 2023. arXiv: 2303.18046 [gr-qc] (cit. on p. 58).

[120] Geraint Pratten et al. "Computationally efficient models for the dominant and subdominant harmonic modes of precessing binary black holes". In: *Phys. Rev. D* 103 (10 May 2021), p. 104056. DOI: 10.1103/PhysRevD.103.104056 (cit. on pp. 58, 82).

[121] Vijay Varma et al. "Surrogate model of hybridized numerical relativity binary black hole waveforms". In: *Phys. Rev. D* 99 (6 Mar. 2019), p. 064045. DOI: 10.1103/PhysRevD.99.064045 (cit. on p. 58).

[122] Eric Thrane and Colm Talbot. "An introduction to Bayesian inference in gravitational-wave astronomy: Parameter estimation, model selection, and hierarchical models". In: *Publications of the Astronomical Society of Australia* 36 (2019), e010. DOI: 10.1017/pasa.2019.2 (cit. on pp. 58–60, 62).

[123] Thibault Damour. "Coalescence of two spinning black holes: An effective one-body approach". In: *Phys. Rev. D* 64 (12 Nov. 2001), p. 124013. DOI: 10.1103/PhysRevD.64.124013 (cit. on p. 59).

[124] Patricia Schmidt, Frank Ohme, and Mark Hannam. “Towards models of gravitational waveforms from generic binaries: II. Modelling precession effects with a single effective precession parameter”. In: *Phys. Rev. D* 91 (2 Jan. 2015), p. 024043. DOI: 10.1103/PhysRevD.91.024043 (cit. on pp. 59, 131).

[125] K. Chatziioannou et al. *Compact binary coalescences: gravitational-wave astronomy with ground-based detectors*. 2024. arXiv: 2409.02037 [gr-qc] (cit. on pp. 59, 61, 63, 130).

[126] Stephen Fairhurst. “Triangulation of gravitational wave sources with a network of detectors”. In: *New Journal of Physics* 11.12 (Dec. 2009), p. 123006. DOI: 10.1088/1367-2630/11/12/123006 (cit. on pp. 59, 77).

[127] Harold Jeffreys. *Theory of Probability*. Oxford University Press, Aug. 1998. ISBN: 9780198503682. DOI: 10.1093/oso/9780198503682.001.0001 (cit. on p. 60).

[128] W. K. Hastings. “Monte Carlo sampling methods using Markov chains and their applications”. In: *Biometrika* 57.1 (Apr. 1970), pp. 97–109. ISSN: 0006-3444. DOI: 10.1093/biomet/57.1.97 (cit. on p. 61).

[129] Nicholas Metropolis et al. “Equation of State Calculations by Fast Computing Machines”. In: *The Journal of Chemical Physics* 21.6 (June 1953), pp. 1087–1092. ISSN: 0021-9606. DOI: 10.1063/1.1699114 (cit. on p. 61).

[130] John Skilling. “Nested Sampling for Bayesian Computations”. In: *Bayesian Statistics 8: Proceedings of the Eighth Valencia International Meeting June 2–6, 2006*. Ed. by J M Bernardo et al. Oxford University Press, July 2007. ISBN: 9780199214655. DOI: 10.1093/oso/9780199214655.003.0019 (cit. on p. 61).

[131] Joshua S Speagle. “dynesty: a dynamic nested sampling package for estimating Bayesian posteriors and evidences”. In: *Monthly Notices of the Royal Astronomical Society* 493.3 (Apr. 2020), pp. 3132–3158. ISSN: 0035-8711. DOI: 10.1093/mnras/staa278 (cit. on pp. 61, 63, 91).

[132] Farhan Feroz et al. “Importance Nested Sampling and the MultiNest Algorithm”. In: *The Open Journal of Astrophysics* 2 (Nov. 2019). DOI: 10.21105/astro.1306.2144 (cit. on p. 63).

[133] F. Feroz, M. P. Hobson, and M. Bridges. “MultiNest: an efficient and robust Bayesian inference tool for cosmology and particle physics”. In: *Monthly Notices of the Royal Astronomical Society* 398.4 (Oct. 2009), pp. 1601–1614. ISSN: 0035-8711. DOI: 10.1111/j.1365-2966.2009.14548.x (cit. on p. 63).

[134] Gregory Ashton et al. “Bilby: A User-friendly Bayesian Inference Library for Gravitational-wave Astronomy”. In: *The Astrophysical Journal Supplement Series* 241.2 (Apr. 2019), p. 27. DOI: 10.3847/1538-4365/ab06fc (cit. on p. 63).

[135] Matthias Bartelmann. “Gravitational lensing”. In: *Classical and Quantum Gravity* 27.23 (Nov. 2010), p. 233001. DOI: 10.1088/0264-9381/27/23/233001 (cit. on p. 65).

[136] Hans C. Ohanian. “On the Focusing of Gravitational Radiation”. In: *International Journal of Theoretical Physics* 9.6 (June 1974), pp. 425–437. ISSN: 1572-9575. DOI: 10.1007/BF01810927 (cit. on p. 65).

[137] Jose María Ezquiaga et al. “Phase effects from strong gravitational lensing of gravitational waves”. In: *Phys. Rev. D* 103 (6 Mar. 2021), p. 064047. DOI: 10.1103/PhysRevD.103.064047. URL: https://link.aps.org/doi/10.1103/PhysRevD.103.064047 (cit. on pp. 65, 73, 76).

[138] Joachim Wambsganss. “Gravitational Lensing in Astronomy”. In: *Living Reviews in Relativity* 1.1 (Dec. 1998), p. 12. ISSN: 1433-8351. DOI: 10.12942/lrr-1998-12 (cit. on p. 65).

[139] Zhiwei Chen and Youjun Lu. “Gravitational Lensing of Gravitational Waves from Astrophysical Sources: Theory, Detection, and Applications”. In: *Research in Astronomy and Astrophysics* 26.6 (Apr. 2026), p. 062001. DOI: 10.1088/1674-4527/ae4bad (cit. on pp. 65, 72, 73, 77, 78).

[140] LIGO Scientific Collaboration. *Gravitational-wave lensing illustration*. https://ligo.org/multimedia/. Accessed: 2026-06-01 (cit. on p. 66).

[141] Anna Liu et al. “Exploring the hidden Universe: a novel phenomenological approach for recovering arbitrary gravitational-wave millilensing configurations”. In: *Monthly Notices of the Royal Astronomical Society* 525.3 (Nov. 2023), pp. 4149–4160. ISSN: 0035-8711. DOI: 10.1093/mnras/stad1302 (cit. on p. 66).

[142] Ryuichi Takahashi and Takashi Nakamura. “Wave Effects in the Gravitational Lensing of Gravitational Waves from Chirping Binaries”. In: *The Astrophysical Journal* 595.2 (Oct. 2003), p. 1039. DOI: 10.1086/377430 (cit. on pp. 66, 72, 74).

[143] Graham P. Smith et al. “Multi-messenger gravitational lensing”. In: *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences* 383.2295 (May 2025), p. 20240134. ISSN: 1364-503X. DOI: 10.1098/rsta.2024.0134 (cit. on p. 67).

[144] The LIGO Scientific Collaboration et al. *GWTC-4.0: An Introduction to Version 4.0 of the Gravitational-Wave Transient Catalog*. 2025. arXiv: 2508.18080 [gr-qc] (cit. on pp. 67, 130).

[145] Ken K. Y. Ng et al. “Precise LIGO lensing rate predictions for binary black holes”. In: *Phys. Rev. D* 97 (2 Jan. 2018), p. 023012. DOI: 10.1103/PhysRevD.97.023012 (cit. on pp. 67, 71).

[146] Bernard F. Schutz. “Determining the Hubble Constant from Gravitational Wave Observations”. In: *Nature* 323.6086 (Sept. 1986), pp. 310–311. ISSN: 1476-4687. DOI: 10.1038/323310a0 (cit. on p. 67).

[147] Otto A Hannuksela et al. “Localizing merging black holes with sub-arcsecond precision using gravitational-wave lensing”. In: *Monthly Notices of the Royal Astronomical Society* 498.3 (Oct. 2020), pp. 3395–3402. ISSN: 0035-8711. DOI: 10.1093/mnras/staa2577 (cit. on pp. 67, 77, 78, 110).

[148] Souvik Jana et al. *Probing the nature of dark matter using strongly lensed gravitational waves from binary black holes*. 2024. arXiv: 2408.05290 [astro-ph.CO]. URL: https://arxiv.org/abs/2408.05290 (cit. on p. 67).

[149] Harsh Narola et al. “How well can modified gravitational wave propagation be constrained with strong lensing?” In: *Phys. Rev. D* 109 (8 Apr. 2024), p. 084064. DOI: 10.1103/PhysRevD.109.084064 (cit. on p. 67).

[150] Liang Dai, Tejaswi Venumadhav, and Kris Sigurdson. “Effect of lensing magnification on the apparent distribution of black hole mergers”. In: *Phys. Rev. D* 95 (4 Feb. 2017), p. 044011. DOI: 10.1103/PhysRevD.95.044011 (cit. on p. 67).

[151] Zhiwei Chen et al. “Enhanced Localization of Dark Lensed Gravitational-wave Events Enables Host Galaxy Identification and Precise Cosmological Inference”. In: *The Astrophysical Journal Letters* 993.2 (Nov. 2025), p. L57. DOI: 10.3847/2041-8213/ae1226 (cit. on p. 67).

[152] Matthias Bartelmann and Matteo Maturi. *Weak gravitational lensing*. 2016. arXiv: 1612.06535 [astro-ph.CO] (cit. on p. 68).

[153] Judit Prat and David Bacon. “Weak gravitational lensing”. In: *Encyclopedia of Astrophysics (First Edition)*. Ed. by Ilya Mandel. First Edition. Oxford: Elsevier, 2026, pp. 508–537. ISBN: 978-0-443-21440-0. DOI: https://doi.org/10.1016/B978-0-443-21439-4.00071-7 (cit. on p. 68).

[154] Prasenjit Saha et al. *Essentials of strong gravitational lensing*. 2024. arXiv: 2401.04165 [astro-ph.CO] (cit. on pp. 68, 70–72, 75).

[155] Michael Sachs. *Gravitational-lensing-angles.png: Diagram of angles involved in gravitational lensing*. Wikimedia Commons. Accessed: 2026-06-12. June 2008. URL: https://commons.wikimedia.org/wiki/File:Gravitational-lensing-angles.png (cit. on p. 69).

[156] W. L. Burke. “Multiple Gravitational Imaging by Distributed Masses”. In: *The Astrophysical Journal Letters* 244 (Feb. 1981), p. L1. DOI: 10.1086/183466 (cit. on p. 69).

[157] Irwin I. Shapiro. “Fourth Test of General Relativity”. In: *Phys. Rev. Lett.* 13 (26 Dec. 1964), pp. 789–791. DOI: 10.1103/PhysRevLett.13.789 (cit. on p. 71).

[158] Oleg Yu Tsupko et al. “An examination of geometrical and potential time delays in gravitational lensing”. In: *Classical and Quantum Gravity* 37.20 (Sept. 2020), p. 205017. DOI: 10.1088/1361-6382/abae86 (cit. on p. 71).

[159] Giovanni Tambalo et al. “Gravitational wave lensing as a probe of halo properties and dark matter”. In: *Phys. Rev. D* 108 (10 Nov. 2023), p. 103529. DOI: 10.1103/PhysRevD.108.103529 (cit. on pp. 72, 75).

[160] R. Kormann, P. Schneider, and M. Bartelmann. “Isothermal elliptical gravitational lens models.” In: *Astronomy and Astrophysics* 284 (Apr. 1994), pp. 285–299 (cit. on p. 73).

[161] Shuo Cao et al. “Limits on the power-law mass and luminosity density profiles of elliptical galaxies from gravitational lensing systems”. In: *Monthly Notices of the Royal Astronomical Society* 461.2 (Sept. 2016), pp. 2192–2199. ISSN: 0035-8711. DOI: 10.1093/mnras/stw932 (cit. on p. 73).

[162] Xiaoyue Cao et al. “Systematic Errors Induced by the Elliptical Power-law model in Galaxy–Galaxy Strong Lens Modeling”. In: *Research in Astronomy and Astrophysics* 22.2 (Feb. 2022), p. 025014. DOI: 10.1088/1674-4527/ac3f2b (cit. on p. 73).

[163] C. M. O’Riordan, S. J. Warren, and D. J. Mortlock. “Galaxy mass profiles from strong lensing II: The elliptical power-law model”. In: *Monthly Notices of the Royal Astronomical Society* 496.3 (Aug. 2020), pp. 3424–3435. DOI: 10.1093/mnras/staa1697 (cit. on p. 73).

[164] Matthew R Gomer and Liliya L R Williams. "Galaxy-lens determination of H0: the effect of the ellipse + shear modelling assumption". In: *Monthly Notices of the Royal Astronomical Society* 504.1 (June 2021), pp. 1340–1354. ISSN: 0035-8711. DOI: 10.1093/mnras/stab930 (cit. on p. 73).

[165] Peter Schneider, Jürgen Ehlers, and Emilio E. Falco. *Gravitational Lenses*. Astronomy and Astrophysics Library. Berlin, Heidelberg: Springer-Verlag, 1992. ISBN: 978-3-540-66506-9. DOI: 10.1007/978-3-662-03758-4 (cit. on p. 74).

[166] The LIGO Scientific Collaboration et al. *GWTC-4.0: Searches for Gravitational-Wave Lensing Signatures*. 2026. arXiv: 2512.16347 [gr-qc] (cit. on pp. 74, 76, 91).

[167] Roger Blandford and Ramesh Narayan. "Fermat's Principle, Caustics, and the Classification of Gravitational Lens Images". In: *The Astrophysical Journal Letters* 310 (Nov. 1986), p. 568. DOI: 10.1086/164709 (cit. on p. 74).

[168] Justin Janquart, Anupreeta More, and Chris Van Den Broeck. "Ordering the confusion: a study of the impact of lens models on gravitational-wave strong lensing detection capabilities". In: *Monthly Notices of the Royal Astronomical Society* 519.2 (Feb. 2023), pp. 2046–2059. ISSN: 0035-8711. DOI: 10.1093/mnras/stac3660 (cit. on p. 76).

[169] Rico K. L. Lo and Ignacio Magaña Hernandez. "Bayesian statistical framework for identifying strongly lensed gravitational-wave signals". In: *Phys. Rev. D* 107 (12 June 2023), p. 123015. DOI: 10.1103/PhysRevD.107.123015 (cit. on pp. 76, 78, 79, 90).

[170] LIGO Scientific Collaboration and Virgo Collaboration and KAGRA Collaboration. *LVK Observing Plans*. https://observing.docs.ligo.org/plan/. Document control number G2002127-v36. 2026 (cit. on p. 77).

[171] Hemantakumar Phurailatpam et al. *ler: LVK (LIGO-Virgo-KAGRA collaboration) event (compact-binary mergers) rate calculator and simulator*. 2026. arXiv: 2407.07526 [astro-ph.IM] (cit. on pp. 81, 83).

[172] Simon Birrer et al. "lenstronomy II: A gravitational lensing software ecosystem". In: *Journal of Open Source Software* 6.62 (2021), p. 3283. DOI: 10.21105/joss.03283. URL: https://doi.org/10.21105/joss.03283 (cit. on p. 81).

[173] LIGO Scientific Collaboration. *A+/O5 strain noise curve projections*. Tech. rep. LIGO-T2500310-v2. LIGO, 2025. URL: https://dcc.ligo.org/T2500310-v2/public (cit. on p. 82).

[174] Marta Colleoni et al. "Fast frequency-domain gravitational waveforms for precessing binaries with a new twist". In: *Phys. Rev. D* 111.10 (2025), p. 104019. DOI: 10.1103/PhysRevD.111.104019 (cit. on p. 82).

[175] R. Abbott et al. "Population of Merging Compact Binaries Inferred Using Gravitational Waves through GWTC-3". In: *Phys. Rev. X* 13 (1 Mar. 2023), p. 011048. DOI: 10.1103/PhysRevX.13.011048 (cit. on p. 83).

[176] Masamune Oguri. "Effect of gravitational lensing on the distribution of gravitational waves from distant binary black hole mergers". In: *Monthly Notices of the Royal Astronomical Society* 480.3 (Nov. 2018), pp. 3842–3855. ISSN: 0035-8711. DOI: 10.1093/mnras/sty2145 (cit. on p. 83).

[177] K. Haris et al. *Identifying strongly lensed gravitational wave signals from binary black hole mergers*. 2018. arXiv: 1807.07062 [gr-qc] (cit. on p. 83).

[178] Yun-Young Choi, Changbom Park, and Michael S. Vogeley. "Internal and Collective Properties of Galaxies in the Sloan Digital Sky Survey". In: *The Astrophysical Journal* 658.2 (Apr. 2007), p. 884. DOI: 10.1086/511060 (cit. on p. 83).

[179] Léon V. E. Koopmans et al. "The Sloan Lens ACS Survey. III. The Structure and Formation of Early-Type Galaxies and Their Evolution since z $\approx$ 1". In: *The Astrophysical Journal* 649.2 (Oct. 2006), p. 599. DOI: 10.1086/505696 (cit. on p. 83).

[180] M. W. Auger et al. "The Sloan Lens ACS Survey. X. Stellar, Dynamical, and Total Mass Correlations of Massive Early-type Galaxies". In: *The Astrophysical Journal* 724.1 (Nov. 2010), p. 511. DOI: 10.1088/0004-637X/724/1/511 (cit. on p. 83).

[181] A. Renske A. C. Wierda et al. "Beyond the Detector Horizon: Forecasting Gravitational-Wave Strong Lensing". In: *The Astrophysical Journal* 921.2 (Nov. 2021), p. 154. DOI: 10.3847/1538-4357/ac1bb4 (cit. on pp. 83, 84).

[182] N. Tessore and R. B. Metcalf. "The elliptical power law profile lens". In: *Astronomy & Astrophysics* 580 (2015), A79. DOI: 10.1051/0004-6361/201526773 (cit. on p. 83).

[183] I. M. Romero-Shaw et al. "Bayesian inference for compact binary coalescences with bilby: validation and application to the first LIGO–Virgo gravitational-wave transient catalogue". In: *Mon. Not. Roy. Astron. Soc.* 499.3 (2020), pp. 3295–3319. DOI: 10.1093/mnras/staa2850 (cit. on p. 90).

[184] Charlie Hoy and Vivien Raymond. "PESummary: The code agnostic Parameter Estimation Summary page builder". In: *SoftwareX* 15 (2021), p. 100765. ISSN: 2352-7110. DOI: https://doi.org/10.1016/j.softx.2021.100765 (cit. on p. 90).

[185] Leo P. Singer and Larry R. Price. "Rapid Bayesian position reconstruction for gravitational-wave transients". In: *Phys. Rev. D* 93 (2016), p. 024013. DOI: 10.1103/PhysRevD.93.024013 (cit. on pp. 90, 108, 110).

[186] LIGO Scientific Collaboration and Virgo Collaboration. *ligo.skymap Documentation*. https://lscsoft.docs.ligo.org/ligo.skymap/. Accessed: 2026-07-14 (cit. on p. 90).

[187] Anna Puecher et al. *Testing general relativity using higher-order modes of gravitational waves from binary black holes*. 2022. arXiv: 2205.09062 [gr-qc] (cit. on p. 90).

[188] Eleanor Hamilton et al. *PhenomXPNR: An improved gravitational wave model linking precessing inspirals and NR-calibrated merger-ringdown*. 2025. arXiv: 2507.02604 [gr-qc] (cit. on p. 91).

[189] LIGO Scientific Collaboration, Virgo Collaboration, and KAGRA Collaboration. *LVK Algorithm Library - LALSuite*. Free software (GPL). 2018. DOI: 10.7935/GT1W-FZ16 (cit. on p. 91).

[190] Brian Bockelman et al. "Principles, technologies, and time: The translational journey of the HTCondor-CE". In: *Journal of Computational Science* 52 (2021). Case Studies in Translational Computer Science, p. 101213. ISSN: 1877-7503. DOI: https://doi.org/10.1016/j.jocs.2020.101213 (cit. on p. 94).

[191] Jason S. C. Poon et al. “Galaxy lens reconstruction based on strongly lensed gravitational waves: similarity transformation degeneracy and mass-sheet degeneracy”. In: *Mon. Not. R. Astron. Soc.* 536.3 (Jan. 2025), pp. 2212–2233. ISSN: 0035-8711. DOI: 10.1093/mnras/stae2660 (cit. on pp. 96, 103).

[192] Hemantakumar Phurailatpam and Otto Akseli Hannuksela. *gwsnr: A Python package for efficient signal-to-noise ratio calculations of gravitational waves.* 2026. arXiv: 2412.09888 [astro-ph.IM] (cit. on p. 101).

[193] Paolo Cremonese, Jose María Ezquiaga, and Vincenzo Salzano. “Breaking the mass-sheet degeneracy with gravitational wave interference in lensed events”. In: *Phys. Rev. D* 104 (2021), p. 023503. DOI: 10.1103/PhysRevD.104.023503 (cit. on p. 103).

[194] Kyungmin Kim, Eungwang Seo, and Chunglee Kim. “Gravitational lensing aided luminosity distance estimation for compact binary coalescences”. In: *Phys. Rev. D* 109 (4 Feb. 2024), p. 043017. DOI: 10.1103/PhysRevD.109.043017 (cit. on p. 103).

[195] K. M. Górski et al. “HEALPix: A Framework for High-Resolution Discretization and Fast Analysis of Data Distributed on the Sphere”. In: *ApJ* 622 (2005), pp. 759–771. DOI: 10.1086/427976 (cit. on p. 106).

[196] B. P. Abbott et al. “Prospects for observing and localizing gravitational-wave transients with Advanced LIGO, Advanced Virgo and KAGRA”. In: *Living Reviews in Relativity* 23.1 (2020), p. 3. DOI: 10.1007/s41114-020-00026-9 (cit. on pp. 107, 111).

Appendix A

Parameters of Compact Binary Coalescences

This appendix provides a reference summary of the parameters used to describe a CBC system with quasi-circular orbits, following the conventions of [125, 144]. The parameters are grouped into mass, spin, tidal, and extrinsic categories; those employed in the injections and PE analyses of this work are indicated with $\star$.

Table A.1: Parameters describing a CBC system with quasi-circular orbits [125, 144]. Parameters used in this work are indicated with $\star$.

Parameter name	Symbol	Definition [Dimensions]
Mass parameters		
Primary/secondary mass	$m_1^\star$, $m_2^\star$	Mass of the more (less) massive component, $m_1 \geq m_2$ [$M_\odot$]
Chirp mass	$\mathcal{M}$	$\mathcal{M} = \eta^{3/5} M$ [$M_\odot$]
Total mass	M	$M = m_1 + m_2$ [$M_\odot$]
Final mass	M_f	Mass of the remnant [$M_\odot$]
Mass ratio	q	$q = m_2/m_1 \leq 1$ [dimensionless]
Symmetric mass ratio	η	$\eta = m_1 m_2 / M^2 \leq 1/4$ [dimensionless]
Radiated energy	$E_{\rm rad}$	$E_{\rm rad} = (M - M_f)c^2$ [energy]
Spin parameters		
Primary/secondary spin vectors	$\mathbf{S}_1$, $\mathbf{S}_2$	Spin angular momentum of the primary/secondary [ang. mom.]
Dimensionless spin magnitudes	$\chi_1^\star$, $\chi_2^\star$	$\chi_{1,2} = c\vert\mathbf{S}_{1,2}\vert/(Gm_{1,2}^2) \leq 1$ [dimensionless]
Remnant spin magnitude	χ_f	$\chi_f = cS_f/(GM_f^2) \leq 1$ where S_f is the magnitude of the remnant's spin angular momentum [dimensionless]
Orbital ang. momentum	$\mathbf{L}$	Instantaneous orbital angular momentum about the center of mass [ang. mom.]
Total ang. momentum	$\mathbf{J}$	$\mathbf{J} = \mathbf{L} + \mathbf{S}_1 + \mathbf{S}_2$ [ang. mom.]
Primary/secondary tilt angles	$\theta_1^\star$, $\theta_2^\star$	Angle between $\mathbf{S}_{1,2}$ and $\mathbf{L}$ [angle]
Spin azimuthal difference	$\phi_{12}^\star$	Angle between $\mathbf{L} \times (\mathbf{S}_1 \times \mathbf{L})$ and $\mathbf{L} \times (\mathbf{S}_2 \times \mathbf{L})$ [angle]

Continued on next page

Table A.1 *(continued)*

Parameter name	Symbol	Definition [Dimensions]
Effective inspiral spin	$\chi_{\rm eff}$	$(m_1\chi_{1z} + m_2\chi_{2z})/M$ [dimensionless]
Effective precession spin	χ_p	Avg. over precession cycles [124] [dimensionless]
Tidal parameters		
Tidal deformabilities	Λ_1, Λ_2	Dimensionless tidal deformability; $\Lambda_{1,2} = 0$ for BHs [dimensionless]
Effective tidal deformability	$\tilde{\Lambda}$	Dominant tidal combination; $\tilde{\Lambda} = 0$ for BBH [dimensionless]
Extrinsic parameters		
Luminosity distance	$D_L^\star$	Distance to the source [Mpc]
Redshift	z	Computable from D_L assuming a cosmological model [dimensionless]
Source inclination	$\theta_{JN}^\star$	Angle between **J** and line of sight **N** [angle]
Orbital inclination	ι	Angle between **L** and **N** [angle]
Coalescence time	$t_c^\star$	GPS time at the geocenter [time]
Coalescence phase	$\phi_c^\star$	Orbital phase at coalescence [angle]
Right ascension	$\alpha^\star$	Celestial longitude [angle]
Declination	$\delta^\star$	Celestial latitude [angle]
Polarization angle	$\psi^\star$	Rotates polarizations in the plane orthogonal to propagation [angle]
Azimuthal spin angle	$\phi_{JL}^\star$	Angle between **J** and **L** projected on the sky plane [angle]

Appendix B

Practical guide to HPC workflows

The motivation for including this appendix comes from the desire to make life easier for future students who will use the same tools, especially since many parts of this tutorial came with their fair share of headaches. It collects the day-to-day commands, configuration patterns, and troubleshooting notes accumulated while running the analyses of this thesis on two HPC clusters. It is not a repetition of the pipeline described in §5.4, but a practical reference for a student facing the same kind of environment for the first time. All pipeline-specific reproducibility details (configuration files, directory layout, job dependency logic) are documented in §5.4 and Appendix §C; the complete code is publicly available at `https://github.com/XxavierGarcia/TFM_reproducibility_material`.

B.1 SSH access and aliases

Both clusters are accessed via SSH (Secure Shell), a protocol that opens an encrypted terminal session on a remote machine: once connected, commands typed locally are executed directly on the cluster, as if sitting in front of it. On a Linux machine, an SSH client is available out of the box:

```
ssh username@picasso.scbi.uma.es
```

(Users connecting from Windows Subsystem for Linux may need to append `-o KexAlgorithms=ecdh-sha2-nistp521` to this and every following `ssh`/`scp` command, since its older default OpenSSH client does not otherwise offer an algorithm accepted by Picasso's server; this flag is omitted from the rest of this tutorial.)

On CIT, the IGWN LDAS head nodes serve as the entry point (LIGO.ORG credentials required):

```
ssh username@ldas-grid.ligo.caltech.edu
```

Typing these every time quickly becomes impractical. Two alternatives exist. The first is to define shell aliases in `~/.bashrc`:

```
alias picasso='ssh username@picasso.scbi.uma.es'
alias caltech='ssh username@ldas-grid.ligo.caltech.edu'
```

Once these lines are added and the shell configuration is reloaded (`source ~/.bashrc` or a new terminal), the full commands above are replaced by simply typing `picasso` or `caltech`.

The second, more robust option is to add entries to `~/.ssh/config`:

```
Host picasso
    HostName picasso.scbi.uma.es
    User username
    IdentityFile ~/.ssh/id_rsa
```

With this configuration, `ssh picasso` and `scp picasso:path/file .` work directly.

B.2 Remote editing with VS Code

Editing files on the cluster through `nano` or `vim` works but is slow for larger scripts. VS Code's Remote - SSH extension provides a full graphical editor connected to the cluster filesystem:

1. Install the *Remote - SSH* extension from the VS Code marketplace.
2. Open the command palette (`Ctrl+Shift+P`), type *Remote-SSH: Connect to Host*, and select the host (or add a new one pointing to `picasso.scbi.uma.es`).
3. Once connected, use *File* → *Open Folder* to navigate to the working directory.

This allows browsing the directory tree, viewing output figures directly, and editing configuration files without leaving the editor.

B.3 Transferring files with scp

`scp` (secure copy) transfers files between any two machines reachable via SSH. The basic syntax is `scp [options] source destination`, where either side can be `host:path`.

```
# Copy a single file from the local machine to the cluster
scp local_file.py username@picasso.scbi.uma.es:~/lenskyloc/

# Copy an entire directory (-r for recursive)
scp -r username@picasso.scbi.uma.es:~/lenskyloc/inj001/result/ .

# Transfer between two clusters (run from CIT)
scp -r username@picasso.scbi.uma.es:~/lenskyloc/pesummary_jsons .
```

On CIT, the Jupyter interface (`https://jupyter.ligo.caltech.edu`) provides an alternative for downloading files: compress them first with `zip -r archive.zip directory/` and download the archive through the browser.

B.4 Persistent terminal sessions with screen

An SSH connection that drops (timeout, laptop closing, network change) kills any process running inside it. The `screen` utility creates a persistent terminal session on the cluster that survives disconnections:

```
screen -S analysis          # create a named session
# ... run any long command ...
# detach: Ctrl+a, then d
screen -ls                  # list active sessions
screen -r analysis          # reattach
exit                        # close the session permanently
```

This is essential for interactive tasks that take more than a few minutes (monitoring logs, running quick post-processing scripts) but are not worth submitting as batch jobs.

B.5 Conda environments

Both clusters use `conda` environments to manage Python dependencies. On Picasso, the analysis environment was created and maintained manually:

```
conda create -n hanabi_bilby260 python=3.10
conda activate hanabi_bilby260
pip install bilby bilby_pipe hanabi lalsuite dynesty ...
```

The activation line is included in every generated SLURM script, so each job runs in the correct environment automatically. On CIT, the `igwn` environment distributed via CVMFS is activated directly instead.

When the home directory quota on Picasso fills up (check with `quota`), the most effective remedy is to clear the conda cache, which accumulates downloaded packages and can grow to several gigabytes:

```
conda clean --all
```

B.6 SLURM: submitting and monitoring jobs on Picasso

Picasso uses the SLURM workload manager. Jobs are submitted with `sbatch`, monitored with `squeue`, and cancelled with `scancel`:

```
sbatch job_script.sh            # submit a job
squeue -u username              # list your running/pending jobs
scancel JOB_ID                  # cancel one job
scancel --user username         # cancel all your jobs
```

A typical SLURM header for a single-image PE job in this thesis (§5.2) looks like:

```
#!/bin/bash
#SBATCH --job-name=lenskyloc_inj_001_img1
#SBATCH --cpus-per-task=128
#SBATCH --mem=110G
#SBATCH --time=48:00:00
#SBATCH --constraint=cal

export MPLCONFIGDIR=$LOCALSCRATCH/$USER
conda activate hanabi_bilby260
```

The `--constraint=cal` flag restricts the job to nodes of a specific hardware class; `MPLCONFIGDIR` is redirected to the node-local scratch space to avoid `matplotlib` cache errors on compute nodes whose home filesystem is read-only.

For dependent jobs (e.g. a joint analysis that must wait for both single-image jobs), SLURM's `--dependency` flag is used:

```
JOB1=$(sbatch img1.sh | awk '{print $4}')
JOB2=$(sbatch img2.sh | awk '{print $4}')
sbatch --dependency=afterok:$JOB1:$JOB2 joint.sh
```

To monitor a running job's progress in real time, inspect the sampler log:

```
tail -f log_data_analysis/*.err
```

If a job is killed by the wall-time limit, it can be resubmitted and will resume from the last `Dynesty` checkpoint (written every 30 minutes by default via `check_point_delta_t=1800`).

B.7 HTCondor: submitting and monitoring jobs on CIT

CIT uses HTCondor instead of SLURM. The analogous commands are:

```
condor_submit file.submit      # submit a job
condor_q                       # list your running/pending jobs
condor_rm JOB_ID               # cancel one job
condor_q -hold                 # diagnose jobs stuck in "held" state
```

One common pitfall: HTCondor requires the executable specified in the `.submit` file to have the Unix execute permission. If a `.sh` script submitted as `executable` produces `errno=13: Permission denied`, the fix is:

```
chmod +x script.sh
```

The permissions can be verified with `ls -l`: the file should show `-rwxr-xr-x` (the `x` flags) rather than `-rw-r--r--`.

B.8 Useful everyday commands

The following commands were used daily throughout this project:

Task	Command
Search for a string inside a file	`grep pattern file`
Compare two files line by line	`diff file1 file2`
Check the size of a file or directory	`du -sh path`
Check sizes of everything in home	`du -sh ~/*`
Jump to the end of a large file	`less file`, then `Shift+G`
View the last/first lines of a log	`tail -n 50 file` / `head -n 50 file`
Check available disk quota (Picasso)	`quota`
See all options of a command	`command --help`

B.9 Troubleshooting summary

The most common issues encountered during this project, and their solutions:

- **Wall-time kills on Picasso:** resubmit the same script; `Dynesty` resumes from the latest checkpoint. This was also the mechanism that limited the last seven joint analyses when the allocation dropped to 24 h (§5.4.3).
- **Matplotlib cache errors on compute nodes:** redirect `MPLCONFIGDIR` to node-local scratch (already included in the SLURM headers above).
- **PESummary failing on Picasso:** the version available in the Picasso environment produced errors; the workaround was to run the summary-page and sky-map generation on CIT instead, where the `igwn` environment handles it correctly (§5.4).
- **HTCondor "held" jobs on CIT:** diagnose with `condor_q -hold`; the most frequent cause was a missing execute permission on the `.sh` script (see above).
- **Home directory quota full on Picasso:** run `conda clean --all` to free cached packages.

B.10 Additional resources

- SLURM documentation: `https://slurm.schedmd.com/documentation.html`
- HTCondor documentation: `https://htcondor.readthedocs.io/en/latest/`

- Picasso user guide (SCBI): https://www.scbi.uma.es/site/
- IGWN Computing Guide (CIT access): https://computing.docs.ligo.org/guide/
- Complete pipeline code: https://github.com/XxavierGarcia/TFM_reproducibility_material

Appendix C

Detailed implementation of the pipeline

This appendix reproduces two representative `.ini` configuration files generated automatically by `Generator_inis_lenskyloc.py` (§5.4) for event `inj001`: the single-image configuration for image 1 (§5.2), and the joint configuration combining both images (§5.3). The complete pipeline code, including both the Picasso (SLURM) and CIT (HTCondor) branches described in §5.4, is publicly available at the GitHub repository listed in Appendix §B.

C.1 Single-image configuration

```
################################################
## Data generation arguments
################################################
trigger-time=1268585121.2749617
gaussian-noise=True
zero-noise=False

#################################################
## Detector arguments
#################################################
detectors=[H1, L1]
duration=4
psd-dict={"H1": "/mnt/home/users/uib54_res/resh000464/noise_curves/
    T2500310-v2_O5a_strain_psd.txt", "L1": "/mnt/home/users/uib54_res/
    resh000464/noise_curves/T2500310-v2_O5a_strain_psd.txt"}
sampling-frequency=2048
reference-frequency=20
minimum-frequency={'H1': 20, 'L1': 20, 'waveform':13.33}
maximum-frequency= {'H1': 896, 'L1': 896}

#################################################
## Injection arguments
#################################################
injection=True
injection-dict={'mass_1': 68.50715378171094, 'mass_2':
    40.35264268122141, 'luminosity_distance': 2889.5612112115464, 'a_1':
    0.4527593183131344, 'a_2': 0.5997569914138758, 'tilt_1':
    1.4148137284549118, 'tilt_2': 1.685673833551925, 'phi_12':
    1.2201031723904865, 'phi_jl': 3.7845775009109737, 'theta_jn':
    0.5351550350716038, 'psi': 2.351204394088696, 'phase':
```

```
    4.596028055035042, 'geocent_time': 1268585121.2749617, 'ra':
    6.094792612884527, 'dec': -0.6855988520635732, 'minimum_frequency':
    20.0, 'reference_frequency': 20.0, 'image_type': 1.0}
injection-waveform-approximant=IMRPhenomXPNR
injection-frequency-domain-source-model=hanabi.lensing.waveform.
    strongly_lensed_BBH_waveform
injection-waveform-arguments={'reference_frequency': 20, '
    minimum_frequency': 20, 'pn_amplitude_order': 0}

################################################
## Job submission arguments
################################################
label=lenskyloc_inj_001_hanabi_single_img1
outdir=/mnt/home/users/uib54_res/resh000464/lenskyloc/lenskyloc_runs/
    inj001/inj001_img1
request-disk=6.0
request-memory=110.0
request-memory-generation=2.0
request-cpus=128
scheduler=slurm
scheduler-args=constraint=cal
extra-lines=export MPLCONFIGDIR=$LOCALSCRATCH/$USER
scheduler-analysis-time=48:00:00

################################################
## Output arguments
################################################
result-format=json

################################################
## Prior arguments
################################################
prior-dict={chirp-mass: bilby.gw.prior.UniformInComponentsChirpMass(
    minimum=31.819, maximum=59.093, name='chirp_mass', boundary=None),
     mass-ratio: bilby.gw.prior.UniformInComponentsMassRatio(minimum
        =0.05, maximum=1.0, name='mass_ratio', latex_label='$q$', unit=
        None, boundary=None),
     mass-1: Constraint(minimum=5.0, maximum=200.0, name='mass_1',
        latex_label='$m_1$', unit=None),
     mass-2: Constraint(minimum=5.0, maximum=200, name='mass_2',
        latex_label='$m_2$', unit=None),
     a-1: Uniform(minimum=0, maximum=0.99, name='a_1', latex_label='$a_1$
        ', unit=None, boundary=None),
     a-2: Uniform(minimum=0, maximum=0.99, name='a_2', latex_label='$a_2$
        ', unit=None, boundary=None),
     tilt-1: Sine(minimum=0, maximum=3.141592653589793, name='tilt_1'),
     tilt-2: Sine(minimum=0, maximum=3.141592653589793, name='tilt_2'),
     phi-12: Uniform(minimum=0, maximum=6.283185307179586, name='phi_12',
         boundary='periodic'),
     phi-jl: Uniform(minimum=0, maximum=6.283185307179586, name='phi_jl',
         boundary='periodic'),
```

```
    luminosity-distance: Uniform(minimum=50, maximum=20000, name='
        luminosity_distance', latex_label='$d_L$', unit='Mpc', boundary=
        None),
    cos-theta-jn: Uniform(minimum=-1, maximum=1, name='cos_theta_jn',
        latex_label='$\cos\theta_{JN}$', unit=None, boundary=None),
    psi: Uniform(minimum=0, maximum=3.141592653589793, name='psi',
        boundary='periodic'),
    phase: Uniform(minimum=0, maximum=6.283185307179586, name='phase',
        boundary='periodic'),
    dec: Cosine(minimum=-1.5707963267948966, maximum=1.5707963267948966,
         name='dec', latex_label='$\mathrm{DEC}$', unit=None, boundary=
        None),
    ra: Uniform(minimum=0, maximum=6.283185307179586, name='ra',
        latex_label='$\mathrm{RA}$', unit=None, boundary='periodic'),
    image_type = 1.0}

##################################################
## Sampler arguments
##################################################
n-parallel=3
sampler-kwargs={'nlive': 1000, 'naccept': 60, 'check_point_plot': True,
   'check_point_delta_t': 1800, 'print_method': 'interval-60', 'sample':
    'acceptance-walk', 'samples': 'acceptance-walk', 'npool': 128}
distance_marginalization=False
time-marginalization=False
phase-marginalization=False

###################################################
## Waveform arguments (used for the analysis)
###################################################
waveform-generator=bilby.gw.waveform_generator.WaveformGenerator
frequency-domain-source-model=hanabi.lensing.waveform.
   strongly_lensed_BBH_waveform
waveform-approximant=IMRPhenomXPNR
enforce-signal-duration=False
catch-waveform-errors=True
pn-amplitude-order=1
```

Listing C.1: Single-image configuration file for image 1 of event `inj001`.

The injection uses `pn_amplitude_order=0` while the analysis uses `pn-amplitude-order=1`. This argument, inherited from the PN waveform families supported by LALSimulation, is not used by the phenomenological `IMRPhenomXPNR` model to construct its amplitude: generating the same waveform with `pn_amplitude_order` set to 0 and to 1, keeping all other parameters fixed, leaves the frequency-domain strain unchanged. This configuration mismatch is therefore a leftover setting from the template with no real effect on the results of this thesis.

C.2 Joint configuration

```
################################################
## Job submission arguments
################################################
label=lenskyloc_inj_001_hanabi_joint
outdir=/mnt/home/users/uib54_res/resh000464/lenskyloc/lenskyloc_runs/
    inj001/inj001_joint
request-disk=6.0
request-memory=180.0
request-memory-generation=2.0
request-cpus=128
scheduler=slurm
scheduler-args=constraint=cal
extra-lines=export MPLCONFIGDIR=$LOCALSCRATCH/$USER
scheduler-analysis-time=71:30:00

################################################
## Output arguments
################################################
result-format=json

################################################
## Sampler arguments
################################################
sampler=dynesty
n-parallel=3
sampler-kwargs={'nlive': 1000, 'naccept': 60, 'check_point_plot': True,
    'check_point_delta_t': 1800, 'print_method': 'interval-60', 'sample':
     'acceptance-walk', 'samples': 'acceptance-walk', 'npool': 128}

################################################
## Joint input arguments
################################################
n-triggers=2
trigger-ini-files=[/mnt/home/users/uib54_res/resh000464/lenskyloc/
    lenskyloc_runs/inj001/single_event_001_image1.ini, /mnt/home/users/
    uib54_res/resh000464/lenskyloc/lenskyloc_runs/inj001/
    single_event_001_image2.ini]

################################################
## Lensing prior argument
################################################
common-parameters=[chirp_mass, mass_ratio, mass_1, mass_2, a_1, a_2,
    tilt_1, tilt_2, phi_12, phi_jl, cos_theta_jn, ra, dec, psi, phase]
lensing-prior-dict={relative-magnification^(1): 1, image_type^(1) =
    hanabi.lensing.prior.DiscreteUniform(name='image_type^(1)', minimum
    =1, N=3), relative_magnification^(2) = 1, image_type^(2) = hanabi.
    lensing.prior.DiscreteUniform(name='image_type^(2)', minimum=1, N=3)}
```

Listing C.2: Joint configuration file for event `inj001`.

Figures C.1 and C.2 lay out, script by script, the two branches of the pipeline introduced in §5.4: the SLURM job generation and submission chain running on Picasso, and the HTCondor post-processing chain running on CIT.

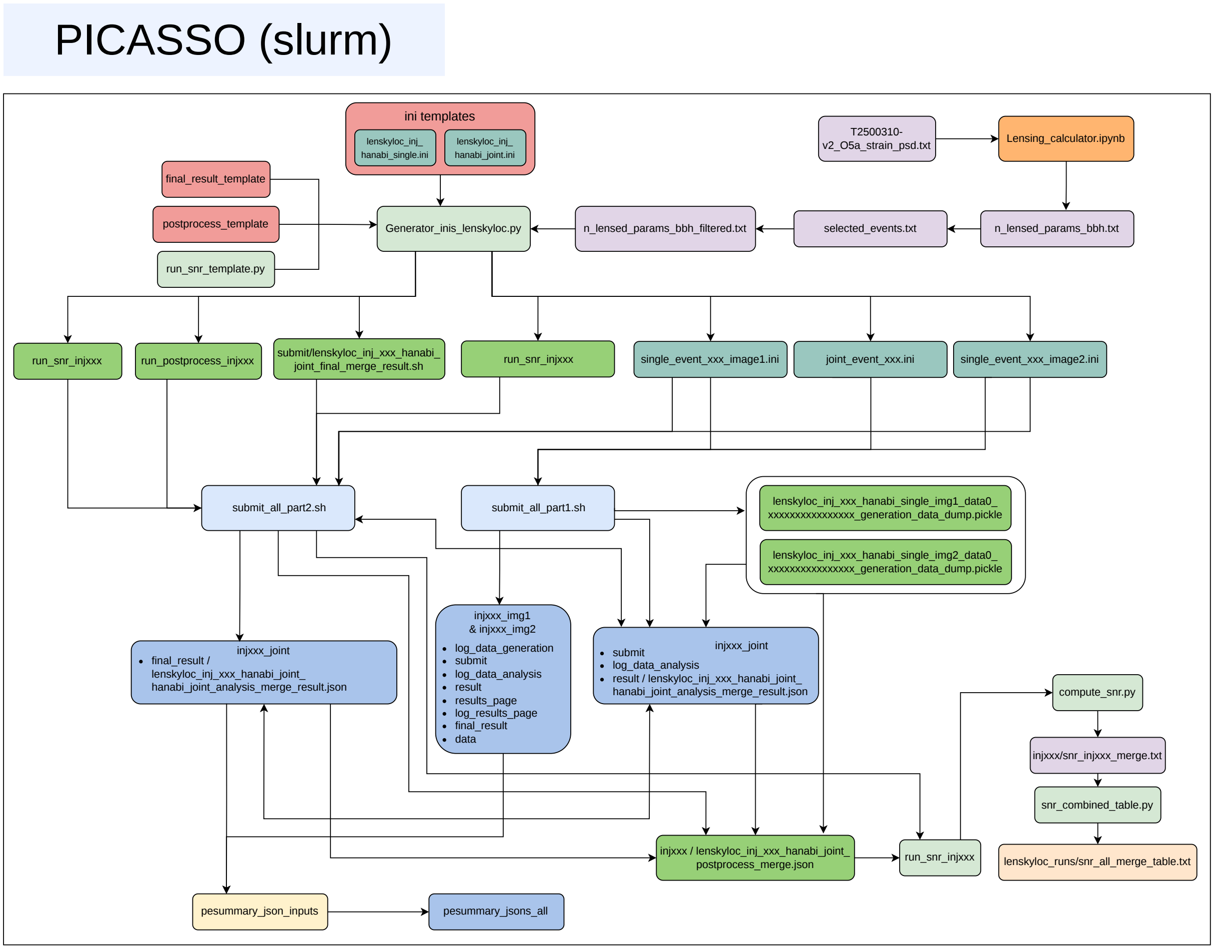

Figure C.1: Detailed implementation of the Picasso (SLURM) branch of the pipeline, referenced in §5.4. Best viewed digitally, zoomed in. An interactive version with clickable links to each script is available at https://github.com/XxavierGarcia/TFM_reproducibility_material/blob/main/diagrama_pipeline.drawio.pdf (download to enable links).

CIT (condor)

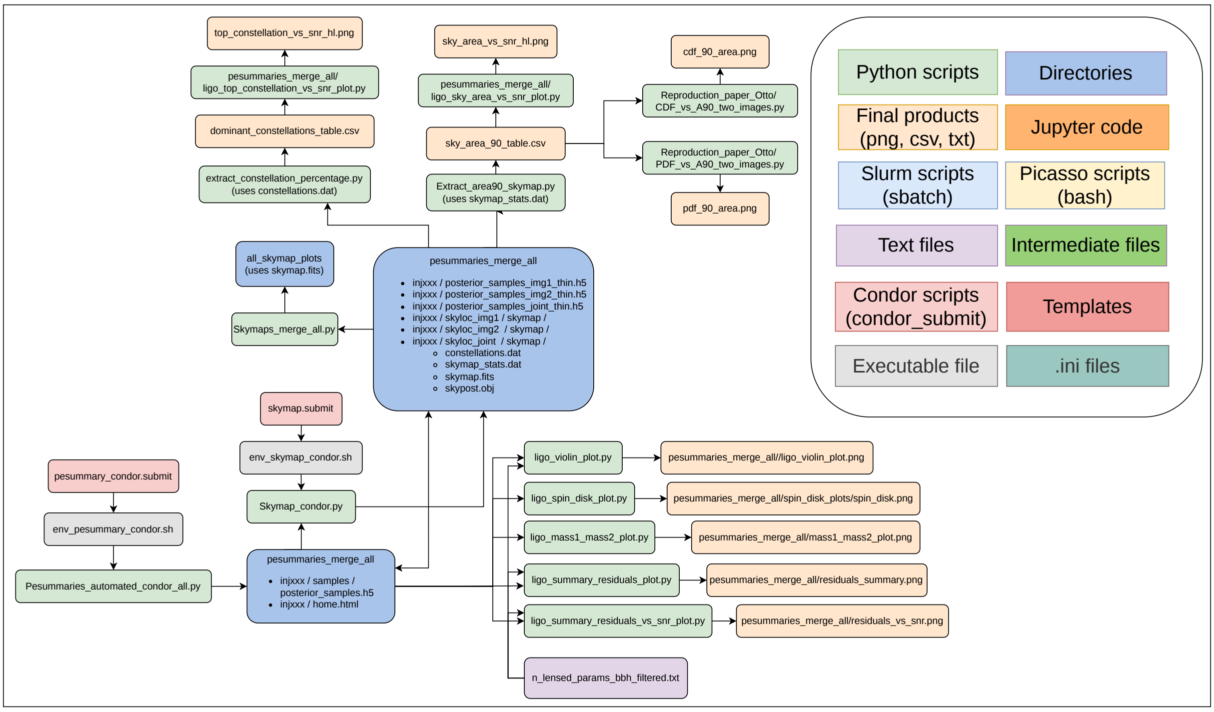

Figure C.2: Detailed implementation of the CIT (HTCondor) post-processing branch of the pipeline, referenced in §5.4 and §6.1. Best viewed digitally, zoomed in. An interactive version with clickable links to each script is available at https://github.com/XxavierGarcia/TFM_reproducibility_material/blob/main/diagrama_pipeline.drawio.pdf (download to enable links).

Appendix D

Full posterior comparison for all events

This appendix collects the recovered posteriors for all 150 analyzed events, complementing the representative examples and population-level residuals of §5.5. For each event, the joint posterior (green, solid) is shown together with the two single-image posteriors (red dash-dot for image 1, blue dotted for image 2) and the injected value (black tick). Figure D.1 shows the chirp mass, mass ratio, and effective luminosity distance; Figure D.2 shows the spin parameters. Within each set the events are ordered by injection index and split across pages.

A closer look at these figures reveals three main patterns. First, the joint analysis produces the largest narrowing on the mass ratio q and the effective spin χ_{eff}. Defining the improvement factor per event and per parameter as

$$\mathrm{IF}_\theta \equiv \frac{\min\left(\Delta_{90}^{(1)},\, \Delta_{90}^{(2)}\right)}{\Delta_{90}^{\mathrm{joint}}}, \tag{D.1}$$

where $\Delta_{90}^{(i)}$ denotes the width of the 90% credible interval (5th to 95th percentile) of the posterior for image i, and $\Delta_{90}^{\mathrm{joint}}$ the same quantity for the joint posterior, values $\mathrm{IF}_\theta > 1$ mean the joint analysis is tighter than the better of the two single-image analyses. The largest improvement factors across the sample are $\mathrm{IF}_\theta = 3.53$ for D_L^{eff} in `inj113`, 2.20 for χ_{eff} in `inj113`, 1.91 for $\mathcal{M}$ in `inj113`, and 1.67 for q in `inj152`. Table D.1 lists the top 25 events ranked by D_L^{eff}; `inj113` stands out as the event where the joint analysis extracts the most information across essentially every parameter simultaneously.

Second, the joint analysis suppresses bimodalities that appear in one or both single-image posteriors of q or χ_{eff} by selecting the mode consistent between the two images: this is visible in `inj030`, `inj055`, `inj069`, `inj091`, `inj101`, `inj119`, and `inj132`. Third, for events already well measured from a single image, typically at high SNR (`inj034`, `inj048`, `inj076`, `inj098`, `inj117`, `inj127`), the joint analysis leaves the posterior essentially unchanged, since the joint and single-image posteriors nearly coincide. The individual spin magnitudes a_1 and a_2 and the precessing spin χ_{p}, by contrast, look very similar to their priors in almost every event, so the joint analysis rarely narrows them by any visible amount: precession leaves a weak signature on most of these signals, and joining the two images cannot compensate for information that the data simply do not contain, a point already made in §5.5 on the population residuals of a_1, a_2, and χ_{p}.

Event	ρ_{inj}	$\mathcal{M}$	q	D_L^{eff}	a_1	a_2	χ_{eff}	χ_{p}
`inj113`	38.9	**1.91**	1.44	**3.53**	**1.81**	1.07	**2.20**	1.24
`inj122`	41.2	**1.55**	1.29	**2.16**	1.25	1.11	**1.50**	1.19
`inj049`	16.7	1.12	0.95	**2.06**	1.00	1.00	1.14	1.00
`inj050`	15.3	1.02	0.95	**1.90**	1.05	0.99	1.06	1.00
`inj135`	17.2	1.36	1.11	**1.87**	0.96	1.00	1.35	1.03
`inj079`	17.7	1.31	1.03	**1.81**	0.98	1.00	1.29	0.96
`inj157`	17.3	1.39	1.04	**1.78**	1.32	0.99	1.25	1.09
`inj021`	20.7	1.08	1.20	**1.75**	1.39	0.99	1.12	1.31
`inj152`	20.5	1.30	**1.67**	**1.70**	1.00	0.97	1.28	1.21
`inj061`	30.8	1.03	0.87	**1.69**	0.95	0.99	1.08	0.99
`inj003`	39.5	1.19	1.34	**1.65**	0.95	1.01	1.38	0.94
`inj068`	22.8	1.41	1.18	**1.63**	1.00	1.01	1.24	0.98
`inj114`	54.0	1.12	1.06	**1.63**	1.12	1.00	1.31	1.19
`inj106`	21.7	1.33	1.09	**1.60**	1.02	0.99	1.22	1.05
`inj134`	53.7	1.40	1.14	**1.59**	0.99	0.92	1.21	1.13
`inj069`	20.1	1.04	0.99	**1.58**	1.00	1.00	1.13	1.01
`inj127`	59.5	0.94	0.79	**1.57**	0.94	0.83	1.13	1.05
`inj004`	49.0	1.21	1.20	**1.55**	1.23	1.11	1.31	1.06
`inj026`	31.5	1.16	1.11	**1.55**	0.92	0.99	1.00	1.02
`inj098`	19.5	1.24	0.97	**1.54**	0.98	0.99	1.21	0.98
`inj028`	16.9	1.26	1.07	**1.52**	1.05	0.99	1.18	1.02
`inj083`	36.6	1.36	1.28	**1.52**	1.05	1.14	1.44	1.11
`inj071`	24.1	0.84	0.97	**1.51**	1.03	1.05	1.16	0.94
`inj085`	26.0	0.80	0.80	**1.51**	1.05	0.98	1.02	1.08
`inj035`	21.5	1.22	0.99	**1.51**	1.06	1.03	1.18	1.07

Table D.1: Improvement factors IF_θ (Eq. (D.1)) for the top 25 events ranked by D_L^{eff}, with the injected pair network SNR ρ_{inj} (the quadrature sum of the two per-image network SNRs computed by LeR, §5.5) shown for context. Values highlighted in bold orange mark $\mathrm{IF}_\theta \geq 1.5$, where the joint analysis narrows the 90% credible interval by at least a factor of 1.5 relative to the better of the two single-image analyses. Values below 1 appear only sporadically for a_1, a_2, and χ_{p}, which remain effectively prior-dominated across the sample, and occasionally for $\mathcal{M}$ or q (e.g. `inj085`, $\mathrm{IF}_q = 0.80$), where a single high-SNR image happens to yield an anomalously narrow posterior that the joint analysis correctly broadens (§5.5).

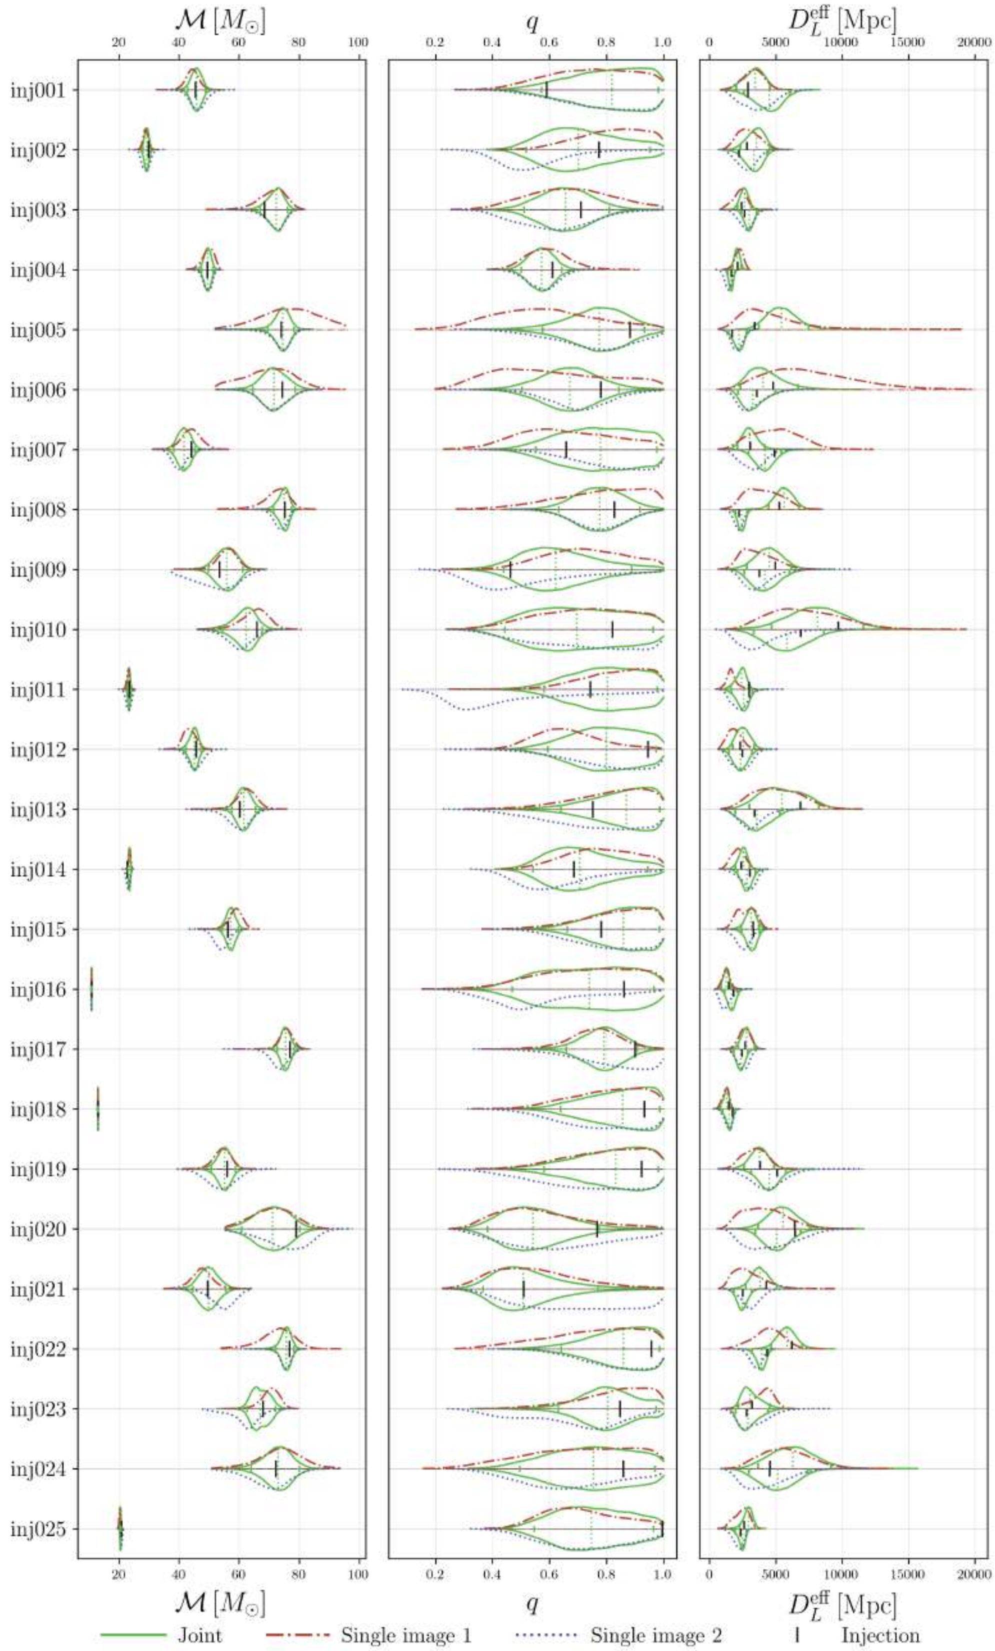

Figure D.1: Recovered posteriors of the chirp mass $\mathcal{M}$, mass ratio q, and effective luminosity distance $D_L^{\rm eff}$, for all 150 events (1 of 6). Joint posterior in green (solid), single-image posteriors in red (dash-dot, image 1) and blue (dotted, image 2); the injected value is the black tick.

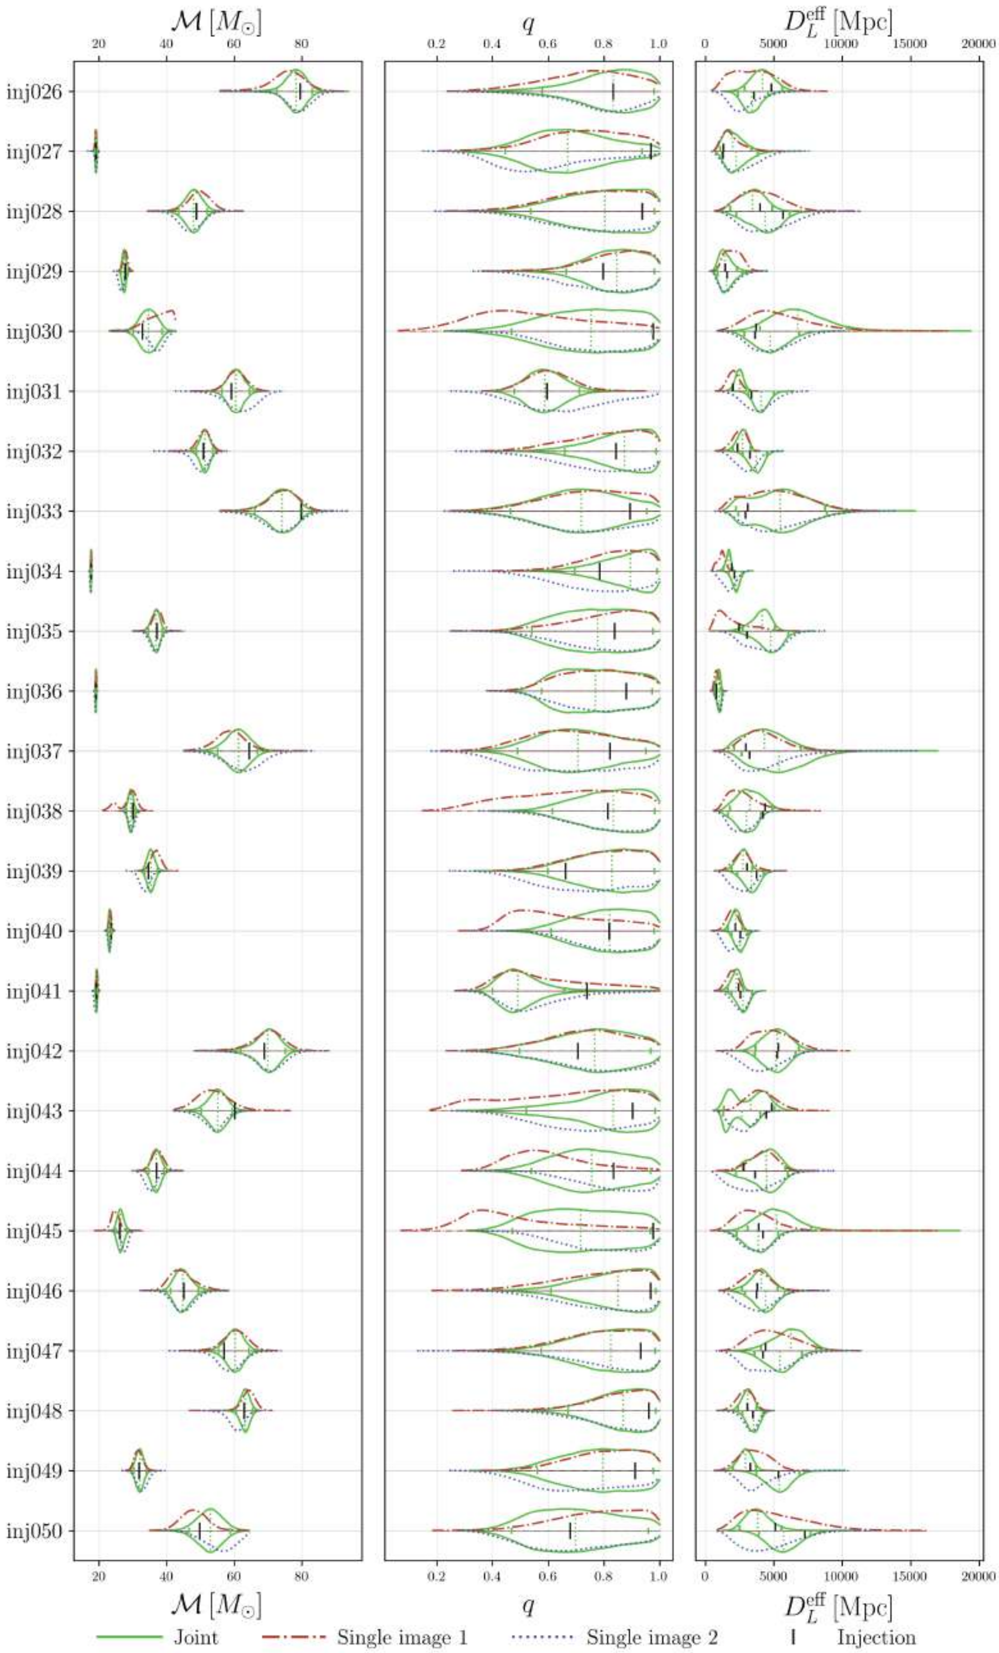

Figure D.1: Recovered posteriors of $\mathcal{M}$, q, and $D_L^{\rm eff}$ (2 of 6).

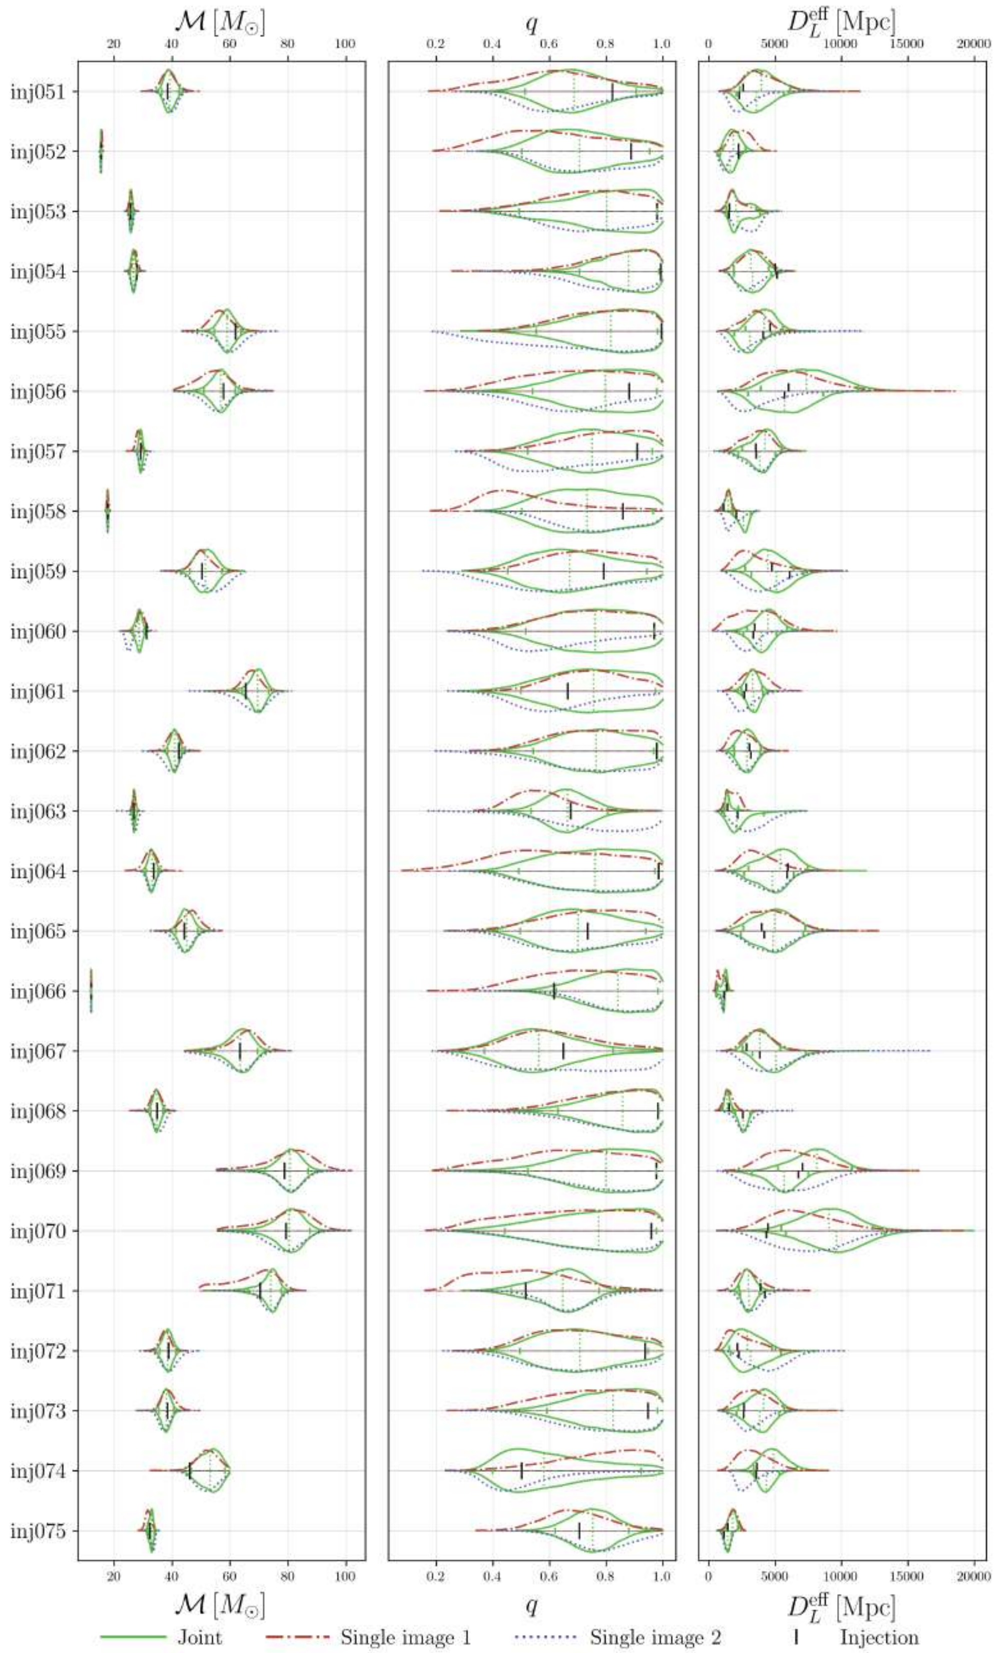

Figure D.1: Recovered posteriors of $\mathcal{M}$, q, and $D_L^{\rm eff}$ (3 of 6).

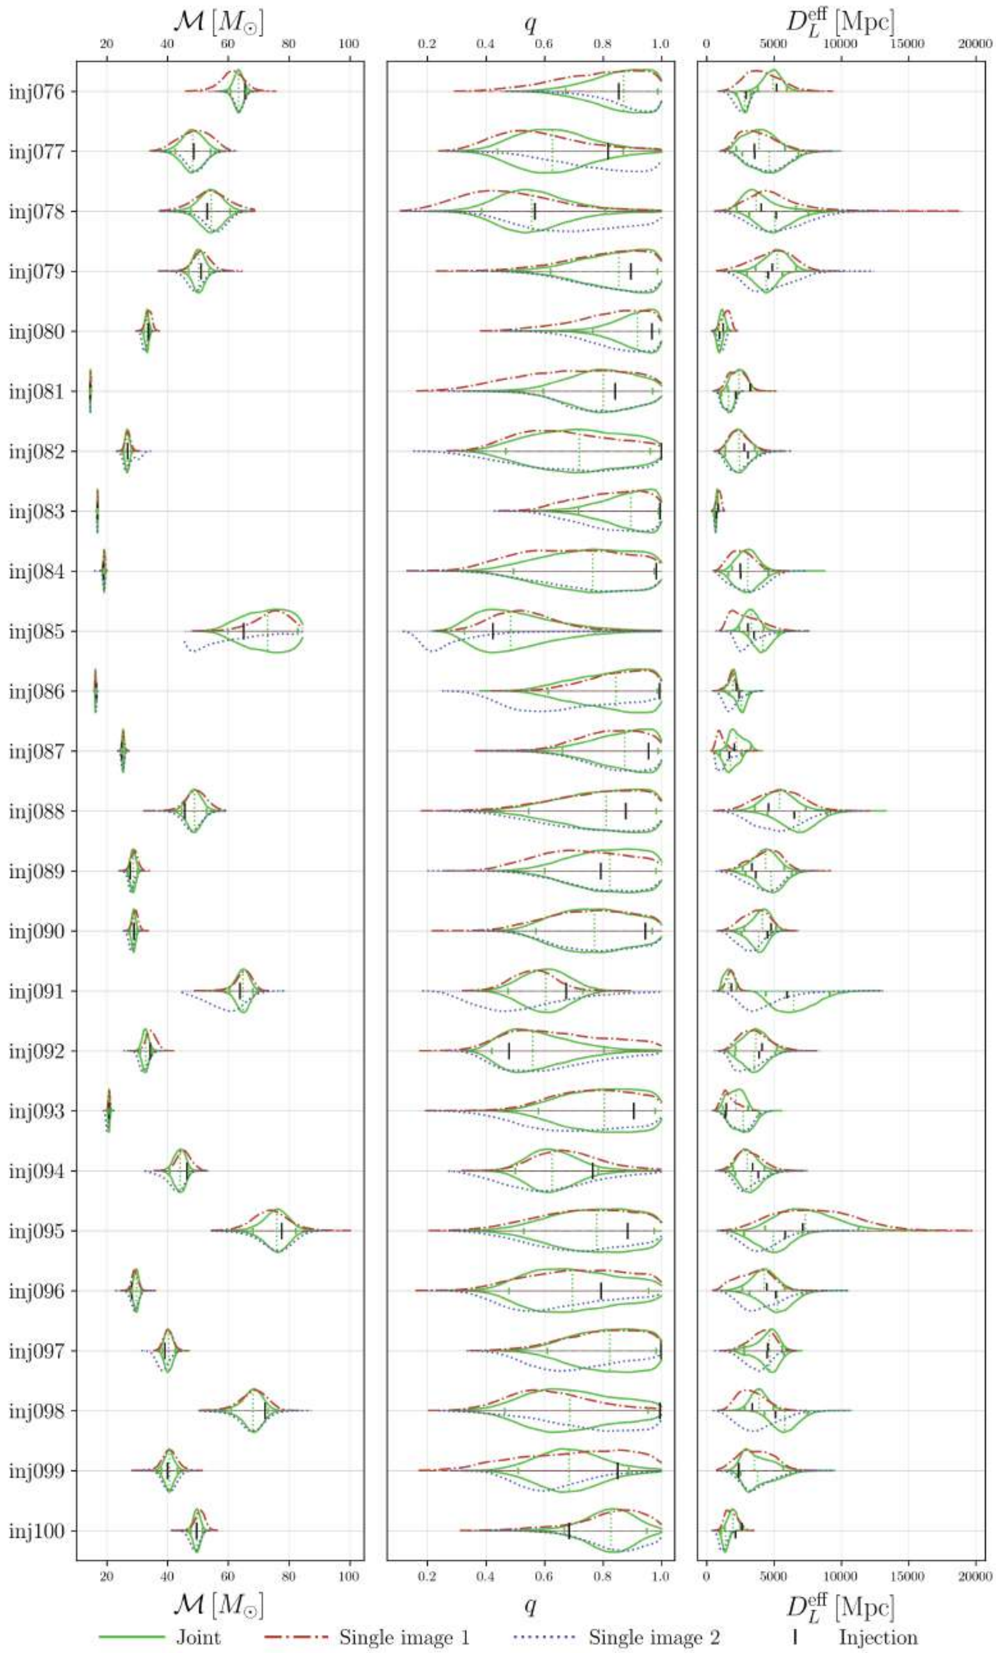

Figure D.1: Recovered posteriors of $\mathcal{M}$, q, and D_L^{eff} (4 of 6).

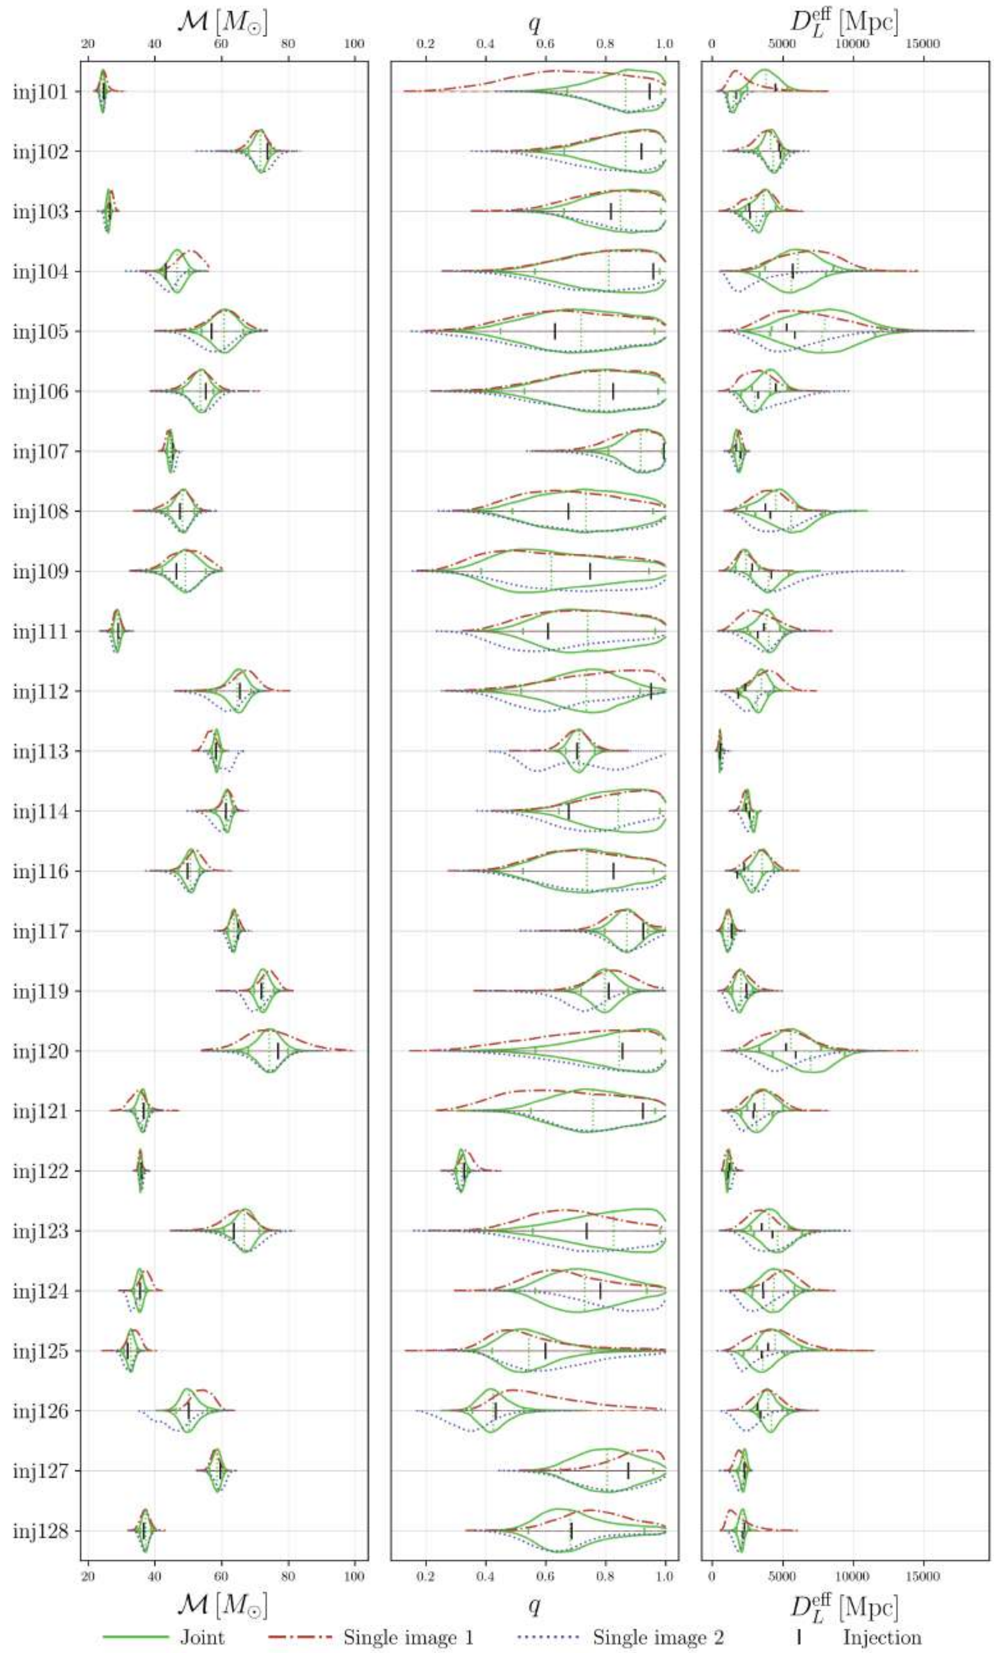

Figure D.1: Recovered posteriors of $\mathcal{M}$, q, and $D_L^{\rm eff}$ (5 of 6).

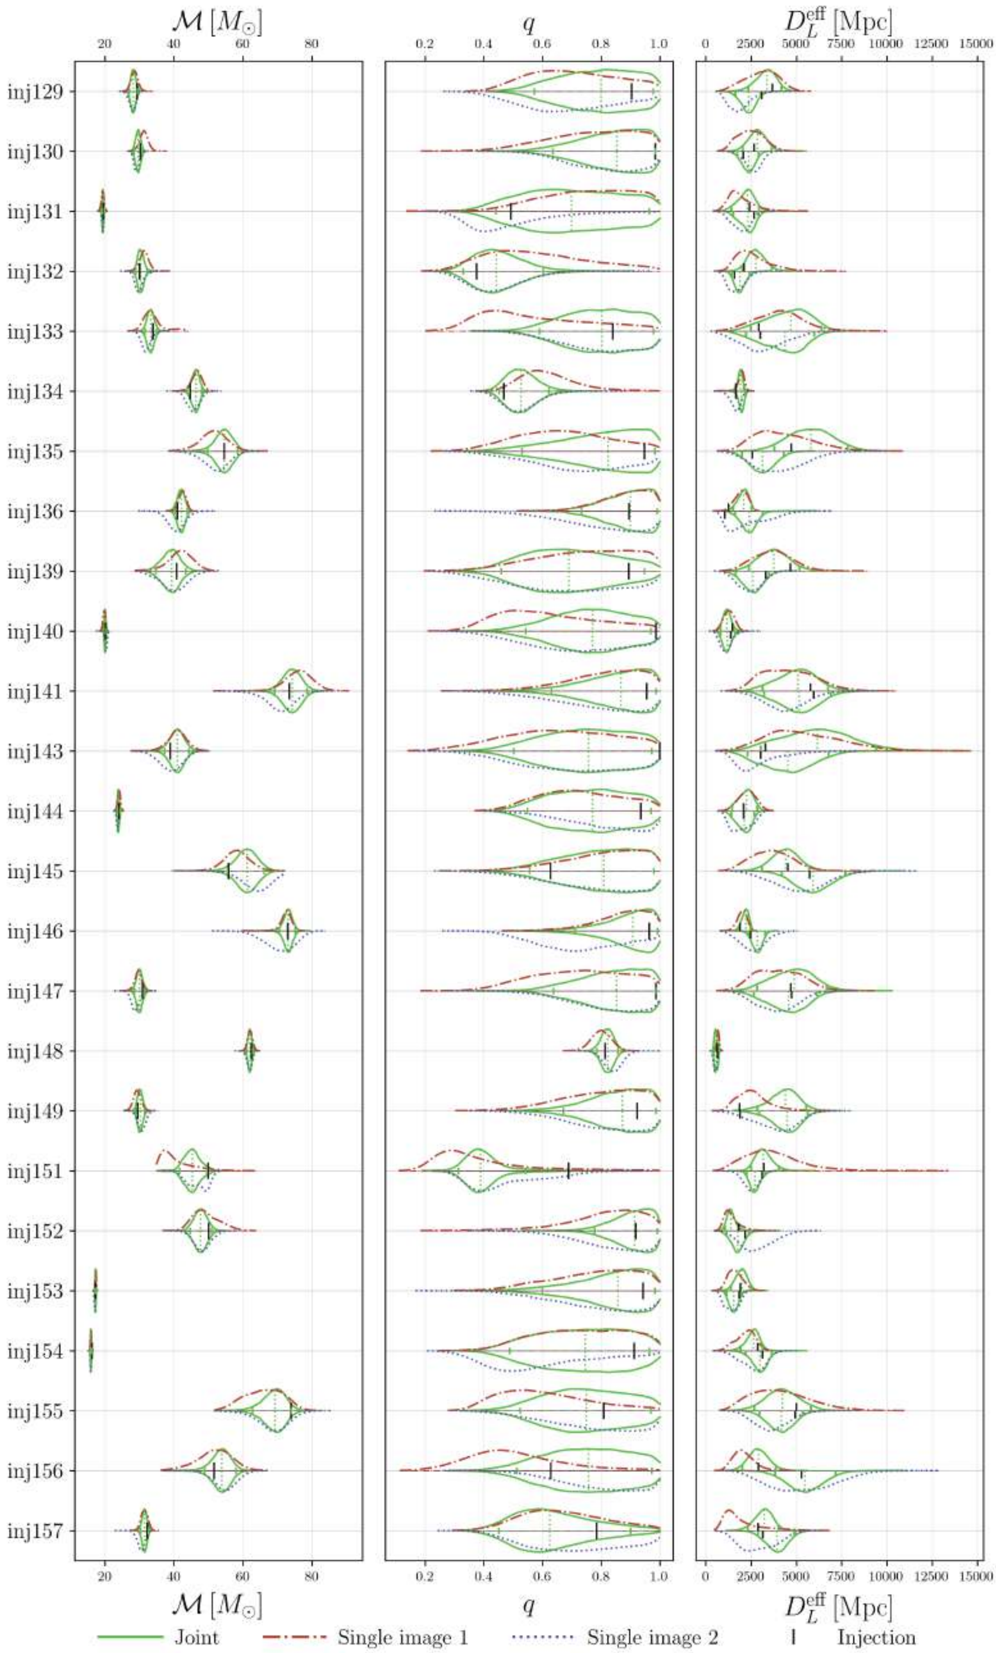

Figure D.1: Recovered posteriors of $\mathcal{M}$, q, and D_L^{eff} (6 of 6).

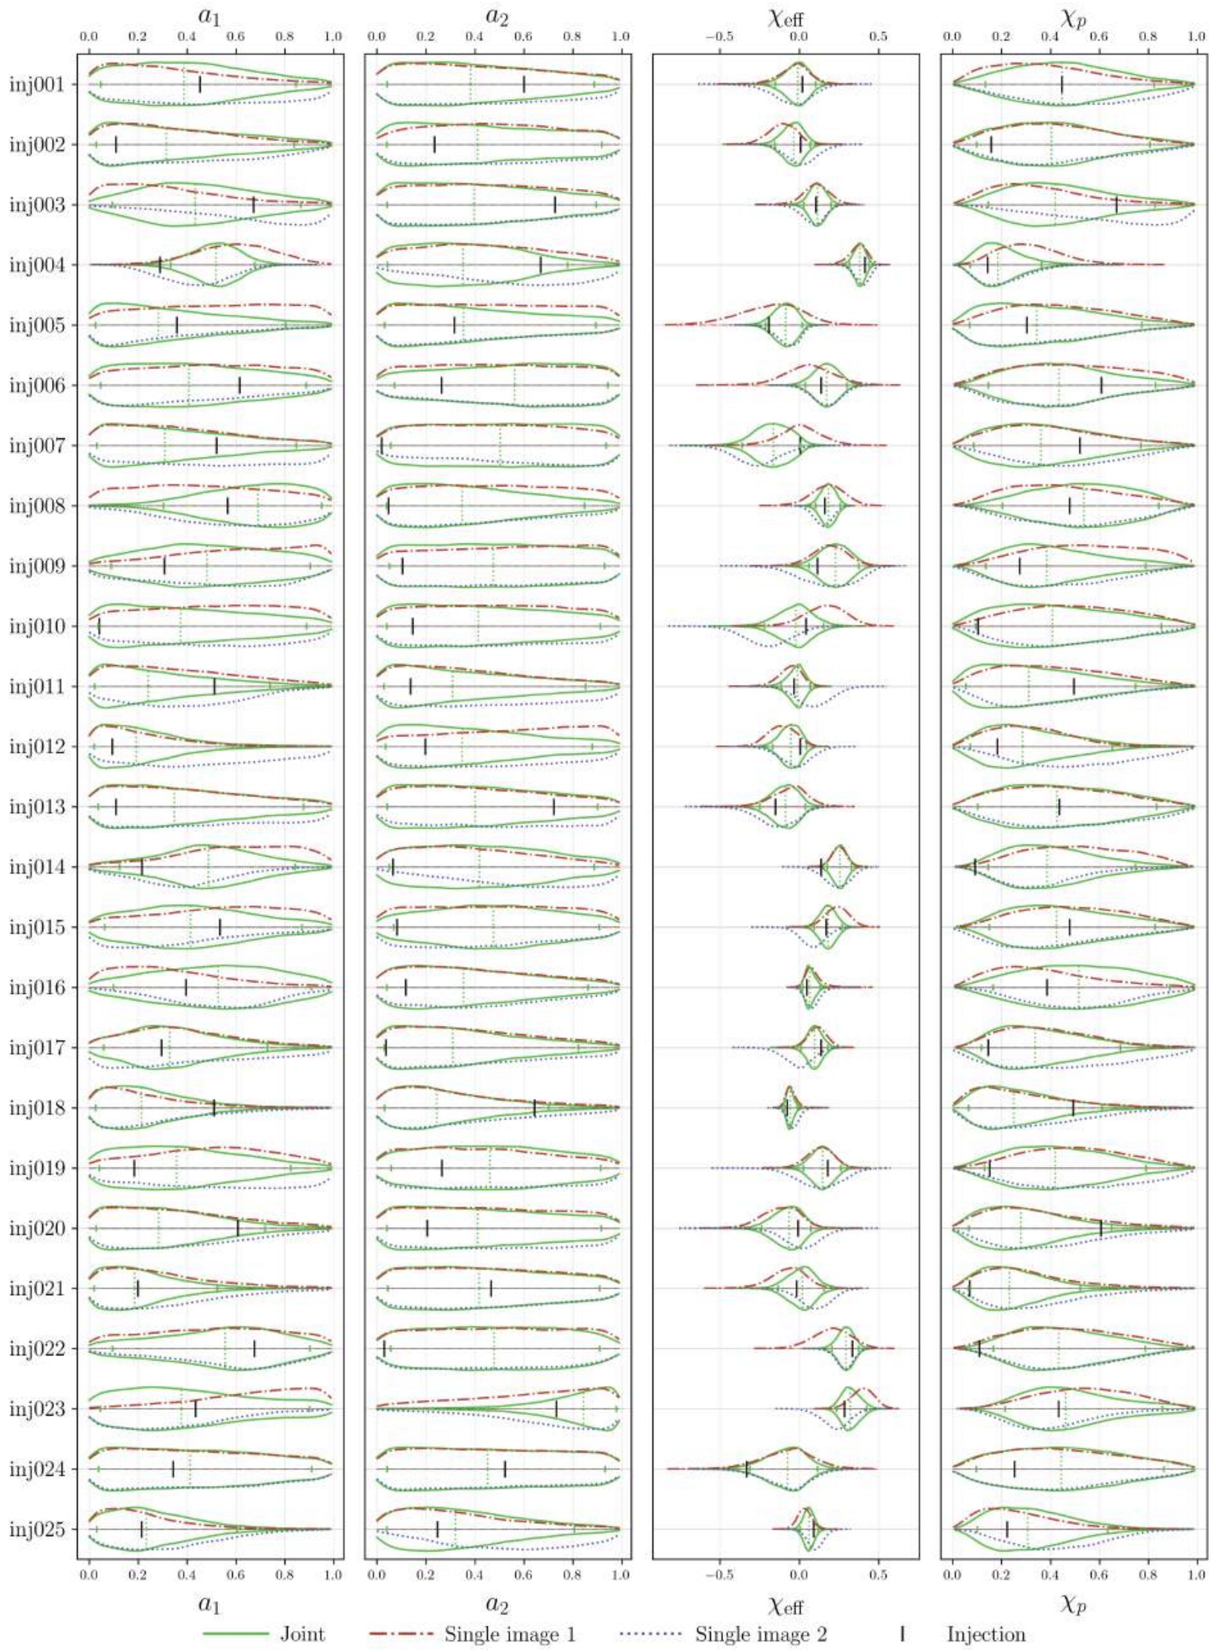

Figure D.2: Recovered posteriors of the spin magnitudes a_1 and a_2, effective spin χ_{eff}, and precessing spin χ_{p}, for all 150 events (1 of 6). Colors as in Figure D.1.

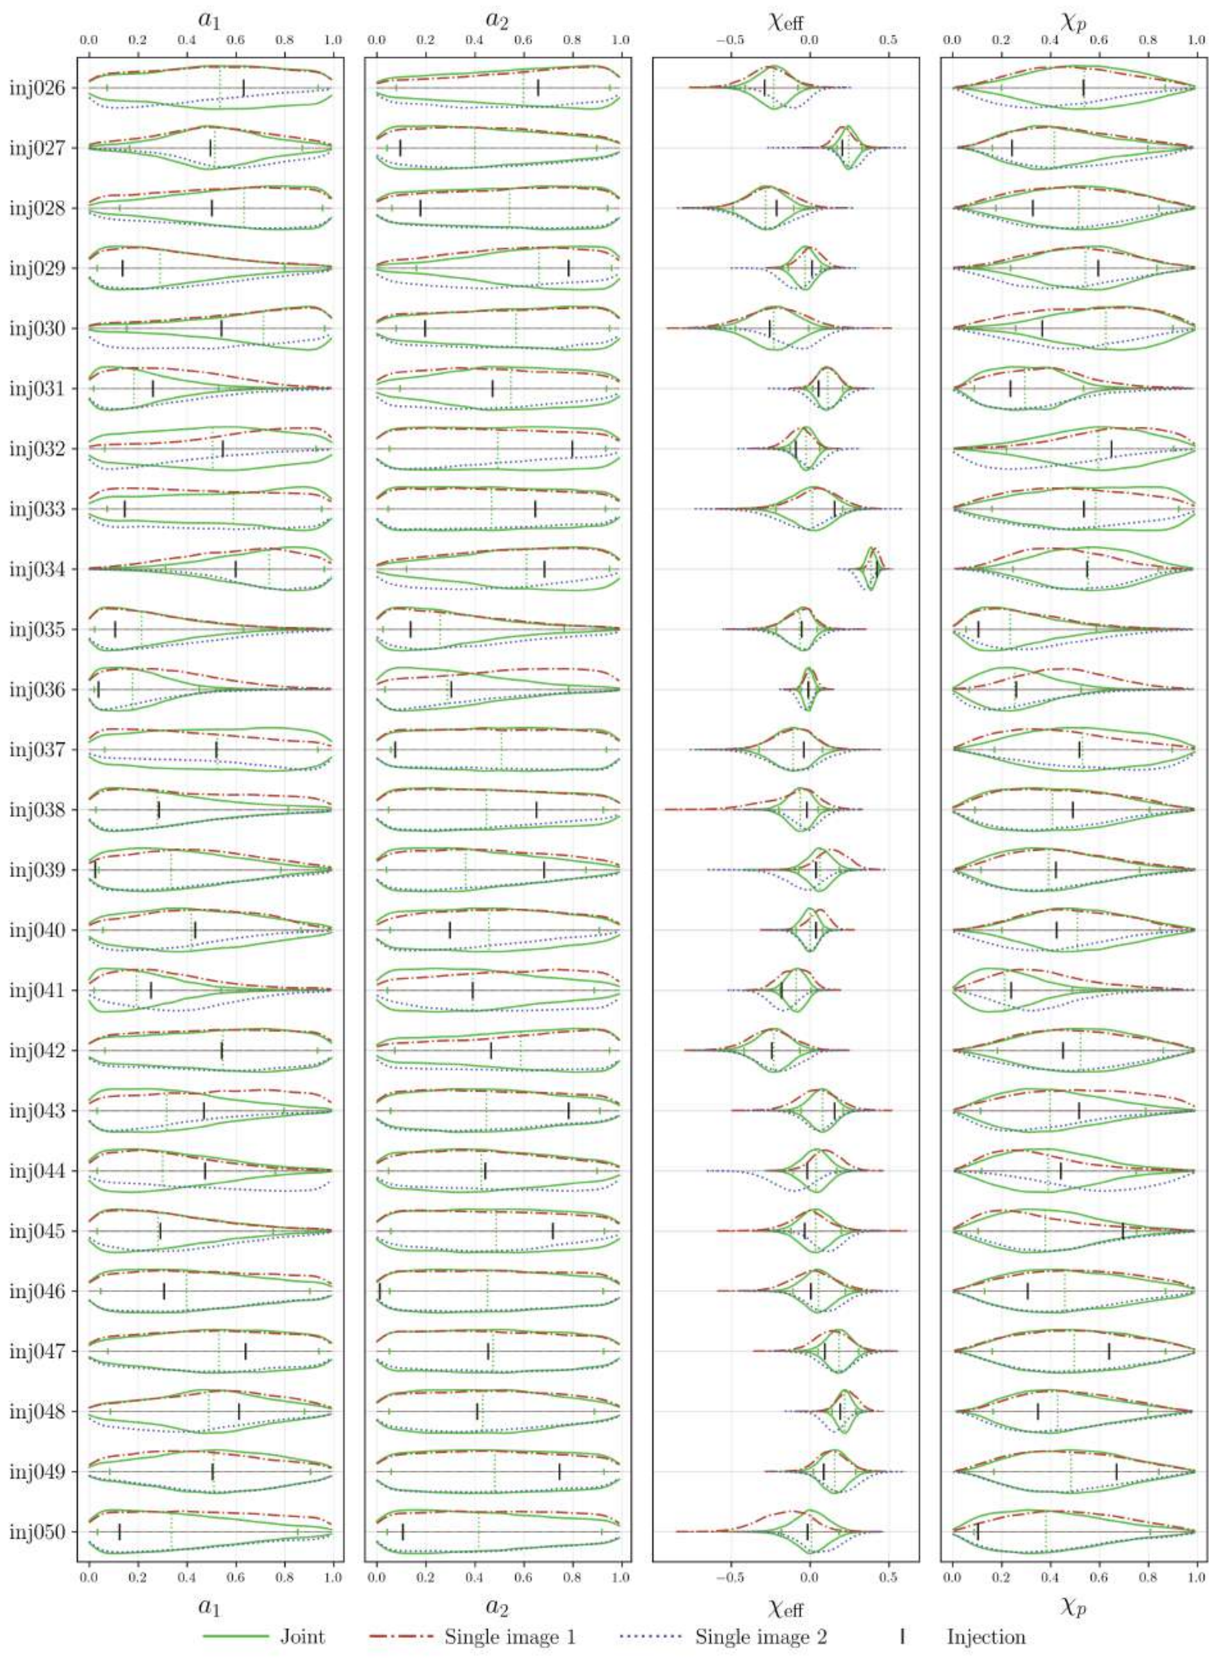

Figure D.2: Recovered posteriors of a_1, a_2, $\chi_{\rm eff}$, and $\chi_{\rm p}$ (2 of 6).

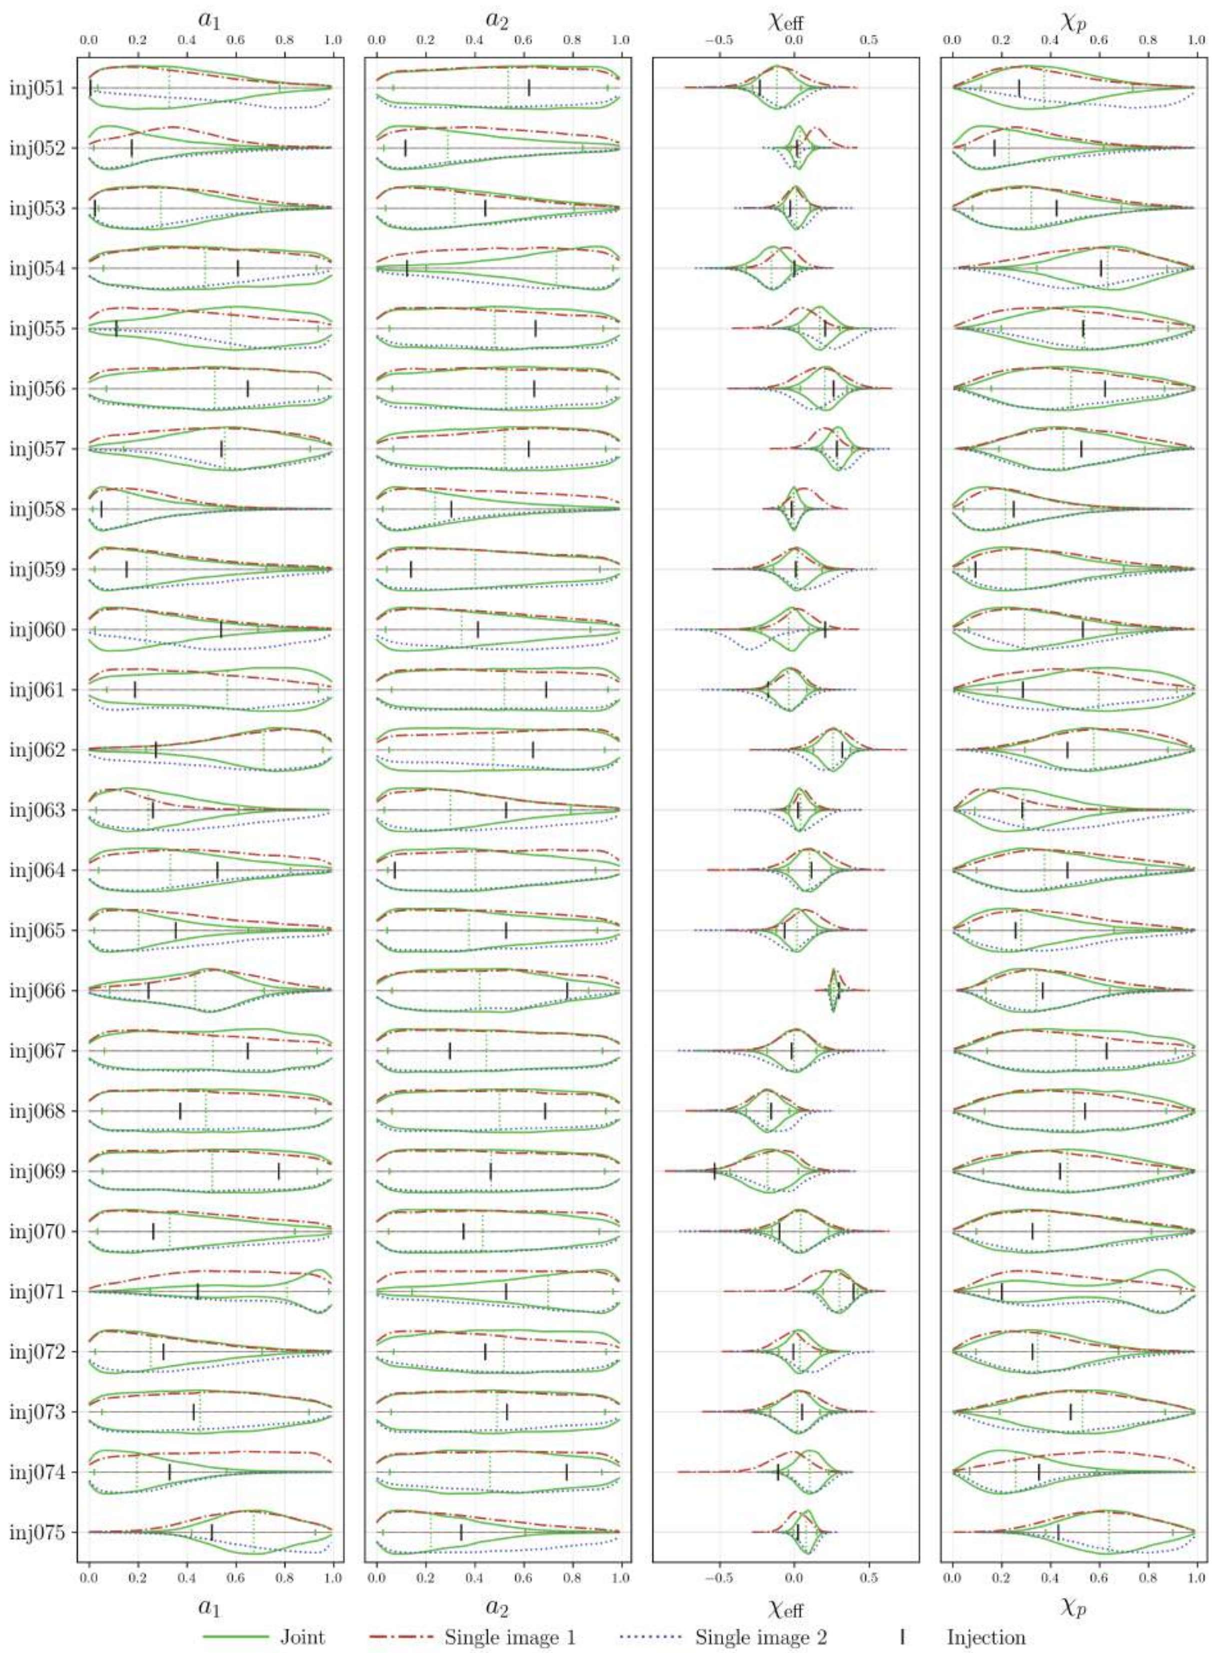

Figure D.2: Recovered posteriors of a_1, a_2, $\chi_{\rm eff}$, and $\chi_{\rm p}$ (3 of 6).

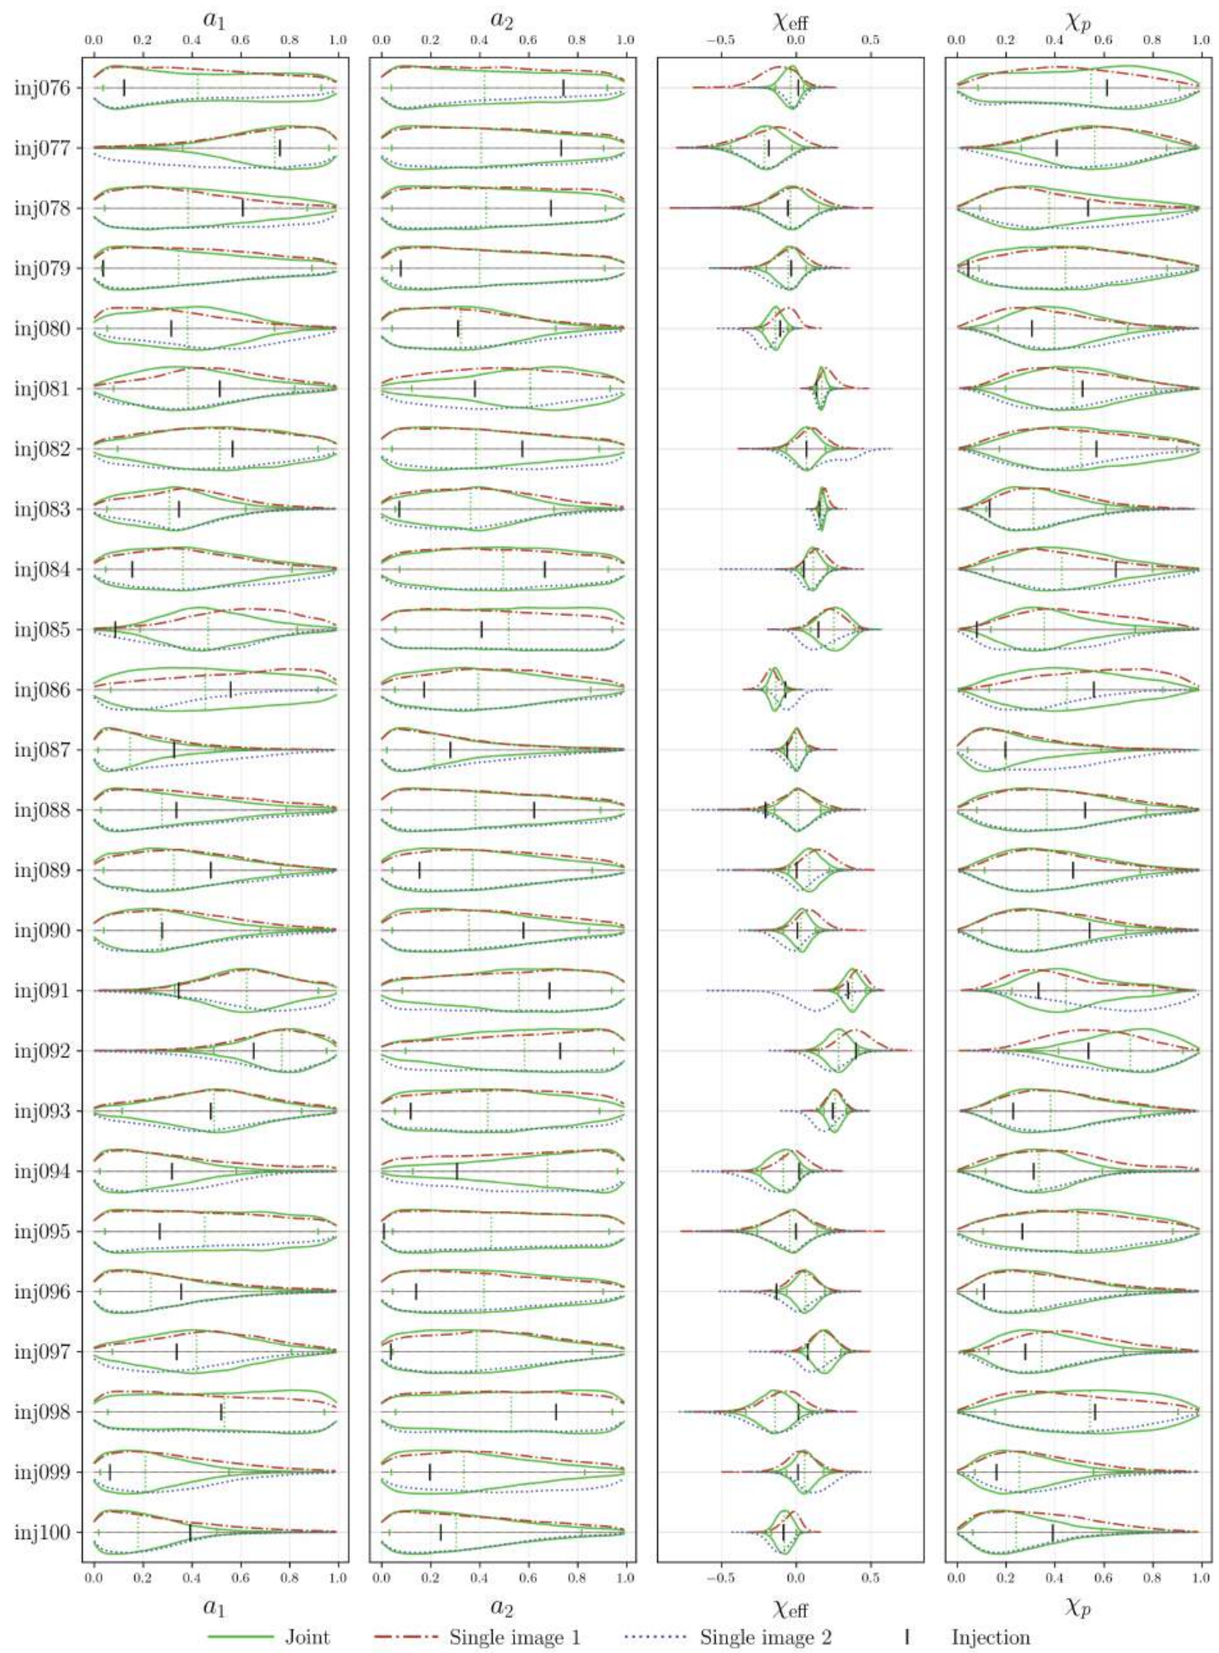

Figure D.2: Recovered posteriors of a_1, a_2, $\chi_{\rm eff}$, and $\chi_{\rm p}$ (4 of 6).

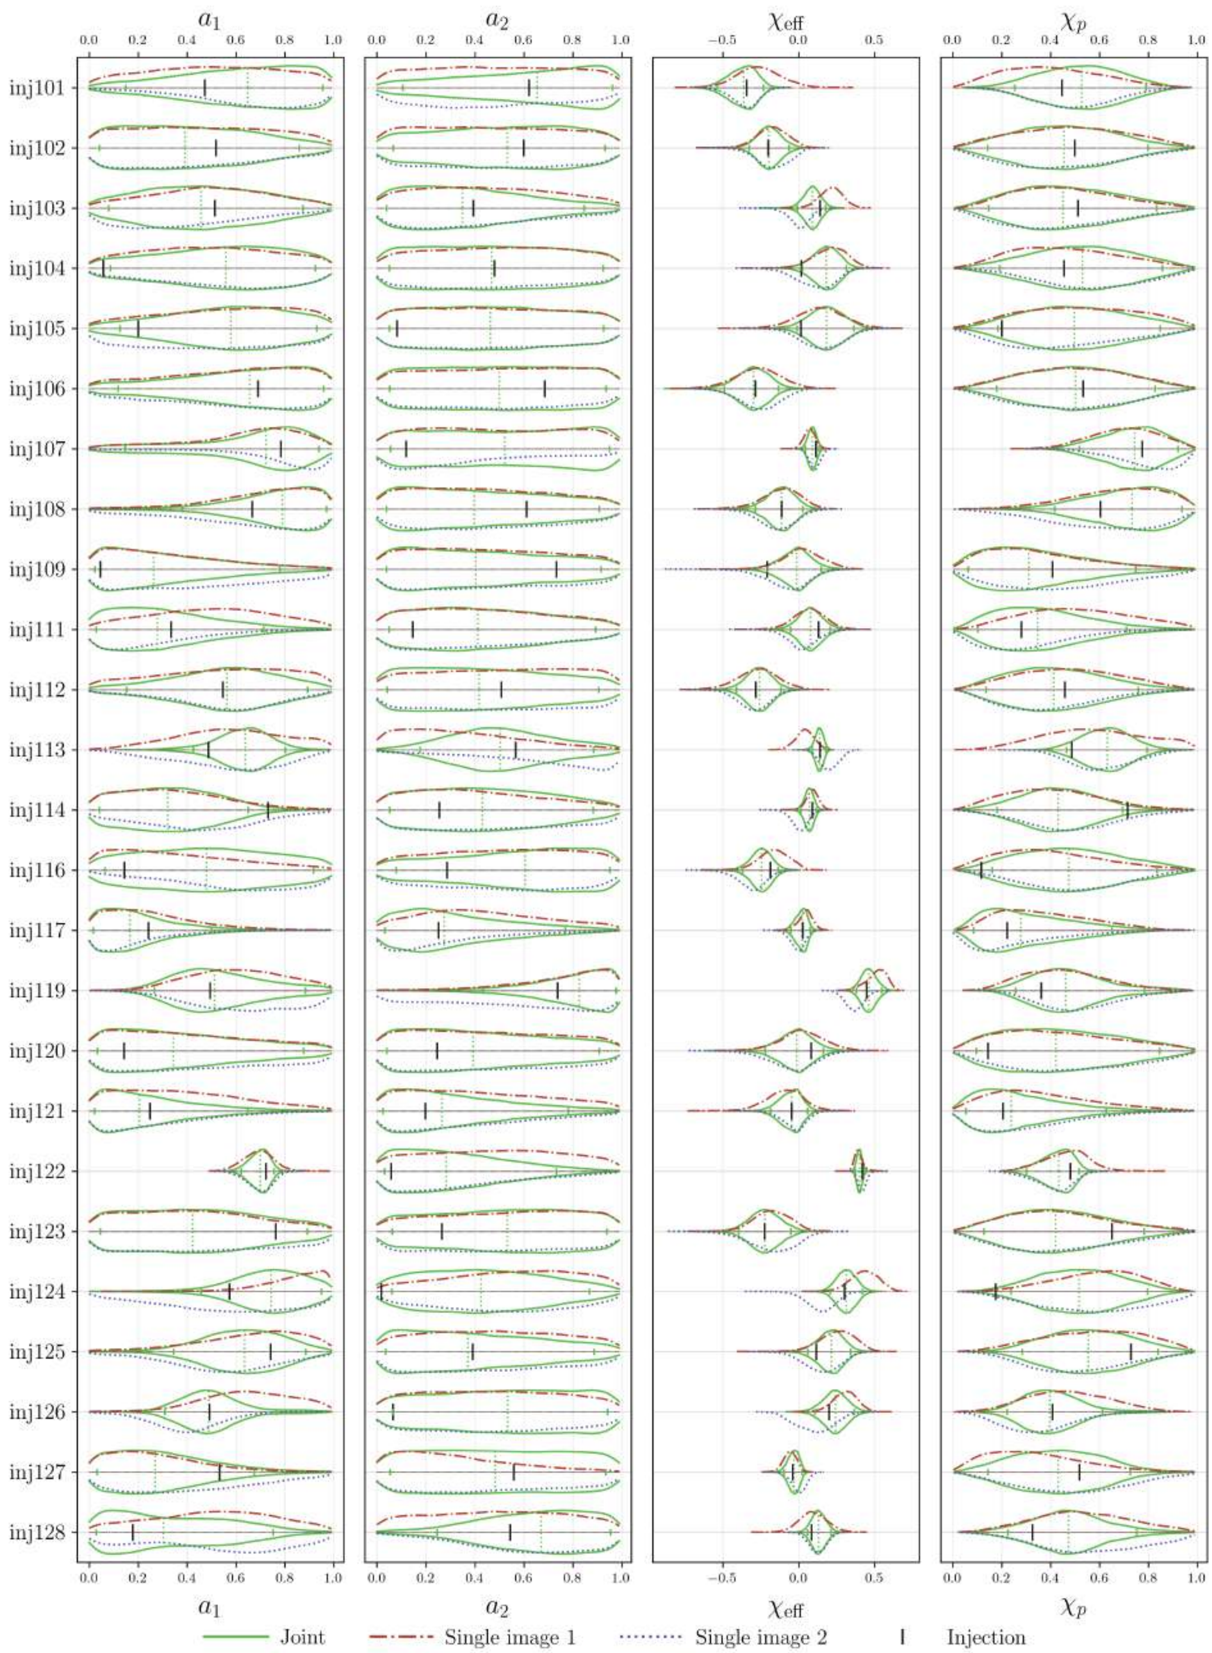

Figure D.2: Recovered posteriors of a_1, a_2, χ_{eff}, and χ_{p} (5 of 6).

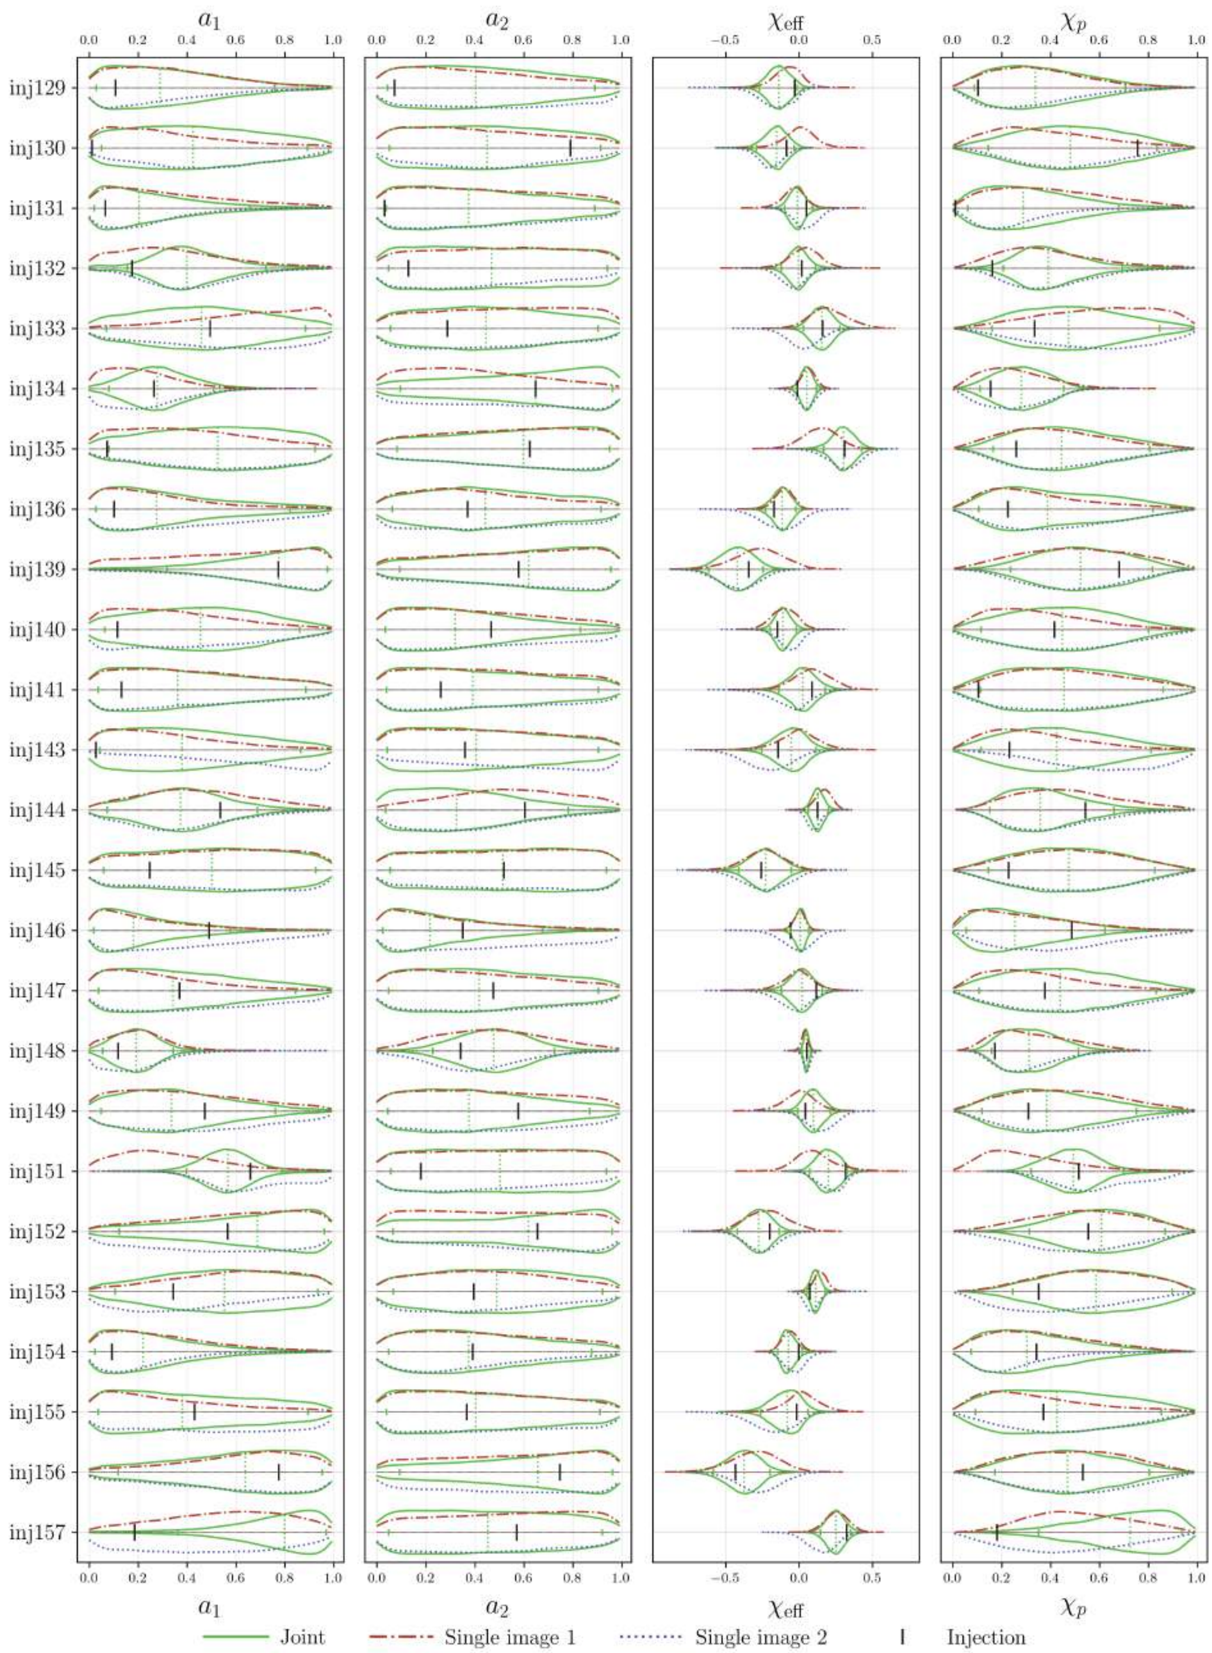

Figure D.2: Recovered posteriors of a_1, a_2, $\chi_{\rm eff}$, and $\chi_{\rm p}$ (6 of 6).